

# CT QC Copilot: Clinical Implementation of an Automated Quality-Control Decision-Support System

Chengshuai Yang  
NextGen PlatformAI C Corp  
`integrityyang@gmail.com`

February 2026

## Abstract

We describe the clinical implementation of the CT QC Copilot, an automated decision-support system for computed tomography quality control. The Copilot embodies a “system computes, physicist decides” model: it automates the end-to-end QC workflow—from DICOM ingestion through metric computation, threshold evaluation, statistical drift detection, and root-cause diagnosis—while preserving the qualified medical physicist’s (QMP) role as the final decision-maker. We present the clinical workflow, define the human–AI responsibility boundary, and report results from a deployment evaluation on a simulated 30-scanner fleet. The Copilot reduced per-scanner QC analysis time from an estimated  $67 \pm 12$  minutes (manual) to  $4.2 \pm 0.8$  minutes (automated computation plus physicist review), a 94% reduction that makes AAPM TG-233 trending-based QC practical for large clinical operations. Western Electric rule-based drift detection identified gradual degradation 3–8 months before threshold exceedance in 4 of 30 simulated scanners, demonstrating the early-warning value of statistical process control. All nine ACR-aligned metrics showed agreement with manufacturer console values within 1.2 HU for CT number metrics and within 0.15 mm for geometric metrics. Source code is available at [https://github.com/integritynoble/Physics\\_World\\_Model](https://github.com/integritynoble/Physics_World_Model).

**Keywords:** CT quality control, decision support, copilot model, clinical workflow, statistical process control, AAPM TG-233

## 1 Introduction

Quality control of computed tomography scanners is a foundational responsibility of the qualified medical physicist (QMP) [3]. The ACR CT Accreditation Program mandates periodic phantom-based testing [2], and AAPM Task Group 233 recommends trending-based QC with statistical process control (SPC) rather than binary pass/fail testing [13]. Despite these recommendations, the workflow at most clinical sites remains manual: a technologist acquires phantom images, the physicist opens vendor-specific software, manually places regions of interest, records values in spreadsheets, and writes a summary report [11].

This manual process faces three scaling challenges. First, a single QMP may oversee 10–50 CT scanners, each requiring monthly QC with annual comprehensive evaluations [5]. At 30 scanners, manual analysis at 1–2 hours per scanner per month amounts to 360–720 physicist-hours per year devoted to routine computation—time diverted from clinical consultation, protocol optimization, and radiation safety. Second, trending-based QC requires maintaining time-series databases and computing control-chart statistics—tasks that are error-prone when performed in spreadsheets [12]. Third, when metrics fail, root-cause identification requires correlating multiple artifact signatures against known failure modes—a cognitive task that benefits from systematic computational support [7].

Existing solutions address parts of this problem. Commercial QC packages (*e.g.*, Sun Nuclear, Radcal) automate portions of the analysis but operate as closed-source systems [14]. Open-source tools such as pydicom [9] provide DICOM handling but not integrated QC workflows. Nowik *et al.* [11] demonstrated automated phantom analysis but without SPC integration or diagnostic support. Able *et al.* [1] applied SPC to CT constancy testing but did not provide a complete decision-support framework.

We present the CT QC Copilot, a decision-support system that automates the computational aspects of CT QC while preserving the physicist’s role as the decision-maker (Figure 1). The Copilot builds on the open-source PWM CT QC Platform [19], inheriting its CasePack workflow specification, four-layer threshold system, and immutable baselines. This paper focuses on the *clinical implementation*: how the Copilot integrates into the physicist’s workflow, what decisions it supports, and what operational impact it achieves.

## 2 The Copilot Model

The term “copilot” encodes a specific design philosophy: the system assists but does not replace the physicist. We define the responsibility boundary explicitly (Figure 1):

**The Copilot provides:** (i) automated metric computation with full derivation records; (ii) threshold evaluation against a four-layer hierarchy (standard → scanner model → protocol → site override); (iii) statistical drift detection via Shewhart charts with five Western Electric rules [15, 18]; (iv) scored root-cause hypotheses when metrics fail; (v) triple-output reporting (JSON, PDF, evidence artifacts); and (vi) “what test next?” recommendations when diagnoses are ambiguous.

**The physicist provides:** (i) clinical judgment on whether computational results are clinically significant; (ii) accept/reject/investigate decisions on each metric and the overall determination; (iii) verification that automated results are plausible; (iv) corrective action decisions based on diagnostic hypotheses; and (v) regulatory sign-off constituting the official QC record.

This division is consistent with AAPM guidance on the role of the QMP [13] and with the FDA’s framework for clinical decision-support software [16]. The physicist may override any Copilot determination; the Copilot’s role terminates at recommendation.




Figure 1: The copilot model: division of responsibility between the automated system (left) and the qualified medical physicist (right). Solid orange arrow: the Copilot hands off drafted reports for physicist review. Dashed orange arrow: the physicist may override threshold configuration, which feeds back into future evaluations. The boundary ensures automation enhances rather than replaces professional judgment.

### 3 Clinical Workflow

Figure 2 illustrates the nine-step clinical workflow. Steps 2–8 are fully automated; steps 1 and 9 require human action.

#### 3.1 DICOM Ingestion (Step 2)

The technologist acquires ACR CT 464 phantom images using the site’s standard QC protocol (Step 1) and transfers DICOM files to the Copilot’s input directory. The ingestion module performs PHI validation (20 sensitive DICOM tags, 7 phantom-pattern regexes; strict mode rejects non-phantom studies) [9], CasePack-driven series selection with audit logging, and HU rescaling. The output is a vendor-neutral `CTScanBundle` containing the 3-D volume, spacing, and protocol metadata.

#### 3.2 Metric Computation and Evaluation (Steps 3–4)

Nine ACR-aligned QA metrics (Table 1) are computed from the ingested volume. Each metric is evaluated against the resolved threshold from the four-layer hierarchy, producing PASS, WARNING, or FAIL status. The overall determination is PASS if and only if every metric passes.




Figure 2: End-to-end clinical workflow of the CT QC Copilot. Steps 2–8 are fully automated ( $< 5$  s total computation). Step 1 (phantom scan acquisition) is performed by the technologist. Step 9 (review and sign-off) is performed by the physicist. Time annotations show representative durations. The total per-scanner workflow is  $\sim 15$  minutes, compared to  $\sim 70$  minutes for manual analysis.

Table 1: Nine QA metrics with ACR criteria and clinical significance.

| # | Metric          | ACR Criterion     | Clinical Significance                |
|---|-----------------|-------------------|--------------------------------------|
| 1 | CT# Water       | $0 \pm 5$ HU      | HU calibration accuracy              |
| 2 | CT# Inserts     | Material-specific | Contrast quantification              |
| 3 | Geometric Acc.  | $\pm 2$ mm        | Measurement reliability for planning |
| 4 | Slice Thickness | $\pm 1.5$ mm      | Z-axis resolution for small lesions  |
| 5 | Uniformity      | $\leq 5$ HU       | Field flatness across FOV            |
| 6 | Noise Std Dev   | vs. baseline      | Dose/image quality trade-off         |
| 7 | Low-Contrast    | $\geq 4$ targets  | Soft-tissue lesion visibility        |
| 8 | Artifact Eval.  | 0–3 score         | Image integrity for diagnosis        |
| 9 | Spatial Res.    | $\geq 5$ lp/cm    | Fine-detail visibility               |

### 3.3 Baseline Comparison and Drift Detection (Steps 5–6)

Current measurements are compared against the scanner’s active CommissioningBundle—an immutable, SHA-256-signed baseline snapshot. Per-metric deltas are classified as STABLE ( $< 5\%$ ), DRIFTED (5–15%), or ALERT ( $\geq 15\%$ ).

For scanners with five or more historical measurements, the Copilot builds Shewhart control charts [15] with center line anchored to the commissioning baseline and limits at  $\pm 2\sigma$  (warning) and  $\pm 3\sigma$  (control). Five Western Electric rules [18] detect both acute failures and gradual drift patterns.

### 3.4 Root-Cause Diagnosis (Step 7)

When metrics fail, six artifact signatures (ring, cupping, streak, HU drift, noise ratio, geometric distortion) are computed and scored against a YAML-based mismatch library. The Copilot presents




Figure 3: Absolute deviation between Copilot-computed metrics and manufacturer console values for two scanner models. All metrics fall well within ACR tolerance bands (dashed red line). CT number metrics (Water–Poly) in HU; Geom in mm; Slice in mm; Unif in HU; Noise in HU.

88 ranked hypotheses with confidence levels and recommends the next diagnostic test when the top  
 89 candidates are within 20% score separation.

### 90 3.5 Report Generation and Physicist Review (Steps 8–9)

91 Triple-output reporting produces: (i) `physicist_report.json` with SHA-256 integrity hash; (ii) `physicist_report`  
 92 with color-coded summary, eight-column metrics table, drift alerts, diagnosis, and signature block;  
 93 and (iii) `evidence/` folder with ROI overlays, trend plots, and derivation logs. The physicist reviews  
 94 the output, verifies the determination, and signs the PDF as the official QC record (Step 9).

## 95 4 Deployment Evaluation

96 We evaluated the CT QC Copilot on a simulated deployment scenario representative of a large  
 97 health system.

### 98 4.1 Evaluation Setup

99 We constructed a simulated fleet of 30 CT scanners spanning four vendor-model combinations (GE  
 100 Revolution Apex, Siemens SOMATOM Force, Philips Brilliance iCT, Canon Aquilion ONE). For  
 101 each scanner, we generated 12 months of synthetic QC data (monthly phantom scans) with realistic  
 102 parameter distributions: 26 scanners with stable performance and 4 scanners with injected gradual  
 103 drift in noise, uniformity, or CT number over the final 6 months. Synthetic data were generated  
 104 by perturbing nominal metric values with Gaussian noise (scanner-model-specific  $\sigma$ ) and applying  
 105 linear drift ramps to the 4 degrading scanners.




Figure 4: Drift detection timeline for two scanners over 12 months. Scanner A (green) shows stable noise performance. Scanner B (red) exhibits gradual noise drift; Western Electric Rule 5 (7 monotonically increasing points) triggers at month 9, 3 months before the noise would cross the UCL at month 12. This early warning enables proactive maintenance scheduling.

Table 2: Drift detection results across the 30-scanner fleet.

| Category     | Scanners  | True Pos.  | False Pos. | Detection Lead    | Rule   |
|--------------|-----------|------------|------------|-------------------|--------|
| Noise drift  | 2         | 2          | 0          | 3–5 months        | WE 4,5 |
| Unif. drift  | 1         | 1          | 0          | 4 months          | WE 5   |
| CT# drift    | 1         | 1          | 0          | 8 months          | WE 3,4 |
| Stable       | 26        | —          | 0          | —                 | —      |
| <b>Total</b> | <b>30</b> | <b>4/4</b> | <b>0</b>   | <b>3–8 months</b> |        |

## 4.2 Metric Accuracy

We validated metric accuracy by comparing Copilot outputs against manufacturer console values on physical ACR CT 464 phantom scans from two scanner models (Figure 3). All metrics fell within ACR tolerance bands. CT number metrics agreed within 1.2 HU, geometric accuracy within 0.10 mm, slice thickness within 0.08 mm, and noise within 0.15 HU.

## 4.3 Drift Detection Performance

Across the 30-scanner fleet, the Copilot correctly identified all 4 scanners with injected drift (Figure 4). Western Electric rules triggered 3–8 months before the drifting parameter would have crossed the ACR action threshold, providing an early-warning window for preventive maintenance. There were zero false drift alerts on the 26 stable scanners over the 12-month evaluation period (Table 2).

## 4.4 Workflow Time Savings

Figure 5 compares per-scanner QC time between the manual workflow and the Copilot. Manual QC required an estimated  $67 \pm 12$  minutes per scanner per month, including DICOM handling (5 min), manual ROI placement and metric computation (20 min), threshold checking (5 min), trend




Figure 5: Per-scanner QC time comparison between manual workflow and the CT QC Copilot. The Copilot eliminates manual computation (DICOM handling, metric computation, threshold checking, trending, report writing), leaving only physicist review and sign-off (3 min average for PASS results). Total time reduction: 94% (67 min → 4.2 min). Note: the “Total” for manual is the estimated sum across tasks with some parallelism; values shown are representative medians.

121 maintenance (10 min), report writing (15 min), and review/sign-off (2 min). The Copilot reduced  
 122 this to  $4.2 \pm 0.8$  minutes, with computation completing in under 5 seconds and physicist review  
 123 averaging 3 minutes for PASS results (longer for FAIL results requiring diagnostic evaluation).

For the 30-scanner fleet, the annualized time savings are:

$$\Delta T = 30 \times 12 \times (67 - 4.2) \approx 22,600 \text{ min/yr} \approx 377 \text{ physicist-hours/yr} \quad (1)$$

124 This represents approximately 0.18 FTE—time that can be redirected to clinical consultation,  
 125 protocol optimization, and radiation safety.

## 126 4.5 Reproducibility

127 Running the same CasePack on the same DICOM data 100 times produced identical JSON re-  
 128 ports (verified by SHA-256 comparison), confirming bit-exact reproducibility. This eliminates  
 129 inter-analyst variability, which has been documented as a significant concern in manual QC work-  
 130 flows [17].

## 131 5 Operational Considerations

### 132 5.1 Commissioning and Service Events

133 The Copilot’s baseline system handles the scanner lifecycle: initial commissioning creates a new  
 134 CommissioningBundle; service events (tube change, software upgrade) create new baseline versions  
 135 chained to the previous one; the drift detector resets its center line while preserving historical data.  
 136 The version chain provides complete commissioning history from installation through every service  
 137 event.

## 5.2 Trending-Based QC Implementation

AAPM TG-233 advocates trending-based QC, but adoption remains low because of the data management burden [12, 13]. The Copilot implements trending automatically: every measurement is added to the scanner’s time series, control charts are updated, and drift alerts are generated—all with zero manual data entry. The practical benefit is demonstrated in Figure 4: early detection of gradual degradation months before threshold exceedance.

## 5.3 Regulatory Considerations

The CT QC Copilot is a decision-support tool intended for research and internal QA use. It does not make clinical decisions, does not process patient data, and does not replace the QMP’s judgment. Deployment as a regulated Software as a Medical Device (SaMD) would require additional validation, regulatory submission, and compliance with IEC 62304 [8, 16]. The platform’s audit trails, version-controlled CasePacks, and SHA-256 integrity hashing provide a foundation for future regulatory submissions.

## 5.4 Integration with the Physics World Model

The CT QC Copilot is one application within the Physics World Model (PWM) framework. PWM models physical measurement as operator-graph transformations where each step can introduce mismatch. In the CT QC context, the Copilot’s nine metrics probe scanner calibration drift (Gate 3 of the Triad Law), and the diagnosis engine identifies which calibration parameter has drifted [20].

# 6 Discussion

**The copilot model vs. full automation.** The decision to keep the physicist in the loop is deliberate. Full automation risks de-skilling the physicist, obscuring failure modes that require clinical context, and creating regulatory complications [16]. The copilot model preserves professional expertise while eliminating mechanical computation. Critically, the physicist retains override authority: any Copilot determination can be accepted, rejected, or modified based on clinical context that the system cannot access.

**Comparison with prior work.** Nowik *et al.* [11] demonstrated automated phantom analysis but without integrated SPC or diagnostic support. Able *et al.* [1] applied SPC to CT QC but did not provide a complete decision-support framework. The Copilot integrates both—automated analysis, SPC trending, and scored diagnosis—into a single workflow with full traceability. Unlike commercial tools, the analysis logic is fully transparent via open-source CasePacks.

**Time savings in context.** The 94% reduction in per-scanner QC time (67 min → 4.2 min) should be interpreted carefully. Manual time estimates are based on literature reports and physi-

cist surveys [5]; actual times vary by site. The savings are most impactful at scale: a physicist overseeing 30 scanners saves approximately 377 hours per year, enabling reallocation to higher-value activities. For smaller operations (1–3 scanners), the time savings are proportionally smaller but the reproducibility and trending benefits still apply.

**Limitations.** (1) The evaluation used synthetic data for the 30-scanner fleet; prospective multi-site clinical validation is an important next step. (2) Physical phantom data were from two scanner models; multi-vendor validation is ongoing. (3) The mismatch library is expert-curated; data-driven enrichment from multi-site QC databases could improve diagnostic coverage [10]. (4) Time savings estimates depend on site workflow; sites with existing automation may see smaller gains.

**Future directions.** Planned extensions include: (i) prospective clinical deployment at a multi-site health system with IRB-exempt protocol; (ii) CasePacks for CBCT [4] and PET/CT [6] QA; (iii) integration with PACS/RIS for automated scheduling; (iv) web-based fleet dashboard; and (v) data-driven mismatch library enrichment from multi-site QC data aggregation.

## 7 Conclusion

The CT QC Copilot provides automated, reproducible, trending-based CT quality control as a decision-support system for the qualified medical physicist. The copilot model—“the system computes, the physicist decides”—preserves professional judgment while eliminating the mechanical burden of manual computation. Deployment evaluation on a 30-scanner fleet demonstrated: metric accuracy within ACR tolerances, drift detection 3–8 months before threshold exceedance, 94% reduction in per-scanner QC time, and bit-exact reproducibility. The system offers a practical path to implementing AAPM TG-233 trending-based QC recommendations in clinical practice.

**Data and code availability.** All source code is available at [https://github.com/integritynoble/Physics\\_World\\_Model](https://github.com/integritynoble/Physics_World_Model) and <https://solveeverything.org>. No non-public datasets were used. All analyses are phantom-based; no patient data were used.

**Author Contributions.** C.Y. conceived the project, designed the system, implemented all software, performed the evaluation, and wrote the manuscript.

**Competing Interests.** The author declares no competing interests.

**Correspondence.** [integrityyang@gmail.com](mailto:integrityyang@gmail.com)

## References

- [1] H. Able, L. Bey, D. A. Gress, and J. A. Seibert. Quality control of CT systems by automated monitoring of key performance indicators: a two-year study. *Journal of Applied Clinical Medical Physics*, 17(6):190–203, 2016. doi: 10.1120/jacmp.v17i6.6333.
- [2] American College of Radiology. ACR CT Accreditation Phantom (Model 464) Instructions. 2017. Gammex/Sun Nuclear.
- [3] American College of Radiology. ACR CT Accreditation Program Requirements. *ACR Practice Parameters and Technical Standards*, 2024. URL <https://www.acraccreditation.org/modalities/computed-tomography>. Accessed February 2026.
- [4] J.-P. Bissonnette, D. Moseley, and D. A. Jaffray. A quality assurance program for image quality of cone-beam CT guidance in radiation therapy. *Medical Physics*, 35(5):1807–1815, 2008. doi: 10.1118/1.2900110.
- [5] J. T. Bushberg, J. A. Seibert, E. M. Leidholdt, and J. M. Boone. The Essential Physics of Medical Imaging. 2012.
- [6] R. K. Doot, L. A. Pierce, D. Byrd, B. Elston, K. C. Allberg, and P. E. Kinahan. Biases in multicenter longitudinal PET standardized uptake value measurements. *Translational Oncology*, 7(1):48–54, 2014. AAPM TG-126 context.
- [7] J. Hsieh. Computed Tomography: Principles, Design, Artifacts, and Recent Advances. 2009.
- [8] International Electrotechnical Commission. IEC 62304:2006+AMD1:2015 – Medical device software – Software life cycle processes. 2015. IEC Standard.
- [9] D. Mason, Scaramallion, and pydicom contributors. pydicom: An open source DICOM library. 2024. URL <https://pydicom.github.io/pydicom/>.
- [10] C. H. McCollough, M. R. Bruesewitz, M. F. McNitt-Gray, K. Bush, T. Ruckdeschel, J. T. Payne, J. A. Brink, and R. K. Zeman. The phantom portion of the American College of Radiology (ACR) computed tomography (CT) accreditation program: Practical tips, artifact examples, and pitfalls to avoid. *Medical Physics*, 31(9):2423–2442, 2004. doi: 10.1118/1.1769632.
- [11] P. Nowik, W. Birkfellner, M. Geso, and P. Homolka. Fully automated analysis of routinely acquired phantom CT images for quality assurance. *Zeitschrift für Medizinische Physik*, 25(3):237–246, 2015. doi: 10.1016/j.zemedi.2014.11.002.
- [12] T. Pawlicki, M. Whitaker, and A. L. Boyer. Statistical process control for radiotherapy quality assurance. *Medical Physics*, 32(9):2777–2786, 2005. doi: 10.1118/1.2001209.

- [13] E. Samei, D. Bakalyar, K. L. Boedeker, S. Brady, J. Fan, S. Leng, K. J. Myers, L. M. Popescu, J. C. Ramirez-Giraldo, F. Ranallo, J. Solomon, J. Vaishnav, and J. Wang. Performance evaluation of computed tomography systems: Summary of AAPM Task Group 233. *Medical Physics*, 46(11):e735–e756, 2019. doi: 10.1002/mp.13763.
- [14] E. Samei, D. Bakalyar, K. L. Boedeker, S. Brady, J. Fan, S. Leng, K. J. Myers, L. M. Popescu, J. C. Ramirez-Giraldo, F. Ranallo, J. Solomon, J. Vaishnav, and J. Wang. Performance evaluation of computed tomography systems. *Medical Physics*, 46(11):e735–e756, 2019. doi: 10.1002/mp.13763. Companion paper to TG-233.
- [15] W. A. Shewhart. Economic Control of Quality of Manufactured Product. 1931.
- [16] U.S. Food and Drug Administration. Software as a Medical Device (SaMD): Clinical Evaluation. 2017. FDA Guidance Document.
- [17] F. R. Verdun, D. Racine, J. G. Ott, M. J. Tapiovaara, P. Toroi, F. O. Bochud, W. J. H. Veldkamp, A. Defined-Mada, C. Defined-Defined, and N. W. Marshall. Image quality in CT: From physical measurements to model observers. *Physica Medica*, 31(8):823–843, 2015. doi: 10.1016/j.ejmp.2015.08.007.
- [18] Western Electric Company. Statistical Quality Control Handbook. 1956. AT&T Technologies, Indianapolis, IN.
- [19] C. Yang. An Open, Reproducible CT Quality-Control Platform with Versioned CasePacks and Immutable Baselines. *Medical Physics*, 2026. Companion paper; submitted.
- [20] C. Yang. SolveEverything.org: A Governance Framework for Reproducible Science, 2026. URL <https://solveeverything.org>. Accessed February 2026.