

# InverseNet: Supplementary Material

Anonymous ECCV submission

Anonymous

## 1 Per-Scene Detailed Results

### 1.1 CASSI Per-Scene PSNR

Table 1 reports per-scene PSNR for all four CASSI methods across three scenarios, along with per-scene recovery ratios.

Table 1: CASSI per-scene PSNR (dB) and recovery ratio  $\rho$  (%). 10 KAIST scenes,  $256 \times 256 \times 28$ . 5-parameter mismatch ( $dx=0.5$  px,  $dy=0.3$  px,  $\theta=0.1^\circ$ ,  $a_1=2.02$ ,  $\alpha=0.15^\circ$ ).

| Scene    | GAP-TV |       |       |        | HDNet |       |       |        | MST-S |       |       |        | MST-L |       |       |              |
|----------|--------|-------|-------|--------|-------|-------|-------|--------|-------|-------|-------|--------|-------|-------|-------|--------------|
|          | I      | II    | III   | $\rho$ | I     | II    | III   | $\rho$ | I     | II    | III   | $\rho$ | I     | II    | III   | $\rho$       |
| Scene 1  | 26.49  | 24.08 | 24.18 | 4.1%   | 34.95 | 24.37 | 24.37 | 0.0%   | 34.73 | 24.03 | 29.69 | 52.9%  | 35.29 | 23.96 | 29.98 | 53.2%        |
| Scene 2  | 24.60  | 21.89 | 22.82 | 34.5%  | 35.65 | 23.26 | 23.26 | 0.0%   | 34.59 | 22.53 | 27.06 | 37.5%  | 36.14 | 22.21 | 28.42 | 44.6%        |
| Scene 3  | 25.96  | 18.62 | 19.68 | 14.4%  | 35.54 | 18.61 | 18.61 | 0.0%   | 34.34 | 16.26 | 22.31 | 33.5%  | 35.66 | 16.09 | 23.57 | 38.2%        |
| Scene 4  | 28.36  | 23.37 | 24.41 | 20.9%  | 41.63 | 23.98 | 23.98 | 0.0%   | 40.75 | 22.19 | 27.92 | 30.9%  | 40.05 | 21.91 | 29.80 | 43.5%        |
| Scene 5  | 23.66  | 20.39 | 21.27 | 26.9%  | 32.56 | 20.22 | 20.22 | 0.0%   | 32.15 | 20.02 | 25.92 | 48.6%  | 32.84 | 20.28 | 26.75 | 51.5%        |
| Scene 6  | 22.34  | 20.39 | 21.00 | 31.2%  | 34.33 | 22.63 | 22.63 | 0.0%   | 33.64 | 22.76 | 27.36 | 42.3%  | 34.56 | 22.37 | 28.67 | 51.7%        |
| Scene 7  | 23.51  | 20.56 | 21.07 | 17.4%  | 33.27 | 20.79 | 20.79 | 0.0%   | 32.56 | 19.94 | 24.93 | 39.6%  | 33.80 | 19.76 | 26.12 | 45.3%        |
| Scene 8  | 22.16  | 20.57 | 21.10 | 33.3%  | 32.26 | 22.73 | 22.73 | 0.0%   | 31.74 | 21.96 | 26.72 | 48.7%  | 32.74 | 21.33 | 27.44 | 53.5%        |
| Scene 9  | 23.03  | 19.14 | 20.61 | 37.8%  | 34.18 | 21.18 | 21.18 | 0.0%   | 33.54 | 20.04 | 25.33 | 39.2%  | 34.37 | 19.75 | 26.71 | 47.6%        |
| Scene 10 | 23.27  | 20.58 | 21.04 | 16.9%  | 32.22 | 21.06 | 21.06 | 0.0%   | 31.79 | 20.19 | 25.59 | 46.5%  | 32.63 | 20.69 | 25.86 | 43.3%        |
| Mean     | 24.34  | 20.96 | 21.72 | 22.5%  | 34.66 | 21.88 | 21.88 | 0.0%   | 33.98 | 20.99 | 26.28 | 40.7%  | 34.81 | 20.83 | 27.33 | <b>46.5%</b> |

### 1.2 CACTI Per-Video PSNR

Table 2 reports per-video PSNR for all four CACTI methods.

### 1.3 SPC Per-Image PSNR

Table 3 reports per-image PSNR for all three SPC methods. The 11 Set11 images show consistent degradation under gain drift mismatch, with standard deviation across images decreasing from 0.95–3.38 dB (Scenario I) to 0.59–0.69 dB (Scenario II), confirming the mismatch-induced performance floor.

Table 2: CACTI per-video PSNR (dB) and recovery ratio  $\rho$  (%). 6 benchmark videos,  $256 \times 256 \times 8$ , 8-parameter mismatch.

| Video       | GAP-TV |       |       |              | PnP-FFDNet |       |       |        | ELP-Unfolding |       |       |        | EfficientSCI |       |       |        |
|-------------|--------|-------|-------|--------------|------------|-------|-------|--------|---------------|-------|-------|--------|--------------|-------|-------|--------|
|             | I      | II    | III   | $\rho$       | I          | II    | III   | $\rho$ | I             | II    | III   | $\rho$ | I            | II    | III   | $\rho$ |
| Kobe        | 26.70  | 18.97 | 26.50 | 97.4%        | 30.02      | 16.00 | 29.20 | 94.2%  | 34.07         | 18.31 | 32.63 | 90.9%  | 35.55        | 18.21 | 32.43 | 82.0%  |
| Traffic     | 20.73  | 13.99 | 20.57 | 97.6%        | 24.06      | 9.95  | 23.37 | 95.1%  | 31.33         | 13.90 | 29.04 | 86.8%  | 32.19        | 13.30 | 27.08 | 72.9%  |
| Runner      | 29.34  | 17.70 | 28.85 | 95.8%        | 32.88      | 13.40 | 31.18 | 91.3%  | 38.14         | 17.00 | 34.06 | 80.7%  | 39.28        | 16.65 | 31.15 | 65.5%  |
| Drop        | 34.22  | 13.64 | 31.43 | 86.4%        | 38.71      | 7.55  | 21.91 | 46.1%  | 40.08         | 13.52 | 25.12 | 43.7%  | 42.36        | 11.63 | 21.95 | 33.6%  |
| Crash       | 24.80  | 14.77 | 24.45 | 96.5%        | 24.81      | 10.11 | 23.32 | 89.9%  | 29.38         | 14.82 | 26.84 | 82.5%  | 30.62        | 14.25 | 25.07 | 66.1%  |
| Aerial      | 25.22  | 16.76 | 24.95 | 96.8%        | 24.56      | 12.79 | 23.77 | 93.3%  | 30.43         | 16.14 | 28.80 | 88.5%  | 31.24        | 15.92 | 27.18 | 73.5%  |
| <b>Mean</b> | 26.75  | 15.81 | 26.01 | <b>93.3%</b> | 29.28      | 11.43 | 25.39 | 78.2%  | 34.09         | 15.47 | 29.40 | 74.8%  | 35.39        | 14.81 | 27.38 | 61.1%  |

Table 3: SPC per-image PSNR (dB) and recovery ratio  $\rho$  (%). 11 Set11 images,  $256 \times 256$ , 25% sampling. Gain drift mismatch ( $\alpha=0.0015$ ,  $\sigma_y=0.03$ ).

| Image       | FISTA-TV |       |       |        | ISTA-Net |       |       |        | HATNet |       |       |              |
|-------------|----------|-------|-------|--------|----------|-------|-------|--------|--------|-------|-------|--------------|
|             | I        | II    | III   | $\rho$ | I        | II    | III   | $\rho$ | I      | II    | III   | $\rho$       |
| Monarch     | 28.04    | 19.22 | 26.15 | 78.6%  | 32.54    | 19.90 | 27.70 | 61.7%  | 30.91  | 20.36 | 29.77 | 89.2%        |
| Parrots     | 27.40    | 18.53 | 26.18 | 86.3%  | 31.42    | 18.99 | 27.82 | 71.0%  | 31.63  | 19.51 | 30.42 | 90.1%        |
| barbara     | 24.48    | 18.49 | 23.79 | 88.4%  | 27.84    | 19.02 | 25.53 | 73.7%  | 30.72  | 19.55 | 29.45 | 88.6%        |
| boats       | 28.79    | 18.86 | 26.85 | 80.4%  | 32.91    | 19.26 | 27.89 | 63.2%  | 31.04  | 19.58 | 29.76 | 88.9%        |
| cameraman   | 26.04    | 18.78 | 24.96 | 85.1%  | 28.61    | 19.31 | 26.16 | 73.7%  | 29.68  | 19.77 | 28.68 | 89.9%        |
| fingerprint | 23.08    | 17.33 | 22.58 | 91.2%  | 28.10    | 18.16 | 25.56 | 74.5%  | 30.11  | 18.84 | 29.15 | 91.5%        |
| flinstones  | 24.64    | 17.18 | 23.59 | 86.0%  | 29.37    | 18.01 | 26.39 | 73.8%  | 29.48  | 18.49 | 28.52 | 91.2%        |
| foreman     | 35.10    | 17.96 | 30.42 | 72.7%  | 38.23    | 18.20 | 29.68 | 57.3%  | 32.80  | 18.34 | 31.31 | 89.7%        |
| house       | 32.20    | 18.90 | 29.20 | 77.4%  | 35.70    | 19.23 | 29.21 | 60.6%  | 32.17  | 19.34 | 30.80 | 89.4%        |
| lena256     | 28.97    | 19.25 | 27.13 | 81.1%  | 32.30    | 19.67 | 27.96 | 65.7%  | 31.22  | 19.96 | 29.89 | 88.2%        |
| peppers256  | 29.95    | 19.11 | 27.51 | 77.5%  | 33.33    | 19.49 | 28.01 | 61.6%  | 31.05  | 19.68 | 29.87 | 89.6%        |
| <b>Mean</b> | 28.06    | 18.51 | 26.21 | 80.7%  | 31.85    | 19.02 | 27.45 | 65.7%  | 30.98  | 19.40 | 29.78 | <b>89.6%</b> |

## 2 Real Hardware Per-Scene Results

### 2.1 CASSI Real Data Per-Scene PSNR

Table 4 reports per-scene PSNR for real CASSI data from the TSA real dataset (Meng et al., ECCV 2020). Reconstructions use the hardware-provided mask (calibrated) or a shifted mask ( $dx=0.5$ ,  $dy=0.3$  px) to simulate operator mismatch. PSNR is computed against reference reconstructions distributed with the dataset.

### 2.2 CACTI Real Data Per-Scene Residuals

Table 5 reports per-scene measurement residuals for real CACTI data from the EfficientSCI dataset (Wang et al., CVPR 2023) (cr=10). Since no ground truth exists, we report the relative measurement residual  $\|y - \Phi\hat{x}\|^2 / \|y\|^2$  (lower is better) and the total variation of reconstructions.

Table 4: CASSI real data per-scene PSNR (dB). 5 real scenes ( $660 \times 660 \times 28$ ). Calibrated uses the hardware mask; mismatched applies subpixel shift ( $dx=0.5$ ,  $dy=0.3$  px).

| Scene       | GAP-TV |       |          | MST-S |       |          | MST-L |       |          |
|-------------|--------|-------|----------|-------|-------|----------|-------|-------|----------|
|             | Cal.   | Mis.  | $\Delta$ | Cal.  | Mis.  | $\Delta$ | Cal.  | Mis.  | $\Delta$ |
| Scene 1     | 23.83  | 24.18 | -0.35    | 19.39 | 20.27 | -0.88    | 18.74 | 18.82 | -0.08    |
| Scene 2     | 21.40  | 21.52 | -0.12    | 20.85 | 20.82 | 0.03     | 20.40 | 20.37 | 0.03     |
| Scene 3     | 20.48  | 20.61 | -0.13    | 18.23 | 18.27 | -0.04    | 18.09 | 18.15 | -0.06    |
| Scene 4     | 21.11  | 21.20 | -0.09    | 16.05 | 16.13 | -0.08    | 15.76 | 15.70 | 0.06     |
| Scene 5     | 21.61  | 21.71 | -0.10    | 16.78 | 16.93 | -0.15    | 16.00 | 15.97 | 0.03     |
| <b>Mean</b> | 21.69  | 21.84 | -0.16    | 18.26 | 18.48 | -0.22    | 17.80 | 17.80 | -0.00    |

### 3 Scenario IV Per-Scene Results

Table 6 reports per-scene/video Scenario IV results for representative methods across all three modalities. Scenario IV applies grid-search calibration over the mismatch parameter space using measurement residual as the objective function (no ground truth required).

#### 3.1 Calibration Convergence

The grid-search calibration in Scenario IV explores mismatch parameters using measurement residual as a blind objective. For CASSI, we search over spatial shifts  $dx, dy \in [-1.0, 1.0]$  px on an  $11 \times 11$  grid (121 evaluations). For CACTI, we search over  $dx, dy \in [-1.0, 1.0]$  px on a  $9 \times 9$  grid (81 evaluations). For SPC, we search over gain drift  $\alpha \in [0, 0.005]$  with 21 grid points. Each evaluation runs a fast inner-loop solver (GAP-TV with 30 iterations for CASSI/CACTI, FISTA-TV with 200 iterations for SPC) and computes the measurement residual. The parameter yielding the lowest residual is selected, and a final high-quality reconstruction is performed.

The total calibration cost measured on a single CPU core:  $\sim 20$  minutes for CASSI (121 evaluations  $\times$  3 scenes),  $\sim 1$  minute for CACTI (81 evaluations  $\times$  2 groups), and  $\sim 3.6$  minutes for SPC (21 evaluations). The CASSI grid search recovered the true parameters with good accuracy (estimated  $dx=0.40$ ,  $dy=0.20$  vs. true  $dx=0.50$ ,  $dy=0.30$ ). CACTI achieved near-perfect recovery (estimated  $dx=0.50$ ,  $dy=0.25$  vs. true  $dx=0.50$ ,  $dy=0.30$ ). For SPC, the measurement residual did not provide a useful signal for gain drift calibration, estimating  $\alpha=0$  regardless of the true value ( $\alpha=0.0015$ ).

Table 5: CACTI real data per-scene measurement residual (lower is better). 4 real scenes ( $512 \times 512$ ,  $cr=10$ ). Calibrated uses the hardware mask; mismatched shifts the mask by ( $dx=0.5$ ,  $dy=0.3$  px).

| Scene        | GAP-TV |        |              | EfficientSCI |        |             |
|--------------|--------|--------|--------------|--------------|--------|-------------|
|              | Cal.   | Mis.   | Ratio        | Cal.         | Mis.   | Ratio       |
| Duomino      | 8e-6   | 8.5e-5 | $10.6\times$ | 0.0217       | 0.0277 | $1.3\times$ |
| Hand         | 7e-6   | 7.7e-5 | $11.0\times$ | 0.0199       | 0.0289 | $1.5\times$ |
| PendulumBall | 3.7e-5 | 3.5e-4 | $9.4\times$  | 0.0334       | 0.0415 | $1.2\times$ |
| WaterBalloon | 1.4e-5 | 1.5e-4 | $10.5\times$ | 0.0204       | 0.0266 | $1.3\times$ |
| <b>Mean</b>  | 1.6e-5 | 1.6e-4 | $10.4\times$ | 0.0238       | 0.0312 | $1.3\times$ |

Table 6: Scenario IV per-scene/video PSNR (dB). Representative methods: GAP-TV (CASSI, CACTI), FISTA-TV (SPC). Recovery percentage =  $(IV-II)/(III-II)\times 100\%$ .

| Modality | Method   | II    | IV    | III   | IV-II | III-II | Recovery |
|----------|----------|-------|-------|-------|-------|--------|----------|
| CASSI    | GAP-TV   | 21.95 | 23.31 | 23.69 | +1.36 | +1.74  | 78%      |
| CACTI    | GAP-TV   | 17.60 | 26.99 | 26.99 | +9.39 | +9.39  | 100%     |
| SPC      | FISTA-TV | 25.47 | 25.47 | 25.43 | +0.00 | -0.04  | 0%       |

## 4 Residual Gap Analysis

?? visualizes the residual gap  $\Delta_{\text{res}} = \text{PSNR}_{\text{I}} - \text{PSNR}_{\text{III}}$  per method, representing unrecoverable losses even with oracle calibration. CASSI exhibits the largest residual gaps (MST-L: 7.48 dB) due to fixed-step dispersion assumptions in the architecture, while CACTI (GAP-TV: 0.74 dB) and SPC (HATNet: 1.20 dB) confirm that spatial and gain-type mismatches are nearly fully recoverable.

## 5 Per-Scene Recovery Ratio Heatmaps

Figure 1 visualizes the recovery ratio  $\rho$  per scene/video per method for all three modalities. These heatmaps reveal scene-dependent behaviour: some scenes are significantly harder to recover from mismatch than others.

## 6 Mismatch Severity Analysis

Figure 2 shows how reconstruction quality varies with mismatch severity (mild, moderate, severe) for one representative method per modality. The moderate




Fig. 1: Residual gap ( $\Delta_{\text{res}} = \text{PSNR}_I - \text{PSNR}_{III}$ ) per method, grouped by modality. CASSI exhibits the largest residual gaps due to dispersion mismatch that oracle mask correction alone cannot address. CACTI and SPC residual gaps are small, confirming high recoverability of spatial and gain-type mismatches.

severity level corresponds to the default mismatch parameters used in the main paper.

#### Mismatch parameters by severity level:

Table 7: Mismatch severity levels for the ablation study.

| Modality | Parameter             | Mild                 | Moderate           | Severe             |
|----------|-----------------------|----------------------|--------------------|--------------------|
| SPC      | $\alpha$ (gain drift) | 0.0005               | 0.0015             | 0.005              |
|          | $\sigma_y$ (noise)    | 0.01                 | 0.03               | 0.10               |
| CASSI    | $dx$ (mask shift, px) | 0.2                  | 0.5                | 1.5                |
|          | $a_1$ (disp. slope)   | 2.01                 | 2.02               | 2.05               |
| CACTI    | all mismatch params   | $0.5 \times$ default | $1 \times$ default | $2 \times$ default |

Table 8: Approximate reconstruction time per scene/image on a single NVIDIA A100 GPU.

|       | Modality Method | Time/scene   | Type          |
|-------|-----------------|--------------|---------------|
| CASSI | GAP-TV          | $\sim 30$ s  | CPU iterative |
|       | HDNet           | $\sim 0.5$ s | GPU inference |
|       | MST-S           | $\sim 0.8$ s | GPU inference |
|       | MST-L           | $\sim 1.2$ s | GPU inference |
| CACTI | GAP-TV          | $\sim 20$ s  | CPU iterative |
|       | PnP-FFDNet      | $\sim 5$ s   | GPU iterative |
|       | ELP-Unfolding   | $\sim 0.3$ s | GPU inference |
|       | EfficientSCI    | $\sim 0.2$ s | GPU inference |
| SPC   | FISTA-TV        | $\sim 15$ s  | CPU iterative |
|       | ISTA-Net        | $\sim 0.1$ s | GPU inference |
|       | HATNet          | $\sim 0.5$ s | GPU inference |

## 7 Implementation Details

### 7.1 Runtime

### 7.2 Hyperparameters

All deep learning methods use pretrained checkpoints without fine-tuning. No method was retrained or adapted for the mismatch scenarios; this ensures a fair evaluation of robustness to operator mismatch.

- **GAP-TV:** 100 iterations (CASSI, CACTI),  $\lambda_{TV} = 0.1$  (CASSI),  $\lambda_{TV} = 1.0$  (CACTI).
- **FISTA-TV:** 500 iterations,  $\lambda = 0.005$  (SPC).
- **PnP-FFDNet:** 50 GAP iterations with FFDNet denoiser ( $\sigma = 25/255$ ).
- **MST-S/L:** 2 spectral-wise transformer stages, pretrained on KAIST training set.
- **HDNet:** Pretrained dual-domain network.
- **ELP-Unfolding:** 8 unfolding stages, pretrained on DAVIS video dataset.
- **EfficientSCI:** Two-stage 3D CNN, pretrained on DAVIS video dataset.
- **ISTA-Net:** 9 learned ISTA layers, pretrained on BSD400 with  $\Phi$  (25% sampling).
- **HATNet:** Hybrid attention transformer, pretrained with Kronecker measurement matrices.

### 7.3 Dataset Generation

**CASSI:** Measurements are generated using a binary random mask (50% fill ratio) with dispersion step  $s = 2$  px/band, yielding  $256 \times 310$  detector images from  $256 \times 256 \times 28$  spectral cubes. Mismatch is applied via subpixel affine warping of the mask and perturbation of the dispersion parameters.

**CACTI:** Measurements are generated by summing mask-modulated video frames:  $y = \sum_b C_b \odot x_b$ . Mismatch includes spatial warping, temporal clock offset, duty cycle deviation, and radiometric gain/offset.

**SPC:** Block-based compressed sensing with  $33 \times 33$  pixel blocks (ISTA-Net, FISTA-TV) or full-image Kronecker-structured measurements (HATNet) at 25% sampling ratio. Mismatch is exponential gain drift applied to measurement rows.

#### 7.4 Code Availability

Code and scripts will be made available upon acceptance. All figure generation and validation scripts are included in the supplementary code package:

- `generate_visual_comparison.py`: Qualitative reconstruction figure
- `generate_scatter_plot.py`: Recovery ratio vs. ideal PSNR scatter plot
- `generate_heatmaps.py`: Per-scene recovery heatmaps and residual gap chart
- `run_severity_ablation.py`: Mismatch severity sweep
- `generate_cassi_figures.py`: CASSI-specific analysis figures
- `generate_cacti_figures.py`: CACTI-specific analysis figures
- `generate_spc_figures.py`: SPC-specific analysis figures
- `validate_cassi_real.py`: Real CASSI hardware validation
- `validate_cacti_real.py`: Real CACTI hardware validation
- `run_scenario_iv.py`: Scenario IV grid-search calibration
- `generate_real_data_figures.py`: Real data comparison figures




Fig. 2: Per-scene/video recovery ratio ( $\rho$ ) heatmaps for CASSI (top, 10 scenes  $\times$  4 methods), CACTI (middle, 6 videos  $\times$  4 methods), and SPC (bottom, 11 images  $\times$  3 methods). Red indicates low recovery; green indicates high recovery. Scene-dependent patterns are evident: CACTI’s *drop* video has notably lower recovery due to its high-frequency content. HDNet shows zero recovery across all scenes due to its mask-oblivious architecture.




Fig. 3: Mismatch severity ablation for three modalities. Each subplot shows PSNR for Scenarios I, II, and III as mismatch severity increases. The shaded red region shows the mismatch gap ( $\Delta_{\text{deg}}$ ); the shaded green region shows the oracle recovery ( $\Delta_{\text{rec}}$ ). As severity increases,  $\Delta_{\text{deg}}$  widens while  $\rho$  generally decreases.