

Ten Primitives and Three Gates: The Universal Structure of Computational Imaging

Chengshuai Yang^{1,*} Xin Yuan²

Abstract

Computational imaging systems routinely underperform in practice because the assumed forward model diverges from the true physics, yet no unified theory explains why reconstructions fail or guarantees that a finite set of operators suffices to describe all imaging modalities. Here we prove two fundamental results. First, the **Finite Primitive Basis Theorem**: every imaging forward model in the Tier-2 operator class $\mathcal{C}_{\text{Tier2}}$ (linear, shift-variant) admits an ε -approximate representation as a typed directed acyclic graph over exactly 10 canonical primitives spanning propagation, interaction, encoding, and detection. Second, the **Triad Decomposition**: every reconstruction failure decomposes into exactly three root causes—information deficiency (**Gate 1**), carrier noise (**Gate 2**), and operator mismatch (**Gate 3**)—with **Gate 3** universally dominant across all validated modalities. Together these results yield a modality-agnostic diagnostic and correction framework. Across seven modalities and five physical carriers—from coded aperture spectral imaging to MRI and CT—autonomous correction recovers +0.76 to +48.25 dB of mismatch-induced degradation without retraining the reconstruction algorithm. Hardware validation on real CASSI and CACTI instruments confirms mismatch dominance ($1.8\times$ to $10.4\times$ residual increase). A held-out closure test on 8 additional modalities—including quantum ghost imaging, THz spectroscopy, and Compton scatter imaging—confirms basis completeness with sublinear, saturating growth.

Introduction

Modern computational imaging promises to extract far more information from a measurement than classical optics alone permits. Coded aperture spectral cameras compress three-dimensional scenes into a single snapshot¹; compressive temporal imagers freeze high-speed video into one exposure²; accelerated MRI scanners reconstruct diagnostic-quality images from a fraction of the acquired k -space data³. Over the past decade, the community has invested enormous effort in improving reconstruction algorithms—progressing from compressed sensing^{4,5} and plug-and-play priors⁶ to deep unrolling networks⁷ and vision transformers⁸—yet these algorithms routinely fail when deployed on real instruments, because

^{*1}NextGen PlatformAI C Corp, USA. ²School of Engineering, Westlake University, Hangzhou, China. ^{*}Correspondence: integrityyang@gmail.com

the forward model assumed by the solver does not match the physics that generated the data.

The scale of this problem is striking. In coded aperture snapshot spectral imaging (CASSI)¹, the state-of-the-art transformer MST-L⁸ achieves 34.81 dB on the KAIST benchmark⁹ when the forward model is perfectly known. A realistic sub-pixel mask perturbation drops MST-L to 20.83 dB—a catastrophic loss of 13.98 dB that erases twice the cumulative gains from a decade of solver development (~ 7 dB from TwIST¹⁰ to MST-L). On real CASSI hardware, the same mask perturbation increases the measurement residual by $1.8\times$; on real CACTI hardware, the increase is $10.4\times$ ². Analogous degradations appear across modalities, from lensless imaging to MRI¹¹ to computed tomography¹².

These failures persist because computational imaging lacks two theoretical foundations. First, a **representation theory**: no formal result guarantees that a finite set of primitive operators suffices to represent all imaging forward models. Without this, every new modality requires bespoke engineering and the field cannot reason about completeness. Second, a **diagnostic theory**: no systematic framework decomposes reconstruction failures into root causes. Calibration errors, noise, and information deficiency are conflated, and failures are addressed *ad hoc*. Calibration methods exist for specific instruments^{13,14}, but they do not generalise; robustness studies perturb individual systems¹⁵, but they lack a unifying formalism.

Here we establish both foundations within Physics World Models (PWM). We prove the **Finite Primitive Basis Theorem** (Theorem 3): every imaging forward model in the operator class $\mathcal{C}_{\text{Tier2}}$ —encompassing all clinical, scientific, and industrial modalities at Tier-2 fidelity—admits an ε -approximate representation as a typed DAG over exactly 10 canonical primitives. We establish the **Triad Decomposition**: every reconstruction failure decomposes into three gates—information deficiency (**Gate 1**), carrier noise (**Gate 2**), and operator mismatch (**Gate 3**)—with **Gate 3** universally dominant. These two results are complementary: the Finite Primitive Basis provides a universal representation for forward models; the Triad Decomposition provides a universal diagnostic law over that representation. Together they yield a practical framework that diagnoses and corrects imaging failures across modalities without modality-specific tuning.

We validate both contributions across seven modalities spanning five physical carriers (CASSI, CACTI, SPC, lensless, ptychography¹⁶, MRI¹⁷, CT¹⁸), with hardware validation on real CASSI and CACTI instruments. A held-out closure test on 8 additional modalities—including quantum ghost imaging, THz time-domain spectroscopy, and Compton scatter imaging—confirms basis completeness. The full formal proof of Theorem 3 is presented in the companion paper¹⁹.

68 The Finite Primitive Basis

69 A foundational question in computational imaging is whether the diversity of imaging forward models—from coded aperture cameras to MRI scanners to electron microscopes—can
70 be captured by a small, fixed set of primitive operators. We answer affirmatively.

72 **The OperatorGraph representation.** Every imaging forward model is encoded as a
73 typed directed acyclic graph (DAG) in which each node wraps a single primitive physical
74 operator and edges define the data flow from source to detector. Every primitive implements
75 both a `forward()` and an `adjoint()` method, with validated adjoint consistency ensuring
76 $\langle H\mathbf{x}, \mathbf{y} \rangle = \langle \mathbf{x}, H^\dagger \mathbf{y} \rangle$ to within numerical precision. Each edge carries tensor shape and
77 dtype metadata, enabling static validation before execution. We call this formalism the
78 OPERATORGRAPH intermediate representation (IR).

79 **Ten canonical primitives.** The OPERATORGRAPH IR defines a library of exactly 10
80 canonical primitives $\mathcal{B} = \{P, M, \Pi, F, C, \Sigma, D, S, W, R\}$:

#	Primitive	Notation	Physical action
1	Propagate	$P(d, \lambda)$	Free-space wave propagation
2	Modulate	$M(\mathbf{m})$	Element-wise multiplication (mask, coil, absorption)
3	Project	$\Pi(\theta)$	Radon line-integral projection
4	Encode	$F(\mathbf{k})$	Fourier-domain encoding (k -space)
5	Convolve	$C(\mathbf{h})$	Spatial convolution (PSF)
6	Accumulate	Σ	Summation over spectral/temporal axis
7	Detect	$D(g, \eta)$	Detector response (5 canonical families)
8	Sample	$S(\Omega)$	Sub-sampling on index set Ω
9	Disperse	$W(\alpha, a)$	Wavelength-dependent spatial shift
10	Scatter	$R(\sigma, \Delta\epsilon)$	Direction change and/or energy shift

81 The Detect nonlinearity η is restricted to five canonical families (linear-intensity, logarithmic,
82 sigmoid, Poisson-rate, coherent-field), each with at most 2 scalar parameters, preventing Detect from becoming a universal approximator (see¹⁹ for formal definitions and
83 non-universality proof).
84

85 **Two-layer architecture.** The framework separates two abstraction layers. *Layer A*
86 (physics-stage families) categorizes carrier physics into four stages: propagation (wave
87 equation), interaction (elastic or scattering), encoding–projection (Radon or Fourier), and
88 detection–readout (integration, sampling, measurement). *Layer B* (primitives) maps each
89 physics-stage instance to one or a finite composition of primitives from $\mathcal{B}$: propagation
90 $\rightarrow \{P, C\}$; elastic interaction $\rightarrow \{M\}$; scattering $\rightarrow \{R\}$; encoding–projection $\rightarrow \{\Pi, F\}$;
91 detection–readout $\rightarrow \{\Sigma, S, W, C, D\}$.

92 **Physics Fidelity Ladder.** The OPERATORGRAPH IR defines a four-tier Physics Fidelity
 93 Ladder: Tier 1 (linear, shift-invariant), Tier 2 (linear, shift-variant), Tier 3 (nonlinear,
 94 ray/wave-based), and Tier 4 (full-wave/Monte Carlo). The Finite Primitive Basis applies to
 95 Tiers 1–2; Tiers 3–4 are handled by refinement sub-DAGs that expand individual primitive
 96 nodes into higher-fidelity sub-graphs while preserving the top-level DAG structure.

97 **Five physical carriers.** The OPERATORGRAPH IR is organized around five carrier fam-
 98 ilies: *photons* (visible, infrared, X-ray, gamma), *electrons* (scanning, transmission, diffrac-
 99 tion), *spins* (nuclear magnetic resonance), *acoustic waves* (ultrasound, photoacoustic), and
 100 *particles* (neutrons, protons, muons). The current library contains 26 validated modality
 101 templates across all five carrier families (Supplementary Table S3).

102 **Theorem statement.**

103 **Definition 1** (Tier-2 Operator Class). $\mathcal{C}_{\text{Tier2}}$ consists of all imaging forward models express-
 104 ible as a finite sequential-parallel composition of linear, shift-variant stages with bounded
 105 operator norm ($\|H_k\| \leq B$) and stage count $K \leq N_{\max}$.

106 **Definition 2** (ε -Approximate Representation). A typed DAG G with nodes $V \subseteq \mathcal{B}$ is an
 107 ε -approximate representation of $H \in \mathcal{C}_{\text{Tier2}}$ if $\sup_{\|\mathbf{x}\| \leq 1} \|H(\mathbf{x}) - H_G(\mathbf{x})\| / \|H(\mathbf{x})\| \leq \varepsilon$ and
 108 $|V| \leq N_{\max}$, $\text{depth}(G) \leq D_{\max}$.

109 **Theorem 3** (Finite Primitive Basis). For every $H \in \mathcal{C}_{\text{Tier2}}$, there exists a typed DAG G
 110 with $V \subseteq \mathcal{B}$ that is an ε -approximate representation of H .

111 *Proof sketch.* By Definition 1, $H = H_K \circ \dots \circ H_1$ with $K \leq N_{\max}$. Each factor is classified
 112 into one of four physics-stage families and realized by primitives: propagation factors $\rightarrow P$
 113 or C (wave-equation solutions); elastic interaction $\rightarrow M$ (exact); scattering $\rightarrow R$; encoding-
 114 projection $\rightarrow \Pi$ or F (exact); detection-readout $\rightarrow$ finite composition from $\{\Sigma, S, W, C, D\}$.
 115 The per-factor approximation errors compose via sub-multiplicativity: $\|H - H_G\| \leq K \cdot$
 116 $\max_k(\varepsilon_k) \cdot B^{K-1} \leq \varepsilon$ for ε_k bounded by the Tier-2 truncation. Full proof with five primitive
 117 realization lemmas in Supplementary Note 10 and ¹⁹. $\square$

118 **Scope of $\mathcal{C}_{\text{Tier2}}$.** Covers: all clinical, scientific, and industrial modalities at Tier 1–2 (coded
 aperture, interferometric, projection, Fourier-encoded, acoustic, scattering, THz). Excludes:
 Tier 3–4 nonlinear models (handled by refinement sub-DAGs), quantum state tomography, rel-
 ativistic regimes.

119 **Extension protocol.** A new primitive is warranted when no DAG over $\mathcal{B}$ achieves $e_{\text{Tier2}} \leq$
 120 ε within the complexity bounds. The extension requires: (1) validated forward/adjoint,
 121 (2) demonstrated representation gap, (3) error reduction below ε , (4) need by ≥ 2 modal-
 122 ities, (5) backward-compatible closure re-test. Scatter (R) was added via this protocol when
 123 Compton imaging produced $e_{\text{Tier2}} = 0.34$ without it.

Held-out closure test. To validate Theorem 3 empirically, we conduct a closure test under a frozen protocol: the primitive library $\mathcal{B}$ (10 primitives), Detect families (5), fidelity threshold ($\varepsilon = 0.01$), and complexity bounds ($N_{\max} = 20$, $D_{\max} = 10$) are all frozen before evaluation. We test 5 Tier-1 held-out modalities (OCT, photoacoustic, SIM, phase-contrast X-ray, electron ptychography) and 3 exotic modalities (quantum ghost imaging, THz-TDS, Compton scatter imaging):

Modality	e_{Tier2}	#Nodes/Depth	Triad transfer	New prim.?
OCT	< 0.01	5 / 3	Y	N
Photoacoustic	< 0.01	4 / 3	Y	N
SIM	< 0.01	4 / 3	Y	N
Phase-contrast X-ray	< 0.01	5 / 4	Y	N
Electron ptycho	< 0.01	4 / 3	Y	N
Ghost imaging	< 0.01	4 / 3	Y	N
THz-TDS	< 0.01	3 / 2	Y	N
Compton scatter	$< 0.01^*$	4 / 3	Y	Y (R)

*With R in library; $e_{\text{Tier2}} > 0.01$ without R .

Seven of eight modalities decompose with the frozen library. Quantum ghost imaging is operator-equivalent to a single-pixel camera at the image-formation level—sharing the canonical DAG $\text{Source} \rightarrow M \rightarrow \Sigma \rightarrow D$ at Tier-2 abstraction despite fundamentally different source physics. THz-TDS requires only the coherent-field Detect family. Compton scatter imaging is the sole modality requiring a new primitive (Scatter, R), which covers 5+ additional modalities (Raman, fluorescence, DOT, Brillouin). The basis grows from 9 to 10—an 11% increase—while modality coverage grows by 19%+.

Basis-growth saturation. Plotting the number of distinct primitives K against the number of registered modalities N reveals clear saturation (Figure 9): 8 of 10 primitives are introduced by the first 10 modalities, Disperse (W) by CASSI-type spectral systems, and Scatter (R) only when Compton/Raman-class modalities enter. The growth is sublinear and saturating: $K = 10$ at $N = 31$, with no new primitive required for the most recent 19 modalities added. This saturation is consistent with Theorem 3: once all four physics-stage families are covered by primitives, new modalities compose existing primitives rather than requiring new ones.

The Triad Decomposition

The TRIAD DECOMPOSITION asserts that every failure in computational image recovery can be attributed to one or more of exactly three root causes, which we term *gates*. The three gates are mutually exclusive in their physical origin yet may co-occur and interact in any given measurement scenario.

Gate 1: Recoverability. **Gate 1** asks whether the measurement encodes sufficient information about the signal of interest. Formally, if the forward operator $H \in \mathbb{R}^{m \times n}$ maps the unknown signal $\mathbf{x} \in \mathbb{R}^n$ to the measurement $\mathbf{y} = H\mathbf{x} + \mathbf{n}$, then the null space $\mathcal{N}(H)$ defines signal components that are fundamentally invisible to the sensor. When $\mathcal{N}(H)$ is large, no solver can recover the missing information. **Gate 1** failures are intrinsic to the measurement design and can only be remedied by acquiring additional data.

Gate 2: Carrier Budget. **Gate 2** asks whether the signal-to-noise ratio (SNR) is sufficient. Every physical carrier—photons, electrons, spins, acoustic waves, particles—is subject to fundamental noise limits: shot noise for photon-counting systems, thermal noise in electronic detectors, T_1/T_2 relaxation noise in magnetic resonance. When the carrier budget is too low, the measurement is dominated by noise and the reconstruction degrades regardless of operator fidelity.

Gate 3: Operator Mismatch. **Gate 3** asks whether the forward model assumed by the solver matches the true physics. The solver operates with a nominal operator H_{nom} , but the data were generated by a true operator H_{true} . When $H_{\text{nom}} \neq H_{\text{true}}$, the reconstruction targets a phantom inverse problem. **Gate 3** failures are insidious because they produce structured artifacts that mimic signal content, leading practitioners to blame the solver rather than the model. Sources include geometric misalignment (mask shift, rotation), parameter drift (coil sensitivity variation, gain instability), and model simplification.

Key finding: Gate 3 dominates. Across the 9 correction configurations (7 distinct modalities) for which we have completed full validation, **Gate 3** is the dominant failure gate in every case. The theoretical basis is established in Supplementary Note 1 (Proposition 2): Gate 3 becomes the binding constraint whenever calibration error exceeds the noise-equivalent resolution—a condition universally satisfied by modern instruments. In CASSI, a sub-pixel mask shift degrades MST-L by 13.98 dB—twice the ~ 7 dB improvement from a decade of solver R&D. In MRI, a 5% coil sensitivity mismatch produces 6.94 dB degradation. In CT, a sub-degree center-of-rotation error produces characteristic ring artifacts. The TRIAD DECOMPOSITION reveals that the imaging community has been optimizing the wrong variable.

Relationship to the Finite Primitive Basis. The two contributions are complementary. The Finite Primitive Basis provides a universal representation: every forward model is a DAG over 10 primitives. The Triad Decomposition provides a universal diagnostic law over that representation. Crucially, the DAG structure makes Gate 3 diagnosis *actionable*: the MismatchAgent can localize the offending primitive node, and the correction pipeline can perturb parameters of that specific node. Without the OperatorGraph representation,

the Triad would be a theoretical framework without a practical mechanism; without the Triad, the OperatorGraph would be a representation without a diagnostic theory.

Falsifiable predictions. The Triad makes two testable predictions. First, for any imaging modality operating above its Gate 1 floor ($\gamma > \gamma_{\min}$) and Gate 2 floor ($\text{SNR} > \text{SNR}_{\min}$), Gate 3 will dominate whenever $\|\delta\theta\|_{\mathbf{J}^\top \mathbf{J}} > (\sigma_n / \|\mathbf{x}\|_\infty) \cdot \kappa(\mathbf{H})$, where $\mathbf{J}$ is the parameter Jacobian and $\kappa(\mathbf{H})$ the condition number. Second, the recovery ratio ρ is bounded by the coherence between the mismatch subspace and the signal prior, predicting that operator-conditioned solvers achieve higher ρ than operator-oblivious solvers—a prediction confirmed across all validated modalities.

TriadReport. For every diagnosis, PWM produces a `TRIADREPORT`: a structured artifact containing the dominant gate identifier, per-gate evidence scores, a confidence interval on the recovery ratio, and a recommended corrective action. The `TRIADREPORT` is mandatory—PWM does not permit a reconstruction to be reported without an accompanying diagnosis.

Consequences: Diagnosis and Correction

The two theoretical results directly imply a practical framework. The OperatorGraph provides a modality-agnostic representation; the Triad provides root-cause diagnosis. PWM implements this through three deterministic agents—no large language model, no learned parameters, no human intervention.

Diagnostic agents. The `RecoverabilityAgent` (Gate 1) computes the effective compression ratio m/n , estimates null-space dimension via randomised SVD, and outputs a recoverability score $s_1 \in [0, 1]$. The `PhotonAgent` (Gate 2) estimates per-pixel carrier budget from measurement statistics and computes the Cramér–Rao lower bound on reconstruction error, outputting a budget score $s_2 \in [0, 1]$. The `MismatchAgent` (Gate 3) compares residual statistics against the expected noise distribution to detect mismatch, then localizes the offending OperatorGraph node by sweeping perturbations, outputting a mismatch score $s_3 \in [0, 1]$ and a pointer to the offending node.

Correction pipeline. When **Gate 3** is dominant, PWM activates a two-stage correction pipeline. **Algorithm 1 (Beam Search)** performs a coarse grid search over the declared mismatch parameter family $\theta = (\theta_1, \dots, \theta_k)$ associated with the offending node, scoring candidates by reconstruction sharpness. **Algorithm 2 (Gradient Refinement)** takes the top candidates and performs continuous optimization via backpropagation through the DAG, combining data-fidelity and regularization losses. Correction operates exclusively on

the forward model, not on the solver: any existing solver—iterative, plug-and-play, or deep unrolling—benefits without modification.

4-Scenario Protocol. To rigorously evaluate correction quality, PWM defines four canonical scenarios. **Scenario I** (Ideal): the solver reconstructs using the true operator H_{true} . **Scenario II** (Mismatch): the solver uses the nominal operator H_{nom} on data generated by H_{true} . **Scenario III** (Corrected): the solver uses the PWM-corrected operator. **Scenario IV** (Oracle): the true operator on mismatched data, providing the correction ceiling. The recovery ratio $\rho = (\text{PSNR}_{\text{III}} - \text{PSNR}_{\text{II}}) / (\text{PSNR}_{\text{I}} - \text{PSNR}_{\text{II}})$ quantifies how much of the mismatch-induced degradation is recovered (see Methods, Equation (5)).

Empirical Validation

We validate both theoretical contributions in two stages: controlled simulation experiments across seven modalities using the 4-Scenario Protocol, and hardware validation on real CASSI and CACTI instruments.

Simulation experiments

Correction results. Supplementary Table S1 summarizes the correction performance across 9 configurations spanning 7 modalities and multiple carrier families. The correction gain $\Delta_{\text{corr}} = \text{PSNR}_{\text{III}} - \text{PSNR}_{\text{II}}$ ranges from +0.54 dB (CASSI Alg 1) to +48.25 dB (MRI single-coil; under clinically realistic multi-coil conditions, the gain is more moderate—see Supplementary Table S13). The validated modalities span photon-domain (CASSI: +0.76/+6.50 dB; CACTI: +10.21 dB, $\rho = 93.3\%$; SPC: +10.38 dB, $\rho = 89.6\%$; Lensless: +3.55 dB), coherent-photon (Ptychography: +7.09 dB), spin (MRI: +48.25 dB), and X-ray (CT: +10.68 dB) modalities.

CASSI deep dive. Under the 5-parameter mismatch validated by InverseNet²⁰ ($dx=0.5$ px, $dy=0.3$ px, $\theta=0.1^\circ$, $a_1=2.02$, $\alpha=0.15^\circ$), all solvers collapse to a narrow Scenario II range of 20.83–21.88 dB regardless of their ideal performance (GAP-TV²¹: 24.34 dB; MST-L⁸: 34.81 dB; HDNet²²: 34.66 dB), confirming operator-driven failure. Under Scenario IV, MST-L achieves $\rho = 0.46$ (+6.50 dB), while HDNet shows 0% recovery due to its mask-oblivious architecture. The ceiling recovery varies by solver, indicating that multi-parameter mismatches are substantially harder to correct than isolated shifts.

CACTI and SPC results. CACTI² exhibits even more severe degradation: EfficientSCI²³ drops from 35.39 to 14.81 dB (−20.58 dB), while GAP-TV recovers to $\rho = 93.3\%$ under Scenario IV. The inverse performance–robustness relationship is stark: the best ideal-condition

method suffers the largest degradation and lowest recovery ratio. SPC²⁴ presents a qualitatively different mismatch (gain drift): HATNet²⁵ recovers to $\rho = 89.6\%$, and blind calibration via total variation recovers 86–92% of the oracle bound (Supplementary Table S7).

Gate binding and sweep validation. **Gate 3** is dominant across all 9 configurations. However, **Gate 1** and **Gate 2** impose fundamental limits: extreme compression produces catastrophic collapse (SPC: 28.3 $\rightarrow$ 14.4 dB from 25% to 1% sampling; CACTI: 25.9 $\rightarrow$ 20.6 dB at CR 64), and noise sweeps reveal steep cliff-edge behaviour (CACTI: 24.8 $\rightarrow$ 10.5 dB from 10,000 to 10 photons). Unlike Gate 3 losses, Gate 1 and Gate 2 losses are information-theoretic and irrecoverable (Supplementary Tables S10–S11).

Zero-shot generalization. Hyperparameters tuned on photon-domain modalities (CASSI, CACTI, SPC) transfer without modification to coherent-photon, spin, and X-ray domains, with comparable correction gains (Figure 6), confirming carrier-agnostic operation enabled by the OperatorGraph abstraction.

Hardware validation

CASSI real data. We reconstruct 5 real TSA scenes¹ under calibrated and perturbed masks. GAP-TV shows a mean residual ratio of $1.8\times$ ($0.00189 \rightarrow 0.00333$); HDNet shows $1.0\times$ (mask-oblivious). MST-S/L show ratios near $1.0\times$ on real data, in contrast to severe synthetic degradation—revealing that real hardware already contains manufacturing imperfections that absorb additional perturbations (Supplementary Table S5).

CACTI real data. GAP-TV shows $10.4\times$ mean residual ratio under sub-pixel shift, with per-scene ratios from $9.4\times$ to $11.0\times$. The order-of-magnitude sensitivity arises because the temporal mask pattern replicates across all compressed frames, amplifying errors multiplicatively.

Autonomous calibration. Grid-search calibration recovers 85% (CASSI, 1,140 s), 100% (CACTI, 60 s), and 86–92% (SPC via TV objective) of the oracle correction. CACTI achieves full recovery because the mismatch manifold is low-dimensional and sensitivity is high.

Simulation-to-hardware gap. The gap depends on baseline calibration quality: CASSI ($1.8\times$ real vs. 3.38 dB synthetic) has complex optics with pre-existing errors; CACTI ($10.4\times$ real, 10.94 dB synthetic) has simpler optics with fewer pre-existing errors. This asymmetry reveals that synthetic studies overestimate marginal impact while underestimating cumulative burden.

Discussion

The central finding is that computational imaging has a universal structure captured by two results: a finite primitive basis for forward models and a tripartite decomposition for failures. Operator mismatch—not solver weakness, not information deficiency, not noise—is the dominant bottleneck across all validated modalities. The implication is direct: the research community should rebalance effort from solver-centric to operator-centric approaches. A single calibration step recovers more reconstruction quality than years of algorithmic innovation.

An instructive analogy is the periodic table of elements. Just as Mendeleev organized known elements by atomic number and predicted gaps, the OPERATORGRAPH organizes imaging modalities by their primitive composition. The analogy is pedagogical, not mathematical: imaging primitives do not have atomic numbers, and the DAG structure is richer than a two-dimensional table. Nevertheless, Theorem 3 implies that the space of imaging forward models, like the space of chemical elements, has a discoverable and finite structure—and for a related reason: the underlying physical interactions are finite.

The coverage of the primitive basis is broader than it may initially appear. Quantum ghost imaging is operator-equivalent to a classical single-pixel camera at the image-formation level, sharing the same DAG despite fundamentally different source physics. THz time-domain imaging requires only the coherent-field Detect family. Of three exotic modalities analyzed, only Compton scatter imaging required a new primitive (Scatter), and that single primitive covers 5+ scattering and fluorescence modalities. This pattern—that exotic modalities decompose into existing primitives, and when they do not, a single new primitive covers a whole family—is the hallmark of a well-chosen basis.

The hardware validation reveals a nuanced picture. On real CASSI, degradation is smaller than predicted ($1.8\times$ vs. 3.38 dB), because as-built masks absorb perturbations. On real CACTI, degradation is severe ($10.4\times$), because temporal compression amplifies errors multiplicatively. This asymmetry is invisible without a unified framework. The simulation-to-hardware gap carries a methodological lesson: synthetic mismatch studies overestimate marginal impact while underestimating cumulative burden.

To demonstrate translational potential, we prototype a CT QC Copilot mapping the Triad gates to clinical failure modes. On a simulated 30-scanner fleet, it detects calibration drift with 100% sensitivity/specificity and reduces per-scanner QC time by 94% (Supplementary Note 7). Consistent with the research findings, Gate 3 dominates.

Several limitations should be noted. First, real-data experiments apply software-simulated perturbations rather than physical displacement; controlled hardware experiments are the natural next step (Supplementary Note 8). Second, Tier 1–2 models may not capture all failure modes visible at Tier 3–4. Third, the correction pipeline is limited to declared mismatch parameter families. Fourth, the CT QC validation uses a simulated fleet.

Looking forward, we envision hardware-in-the-loop validation across additional instru-

ments, real-time adaptive calibration during acquisition, prospective clinical deployment of the CT QC Copilot, and scaling the OperatorGraph library to compile a comprehensive atlas of imaging failure modes.

Acknowledgements. We thank the open-source computational imaging community for making reconstruction code and benchmark datasets publicly available. We acknowledge discussions with David J. Brady regarding controlled hardware validation protocols for CASSI and CACTI instruments, and with Steve B. Jiang regarding clinical CT quality assurance validation. This work was supported by NextGen PlatformAI C Corp.

Author Contributions. C.Y. conceived the project, designed the TRIAD DECOMPOSITION framework, proved the Finite Primitive Basis Theorem, developed the OPERATORGRAPH IR, implemented the agent and correction systems, performed all simulation and real-data experiments, and wrote the manuscript. X.Y. contributed domain expertise on the CASSI and CACTI forward models and reconstruction algorithms (GAP-TV, EfficientSCI), validated the mismatch parameter specifications, and edited the manuscript.

Competing Interests. C.Y. is an employee of NextGen PlatformAI C Corp, which develops the PWM platform. The authors declare no other competing interests.

Data Availability. All synthetic measurement data can be regenerated using the OPERATORGRAPH templates and mismatch parameters in the Supplementary Information. The KAIST hyperspectral dataset⁹ and TSA real-data scenes used for CASSI experiments are publicly available. CACTI real-data scenes are available from the EfficientSCI repository²³.

Code Availability. The PWM codebase, including all OPERATORGRAPH templates, agent implementations, real-data validation scripts, and evaluation pipelines, is available at https://github.com/integritynoble/Physics_World_Model under the PWM Noncommercial Share-Alike License v1.0 (see LICENSE in the repository).

Correspondence. Correspondence and requests for materials should be addressed to C.Y. (integrityyang@gmail.com).

Online Methods

OperatorGraph Specification

Formal definition. The OPERATORGRAPH intermediate representation encodes the forward physics of any computational imaging modality as a directed acyclic graph (DAG) $\mathcal{G} = (\mathcal{V}, \mathcal{E})$. Each node $v_i \in \mathcal{V}$ wraps a *primitive operator* and implements two entry points:

352 `forward( $x$ )  $\rightarrow$   $y$` and `adjoint( $y$ )  $\rightarrow$   $x$` , the latter defined only when the primitive is lin-
 353 ear. Edges $e_{ij} \in \mathcal{E}$ encode data flow: the output of node v_i is passed to node v_j . Each
 354 node additionally exposes a set of learnable parameters θ_i that may be perturbed during
 355 mismatch simulation or optimized during calibration, as well as read-only metadata flags
 356 (`is_linear`, `is_stochastic`, `is_differentiable`). The graph is stored as a declarative
 357 YAML specification (`OperatorGraphSpec`) and compiled to an executable `GraphOperator`
 358 object by the `GraphCompiler`.

359 **Node types.** Primitive operators fall into two categories:

- 360 • **Linear operators.** Convolution (`conv2d`), mask modulation (`mask_modulate`), sub-
 361 pixel shift (`subpixel_shift_2d`), Radon transform (`radon_fanbeam`), Fourier encod-
 362 ing (`fourier_encode`), spectral dispersion (`spectral_disperse`), Fresnel propaga-
 363 tion (`fresnel_propagate`), random projection (`random_project`), and structured il-
 364 lumination (`sim_modulate`). Each implements both `forward()` and `adjoint()`.
- 365 • **Nonlinear operators.** Squared magnitude (`magnitude_sq`), Poisson–Gaussian noise
 366 (`poisson_gaussian`), saturation clipping (`saturation_clip`), phase retrieval nonlin-
 367 earity (`phase_abs`), and detector quantization (`quantize`). These set `is_linear` =
 368 `False` and raise `NotImplementedError` on `adjoint()`, except where a well-defined
 369 pseudo-adjoint exists (*e.g.*, the identity adjoint for magnitude-squared in Gerchberg–
 370 Saxton-type algorithms).

Adjoint validation. Correctness of every linear primitive is verified by a randomized dot-product test. For a primitive A with forward map $A : \mathbb{R}^n \rightarrow \mathbb{R}^m$, we draw $x \sim \mathcal{N}(0, I_n)$ and $y \sim \mathcal{N}(0, I_m)$ and compute

$$\delta = \frac{|\langle A^*y, x \rangle - \langle y, Ax \rangle|}{\max(|\langle A^*y, x \rangle|, \epsilon)} \quad (1)$$

371 where $\epsilon = 10^{-12}$ guards against division by zero. The test is repeated $n_{\text{trials}} = 5$ times
 372 with independent random draws; the primitive passes if $\delta_{\text{max}} < 10^{-6}$. At the graph level, a
 373 compiled `GraphOperator` composed entirely of linear nodes executes the same test over the
 374 composed forward–adjoint chain. A `GraphAdjointCheckReport` records n_{trials} , δ_{max} , and $\bar{\delta}$
 375 for audit. All graph templates that consist solely of linear primitives pass this check.

376 **Graph compilation.** The compiler executes a four-stage pipeline:

- 377 1. **Validate.** Confirm acyclicity via topological sort (Kahn’s algorithm), verify that ev-
 378 ery `primitive_id` exists in the global `PRIMITIVE_REGISTRY`, reject duplicate `node_id`
 379 values, and optionally verify shape compatibility along edges when a `canonical_chain`
 380 metadata flag is set.

- 381 2. **Bind.** Instantiate each primitive with its parameter dictionary θ_i .
- 382 3. **Plan forward.** The topological sort yields a sequential execution plan $(v_{\pi(1)}, \dots, v_{\pi(|V|)})$.
- 383 4. **Plan adjoint.** For graphs where `all_linear = True`, the adjoint plan reverses the
 384 topological order and applies each node’s individual adjoint in sequence, implementing
 385 the chain rule $A^* = A_1^* \circ \dots \circ A_{|V|}^*$ for a composition $A = A_{|V|} \circ \dots \circ A_1$. For
 386 graphs containing nonlinear nodes, the adjoint plan is not generated, and any call to
 387 `adjoint()` raises `NotImplementedError` at runtime.

388 The compiled `GraphOperator` is serializable to JSON and hashable via SHA-256 for prove-
 389 nance tracking in `RunBundle` manifests.

390 **Template library.** The `graph_templates.yaml` registry contains templates organized
 391 across 26 registered modalities (7 with full end-to-end correction validation, 1 with Scenario I
 392 baseline, 18 with template-level validation), grouped by physical carrier:

- 393 • **Photons (optical and X-ray):** CASSI, SPC, CACTI, structured illumination
 394 microscopy (SIM), confocal, light-sheet, holography, ptychography, Fourier ptycho-
 395 graphic microscopy (FPM), optical coherence tomography (OCT), lensless imaging,
 396 light field, integral imaging, neural radiance fields (NeRF), Gaussian splatting, fluo-
 397 rescence lifetime imaging (FLIM), diffuse optical tomography (DOT), phase retrieval,
 398 X-ray computed tomography (CT), and cone-beam CT (CBCT).
- 399 • **Electrons:** Electron diffraction, electron backscatter diffraction (EBSD), electron
 400 energy loss spectroscopy (EELS), and electron holography.
- 401 • **Spins (MRI):** Functional MRI (fMRI), diffusion-weighted MRI (DW-MRI), and
 402 magnetic resonance spectroscopy (MRS).
- 403 • **Acoustic:** Ultrasound B-mode, Doppler ultrasound, shear-wave elastography, sonar,
 404 and photoacoustic tomography (combines optical excitation with acoustic detection).
- 405 • **Particles:** Neutron tomography, proton radiography, and muon tomography.

406 **Physics Fidelity Ladder.** Each template is parameterized by a fidelity tier that controls
 407 the degree of physical realism in the simulated forward model:

- 408 **Tier 1 (Linear, shift-invariant):** The forward model is a linear, spatially uniform operator—
 409 the simplest approximation, suitable for initial diagnostics and rapid prototyping.
- 410 **Tier 2 (Linear, shift-variant):** Spatially varying operator parameters (e.g. non-uniform
 411 illumination, position-dependent PSF, multi-coil sensitivity maps in MRI). Adds a
 412 modality-appropriate noise model (Poisson shot noise plus Gaussian read noise for
 413 photon-counting modalities, Rician noise for MRI, Poisson for CT).

414 **Tier 3 (Nonlinear, ray/wave-based):** Includes nonlinear effects such as wavefront cur-
 415 vature, diffraction, and scattering. Perturbation families and ranges are specified in
 416 `mismatch_db.yaml`.

417 **Tier 4 (Full-wave / Monte Carlo):** Complete physical simulation including wave-optical
 418 propagation, spatially varying aberrations, detector nonlinearities, and environmen-
 419 tal drift. Currently implemented for holography and ptychography; other modalities
 420 degrade gracefully to Tier 3.

421 **Triad Decomposition Formalization**

422 The TRIAD DECOMPOSITION asserts that the quality of any computational imaging re-
 423 construction is bounded by three fundamental gates. Rather than a qualitative guideline,
 424 PWM quantifies each gate numerically and uses the resulting scores to diagnose the domi-
 425 nant bottleneck in any imaging configuration.

426 **Gate 1 (Recoverability).** Recoverability measures the information-theoretic capacity
 427 of the sensing geometry. We quantify it via the *effective compression ratio* $r = m/n$, where
 428 m is the number of independent measurements and n the dimension of the signal. The
 429 `compression_db.yaml` registry (1,186 lines) stores, for each modality, a lookup table map-
 430 ping compression ratio to expected reconstruction PSNR under ideal conditions, obtained
 431 from calibration experiments or published benchmarks. Each entry carries a **provenance**
 432 field citing the source (paper DOI, internal experiment ID, or theoretical formula). Addi-
 433 tional recoverability indicators include the effective rank of the measurement matrix (esti-
 434 mated via randomized SVD for large operators), the dimension of the null space, and the
 435 restricted isometry property (RIP) constant where analytically tractable (*e.g.*, for Gaussian
 436 random projections in SPC).

437 **Gate 2 (Carrier Budget).** The carrier budget quantifies the signal-to-noise ratio (SNR)
 438 of the measurement channel. The **PhotonAgent** consumes the `photon_db.yaml` registry
 439 (624 lines) which stores, per modality, a deterministic photon model parameterized by
 440 source power, quantum efficiency, exposure time, and detector characteristics. The agent
 441 classifies the noise regime into one of three categories: *shot-limited* (Poisson-dominated,
 442 $\text{SNR} \propto \sqrt{N_{\text{photon}}}$), *read-limited* (Gaussian read noise dominates, $\text{SNR} \propto N_{\text{photon}}/\sigma_{\text{read}}$),
 443 and *dark-current-limited* (long exposures where dark current accumulation dominates). The
 444 output is a **PhotonReport** containing the estimated SNR in decibels, the noise regime
 445 classification, per-element photon count, and a feasibility verdict (**sufficient**, **marginal**,
 446 or **insufficient**).

447 **Gate 3 (Operator Mismatch).** Operator mismatch quantifies the discrepancy between
 448 the assumed forward model H_{nom} and the true physical operator H_{true} . The **MismatchAgent**

449 consults `mismatch_db.yaml` (797 lines) which catalogs, for each modality, the set of mis-
 450 match parameters (spatial shifts, rotational offsets, dispersion errors, PSF deviations, coil
 451 sensitivity errors, center-of-rotation offsets, *etc.*), their typical ranges, and available cor-
 452 rection methods. The mismatch severity score $s \in [0, 1]$ is computed as the normalized ℓ_2
 453 distance $\|\boldsymbol{\theta}_{\text{true}} - \boldsymbol{\theta}_{\text{nom}}\| / \|\boldsymbol{\theta}_{\text{range}}\|$, where $\boldsymbol{\theta}_{\text{range}}$ is the per-parameter dynamic range from the
 454 registry. Sensitivity analysis $\partial \text{PSNR} / \partial \theta_k$ is estimated via finite differences on the forward
 455 model. The output is a **MismatchReport** containing the severity score, the dominant mis-
 456 match parameter, the recommended correction method, and the expected PSNR gain from
 457 correction.

Gate binding determination. Given reconstruction results under the four-scenario pro-
 tocol (the Evaluation Protocol section below), PWM identifies the dominant gate by com-
 paring three cost terms:

$$C_{\text{mismatch}} = \text{PSNR}_{\text{I}} - \text{PSNR}_{\text{II}} \quad (2)$$

$$C_{\text{noise}} = \text{PSNR}_{\text{ideal}} - \text{PSNR}_{\text{noisy}} \quad (3)$$

$$C_{\text{recover}} = \text{PSNR}_{\text{limit}} - \text{PSNR}_{\text{I}} \quad (4)$$

458 where PSNR_{I} is the reconstruction PSNR under Scenario I (ideal operator), PSNR_{II} under
 459 Scenario II (mismatched operator), $\text{PSNR}_{\text{noisy}}$ under the corresponding noisy condition,
 460 and $\text{PSNR}_{\text{limit}}$ is the theoretical upper bound from the compression table. The dominant
 461 gate is $\arg \max_g C_g$.

462 **TriadReport schema.** The analysis output is a Pydantic-validated `TRIADREPORT` com-
 463 prising: `dominant_gate` (enum: `recoverability`, `carrier_budget`, `operator_mismatch`),
 464 `evidence_scores` (three floats, one per gate), `confidence_interval` (float, 95% CI width
 465 from bootstrap), `recommended_action` (string, *e.g.* “increase compression ratio” or “ap-
 466 ply mismatch correction”), and `parameter_sensitivities` (dictionary mapping each mis-
 467 match parameter name to its $\partial \text{PSNR} / \partial \theta_k$ value).

Recovery ratio. We define the *recovery ratio*

$$\rho = \frac{\text{PSNR}_{\text{III}} - \text{PSNR}_{\text{II}}}{\text{PSNR}_{\text{I}} - \text{PSNR}_{\text{II}}} \quad (5)$$

468 which lies in $[0, 1]$ under standard convexity conditions (see Supplementary Note 1 for
 469 formal analysis; values $\rho > 1$ are possible when the corrected operator provides beneficial
 470 regularization). $\rho = 0$ indicates that calibration yields no benefit (mismatch is not the
 471 bottleneck), while $\rho = 1$ indicates that calibration fully closes the mismatch gap.

472 Agent System Architecture

473 The PWM agent system comprises 6 specialist agents, 1 optional hybrid agent, and 8
474 support classes totalling 10,545 lines of Python. All agents execute deterministically; no
475 large language model (LLM) is required for pipeline operation.

476 **PlanAgent.** The orchestrator agent. Given a user prompt or a structured `ExperimentSpec`,
477 `PlanAgent` parses the intent (`simulate`, `operator_correction`, or `auto`), maps the re-
478 quested modality to its canonical key via the `modalities.yaml` registry (which contains 64
479 modality entries with keywords, forward model equations, and default solvers), builds an
480 `ImagingSystem` contract, and dispatches to the appropriate sub-agents. When the mode is
481 `auto`, `PlanAgent` inspects the available data and operator specification to determine whether
482 simulation or operator correction is more appropriate.

483 **PhotonAgent.** Computes SNR feasibility deterministically from the `photon_db.yaml`
484 registry. For each modality and photon-level tier (`bright`, `standard`, `low_light`), the
485 agent evaluates the photon budget by combining source power, quantum efficiency, ex-
486 posure time, and noise model parameters. The output `PhotonReport` is a strict Pydan-
487 tic model containing `noise_regime` (enum), `snr_db` (float), `feasibility` (enum), and
488 `per_element_photons` (float).

489 **RecoverabilityAgent.** A table-driven agent that consults `compression_db.yaml` (1,186
490 lines) to map the modality and compression ratio to an expected PSNR range. Each table
491 entry includes provenance metadata citing the original source. The output `RecoverabilityReport`
492 contains `compression_ratio`, `psnr_prediction`, `feasibility`, and `null_space_dim` where
493 available.

494 **MismatchAgent.** Scores the mismatch severity for a given imaging configuration us-
495 ing `mismatch_db.yaml` (797 lines). For each modality, the database enumerates the rel-
496 evant mismatch parameters, their physical units, typical perturbation ranges, and avail-
497 able correction algorithms. The output `MismatchReport` includes `severity` (float, 0–1),
498 `correction_method` (string), `expected_gain_db` (float), and `dominant_parameter` (string).

499 **AnalysisAgent.** The bottleneck classifier. It receives reports from the Photon, Recover-
500 ability, and Mismatch agents, computes the gate costs (Equations (2) to (4)), identifies the
501 dominant gate, and generates actionable suggestions. The `AnalysisAgent` also computes
502 the recovery ratio ρ and its bootstrap confidence interval.

503 **AgentNegotiator.** Implements a cross-agent veto protocol. Before reconstruction is au-
504 thorized, the negotiator inspects all three upstream reports and applies three veto con-
505 ditions: (1) low photon budget combined with aggressive compression (C_{noise} and C_{recover}

both large); (2) severe mismatch (severity > 0.7) without a planned correction step; (3) joint probability below the floor threshold ($p_{\text{joint}} < 0.15$), indicating that all three subsystems are simultaneously marginal. When any veto fires, reconstruction halts with an actionable explanation and suggested remediation.

HybridAgent. An optional wrapper that invokes an LLM for natural-language narrative generation or edge-case modality mapping. All quantitative decisions remain on the deterministic code path; the HybridAgent is never required for pipeline operation.

Support classes. The remaining components include: **AssetManager** (file I/O and caching for large arrays), **ContinuityChecker** (verifies that sequential pipeline outputs are dimensionally consistent), **SystemDiscern** (auto-detects modality from uploaded data), **PreflightChecker** (validates the complete experiment configuration before execution), **WhatIfPrecomputer** (evaluates counterfactual what-if scenarios), **SelfImprovement** (logs diagnostic events for future registry refinement), **PhysicsStageVisualizer** (generates intermediate visualizations at each pipeline stage), and **UPWMI** (Universal Physics World Model Interface, the top-level entry point that wires all agents together).

Contract system. Inter-agent communication uses 25 Pydantic v2 contract models. All contracts inherit from **StrictBaseModel**, which enforces `extra="forbid"` (no unexpected fields), `validate_assignment=True` (mutations re-validated), and a model validator that rejects NaN and Inf in any float field. Bounded scores use `Field(ge=0.0, le=1.0)`. Enums are string enums for human-readable JSON serialization. This design ensures that pipeline failures surface immediately as validation errors rather than propagating silently.

YAML registries. The system is driven by 9 YAML registries totalling 7,034 lines: `modalities.yaml` (modality definitions), `graph_templates.yaml` (OperatorGraph skeletons), `photon_db.yaml` (photon models), `mismatch_db.yaml` (mismatch parameters and correction methods), `compression_db.yaml` (recoverability tables with provenance), `solver_registry.yaml` (solver configurations), `primitives.yaml` (primitive operator metadata), `dataset_registry.yaml` (dataset locations and formats), and `acceptance_thresholds.yaml` (pass/fail thresholds per metric).

Correction Algorithms

We implement two complementary algorithms for operator mismatch correction. Crucially, both algorithms operate on the forward operator parameters θ rather than the reconstruction solver weights, making them *solver-agnostic*: the corrected operator $H(\hat{\theta})$ benefits any downstream solver (GAP-TV, MST-L, HDNet²², CST, *etc.*) without retraining.

539 **Algorithm 1: Hierarchical Beam Search.** The coarse correction phase employs a
 540 hierarchical search strategy to rapidly explore the mismatch parameter space. For CASSI,
 541 the five-parameter mismatch model comprises mask affine parameters (spatial shifts dx , dy
 542 and rotation θ) and dispersion parameters (slope a_1 and axis angle α); an optional sixth
 543 parameter, PSF width σ_{psf} , is available but not used in the primary experiments. The
 544 algorithm proceeds as follows:

- 545 1. **1D sweeps.** Each parameter is swept independently over its full range while holding
 546 others at nominal values. This produces five 1D cost curves from which coarse optima
 547 are extracted.
- 548 2. **3D beam search.** The mask affine subspace (dx, dy, θ) is searched over a $5 \times 5 \times 5$
 549 grid centered on the 1D optima. The top- k ($k = 5$) candidates by reconstruction
 550 PSNR are retained.
- 551 3. **2D beam search.** For each retained mask candidate, the dispersion subspace (a_1, α)
 552 is searched over a 5×7 grid. The joint top- k candidates are retained.
- 553 4. **Coordinate descent refinement.** Three rounds of univariate refinement on each
 554 parameter, shrinking the search interval by factor 2 at each round, produce the final
 555 estimate $\hat{\theta}_{\text{Alg1}}$.

556 Total runtime is approximately 300 seconds per scene on a single GPU. Accuracy is
 557 ± 0.1 – 0.2 pixels for spatial parameters and $\pm 0.05^\circ$ for angular parameters.

558 **Algorithm 2: Joint Gradient Refinement.** The fine correction phase uses a differ-
 559 entiable forward model to jointly optimize all mismatch parameters via gradient descent.
 560 The key components are:

- 561 1. **Differentiable mask warp.** The binary mask is warped by a continuous affine
 562 transformation using bilinear interpolation, implemented as a custom PyTorch module
 563 (`DifferentiableMaskWarpFixed`). The mask values are passed through a straight-
 564 through estimator (STE) to maintain binary structure while permitting gradient flow.
- 565 2. **Differentiable forward model.** The CASSI forward model $y = \text{CASSI}(x; \theta)$ is
 566 implemented as a differentiable PyTorch module (`DifferentiableCassiForwardSTE`)
 567 that accepts mismatch parameters as differentiable inputs.
- 568 3. **GPU grid initialization.** A full-range 3D grid search over (dx, dy, θ) with $9 \times 9 \times 7 =$
 569 567 points provides diverse starting candidates. The top 9 candidates seed multi-start
 570 gradient refinement.

571 **4. Staged gradient refinement.** Each of the 9 candidates is refined using Adam
 572 optimization (learning rate 10^{-2} , decaying to 10^{-3}) for 200 steps. For each candidate,
 573 4 random restarts with jittered initialization guard against local minima. The loss
 574 function is the negative PSNR computed via an unrolled K -iteration differentiable
 575 GAP-TV solver (`DifferentiableGAPTV`, $K = 10$ unrolled iterations).

576 Total runtime for Algorithm 2 is approximately 3,200 seconds (200 steps $\times$ 4 restarts $\times$
 577 9 candidates with early stopping). Accuracy improves to ± 0.05 – 0.1 pixels, a 3 – $5\times$ improve-
 578 ment over Algorithm 1. The two algorithms are used sequentially in practice: Algorithm 1
 579 provides a warm start, and Algorithm 2 refines to sub-pixel precision.

580 Evaluation Protocol

581 **Four-Scenario Protocol.** We evaluate every modality under four standardized scenarios
 582 that isolate different sources of quality degradation:

583 **Scenario I (Ideal):** $\mathbf{y}_{\text{obs}} = H_{\text{true}} \mathbf{x}_{\text{gt}}$; reconstruct with H_{true} . In this scenario the system
 584 is perfectly calibrated ($H_{\text{true}} = H_{\text{nom}}$), so the operator used for reconstruction matches
 585 the one that generated the data. This yields the oracle upper bound on reconstruction
 586 quality, limited only by the sensing geometry and solver convergence.

587 **Scenario II (Mismatch):** $\mathbf{y}_{\text{obs}} = H_{\text{true}} \mathbf{x}_{\text{gt}}$; reconstruct with H_{nom} ($H_{\text{nom}} \neq H_{\text{true}}$). This
 588 is the standard operating condition in practice: the measurement is generated by the
 589 true physics, but the reconstruction uses a nominal (potentially mismatched) forward
 590 model.

591 **Scenario III (Corrected):** $\mathbf{y}_{\text{obs}} = H_{\text{true}} \mathbf{x}_{\text{gt}}$; reconstruct with $\hat{H} = H(\hat{\theta})$ where $\hat{\theta}$ is
 592 estimated by Algorithms 1 and 2. This quantifies the benefit of mismatch calibration.

593 **Scenario IV (Oracle Mask):** Same measurements as Scenario II ($\mathbf{y}_{\text{obs}} = H_{\text{true}} \mathbf{x}_{\text{gt}}$ with
 594 $H_{\text{true}} \neq H_{\text{nom}}$); reconstruct with H_{true} instead of H_{nom} . Provides the correction
 595 ceiling: the best reconstruction achievable when the true operator is known exactly,
 596 applied to data that were sensed by the mismatched system. The gap between Sce-
 597 nario IV and Scenario I reveals the irreducible loss from the degraded sensing config-
 598 uration itself (e.g., a shifted mask pattern is suboptimal even when perfectly known).

599 **Metrics.** Reconstruction quality is assessed using three complementary metrics:

- 600 • **PSNR** (peak signal-to-noise ratio, in dB): the primary metric, computed per scene
 601 and averaged. For signals normalized to $[0, 1]$, $\text{PSNR} = 10 \log_{10}(1/\text{MSE})$. For SPC
 602 data normalized to $[0, 255]$, the peak value is 255.

- **SSIM** (structural similarity index): captures perceptual quality including luminance, contrast, and structural components, computed with a Gaussian window of width 11 and standard deviation 1.5.
- **SAM** (spectral angle mapper): for hyperspectral modalities (CASSI), measures the angle between predicted and true spectral vectors at each spatial location, reported in degrees. Lower is better.

Datasets.

- **CASSI**: 10 scenes from the KAIST dataset⁹, each a $256 \times 256 \times 28$ spectral cube (28 spectral bands from 450 nm to 650 nm). Data range $[0, 1]$.
- **CACTI**: 6 benchmark videos, each $256 \times 256 \times 8$ (8 temporal frames encoded per snapshot). Data range $[0, 1]$.
- **SPC**: 11 natural images from the Set11 benchmark, each 256×256 grayscale. Data range $[0, 255]$.

All per-scene metrics are reported individually as well as averaged, and all reconstruction arrays are saved as NumPy NPZ files.

Experimental Details

Hardware. All experiments are conducted on a single NVIDIA GPU. Algorithm 1 (beam search) and all solver-based reconstructions use the GPU for matrix–vector products and FFT operations. Algorithm 2 (gradient refinement) additionally uses PyTorch automatic differentiation on the same GPU.

CASSI configuration. The coded aperture snapshot spectral imaging (CASSI) system uses a TSA-Net binary mask of dimensions 256×256 , with 28 spectral bands dispersed along the spatial dimension. The five-parameter mismatch model $\theta = (dx, dy, \theta, a_1, \alpha)$ describes: mask spatial shift in x (dx , pixels), mask spatial shift in y (dy , pixels), mask rotation angle (θ , degrees), dispersion slope (a_1 , pixels per band), and dispersion axis angle (α , degrees). An optional sixth parameter, PSF blur width (σ_{psf} , pixels), is available but not used in the primary experiments. These mismatch parameter values were determined through systematic characterization of realistic CASSI assembly errors (Supplementary Note 9). The true mismatch parameters are $\theta_{\text{true}} = (dx = 0.5 \text{ px}, dy = 0.3 \text{ px}, \theta = 0.1^\circ, a_1 = 2.02, \alpha = 0.15^\circ)$. Solvers evaluated include TwIST¹⁰, GAP-TV²¹, DGSMP²⁶, MST-L⁸, and CST-L²⁷, all of which receive the same operator and differ only in their reconstruction algorithm. The supplementary per-scene analysis additionally includes DeSCI²⁸ and HDNet²².

635 **CACTI configuration.** The coded aperture compressive temporal imaging system uses
 636 binary temporal masks of dimensions 256×256 , encoding 8 video frames into a single
 637 snapshot measurement. Mismatch is parameterized as a temporal mask timing offset (sub-
 638 frame shift). The default solver is EfficientSCI²³.

639 **SPC configuration.** The single-pixel camera uses random binary measurement patterns
 640 at three compression ratios: 10%, 25%, and 50% ($r = m/n \in \{0.10, 0.25, 0.50\}$). Mismatch
 641 is modeled as an exponential gain drift ($g_i = \exp(-\alpha \cdot i)$) on the measurement matrix. The
 642 default solver is FISTA-TV with total-variation regularization.

643 **MRI configuration.** Cartesian k -space sampling with $4\times$ acceleration (25% of k -space
 644 lines acquired). Mismatch is parameterized as a 5% multiplicative error in the coil sensitivity
 645 maps used for parallel imaging reconstruction. The default solver is SENSE¹⁷ with ℓ_1 -
 646 wavelet regularization.

647 **CT configuration.** Fan-beam geometry with 180 projections over 180° . Mismatch is
 648 modeled as a center-of-rotation (CoR) offset, which produces characteristic arc artifacts in
 649 the reconstruction. The default solver is filtered back-projection (FBP)¹⁸ with a Ram-Lak
 650 filter, supplemented by iterative SART for comparison.

651 **CASSI real-data configuration.** The TSA real hyperspectral dataset¹ consists of 5
 652 scenes at 660×660 spatial resolution with 28 spectral bands and mask-shift step 2. Four
 653 solvers are evaluated: GAP-TV (200 iterations), HDNet (pre-trained checkpoint, full spatial
 654 resolution), MST-S and MST-L (pre-trained checkpoints, centre-cropped to 256×256 due
 655 to hardcoded spatial assumptions in the model architecture). The coded aperture mask
 656 is perturbed by $dx = 0.5$ px, $dy = 0.3$ px to simulate assembly-induced mismatch. No
 657 ground truth is available; quality is assessed via the normalised measurement residual $r =$
 658 $\|\mathbf{y} - H\hat{\mathbf{x}}\|^2 / \|\mathbf{y}\|^2$.

659 **CACTI real-data configuration.** The EfficientSCI real temporal dataset²³ consists of
 660 4 dynamic scenes (duomino, hand, pendulumBall, waterBalloon) at 512×512 with compres-
 661 sion ratio 10. The real mask is stored separately from the measurement data. Two solvers
 662 are evaluated: GAP-TV (50 iterations) and PnP-FFDNet (50 iterations with FFDNet de-
 663 noiser). Mismatch is induced by shifting the mask by $dx = 0.5$ px, $dy = 0.3$ px. Quality is
 664 assessed via the normalised measurement residual and total variation of the reconstruction.

665 **Controlled hardware experiment protocol.** The software-perturbation protocol above
 666 applies calibrated mask shifts to existing real measurements. A full hardware-in-the-loop
 667 validation requires physically displacing the coded aperture mask and re-acquiring data.
 668 The protocol proceeds as follows: (i) acquire a baseline dataset with the mask at its

factory-calibrated position; (ii) physically translate the mask by a known displacement ($\Delta x \in \{0.25, 0.5, 1.0\}$ px equivalent, verified by micrometer stage) and re-acquire under identical illumination; (iii) reconstruct both datasets with the factory mask specification and compute the PSNR degradation and measurement residual; (iv) apply PWM autonomous calibration and measure recovery. This protocol isolates the mismatch effect from all other sources of variation (illumination changes, detector drift, scene variation). Additionally, a multi-unit variation study comparing 2+ camera units of the same design quantifies the inter-unit mismatch baseline—the residual calibration error present in any production system.

Clinical CT phantom configuration. For clinical translation, PWM is evaluated on CT quality assurance using the ACR CT accreditation phantom (Gammex 464)¹². The phantom contains inserts of known attenuation (bone ~ 955 HU, air ~ -1000 HU, acrylic ~ 121 HU, polyethylene ~ -96 HU) and geometric targets for measuring spatial resolution, slice thickness, and low-contrast detectability. Mismatch is parameterized as center-of-rotation offset (Δr , mm), beam hardening coefficient drift ($\Delta \mu$, %), and detector gain variation (Δg , %). Nine ACR-aligned metrics are computed automatically: CT number accuracy (5 materials), geometric accuracy (± 2 mm tolerance), slice thickness (± 1.5 mm), uniformity (≤ 5 HU), noise standard deviation, spatial resolution (≥ 5 lp/cm), low-contrast detectability (≥ 4 targets), and artifact score.

Clinical MRI validation configuration. For MRI clinical validation, PWM processes multi-coil k -space data from public datasets (fastMRI¹¹). Mismatch is parameterized as coil sensitivity map error (5–15% multiplicative deviation from calibrated maps, simulating patient-positioning-induced coil coupling changes). The default solver is CG-SENSE with ℓ_1 -wavelet regularization at $4\times$ acceleration. Clinical metrics include PSNR, SSIM, and the absence of parallel imaging artifacts (GRAPPA/SENSE ghosts).

Statistical Analysis

Per-scene reporting. All metrics are reported per scene, not merely as dataset averages. This enables identification of scene-dependent failure modes (*e.g.*, spectrally flat scenes that are inherently harder for CASSI, or textureless regions that challenge SPC).

Summary statistics. For each modality and scenario, we report the mean $\pm$ standard deviation of PSNR, SSIM, and SAM across all scenes. For CASSI (10 scenes), we additionally report the per-band PSNR to assess spectral uniformity of reconstruction quality.

Recovery ratio confidence intervals. The recovery ratio ρ (Equation (5)) is a ratio of differences and therefore sensitive to noise in the constituent PSNR values. We compute

95% confidence intervals via the bootstrap percentile method with $B = 1,000$ resamples. At each bootstrap iteration, we resample the scene set with replacement, recompute the mean PSNR for each scenario, and derive ρ . The 2.5th and 97.5th percentiles of the bootstrap distribution define the 95% CI.

Parameter recovery accuracy. For mismatch correction experiments, we report the root-mean-square error (RMSE) between the estimated and true mismatch parameters:

$$\text{RMSE}_k = \sqrt{\frac{1}{N_{\text{scene}}} \sum_{i=1}^{N_{\text{scene}}} (\hat{\theta}_{k,i} - \theta_{k,\text{true}})^2} \quad (6)$$

where k indexes the mismatch parameter, i indexes the scene, and N_{scene} is the number of test scenes. Uncertainty in the RMSE is estimated via bootstrap ($B = 1,000$).

Ablation significance. Ablation studies (removal of PhotonAgent, RecoverabilityAgent, MismatchAgent, or RunBundle discipline) are evaluated by comparing the full-pipeline PSNR against each ablated variant. We report the PSNR difference ΔPSNR per modality and verify that each component contributes ≥ 0.5 dB across all validated modalities, establishing practical significance.

Code and Data Availability

Source code. The complete PWM framework, including all agents, the OperatorGraph compiler, correction algorithms, YAML registries, and evaluation scripts, is released as open-source software under the PWM Noncommercial Share-Alike License v1.0 at https://github.com/integritynoble/Physics_World_Model. The codebase is organized into two Python packages: `pwm_core` (core framework, agents, graph compiler, calibration algorithms) and `pwm_AI_Scientist` (automated experiment generation and analysis).

Reconstruction data. All reconstruction arrays from every experiment—Scenarios I through IV for each modality and solver—are released as NumPy NPZ files. Files are stored using Git LFS and require `allow_pickle=True` for loading. Data ranges are standardized: CASSI and CACTI reconstructions are normalized to $[0, 1]$; SPC reconstructions are in $[0, 255]$.

Experiment manifests. Every experiment is recorded in a RunBundle v0.3.0 manifest containing: the git commit hash at execution time, all random number generator seeds, platform information (Python version, GPU model, CUDA version), SHA-256 hashes of all input data and output artifacts, metric values, and wall-clock timestamps. These manifests enable exact reproduction of every reported result.

731 **Registry data.** All 9 YAML registries (7,034 lines total) that drive the agent system—
 732 including modality definitions, graph templates, photon models, mismatch databases, com-
 733 pression tables, solver configurations, primitive specifications, dataset paths, and acceptance
 734 thresholds—are publicly available in the repository under `packages/pwm_core/contrib/`.
 735 The `ExperimentSpec` JSON schemas used for pipeline input validation are included along-
 736 side worked examples in `examples/`.

737 References

- 738 [1] Ashwin A. Wagadarikar, Renu John, Rebecca Willett, and David J. Brady. Single
 739 disperser design for coded aperture snapshot spectral imaging. *Applied Optics*, 47(10):
 740 B44–B51, 2008. doi: 10.1364/AO.47.000B44.
- 741 [2] Patrick Llull, Xuejun Liao, Xin Yuan, Jianbo Yang, David Kittle, Lawrence Carin,
 742 Guillermo Sapiro, and David J. Brady. Coded aperture compressive temporal imaging.
 743 *Optics Express*, 21(9):10526–10545, 2013. doi: 10.1364/OE.21.010526.
- 744 [3] Michael Lustig, David Donoho, and John M. Pauly. Sparse MRI: The application of
 745 compressed sensing for rapid MR imaging. *Magnetic Resonance in Medicine*, 58(6):
 746 1182–1195, 2007. doi: 10.1002/mrm.21391.
- 747 [4] Emmanuel J. Candès and Michael B. Wakin. An introduction to compressive sampling.
 748 *IEEE Signal Processing Magazine*, 25(2):21–30, 2008. doi: 10.1109/MSP.2007.914731.
- 749 [5] David L. Donoho. Compressed sensing. *IEEE Transactions on Information Theory*,
 750 52(4):1289–1306, 2006. doi: 10.1109/TIT.2006.871582.
- 751 [6] Singanallur V. Venkatakrishnan, Charles A. Bouman, and Brendt Wohlberg. Plug-
 752 and-play priors for model based reconstruction. In *Proceedings of the IEEE Global*
 753 *Conference on Signal and Information Processing (GlobalSIP)*, pages 945–948, 2013.
 754 doi: 10.1109/GlobalSIP.2013.6737048.
- 755 [7] Vishal Monga, Yuelong Li, and Yonina C. Eldar. Algorithm unrolling: Interpretable,
 756 efficient deep learning for signal and image processing. *IEEE Signal Processing Maga-*
 757 *zine*, 38(2):18–44, 2021. doi: 10.1109/MSP.2020.3016905.
- 758 [8] Yuanhao Cai, Jing Lin, Xiaowan Hu, Haoqian Wang, Xin Yuan, Yulun Zhang, Radu
 759 Timofte, and Luc Van Gool. Mask-guided spectral-wise transformer for efficient hy-
 760 perspectral image reconstruction. In *Proceedings of the IEEE/CVF Conference on*
 761 *Computer Vision and Pattern Recognition (CVPR)*, pages 17502–17511, 2022.
- 762 [9] Inchang Choi, Daniel S. Jeon, Giljoo Nam, Diego Gutierrez, and Min H. Kim. High-
 763 quality hyperspectral reconstruction using a spectral prior. *ACM Transactions on*

- 764 *Graphics (Proceedings of SIGGRAPH Asia)*, 36(6):218:1–218:13, 2017. doi: 10.1145/
765 3130800.3130810.
- 766 [10] José M. Bioucas-Dias and Mário A. T. Figueiredo. A new TwIST: Two-step iterative
767 shrinkage/thresholding algorithms for image restoration. *IEEE Transactions on Image*
768 *Processing*, 16(12):2992–3004, 2007. doi: 10.1109/TIP.2007.909319.
- 769 [11] Jure Zbontar, Florian Knoll, Anuroop Sriram, Tullie Murrell, Zhengnan Huang,
770 Matthew J. Muckley, Aaron Defazio, Ruben Stern, Patricia Johnson, Mary Bruno,
771 Marc Parente, Krzysztof J. Geras, Joe Katsnelson, Hersh Chandarana, Zizhao Zhang,
772 Michal Drozdal, Adriana Romero, Michael Rabbat, Pascal Vincent, Nafissa Yakubova,
773 James Pinkerton, Duo Wang, Erich Owens, C. Lawrence Zitnick, Michael P. Recht,
774 Daniel K. Sodickson, and Yvonne W. Lui. fastMRI: An open dataset and benchmarks
775 for accelerated MRI. *arXiv preprint arXiv:1811.08839*, 2018.
- 776 [12] Hu Chen, Yi Zhang, Mannudeep K. Kalra, Feng Lin, Yang Chen, Peixi Liao, Jiliu
777 Zhou, and Ge Wang. Low-dose CT with a residual encoder-decoder convolutional
778 neural network. *IEEE Transactions on Medical Imaging*, 36(12):2524–2535, 2017. doi:
779 10.1109/TMI.2017.2715284.
- 780 [13] Martin Uecker, Peng Lai, Mark J. Murphy, Patrick Virtue, Michael Elad, John M.
781 Pauly, Shreyas S. Vasanawala, and Michael Lustig. ESPIRiT — an eigenvalue approach
782 to autocalibrating parallel MRI: where SENSE meets GRAPPA. *Magnetic Resonance*
783 *in Medicine*, 71(3):990–1001, 2014. doi: 10.1002/mrm.24751.
- 784 [14] Andrew M. Maiden and John M. Rodenburg. An improved ptychographical phase
785 retrieval algorithm for diffractive imaging. *Ultramicroscopy*, 109(10):1256–1262, 2009.
786 doi: 10.1016/j.ultramic.2009.05.012.
- 787 [15] Vegard Antun, Francesco Renna, Clarice Poon, Ben Adcock, and Anders C. Hansen.
788 On instabilities of deep learning in image reconstruction and the potential costs of
789 AI. *Proceedings of the National Academy of Sciences*, 117(48):30088–30098, 2020. doi:
790 10.1073/pnas.1907377117.
- 791 [16] J. M. Rodenburg and H. M. L. Faulkner. A phase retrieval algorithm for shifting
792 illumination. *Applied Physics Letters*, 85(20):4795–4797, 2004. doi: 10.1063/1.1823034.
- 793 [17] Klaas P. Pruessmann, Markus Weiger, Markus B. Scheidegger, and Peter Boesiger.
794 SENSE: Sensitivity encoding for fast MRI. *Magnetic Resonance in Medicine*, 42(5):
795 952–962, 1999. doi: 10.1002/(SICI)1522-2594(199911)42:5<952::AID-MRM16>3.0.
796 CO;2-S.
- 797 [18] L. A. Feldkamp, L. C. Davis, and J. W. Kress. Practical cone-beam algorithm. *Journal*
798 *of the Optical Society of America A*, 1(6):612–619, 1984. doi: 10.1364/JOSAA.1.000612.

- [19] Chengshuai Yang. The finite primitive basis theorem for computational imaging: Formal foundations of the OperatorGraph representation. *Companion paper*, 2026. Full formal definitions, primitive semantics, typed DAG denotation, and constructive proof.
- [20] Chengshuai Yang. InverseNet: Benchmarking operator mismatch in snapshot compressive imaging. Technical report, NextGen PlatformAI C Corp, 2026.
- [21] Xin Yuan. Generalized alternating projection based total variation minimization for compressive sensing. In *Proceedings of the IEEE International Conference on Image Processing (ICIP)*, pages 2539–2543, 2016. doi: 10.1109/ICIP.2016.7532817.
- [22] Xiaowan Hu, Yuanhao Cai, Jing Lin, Haoqian Wang, Xin Yuan, Yulun Zhang, Radu Timofte, and Luc Van Gool. HDNet: High-resolution dual-domain learning for spectral compressive imaging. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 17542–17551, 2022.
- [23] Lishun Wang, Miao Cao, and Xin Yuan. EfficientSCI: Densely connected network with space-time factorization for large-scale video snapshot compressive imaging. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 18477–18486, 2023.
- [24] Marco F. Duarte, Mark A. Davenport, Dharmpal Takhar, Jason N. Laska, Ting Sun, Kevin F. Kelly, and Richard G. Baraniuk. Single-pixel imaging via compressive sampling. *IEEE Signal Processing Magazine*, 25(2):83–91, 2008. doi: 10.1109/MSP.2007.914730.
- [25] Yichen Qu, Lishun Wang, Ao Jiang, and Xin Yuan. HATNet: Hierarchical attention transformer network for single-pixel imaging. In *Proceedings of the AAAI Conference on Artificial Intelligence*, 2024.
- [26] Tao Huang, Weisheng Dong, Xin Yuan, Jinjian Wu, and Guangming Shi. Deep gaussian scale mixture prior for spectral compressive imaging. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 16216–16225, 2021.
- [27] Yuanhao Cai, Jing Lin, Xiaowan Hu, Haoqian Wang, Xin Yuan, Yulun Zhang, Radu Timofte, and Luc Van Gool. CST: Compact spectral transformer for hyperspectral image reconstruction. In *Proceedings of the European Conference on Computer Vision (ECCV)*, 2022.
- [28] Yang Liu, Xin Yuan, Jinli Suo, David J. Brady, and Qionghai Dai. Rank minimization for snapshot compressive imaging. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 41(12):2990–3006, 2019.




Figure 1: **PWM overview.** The Physics World Models pipeline. **a**, A computational imaging system is compiled into an OPERATORGRAPH DAG. **b**, The TRIAD DECOMPOSITION diagnostic agents evaluate each gate. **c**, The dominant gate is identified and a TRIADREPORT is produced. **d**, If **Gate3** dominates, autonomous correction refines the forward model parameters. **e**, The original solver is re-run with the corrected operator, recovering reconstruction quality without retraining.




Figure 2: **OperatorGraph IR and Physics Fidelity Ladder.** **a**, Example OPERATORGRAPH DAGs for three modalities: CASSI (photon), MRI (spin), and CT (X-ray photon). Each node wraps a primitive operator; edges define data flow. **b**, The Physics Fidelity Ladder. Tier 1: linear shift-invariant. Tier 2: linear shift-variant. Tier 3: nonlinear ray/wave-based. Tier 4: full-wave/Monte Carlo. **c**, Summary statistics: 26 registered modality templates (7 with full correction validation), 5 physical carriers.




Figure 3: **Triad Decomposition structure and gate binding.** **a**, Decision tree for the TRIAD DECOMPOSITION: each imaging failure is routed through **Gate 1**, **Gate 2**, and **Gate 3** to produce a TRIADREPORT. **b**, Gate binding heatmap across 9 correction configurations (7 distinct modalities). Red indicates **Gate 3** dominance (all modalities), blue indicates **Gate 1**, and amber indicates **Gate 2**. **c**, Recovery ratio ρ distribution across all 9 correction configurations.




Figure 4: **Correction results across 9 validated configurations.** Bar chart showing correction gain Δ_{corr} (dB) for each of the 9 correction configurations (7 distinct modalities), grouped by carrier family. Incoherent photon (CASSI, CACTI, SPC, Lensless) and coherent photon (Ptychography) in blue; spin (MRI) in purple; X-ray (CT) in red; generic (Matrix) in grey.




Figure 5: **CASSI and CACTI deep dive.** **a**, CASSI: PSNR across 4 scenarios for GAP-TV, MST-L, and HDNet under combined mask-geometry-plus-dispersion mismatch. The uniform collapse under Scenario II (range 20.83–21.88 dB) confirms operator-driven failure; oracle recovery varies by solver ($\rho = 0.22$ –0.46). **b**, CACTI: four methods across 4 scenarios, showing up to 20.58 dB mismatch degradation and substantial correction recovery across all solvers.




Figure 6: **Zero-shot generalization across carrier families.** Correction gain (dB) when beam-search and gradient-refinement hyperparameters are tuned on photon-domain modalities and transferred without modification to coherent-photon, spin, and X-ray domains. Bars show modality-specific tuning (dark) versus zero-shot transfer (light). Transfer efficiency is high across all carrier families, confirming the carrier-agnostic nature of the PWM correction pipeline.




Figure 7: **Hardware validation on real CASSI and CACTI instruments.** **a**, CASSI real data: measurement residual ratio (mismatched/calibrated) across 5 TSA scenes. GAP-TV shows 1.8 $\times$ mean ratio. **b**, CACTI real data: residual ratio across 4 scenes. GAP-TV shows 10.4 $\times$ mean ratio; PnP-FFDNet shows 2.0 $\times$. **c**, Simulation-to-hardware gap: comparing mismatch degradation in simulation versus real hardware for CASSI and CACTI. **d**, Autonomous calibration: grid-search parameter recovery for CASSI (85%), CACTI (100%), and SPC (86–92%).




Figure 8: **Visual reconstruction comparison across three modalities.** Each row shows one modality (CASSI, CACTI, MRI); columns from left to right: ground truth, Scenario I (ideal operator), Scenario II (mismatched operator—note structured artifacts), Scenario III (PWM-corrected operator), and error map ($|Sc. II - GT|$). Mismatch produces severe structured artifacts (column 3) that are qualitatively distinct from noise. PWM correction (column 4) substantially reduces these artifacts.




Figure 9: **Basis-growth saturation and primitive decomposition.** **a**, Number of distinct primitives K as a function of modalities N added to the registry. The curve saturates at $K = 10$ for $N \geq 27$; annotated points mark the introduction of each primitive. **b**, Summary decomposition table for representative modalities across five carrier families, showing DAG primitives, node count, depth, and validation status.