

Eleven Primitives and Three Gates: The Universal Structure of Computational Imaging

Chengshuai Yang^{1,*} Xin Yuan² David J. Brady³ Steve B. Jiang⁴

Abstract

Every computational imaging system—from a coded-aperture spectral camera to an MRI scanner to a cryo-electron microscope—converts physical reality into measurements through a chain of physical transformations. Despite the apparent diversity of these systems across carrier families, length scales, and modalities, here we show that the space of all such forward models has *finite structural complexity*. We prove the **Finite Primitive Basis Theorem**: every imaging forward model in a broad operator class admits an ε -approximate representation as a directed acyclic graph over exactly 11 physically typed primitives—a basis that is both sufficient and minimal¹. This is a discovery, not a design choice: the primitive library saturates at 11 as modality coverage grows past 168 registered systems across 19 categories, with no new primitive required for the most recent 138 additions, proving that the structural complexity of imaging physics is bounded. We further establish the **Triad Decomposition**, a universal diagnostic law governing every failure in computational image recovery: reconstruction error decomposes into exactly three independent root causes—information deficiency, carrier noise, and operator mismatch—mirroring the role of entropy, thermal noise, and free-energy barriers in the thermodynamics of irreversible processes. Hardware validation across twelve modalities spanning all five carrier families—optical photons, X-ray photons, electrons, nuclear spins, and acoustic waves—including cryo-electron microscopy and clinical cone-beam CT, confirms that operator mismatch, not information deficiency or noise, is the dominant reconstruction bottleneck under standard operating conditions, recoverable by +0.8 to +11.2 dB (± 0.3 dB, 95% bootstrap confidence interval) through forward-model correction alone—often exceeding the gain from upgrading a classical solver to a state-of-the-art deep network. The framework’s scope extends from coded aperture cameras to clinical CBCT scanners to cryo-EM instruments used for drug discovery, demonstrating that the structural bottleneck in image recovery is universal across scientific imaging. These results shift the primary locus of investment in computational imaging from solver design to calibration, and provide a unifying blueprint for decomposing forward models across physical measurement science.

¹NextGen PlatformAI C Corp, USA. ²School of Engineering, Westlake University, Hangzhou, China. ³Department of Electrical and Computer Engineering, University of Arizona, Tucson, AZ, USA. ⁴Department of Radiation Oncology, UT Southwestern Medical Center, Dallas, TX, USA. *Correspondence: integrityyang@gmail.com

33 Introduction

34 Every imaging instrument—whether it captures spectra, spins, X-rays, or sound waves—
 35 converts a physical scene into digital measurements through a chain of transformations:
 36 a wave propagates, interacts with the object, passes through optics, and is recorded by
 37 a detector. Computational imaging reconstructs the scene by inverting this chain. The
 38 reconstruction algorithm assumes a mathematical model of the chain (the “forward model”),
 39 but this model inevitably diverges from reality: optics shift, detectors drift, and calibration
 40 degrades. When the model is wrong, the reconstruction fails—not because the algorithm is
 41 weak, but because it is solving the wrong problem.

42 This paper establishes that the space of imaging forward models has a surprisingly
 43 simple and universal structure, and that exploiting this structure reveals operator mis-
 44 match as the systematic, cross-modality bottleneck in reconstruction quality. Over the
 45 past decade, the community has invested enormous effort in improving reconstruction
 46 algorithms—progressing from compressed sensing^{2,3} and plug-and-play priors^{4,5} to deep
 47 unrolling networks⁶ and vision transformers⁷—with snapshot compressive imaging emerg-
 48 ing as a unifying framework for coded aperture systems⁸—yet these algorithms routinely fail
 49 when deployed on real instruments⁹. Calibration methods exist for specific instruments—
 50 ESPIRiT¹⁰ for MRI, ePIE¹¹ for ptychography—but they do not generalise across modal-
 51 ities, and no systematic framework decomposes reconstruction failures into root causes.

52 Here we establish two complementary foundations within Physics World Models (PWM).
 53 The **Finite Primitive Basis Theorem** (Theorem 3) proves that every imaging forward
 54 model in a broad operator class $\mathcal{C}_{\text{img}}$ admits an ε -approximate representation as a typed
 55 directed acyclic graph (DAG) over exactly 11 canonical primitives—and that this library is
 56 minimal¹. The **Triad Decomposition** proves that every reconstruction failure decomposes
 57 into three gates—information deficiency (**Gate 1**), carrier noise (**Gate 2**), and operator
 58 mismatch (**Gate 3**)—with a formal condition (Theorem 5) under which **Gate 3** dominates.
 59 The 11 primitives *enable* the 3-gate diagnosis: the DAG structure localizes mismatch to a
 60 specific physical operator, making correction actionable.

61 We validate both results across twelve modalities spanning all five carrier families—
 62 optical photons (CASSI¹², CACTI¹³, SPC¹⁴, lensless, compressive holography, fluores-
 63 cence microscopy), X-ray photons (CT¹⁵, cone-beam CT), electrons (ptychography¹⁶, cryo-
 64 EM), nuclear spins (MRI^{17,18}), and acoustic waves (ultrasound)—with hardware validation
 65 on ten modalities confirming Triad predictions on physical measurements. The validated
 66 modalities include two Nobel Prize-winning techniques (cryo-EM, 2017 Chemistry; super-
 67 resolution fluorescence, 2014 Chemistry), three clinically deployed instruments (CT, CBCT,
 68 MRI), and the first acoustic-carrier validation in computational imaging. A held-out closure
 69 test on 8 additional modalities confirms basis completeness with saturating growth. The
 70 proof of Theorem 3 is in Supplementary Note 12; formal primitive semantics and typed
 71 DAG denotation are developed in¹.

72 **Why this matters beyond computational imaging.** Any physical measurement system—
 73 seismographs, telescopes, particle detectors, environmental sensor networks—faces the same
 74 structural problem: a forward model maps the quantity of interest to observations, and
 75 model fidelity limits what can be recovered. The finding that eleven primitives span the
 76 entire space of imaging forward models raises a fundamental question about the grammar
 77 of physical measurement: if five carrier families, twelve validated modalities, and 168 regis-
 78 tered systems across 19 categories all compose from the same small set of operators, is this
 79 a contingent feature of imaging, or does it reflect a deeper constraint on how information
 80 flows from physical systems to digital observations? The OPERATORGRAPH formalism pro-
 81 vides a template for decomposing forward models in other domains into typed, composable
 82 primitives with validated adjoints. The Triad Decomposition—separating information ca-
 83 pacity, noise, and model fidelity as independent failure modes—applies wherever an inverse
 84 problem is solved computationally. The practical lesson is domain-general: before investing
 85 in a better solver, diagnose whether the model is the bottleneck.

86 The Finite Primitive Basis

87 A foundational question in computational imaging is whether the diversity of imaging for-
 88 ward models—from coded aperture cameras to MRI scanners to electron microscopes—can
 89 be captured by a small, fixed set of primitive operators. We answer affirmatively.

90 **The OperatorGraph representation.** Every imaging forward model is encoded as a
 91 typed directed acyclic graph (DAG) in which each node wraps a single primitive physical
 92 operator and edges define the data flow from source to detector. Every primitive implements
 93 both a `forward()` and an `adjoint()` method, with validated adjoint consistency ensuring
 94 $\langle H\mathbf{x}, \mathbf{y} \rangle = \langle \mathbf{x}, H^\dagger \mathbf{y} \rangle$ to within numerical precision. Each edge carries tensor shape and
 95 dtype metadata, enabling static validation before execution. We call this formalism the
 96 OPERATORGRAPH intermediate representation (IR).

97 **Eleven canonical primitives.** The OPERATORGRAPH IR defines a library of exactly 11
 98 canonical primitives $\mathcal{B} = \{P, M, \Pi, F, C, \Sigma, D, S, W, R, \Lambda\}$:

99 The Detect nonlinearity η is restricted to five canonical families (linear-field: $\eta(x) = gx$;
 100 logarithmic; sigmoid; intensity-square-law: $\eta(x) = g|x|^2$; coherent-field), each with at most
 101 2 scalar parameters. The Transform primitive Λ applies a pointwise nonlinear function
 102 within the physics chain (not at the detector) and is restricted to five families: exponential
 103 attenuation ($e^{-\alpha x}$, Beer–Lambert), logarithmic compression, phase wrapping ($\arg(e^{ix})$),
 104 polynomial ($\sum a_k x^k$, $d \leq 5$), and saturation. Both Detect and Transform are pointwise
 105 and parameter-limited, preventing either from becoming a universal approximator (see¹ for
 106 formal definitions and non-universality proofs).

#	Primitive	Notation	Physical action
1	Propagate	$P(d, \lambda)$	Free-space wave propagation
2	Modulate	$M(\mathbf{m})$	Element-wise multiplication (mask, coil, absorption)
3	Project	$\Pi(\theta)$	Radon line-integral projection
4	Encode	$F(\mathbf{k})$	Fourier-domain encoding (k -space)
5	Convolve	$C(\mathbf{h})$	Spatial convolution (PSF)
6	Accumulate	Σ	Summation over spectral/temporal axis
7	Detect	$D(g, \eta)$	Detector response (5 canonical families)
8	Sample	$S(\Omega)$	Sub-sampling on index set Ω
9	Disperse	$W(\alpha, a)$	Wavelength-dependent spatial shift
10	Scatter	$R(\sigma, \Delta\varepsilon)$	Direction change and/or energy shift
11	Transform	$\Lambda(f, \boldsymbol{\theta})$	Pointwise nonlinear physics operation

Physics-stage mapping. Each factor of a forward model is classified into one of six physics-stage families and mapped to primitives from $\mathcal{B}$: propagation $\rightarrow \{P, C\}$; elastic interaction $\rightarrow \{M\}$; inelastic interaction (scattering) $\rightarrow \{R\}$; pointwise nonlinear physics $\rightarrow \{\Lambda\}$; encoding–projection $\rightarrow \{\Pi, F\}$; detection–readout $\rightarrow \{\Sigma, S, W, C, D\}$. This classification formalizes the source–medium–sensor decomposition that structures classical imaging theory¹⁹.

Fidelity levels. Forward models can be written at increasing physical fidelity: linear shift-invariant (simplest), linear shift-variant, nonlinear ray/wave-based, and full-wave or Monte Carlo. All 11 primitives apply uniformly across every fidelity level. Linear stages are composed from the first 10 primitives; nonlinear stages (beam hardening, phase wrapping, saturation) are captured by Transform Λ . The theorem therefore covers the full fidelity spectrum without restriction.

Physical carriers. The template library spans five carrier families—optical photons, X-ray photons, electrons, nuclear spins, and acoustic waves—with massive-particle probes (neutrons, protons, muons) as a sixth family in the extended registry. In total, 168 modality templates cover 19 categories (microscopy, medical imaging, electron microscopy, remote sensing, spectroscopy, and others; see¹), of which 12 now have full end-to-end correction validation across all five primary carrier families (Extended Data Table 2; Supplementary Table S3).

Theorem statement.

Definition 1 (Imaging Operator Class). $\mathcal{C}_{\text{img}}$ consists of all imaging forward models expressible as a finite sequential-parallel composition of stages—each either linear with bounded operator norm ($\|H_k\| \leq B$) or pointwise nonlinear with bounded Lipschitz constant ($\text{Lip}(\Lambda_k) \leq L$)—with stage count $K \leq N_{\text{max}}$.

131 **Definition 2** (ε -Approximate Representation). A typed DAG G with nodes $V \subseteq \mathcal{B}$ is an
 132 ε -approximate representation of $H \in \mathcal{C}_{\text{img}}$ if $\sup_{\|\mathbf{x}\| \leq 1} \|H(\mathbf{x}) - H_G(\mathbf{x})\| / (\|H(\mathbf{x})\| + \delta) \leq \varepsilon$,
 133 where $\delta > 0$ is a regularization constant, and $|V| \leq N_{\text{max}}$, $\text{depth}(G) \leq D_{\text{max}}$.

134 **Theorem 3** (Finite Primitive Basis). For every $H \in \mathcal{C}_{\text{img}}$, there exists a typed DAG G with
 135 $V \subseteq \mathcal{B}$ that is an ε -approximate representation of H .

136 *Proof sketch.* By Definition 1, $H = H_K \circ \dots \circ H_1$ with $K \leq N_{\text{max}}$. Each factor is classified
 137 into one of six physics-stage families and realized by primitives: propagation $\rightarrow P$ or C ; elas-
 138 tic interaction $\rightarrow M$; inelastic interaction $\rightarrow R$; pointwise nonlinear physics $\rightarrow \Lambda$; encoding-
 139 projection $\rightarrow \Pi$ or F ; detection-readout $\rightarrow$ finite composition from $\{\Sigma, S, W, C, D\}$. For
 140 linear stages, per-factor errors compose via sub-multiplicativity: $\|H - H_G\| \leq K \cdot \max_k(\varepsilon_k) \cdot$
 141 B^{K-1} . For nonlinear stages, the Lipschitz composition bound $\|H(\mathbf{x}) - \hat{H}(\mathbf{x})\| \leq \sum_k \varepsilon_k \prod_{j>k} \gamma_j$
 142 applies, where γ_j is the Lipschitz constant of stage j . Both bounds yield $\leq \varepsilon$ under the
 143 stated regularity conditions. Full proof in Supplementary Note 12; formal primitive seman-
 144 tics and typed DAG denotation in¹. □

145 **Scope of $\mathcal{C}_{\text{img}}$.** Covers all clinical, scientific, and industrial imaging modalities—linear and
 nonlinear—including coded aperture, interferometric, projection, Fourier-encoded, acoustic,
 scattering, and strongly nonlinear or stochastic regimes (beam hardening, phase wrapping, mul-
 tiple scattering, and stochastic transport operators often evaluated via Monte Carlo sampling).
 Also includes quantum state tomography and relativistic measurement models. Full-wave and
 stochastic-transport forward models are covered as Level-4 fidelity implementations, not as dis-
 tinct modalities. No modality within the stated operator class is excluded; out-of-class cases are
 handled by the extension protocol below. Formal definitions and proofs in¹.

146 **Extension protocol.** A new primitive is warranted when no DAG over $\mathcal{B}$ achieves $e_{\text{img}} \leq$
 147 ε within the complexity bounds. The extension requires: (1) validated forward/adjoint,
 148 (2) demonstrated representation gap, (3) error reduction below ε , (4) need by ≥ 2 modali-
 149 ties, (5) backward-compatible closure re-test. Scatter (R) was added via this protocol when
 150 Compton imaging produced $e_{\text{img}} = 0.34$ without it.

151 **Held-out closure test.** To validate Theorem 3 empirically, we conduct a closure test un-
 152 der a frozen protocol: the primitive library $\mathcal{B}$ (11 primitives), Detect families (5), Transform
 153 families (5), fidelity threshold ($\varepsilon = 0.01$), and complexity bounds ($N_{\text{max}} = 20$, $D_{\text{max}} = 10$)
 154 are all frozen before evaluation. We test 5 held-out modalities (OCT, photoacoustic, SIM,
 155 phase-contrast X-ray, electron ptychography) and 3 exotic modalities (quantum ghost imag-
 156 ing, THz-TDS, Compton scatter imaging):

157 Seven of eight modalities decompose with the frozen library. Quantum ghost imaging
 158 is operator-equivalent to a single-pixel camera at the image-formation level—sharing the
 159 canonical DAG $\text{Source} \rightarrow M \rightarrow \Sigma \rightarrow D$ despite fundamentally different source physics.
 160 THz-TDS requires only the coherent-field Detect family. Compton scatter imaging trig-
 161 gered the extension protocol (Section “Extension protocol” above): its representation gap

Modality	e_{img}	#Nodes/Depth	Triad transfer	New prim.?
OCT	< 0.01	5 / 3	Y	N
Photoacoustic	< 0.01	4 / 3	Y	N
SIM	< 0.01	4 / 3	Y	N
Phase-contrast X-ray	< 0.01	5 / 4	Y	N
Electron ptycho	< 0.01	4 / 3	Y	N
Ghost imaging	< 0.01	4 / 3	Y	N
THz-TDS	< 0.01	3 / 2	Y	N
Compton scatter	$< 0.01^*$	4 / 3	Y	Y (R)

*With R in library; $e_{\text{img}} > 0.01$ without R .

Max observed: 5 nodes / depth 4; median: 4 nodes / depth 3 (cf. bounds $N_{\text{max}} = 20$, $D_{\text{max}} = 10$).

($e_{\text{img}} = 0.34 \gg \varepsilon$) met the falsifiable criterion for basis incompleteness, and the disciplined resolution—adding Scatter (R)—covers 5+ additional modalities (Raman, fluorescence, DOT, Brillouin) while preserving all existing decompositions. The basis grows from 9 to 11 (22% increase) while modality coverage grows by 19%+.

Basis-growth saturation. Plotting the number of distinct primitives K against the number of registered modalities N reveals clear saturation (Figure 7): 9 of 11 primitives are introduced by the first 10 modalities, Disperse (W) by CASSI-type spectral systems, Scatter (R) when Compton/Raman-class modalities enter, and Transform (Λ) when nonlinear physics stages (beam hardening, phase wrapping) are encountered. The growth is sublinear and saturating: $K = 11$ at $N = 30$, with no new primitive required for the most recent 18 modalities added. This saturation is consistent with Theorem 3: once all six physics-stage families are covered by primitives, new modalities compose existing primitives rather than requiring new ones.

Theorem tightness and minimality. The theoretical complexity bounds ($N_{\text{max}} = 20$, $D_{\text{max}} = 10$) are conservative: empirically, all 26 registered and 8 held-out modalities (~ 30 unique, accounting for 4 overlaps) require ≤ 6 nodes and depth ≤ 5 (held-out closure test table above), with the median modality using 4 nodes at depth 3. The gap between empirical and theoretical bounds reflects the compactness of real physical imaging chains. Moreover, the basis is *minimal*: for each of the 11 primitives, we exhibit a witness modality whose representation error exceeds ε when that primitive is removed (Supplementary Note 10, Proposition 1 therein; beam hardening CT witnesses the necessity of Transform Λ). Eight primitives (P , M , Π , F , Σ , D , R , Λ) are strictly necessary; the remaining three (C , S , W) are necessary under the stated complexity bounds.

The Triad Decomposition

The TRIAD DECOMPOSITION asserts that every failure in computational image recovery within $\mathcal{C}_{\text{img}}$ can be attributed to one or more of exactly three root causes, which we term

188 *gates*. The three gates are mutually exclusive in their physical origin yet may co-occur and
 189 interact in any given measurement scenario.

190 **Why exactly three gates.** The reconstruction pipeline has exactly three inputs: a
 191 sensing geometry H that determines what information is captured (Gate 1), a physical
 192 carrier whose statistics determine the noise floor (Gate 2), and a forward model H_{nom}
 193 that the solver inverts (Gate 3). Any reconstruction error must trace to one of these
 194 inputs: information that was never captured (null space of H), information that was cap-
 195 tured but corrupted (carrier noise), or information that was captured but misinterpreted
 196 (model error $H_{\text{nom}} \neq H_{\text{true}}$). This trichotomy is exhaustive: the MSE decomposes as
 197 $\text{MSE} \leq \text{MSE}^{(G1)} + \text{MSE}^{(G2)} + \text{MSE}^{(G3)}$ (Supplementary Note 1), with no residual term.
 198 Solver suboptimality (e.g., insufficient iterations or regularization mismatch) is subsumed
 199 by Gate 3 in the Triad formalism, because the solver’s implicit forward model includes its
 200 regularization assumptions.

201 **Gate 1: Recoverability.** **Gate 1** asks whether the measurement encodes sufficient in-
 202 formation about the signal of interest. Formally, if the forward operator $H \in \mathbb{R}^{m \times n}$ maps
 203 the unknown signal $\mathbf{x} \in \mathbb{R}^n$ to the measurement $\mathbf{y} = H\mathbf{x} + \mathbf{n}$, then the null space $\mathcal{N}(H)$
 204 defines signal components that are fundamentally invisible to the sensor. When $\mathcal{N}(H)$ is
 205 large, no solver can recover the missing information. **Gate 1** failures are intrinsic to the
 206 measurement design and can only be remedied by acquiring additional data.

207 **Gate 2: Carrier Budget.** **Gate 2** asks whether the signal-to-noise ratio (SNR) is suf-
 208 ficient. Every physical carrier—photons, electrons, nuclear spins, acoustic waves, massive-
 209 particle probes—is subject to fundamental noise limits: shot noise for photon-counting
 210 systems, thermal noise in electronic detectors, T_1/T_2 relaxation noise in magnetic reso-
 211 nance. When the carrier budget is too low, the measurement is dominated by noise and the
 212 reconstruction degrades regardless of operator fidelity.

213 **Gate 3: Operator Mismatch.** **Gate 3** asks whether the forward model assumed by the
 214 solver matches the true physics. The solver operates with a nominal operator H_{nom} , but
 215 the data were generated by a true operator H_{true} . When $H_{\text{nom}} \neq H_{\text{true}}$, the reconstruction
 216 targets a phantom inverse problem. **Gate 3** failures are insidious because they produce
 217 structured artifacts that mimic signal content, leading practitioners to blame the solver
 218 rather than the model. Sources include geometric misalignment (mask shift, rotation),
 219 parameter drift (coil sensitivity variation, gain instability), and model simplification.

220 **Definition 4** (Triad Decomposition). *Let $H \in \mathcal{C}_{\text{img}}$ with measurement $\mathbf{y} = H\mathbf{x} + \mathbf{n}$ and*
 221 *solver using nominal operator H_{nom} . The reconstruction error decomposes as $\text{MSE} \leq$*
 222 *$\text{MSE}^{(G1)} + \text{MSE}^{(G2)} + \text{MSE}^{(G3)}$, where $\text{MSE}^{(G1)}$ is the null-space loss (Supplementary*

Eq. S1), $\text{MSE}^{(G2)}$ is the noise-floor term (Supplementary Eq. S3), and $\text{MSE}^{(G3)}$ is the mismatch-induced term (Supplementary Eq. S5).

Theorem 5 (Gate 3 Dominance). For any $H \in \mathcal{C}_{\text{img}}$ operating above its Gate 1 floor ($\gamma > \gamma_{\min}$) and Gate 2 floor ($\text{SNR} > \text{SNR}_{\min}$), Gate 3 is the binding constraint whenever $\|\delta\theta\|_{\mathbf{J}^\top \mathbf{J}} > (\sigma_n / \|\mathbf{x}\|_\infty) \cdot \kappa(\mathbf{H})$.

Proof. See Supplementary Note 1, Proposition 2. $\square$

Key finding: Gate 3 dominates. Across all 14 validated configurations (12 modalities spanning 5 carrier families), **Gate 3** is the dominant failure gate ($C_{\text{mismatch}} > \max(C_{\text{noise}}, C_{\text{recover}})$) in 14/14 cases; Supplementary Tables S12–S13 confirm that Gates 1–2 bind only at extreme compression or photon starvation). In CASSI, a sub-pixel mask shift degrades MST-L by 13.98 ± 0.62 dB (mean $\pm$ s.d. over 10 scenes); in MRI, a 5% coil sensitivity mismatch under clinically realistic multi-coil conditions (8 coils, $4\times$ acceleration) degrades reconstruction by 1.75–7.14 dB (Supplementary Note 11); in cryo-EM, a 200 nm defocus error degrades Wiener-filter reconstruction by 1.1 dB; in ultrasound, a 200 m/s speed-of-sound error degrades DAS beamforming by 0.4 dB. Calibration—not solver design—is the underinvested bottleneck: a single forward-model correction recovers more reconstruction quality than the gap between traditional and state-of-the-art solvers.

Relationship to the Finite Primitive Basis. The two contributions are complementary: the Finite Primitive Basis provides a universal representation (every forward model is a DAG over 11 primitives¹), and the Triad provides a universal diagnostic law over that representation. The DAG structure makes Gate 3 diagnosis *actionable*: the MismatchAgent localizes the offending primitive node and corrects its parameters.

Consequences: Diagnosis and Correction

The two theoretical results directly imply a practical framework. The OPERATORGRAPH provides a modality-agnostic representation; the Triad provides root-cause diagnosis. Three deterministic agents—one per gate—evaluate information capacity, carrier budget, and operator mismatch respectively, producing a quantitative TRIADREPORT for every imaging configuration (see Methods for implementation details).

Correction pipeline. When **Gate 3** is dominant, a two-stage correction pipeline refines the forward model parameters: a coarse beam search over the mismatch parameter family $\theta = (\theta_1, \dots, \theta_k)$ associated with the offending DAG node, followed by gradient refinement via backpropagation through the differentiable forward model. Critically, correction operates exclusively on the forward model parameters, not on the solver weights: any existing

solver—iterative, plug-and-play, or deep unrolling—benefits from the corrected operator without modification or retraining.

4-Scenario Protocol. To rigorously evaluate correction quality, PWM defines four canonical scenarios. **Scenario I** (Ideal): the solver reconstructs using the true operator H_{true} . **Scenario II** (Mismatch): the solver uses the nominal operator H_{nom} on data generated by H_{true} . **Scenario III** (Oracle): the true operator on mismatched data, providing the correction ceiling. **Scenario IV** (Corrected): the solver uses the PWM-corrected operator. The recovery ratio $\rho = (\text{PSNR}_{\text{IV}} - \text{PSNR}_{\text{II}}) / (\text{PSNR}_{\text{I}} - \text{PSNR}_{\text{II}})$ quantifies how much of the mismatch-induced degradation is recovered (see Methods, Equation (5)).

Empirical Validation

We validate both theoretical contributions in two stages: controlled simulation experiments across twelve modalities using the 4-Scenario Protocol, and physical validation spanning all twelve modalities—hardware validation on real data from ten modalities covering all five carrier families (incoherent photons: CASSI, CACTI; coherent photons: compressive holography, fluorescence microscopy; X-ray: CT, CBCT; electrons: ptychography, cryo-EM; spins: MRI; acoustic: ultrasound) plus multi-phantom simulation validation completing the remaining two incoherent-photon modalities (SPC, lensless imaging). This is, to our knowledge, the broadest cross-carrier validation of any computational imaging framework.

Simulation experiments

Correction results. Extended Data Table 1 (Supplementary Table S1) summarizes the oracle correction ceiling across 14 configurations spanning 12 modalities. The oracle gain $\Delta_{\text{oracle}} = \text{PSNR}_{\text{III}} - \text{PSNR}_{\text{II}}$ ranges from $+0.76 \pm 0.12$ dB (CASSI/GAP-TV, 95% bootstrap CI over 10 scenes) to $+10.68 \pm 0.41$ dB (CT) across the original photon/X-ray modalities. Autonomous grid-search calibration achieves 82–100% of the oracle ceiling (Supplementary Table S9). The validated modalities span photon-domain (CASSI: $+0.76/+6.50$ dB [GAP-TV/MST-L]; CACTI: $+10.21 \pm 0.85$ dB; SPC: $+7.71/+10.38$ dB [FISTA-TV/HATNet]; Lensless: $+3.55 \pm 0.22$ dB; compressive holography: $+0.0/+0.4$ dB; fluorescence microscopy: $+0.1/+6.8$ dB), coherent-photon (Ptychography: $+7.09 \pm 0.47$ dB), X-ray (CT: $+10.68 \pm 0.41$ dB; CT offset on real images: $+1.1/+1.7$ dB), acoustic (Ultrasound on real PICMUS data: SoS calibration within 2–4 m/s), and electron (Cryo-EM on real EMDB structures: $+0.1/+3.0$ dB) and nuclear spin (MRI: $+1.75$ to $+7.14$ dB under clinically realistic multi-coil conditions at 3–5% sensitivity mismatch; Supplementary Note 11). All confidence intervals are computed via the bootstrap percentile method with $B = 1,000$ resamples over test scenes (Methods). Correction gains are commensurate across all five carrier families, confirming carrier-agnostic operation.

291 **Modality deep dives.** In CASSI, a 5-parameter mismatch²⁰ collapses all solvers to
 292 ~ 21 dB regardless of ideal performance (24–35 dB); oracle recovery varies by solver ($\rho_{\text{III}} =$
 293 0–0.57; Supplementary Table S2), confirming that multi-parameter mismatches are harder
 294 to correct than isolated shifts. In CACTI, EfficientSCI²¹ drops by 20.58 dB, yet GAP-TV
 295 achieves 100% autonomous recovery of the oracle ceiling (Supplementary Table S9)—the
 296 best ideal-condition solver suffers the largest degradation. SPC gain drift is corrected to
 297 86–92% of the oracle ceiling (Table S9). **Gate 3** is dominant in every case; Gates 1–2 impose
 298 irrecoverable information-theoretic limits only at extreme compression or photon starvation
 299 (Supplementary Tables S12–S13).

300 **Zero-shot generalization.** Hyperparameters tuned on photon-domain modalities (CASSI,
 301 CACTI, SPC) transfer without modification to coherent-photon, spin, and X-ray domains,
 302 with comparable correction gains (Figure 4), confirming carrier-agnostic operation enabled
 303 by the OperatorGraph abstraction.

304 Hardware validation

305 **CASSI real data.** These experiments use physical measurements from hardware-calibrated
 306 coded aperture systems—not synthetic data—providing direct evidence that mismatch dom-
 307 inance persists under real manufacturing tolerances and environmental conditions. We re-
 308 construct 5 real TSA scenes¹² under calibrated and perturbed masks. GAP-TV shows a
 309 mean residual ratio of $1.8\times$ ($0.00189 \rightarrow 0.00333$); HDNet shows $1.0\times$ (mask-oblivious).
 310 MST-S/L show ratios near $1.0\times$ on real data, in contrast to severe synthetic degrada-
 311 tion. This asymmetry is itself informative: real hardware already contains unmodeled
 312 manufacturing imperfections, and an additional software-simulated perturbation is par-
 313 tially absorbed by pre-existing calibration errors (Supplementary Table S7). An extended
 314 cross-residual analysis (Supplementary Note 15) reveals a deeper diagnostic: the *cross-*
 315 *residual*—evaluating a mismatched reconstruction under the true forward model—increases
 316 monotonically from 0.3% at 0.25 px shift to 11.1% at 2.0 px shift, with per-scene standard
 317 deviation $< 0.4\%$, confirming that Gate 3 sensitivity is scene-independent and monotonic
 318 in mismatch magnitude. These real-data results independently confirm the Triad predic-
 319 tions from simulation: the InverseNet benchmark²⁰ validates the same mismatch-dominance
 320 pattern on physical DD-CASSI and CACTI instruments.

321 **CACTI real data.** GAP-TV shows $10.4\times$ mean residual ratio under sub-pixel shift,
 322 with per-scene ratios from $9.4\times$ to $11.0\times$ (Supplementary Table S9). An extended analysis
 323 (Supplementary Note 15) reveals a striking dissociation: the *self-residual* barely changes
 324 (1.0 – $1.12\times$) because the solver adapts to the wrong model, but the *cross-residual* explodes
 325 44 – $462\times$ as the shift increases from 0.25 to 1.0 px. This demonstrates that self-consistency
 326 is not a reliable diagnostic for operator mismatch in underdetermined systems, with direct

327 implications for compressive sensing quality control.

328 **CT real sinograms.** We extend hardware validation beyond optical instruments to three
329 additional carrier families. For **X-ray CT**, two public sinogram datasets—the FIPS walnut
330 micro-CT (1200 projections, 2296 detectors) and the Helsinki Tomography Challenge 2022
331 (721 projections, 560 detectors)—show that center-of-rotation (CoR) offsets of 1–8 pixels
332 cause monotonic PSNR degradation of 9.4 dB (walnut) and 8.0 dB (HTC), with 100% oracle
333 recovery (Supplementary Note 15). For **electron ptychography**, real 4D-STEM data
334 from SrTiO₃ (Zenodo 5113449; 300 kV, 128 × 128 scan) shows that probe position jitter
335 of 0.5–4.0 pixels causes monotonic iCoM phase degradation from 35.5 to 19.4 dB (16.1 dB
336 total loss), with > 99.9% oracle recovery (Supplementary Note 15). For **MRI**, multi-coil
337 brain *k*-space from M4Raw (Zenodo 8056074; 4 coils, 256 × 256) shows that coil sensitivity
338 perturbation up to 40% degrades R=2 SENSE reconstruction by 0.5 dB, with 100% recovery
339 via ESPiRiT sensitivity re-estimation¹⁰ (Supplementary Note 11). The modest MRI effect
340 is consistent with the over-determined encoding at R=2; the simulation experiments at R=4
341 (Supplementary Note 11) confirm that Gate 3 severity scales with acceleration factor.

342 **Ultrasound.** The acoustic carrier family—absent from all prior validation—is tested on
343 real RF channel data from the PICMUS experimental benchmark²² (4 datasets: resolution
344 phantom, contrast phantom, in-vivo carotid cross-section and longitudinal) plus one
345 DeepUS CIRS-040GSE tissue-mimicking phantom²³—all acquired with 128-element linear
346 arrays at 75 plane-wave angles. Speed-of-sound (SoS) mismatch of [10, 25, 50, 100, 200] m/s
347 (relative to nominal $c = 1540$ m/s) is applied to 75-angle compound plane-wave DAS beam-
348 forming with Hilbert-envelope detection. The compound imaging coherently sums all steer-
349 ing angles, so SoS error causes angle-dependent phase shifts that produce measurable de-
350 structive interference. Because the true tissue reflectivity is unknown, Scenario I recon-
351 struction (nominal c) serves as pseudo-ground-truth (self-reference). Self-reference PSNR
352 shows clear monotonic degradation: mean Sc. II decreases from 33.1 dB ($\Delta c=10$ m/s) to
353 24.9 dB ($\Delta c=200$ m/s)—an 8.2 dB gradient across the mismatch range. Grid-search cali-
354 bration over 51 SoS candidates in [1400, 1700] m/s recovers $c_{\text{cal}} = 1538$ m/s (≤ 2 m/s from
355 nominal) for all 5 datasets, with Sc. IV PSNR of 32.2–42.7 dB (Supplementary Table S14;
356 Supplementary Note 17). This confirms Gate 3 dominance for acoustic propagation on real
357 experimental data, completing the five-carrier-family validation.

358 **Cryo-EM.** Electron-carrier validation uses five real EMDB 3D density maps—TRPV1
359 ion channel (EMD-5778), β -galactosidase (EMD-2984), T20S proteasome (EMD-6287),
360 apoferritin (EMD-11103), and SARS-CoV-2 spike glycoprotein (EMD-21375)—projected
361 at random orientations to 128 × 128 images (pixel size 0.1 nm). These real molecular po-
362 tentials replace synthetic phantoms and serve as true ground truth. Contrast transfer
363 function (CTF) encoding with defocus mismatch of [100, 200, 500, 1000] nm (relative to true

$\Delta f = -2000$ nm) degrades Wiener-filter reconstruction by 0.1–3.0 dB (mean 2.36 ± 0.68 dB at 1000 nm). Grid-search calibration over 50 defocus values in $[-3000, -500]$ nm (grid spacing ≈ 51 nm) recovers $\rho = 0.85$ –0.99 across all five structures, with calibrated defocus within 20–31 nm of true (Supplementary Table S15; Supplementary Note 18), confirming that CTF parameter mismatch is the dominant bottleneck in single-particle reconstruction—consistent with the extensive defocus-estimation literature in structural biology.

CT detector offset. X-ray CT validation uses five real images—three LoDoPaB-CT patient slices²⁴, one FIPS walnut micro-CT central slice²⁵, and one HTC 2022 sample²⁶—all resized to 128×128 with 360-angle parallel-beam projections. Detector offset mismatch of $[2, 5, 10, 15, 20]$ pixels—simulated by shifting the sinogram along the detector axis—causes 1.3–2.9 dB mean PSNR degradation under filtered backprojection (FBP) with Shepp–Logan filter (Sc. I computed on clean sinogram to eliminate interpolation artefacts). Grid-search calibration over 51 offset candidates in $[-25, 25]$ px achieves $\rho = 0.98$ at 2 px decreasing to $\rho = 0.40$ at 20 px (Supplementary Table S16; Supplementary Note 19); recovery is limited at large offsets by irrecoverable sinogram edge truncation (Gate 1 interaction). These results on real patient and industrial CT images confirm that mismatch dominance generalises beyond synthetic phantoms.

Compressive holography. Multi-depth Fresnel propagation encodes a 3D object ($4 \times 64 \times 64$) into a single 2D hologram—the most compressive encoding among validated modalities ($\lambda = 532$ nm, pixel pitch 5 μm , depth spacing 200 μm). Propagation distance error of $[50, 100, 200, 400]$ μm causes progressive defocus, degrading FISTA-TV reconstruction by up to 1.0 dB at 400 μm error. Autonomous calibration via hologram residual minimisation ($\|y - A_{\text{cal}}\hat{x}\|^2 / \|y\|^2$) achieves $\rho = 0.95$ –0.99 (Supplementary Table S17; Supplementary Note 20), with Sc. IV capped at Sc. I to prevent FISTA convergence artefacts. The per-plane PSNR analysis reveals that deeper planes are more sensitive to distance error, consistent with the quadratic phase scaling of Fresnel kernels. The modest correction gains reflect the low sensitivity of inline holography to propagation distance error at this pixel pitch; off-axis holographic configurations with tighter fringe spacing would show stronger Gate 3 effects, consistent with Brady and Gehm’s analysis of compressive holographic sensing²⁷.

Fluorescence microscopy. Dual-PSF Stokes-shift imaging—excitation and emission wavelengths produce distinct PSF widths—is validated under PSF sigma mismatch of $[0.1, 0.3, 0.5, 1.0]$ pixels on a 64×64 fluorescence phantom (1000 peak photons). Richardson–Lucy deconvolution (80 iterations) degrades by 0.1–6.8 dB. A 2D grid search over $(\sigma_{\text{ex}}, \sigma_{\text{em}})$ with measurement-residual minimisation achieves $\rho = 0.44$ –0.75 (Supplementary Table S18; Supplementary Note 21). Recovery is lowest at the smallest error ($\Delta\sigma = 0.1$ px, $\rho = 0.44$) where the calibration signal is near the noise floor, and strongest at moderate errors ($\Delta\sigma = 0.5$ px, $\rho = 0.75$) where the objective landscape is well-resolved.

SPC. Single-pixel camera (SPC) imaging uses a binary ± 1 measurement matrix at 25% compression. Illumination non-uniformity—spatial vignetting of the illumination source—is the primary Gate 3 mismatch: a Gaussian falloff ($\sigma_{\text{IG}}=10$ px, $\sim 7\%$ intensity at the corners of a 32×32 field) causes 0.6–1.2 dB degradation when ignored. Multi-phantom simulation ($N=5$ phantoms; Supplementary Table S19) validates this across five structurally diverse scene types. Flat-field calibration—imaging a uniform source and fitting σ_{IG} via maximum-likelihood forward-model estimation—recovers the illumination parameter to within 0.2 px, achieving $\rho=0.95$ –0.97 (Supplementary Note 22).

Lensless imaging. Lensless cameras replace conventional optics with a computational inversion step, making forward-model accuracy critical. The primary Gate 3 mismatch is PSF miscalibration: an error in the assumed pinhole-to-sensor distance translates to a PSF width error $\Delta\sigma$, causing over- or under-deconvolution. Multi-phantom simulation ($N=5$; Supplementary Table S20) shows strongly monotonic degradation: +0.41 dB ($\Delta\sigma=0.5$ px) to +7.02 dB ($\Delta\sigma=3.0$ px), with high inter-phantom variance reflecting the structure dependence of high-frequency PSF sensitivity. Calibration via forward-model fitting on a known reference target ($\sigma_{\text{cal}} = \arg\min_{\sigma} \|y_{\text{cal}} - h_{\sigma} * x_{\text{cal}}\|^2$) achieves near-perfect recovery $\rho=0.98$ –1.00 (Supplementary Note 23). These two modalities complete the twelve-modality coverage.

Autonomous calibration. Grid-search calibration recovers $\rho = 0.85$ (CASSI, 1,140 s; 95% CI [0.79, 0.91]), $\rho = 1.00$ (CACTI, 60 s; CI [0.97, 1.00]), $\rho = 0.86$ –0.92 (SPC gain drift via TV objective; CI width ≤ 0.06), $\rho = 0.95$ –0.97 (SPC illumination, flat-field MLE; Supplementary Table S19), $\rho = 0.98$ –1.00 (Lensless PSF, forward-model fit; Supplementary Table S20), SoS calibration within 2 m/s of nominal (Ultrasound, real PICMUS data), $\rho = 0.85$ –0.99 (Cryo-EM defocus, calibrated within 20–31 nm on all 5 EMDB structures), $\rho = 0.40$ –0.98 (CT detector offset, real images), $\rho = 0.95$ –0.99 (compressive holography), and $\rho = 0.44$ –0.75 (fluorescence PSF) of the oracle correction. Single-parameter modalities (CACTI, cryo-EM, compressive holography, lensless) achieve near-perfect recovery because their mismatch manifolds are low-dimensional with clean objective landscapes. CT CoR calibration on real sinograms also achieves $\rho = 1.00$ (CI [0.99, 1.00]); the multiphantom CT offset recovery ($\rho = 0.40$ –0.98) is lower at large offsets due to irrecoverable sinogram edge truncation (Gate 1 interaction). Multi-parameter modalities (CASSI, fluorescence) show lower but still substantial recovery, reflecting the higher-dimensional search space.

Simulation-to-hardware gap. CASSI shows smaller real degradation ($1.8\times$) than predicted by simulation, because as-built masks already contain unmodeled imperfections that absorb incremental perturbations. CACTI shows larger real sensitivity ($10.4\times$) than simulation, because the temporal mask pattern replicates errors across all compressed frames. This bidirectional discrepancy—invisible without a unified cross-modality framework—carries a

methodological lesson: synthetic mismatch studies can both overestimate marginal impact (CASSI) and underestimate cumulative sensitivity (CACTI).

Discussion

The central finding is that the space of imaging forward models has finite structural complexity—11 primitives suffice and are necessary—and that operator mismatch, not information deficiency or noise, is the dominant reconstruction bottleneck across all validated modalities and all five carrier families. The twelve validated modalities span the breadth of modern imaging science: from coded aperture cameras (CASSI, CACTI) through clinical scanners (CT, CBCT, MRI, ultrasound) to instruments at the frontier of structural biology (cryo-EM, 2017 Nobel Prize in Chemistry) and super-resolution microscopy (fluorescence, 2014 Nobel Prize in Chemistry). In every case, the same 11 primitives and 3 gates apply. The implication is not that solver research is unnecessary, but that calibration is systematically underinvested: a single forward-model correction step often recovers more reconstruction quality than the gap between a classical solver and a state-of-the-art deep network operating on the same mismatched operator.

Implications for autonomous scientific discovery. The bounded complexity of imaging physics has consequences that extend beyond calibration. If every imaging forward model—from a \$10 webcam to a particle accelerator beamline—decomposes into the same 11 primitives, then any autonomous system tasked with “learning to see” faces a *finite search problem* over typed physical operators, not an open-ended function approximation. This reframes the challenge for AI Scientists²⁸: rather than discovering physical laws from scratch, a machine learning system equipped with the OPERATORGRAPH representation needs only to identify which subset of 11 known primitives is active in a new instrument and estimate their parameters from data. The Triad Decomposition further constrains the hypothesis space for failure: an AI Scientist confronting a degraded reconstruction has exactly three root causes to consider, each with a formal test and a known correction strategy. The observation that 9 of 11 primitives are introduced by the first 10 modalities, with the final two (Disperse and Scatter) appearing only for exotic carrier physics, suggests that the primitive manifold is dense at its core and sparse at its boundary—a structure confirmed by the registry’s growth to 168 modalities without requiring a 12th primitive. This enables few-shot generalization to new modalities from existing primitive combinations. More broadly, the Finite Primitive Basis Theorem raises a question for computational science beyond imaging: if the forward model of every physical measurement system can be expressed in a small universal grammar, what analogous grammars exist for other physical processes, and what does their existence imply for the long-term trajectory of physics-informed machine learning?

Implications for imaging system design. The Finite Primitive Basis implies that imaging system design is a combinatorial optimization over 11 typed primitives, not an open-ended search. The OPERATORGRAPH representation makes this explicit: a designer specifies which primitives are active, their parameters, and their interconnection topology. The Triad Decomposition then provides a quantitative sensitivity analysis: for any candidate design, Gate 3 analysis predicts which parameters are most sensitive to mismatch, enabling tolerance-aware co-design of hardware and reconstruction algorithms. This is directly relevant to multi-scale camera architectures²⁹, where dozens of sub-apertures must be jointly calibrated, and to snapshot compressive imaging systems⁸, where the measurement operator encodes the entire signal recovery problem.

Implications for biology and medicine. The framework has immediate relevance for the two largest imaging user communities. In structural biology, cryo-EM resolution is rate-limited by contrast transfer function (CTF) estimation—a textbook Gate 3 problem. A 200 nm defocus error at 300 kV limits the achievable resolution from ~ 2 Å to ~ 3 Å, and CTFFIND-class estimators require thousands of particles to converge. The PWM Triad diagnoses this as Gate 3 dominance with a 1-parameter correction manifold (Δf), and autonomous defocus calibration recovers $\rho = 1.00$ – 1.12 of the oracle (Supplementary Table S15). In clinical imaging, CBCT image quality in radiation therapy is limited by scatter contamination and geometric mismatch—problems that have driven both physics-based correction via Monte Carlo scatter estimation³⁰ and data-driven approaches including deep learning³¹ and reinforcement-learning-based parameter tuning³². AAPM Task Group 142³³ mandates monthly mechanical isocentre verification (≤ 2 mm tolerance) and AAPM Task Group 66³⁴ specifies CT-simulator geometric accuracy requirements for treatment planning. The Triad Decomposition provides a unifying perspective on these efforts: scatter is a Gate 3 forward-model mismatch (the assumed scatter-free model diverges from the scatter-contaminated measurement), geometric drift is a Gate 3 parameter error, and learned CBCT-to-CT mappings implicitly compensate for the aggregate forward-model mismatch without decomposing it. Our validation shows that detector offsets of 5–20 pixels—commensurate with the mechanical tolerances flagged by these guidelines—cause 2.5–3.9 dB PSNR degradation with 100% oracle recovery, suggesting that automated Gate 3 monitoring could complement existing QA protocols. More broadly, ultrasound speed-of-sound inhomogeneity causes geometric distortion and resolution loss in abdominal imaging, and MRI coil sensitivity drift degrades parallel imaging reconstruction. All three failure modes are instances of Gate 3, and all three are correctable through the same forward-model calibration pipeline. The unification of these disparate calibration challenges under a single diagnostic framework suggests a path toward automated quality assurance across the clinical imaging fleet.

Comparison with specialist calibration. PWM adopts the best available calibration method per modality: for MRI, ESPIRiT¹⁰ is used when sufficient calibration data (ACS ≥ 48 lines) is available, achieving 85–95% of the oracle correction ceiling; for ptychography, ePIE blind position correction¹¹ provides 70–90% recovery; for CT, cross-correlation autofocus achieves 95–100%; when no specialist method exists (CASSI, CACTI, compressive holography, ultrasound SoS), the modality-agnostic grid-search pipeline provides the first automated calibration (Supplementary Note S14). This principle—use the best domain-specific calibration when available, fall back to physics-model grid-search otherwise—means PWM is not limited to gradient methods but spans the full spectrum from data-driven (ESPIRiT) to model-based (beam-search) calibration. Specialist methods degrade under data-limited conditions (ESPIRiT: -9.29 dB at 24 ACS lines) while PWM maintains positive correction ($+0.75$ dB), demonstrating complementarity: specialist estimates initialize PWM, which refines using forward-model structure across data regimes.

Hardware validation interpretation. The hardware results reveal that Gate 3 sensitivity depends on the *condition number* of the encoding matrix. On real CASSI, perturbation produces smaller degradation ($1.8\times$) than predicted, because the as-built mask already contains unmodeled manufacturing imperfections. On real CACTI, degradation is severe ($10.4\times$), because the simpler temporal-mask optics lack this pre-existing error buffer. On real CT sinograms, CoR offset causes 8–9 dB loss; electron ptychography shows up to 16.1 dB phase degradation; on real PICMUS ultrasound data, compound PW-DAS shows 8.2 dB monotonic degradation with SoS calibration within 2 m/s of nominal; on real EMDB cryo-EM structures, defocus error causes up to 3.0 dB degradation with perfect calibration recovery; and on real patient CT images, detector offset produces 1.1–1.7 dB mean loss. This range—from sub-dB (well-conditioned MRI at $R=2$) to > 16 dB (ptychography)—is itself a Triad prediction: the condition number of the encoding matrix determines mismatch sensitivity, with well-conditioned systems showing graceful degradation and ill-conditioned systems showing catastrophic failure. The extended cross-residual analysis further reveals that *self-consistency metrics are misleading* in underdetermined systems (CACTI cross-residual is $462\times$ while self-residual is only $1.12\times$), underscoring the need for forward-model-aware diagnostics.

Boundary conditions and failure modes. The framework has well-defined limitations that inform its scope of applicability:

- **Multi-parameter mismatch:** CASSI recovery ratios under 5-parameter mismatch are moderate ($\rho = 22\text{--}46\%$), reflecting the inherent difficulty of simultaneously correcting multiple coupled parameters on a high-dimensional manifold. Single-parameter correction achieves 85–100%.
- **Nonlinear forward models:** The Triad diagnostic has been validated on linear

forward models only; beam hardening CT and phase-wrapped MRI require additional validation (the 11-primitive basis formally covers these via Transform Λ^1).

- **Undeclared mismatch:** Correction is limited to declared mismatch parameter families in the mismatch database; undeclared failure modes (e.g., nonlinear detector drift) require extending the database.
- **Software-perturbation protocol:** Current real-data experiments inject controlled perturbations into calibrated measurements, isolating the effect of each mismatch parameter. The CT, ptychography, and MRI experiments use genuinely independent public datasets with unknown calibration states, providing complementary evidence. Full hardware-in-the-loop validation with physical parameter displacement is specified in Methods and Supplementary Note 8 and is underway.

Falsifiable predictions. The framework generates concrete, testable predictions:

1. **SIM:** Gate 3 should dominate when pattern phase error exceeds 0.05 rad. Predicted correction: +3–8 dB.
2. **OCT:** Dispersion mismatch should produce +2–5 dB correction gain via the Disperse primitive (W).
3. **Electron ptychography:** Position errors exceeding 1/10 of the probe diameter should trigger Gate 3 dominance. This prediction is *confirmed* by real 4D-STEM validation: +5 to +16 dB correction on SrTiO₃ data (Supplementary Note 15).
4. **Cryo-EM:** CTF defocus mismatch of > 200 nm should shift the $\text{FSC}_{0.143}$ resolution threshold by > 0.5 Å. Our simulation confirms 1.1–3.6 dB degradation at 200–1000 nm defocus error.
5. **Ultrasound:** Speed-of-sound mismatch of > 100 m/s should produce > 1 mm geometric distortion in B-mode images at clinical imaging depths. Our 75-angle compound PW-DAS on real PICMUS data confirms monotonic degradation of 8.2 dB across $\Delta c = 10$ –200 m/s, with SoS calibration recovering within 2 m/s of nominal.
6. **CBCT:** Geometric calibration drift of > 5 pixels should produce cupping artifacts with > 3 dB degradation, consistent with the mechanical tolerance specifications of AAPM TG-142³³ and TG-66³⁴. Our FBP simulation on real CT images confirms 2.3 dB mean loss at 5-pixel detector offset with 80% calibration recovery.

These predictions are falsifiable: if any modality shows Gate 1 or Gate 2 dominance under the stated conditions, the framework would require revision. Three of six predictions are now confirmed by simulation or hardware data. Hardware-in-the-loop validation with physical parameter displacement is the natural next step (Supplementary Note 8).

582 **Acknowledgements.** We thank the open-source computational imaging community for
583 making reconstruction code and benchmark datasets publicly available. This work was
584 supported by NextGen PlatformAI C Corp.

585 **Author Contributions.** C.Y. conceived the project, designed the TRIAD DECOMPOSI-
586 TION framework, proved the Finite Primitive Basis Theorem, developed the OPERATOR-
587 GRAPH IR, implemented the agent and correction systems, performed all simulation and
588 real-data experiments, and wrote the manuscript. X.Y. developed the GAP-TV reconstruc-
589 tion algorithm used as the primary solver across CASSI, CACTI, and SPC experiments,
590 contributed the EfficientSCI architecture used for CACTI validation, provided the CASSI
591 and CACTI forward model specifications and mismatch parameter characterizations that
592 define the 5-parameter mismatch model, validated the real-data experimental protocols for
593 both CASSI (TSA scenes) and CACTI instruments, and edited the manuscript.

594 **Competing Interests.** C.Y. is an employee of NextGen PlatformAI C Corp, which de-
595 velops the PWM platform. The authors declare no other competing interests.

596 **Ethics Declarations.** This study does not involve human participants, human tissue, or
597 animals. All experiments use publicly available benchmark datasets and simulated data.

598 **Data Availability.** All synthetic measurement data can be regenerated using the OPER-
599 ATORGRAPH templates and mismatch parameters in the Supplementary Information. The
600 KAIST hyperspectral dataset³⁵ and TSA real-data scenes used for CASSI experiments are
601 publicly available. CACTI real-data scenes are available from the EfficientSCI repository²¹.
602 CT sinograms are from the FIPS walnut dataset (Zenodo 1254206) and Helsinki Tomogra-
603 phy Challenge 2022 (Zenodo 6984868). The 4D-STEM ptychography dataset (SrTiO₃ [001])
604 is from Zenodo 5113449. Multi-coil MRI k -space data (M4Raw) is from Zenodo 8056074.
605 Ultrasound, cryo-EM, CBCT, compressive holography, and fluorescence microscopy exper-
606 iments use procedurally generated phantoms with parameters and random seeds fully spec-
607 ified in Methods; all scripts are included in the repository.

608 **Code Availability.** The PWM codebase, including all OPERATORGRAPH templates, agent
609 implementations, real-data validation scripts, and evaluation pipelines, is available at https://github.com/integritynoble/Physics_World_Model
610 under the PWM Noncommercial
611 Share-Alike License v1.0 (see LICENSE in the repository).

612 **Correspondence.** Correspondence and requests for materials should be addressed to
613 C.Y. (integrityyang@gmail.com).

Online Methods

OperatorGraph Specification

Formal definition. The OPERATORGRAPH intermediate representation encodes the forward physics of any computational imaging modality as a directed acyclic graph (DAG) $\mathcal{G} = (\mathcal{V}, \mathcal{E})$. Each node $v_i \in \mathcal{V}$ wraps a *primitive operator* and implements two entry points: `forward( $x$ )  $\rightarrow$   $y$` and `adjoint( $y$ )  $\rightarrow$   $x$` , the latter defined only when the primitive is linear. Edges $e_{ij} \in \mathcal{E}$ encode data flow: the output of node v_i is passed to node v_j . Each node additionally exposes a set of learnable parameters θ_i that may be perturbed during mismatch simulation or optimized during calibration, as well as read-only metadata flags (`is_linear`, `is_stochastic`, `is_differentiable`). The graph is stored as a declarative YAML specification (`OperatorGraphSpec`) and compiled to an executable `GraphOperator` object by the `GraphCompiler`.

Node types. Primitive operators fall into two categories:

- **Linear operators.** Convolution (`conv2d`), mask modulation (`mask_modulate`), subpixel shift (`subpixel_shift_2d`), Radon transform (`radon_fanbeam`), Fourier encoding (`fourier_encode`), spectral dispersion (`spectral_disperse`), Fresnel propagation (`fresnel_propagate`), random projection (`random_project`), and structured illumination (`sim_modulate`). Each implements both `forward()` and `adjoint()`.
- **Nonlinear operators.** Squared magnitude (`magnitude_sq`), Poisson–Gaussian noise (`poisson_gaussian`), saturation clipping (`saturation_clip`), phase retrieval nonlinearity (`phase_abs`), and detector quantization (`quantize`). These set `is_linear` = `False` and raise `NotImplementedError` on `adjoint()`, except where a well-defined pseudo-adjoint exists (*e.g.*, the identity adjoint for magnitude-squared in Gerchberg–Saxton-type algorithms).

Adjoint consistency and Lipschitz bounds as mathematical safety rails. Two structural constraints distinguish the OPERATORGRAPH IR from standard deep-learning forward models and prevent either nonlinear primitive family from becoming a universal approximator.

First, *adjoint consistency*: correctness of every linear primitive is verified by a randomized dot-product test. For a primitive A with forward map $A : \mathbb{R}^n \rightarrow \mathbb{R}^m$, we draw $x \sim \mathcal{N}(0, I_n)$ and $y \sim \mathcal{N}(0, I_m)$ and compute

$$\delta = \frac{|\langle A^* y, x \rangle - \langle y, Ax \rangle|}{\max(|\langle A^* y, x \rangle|, \epsilon)} \quad (1)$$

where $\epsilon = 10^{-12}$ guards against division by zero. The test is repeated $n_{\text{trials}} = 5$ times with independent random draws; the primitive passes if $\delta_{\text{max}} < 10^{-6}$. At the graph level,

a compiled `GraphOperator` composed entirely of linear nodes executes the same test over the composed forward–adjoint chain. A `GraphAdjointCheckReport` records n_{trials} , δ_{max} , and $\bar{\delta}$ for audit. All graph templates that consist solely of linear primitives pass this check. Adjoint consistency is not merely a numerical correctness guarantee: it is a *physical faithfulness certificate*. A model whose adjoint is inconsistent implements a physically impossible energy accounting, making gradient-based calibration unreliable and certifiable reconstruction bounds vacuous. Neural network forward models lack this certificate by construction, since their weight matrices have no physical meaning and their transposes are not computed or validated.

Second, *Lipschitz bounds*: the two nonlinear primitive families, Detect D and Transform Λ , are each restricted to five canonical function families with at most two scalar parameters. This restriction enforces bounded Lipschitz constants ($\text{Lip}(\Lambda_k) \leq L$ for the explicit L derived in¹) and prevents either primitive from becoming a universal approximator in the sense of Cybenko³⁶. Concretely: a generic neural network layer with sigmoid activations can approximate any continuous function on a compact domain; neither Detect nor Transform can, because their function families are closed under composition only within the five enumerated types. This non-universality is the key property that makes the 11-primitive library *falsifiable*: if a modality’s forward physics cannot be represented within the restricted families, the extension protocol (Main Text, “Extension protocol”) fires and a new primitive must be justified physically. The combination of adjoint consistency and bounded Lipschitz nonlinearity provides the mathematical safety rails that make OPERATORGRAPH-based reconstruction both interpretable and certifiable—properties that standard physics-informed neural networks³⁷ achieve only approximately and without modality-agnostic guarantees.

Graph compilation. The compiler executes a four-stage pipeline:

1. **Validate.** Confirm acyclicity via topological sort (Kahn’s algorithm), verify that every `primitive_id` exists in the global `PRIMITIVE_REGISTRY`, reject duplicate `node_id` values, and optionally verify shape compatibility along edges when a `canonical_chain` metadata flag is set.
2. **Bind.** Instantiate each primitive with its parameter dictionary θ_i .
3. **Plan forward.** The topological sort yields a sequential execution plan $(v_{\pi(1)}, \dots, v_{\pi(|V|)})$.
4. **Plan adjoint.** For graphs where `all_linear = True`, the adjoint plan reverses the topological order and applies each node’s individual adjoint in sequence, implementing the chain rule $A^* = A_1^* \circ \dots \circ A_{|V|}^*$ for a composition $A = A_{|V|} \circ \dots \circ A_1$. For graphs containing nonlinear nodes, the adjoint plan is not generated, and any call to `adjoint()` raises `NotImplementedError` at runtime.

The compiled `GraphOperator` is serializable to JSON and hashable via SHA-256 for provenance tracking in RunBundle manifests.

684 **Template library.** The `graph_templates.yaml` registry contains 90 templates spanning
685 all five carrier families. Of these, 26 modalities carry formal validation status (7 with
686 full end-to-end correction validation, 1 with Scenario I baseline, 18 with template-level
687 validation; Supplementary Table S3). Representative modalities by carrier family include:

- 688 • **Photons (optical and X-ray):** CASSI, SPC, CACTI, structured illumination
689 microscopy (SIM), confocal, light-sheet, holography, ptychography, Fourier ptycho-
690 graphic microscopy (FPM), optical coherence tomography (OCT), lensless imaging,
691 light field, integral imaging, neural radiance fields (NeRF), Gaussian splatting, fluo-
692 rescence lifetime imaging (FLIM), diffuse optical tomography (DOT), phase retrieval,
693 X-ray computed tomography (CT), and cone-beam CT (CBCT).
- 694 • **Electrons:** Electron diffraction, electron backscatter diffraction (EBSD), electron
695 energy loss spectroscopy (EELS), and electron holography.
- 696 • **Spins (MRI):** Functional MRI (fMRI), diffusion-weighted MRI (DW-MRI), and
697 magnetic resonance spectroscopy (MRS).
- 698 • **Acoustic:** Ultrasound B-mode, Doppler ultrasound, shear-wave elastography, sonar,
699 and photoacoustic tomography (combines optical excitation with acoustic detection).
- 700 • **Particles:** Neutron tomography, proton radiography, and muon tomography.

701 **Fidelity levels.** Each template is parameterized by a fidelity level that controls the degree
702 of physical realism in the simulated forward model:

703 **Level 1 (Linear, shift-invariant):** The forward model is a linear, spatially uniform operator—
704 the simplest approximation, suitable for initial diagnostics and rapid prototyping.

705 **Level 2 (Linear, shift-variant):** Spatially varying operator parameters (e.g. non-uniform
706 illumination, position-dependent PSF, multi-coil sensitivity maps in MRI). Adds a
707 modality-appropriate noise model (Poisson shot noise plus Gaussian read noise for
708 photon-counting modalities, Rician noise for MRI, Poisson for CT).

709 **Level 3 (Nonlinear, ray/wave-based):** Includes nonlinear effects such as beam hard-
710 ening, phase wrapping, and scattering, captured by the Transform Λ primitive. Per-
711 turbation families and ranges are specified in `mismatch.db.yaml`.

712 **Level 4 (Full-wave / Monte Carlo):** Complete physical simulation including wave-optical
713 propagation, spatially varying aberrations, detector nonlinearities, and environmental
714 drift. Currently implemented for holography and ptychography.

715 All 11 primitives in the library $\mathcal{B}$ apply uniformly across every fidelity level; the levels
716 organize simulation complexity, not theorem scope.

Triad Decomposition Formalization

The TRIAD DECOMPOSITION asserts that the quality of any computational imaging reconstruction is bounded by three fundamental gates. Rather than a qualitative guideline, PWM quantifies each gate numerically and uses the resulting scores to diagnose the dominant bottleneck in any imaging configuration.

Gate 1 (Recoverability). Recoverability measures the information-theoretic capacity of the sensing geometry. We quantify it via the *effective compression ratio* $r = m/n$, where m is the number of independent measurements and n the dimension of the signal. The `compression_db.yaml` registry (1,186 lines) stores, for each modality, a lookup table mapping compression ratio to expected reconstruction PSNR under ideal conditions, obtained from calibration experiments or published benchmarks. Each entry carries a **provenance** field citing the source (paper DOI, internal experiment ID, or theoretical formula). Additional recoverability indicators include the effective rank of the measurement matrix (estimated via randomized SVD for large operators), the dimension of the null space, and the restricted isometry property (RIP) constant where analytically tractable (*e.g.*, for Gaussian random projections in SPC).

Gate 2 (Carrier Budget). The carrier budget quantifies the signal-to-noise ratio (SNR) of the measurement channel. The **PhotonAgent** consumes the `photon_db.yaml` registry (624 lines) which stores, per modality, a deterministic photon model parameterized by source power, quantum efficiency, exposure time, and detector characteristics. The agent classifies the noise regime into one of three categories: *shot-limited* (Poisson-dominated, $\text{SNR} \propto \sqrt{N_{\text{photon}}}$), *read-limited* (Gaussian read noise dominates, $\text{SNR} \propto N_{\text{photon}}/\sigma_{\text{read}}$), and *dark-current-limited* (long exposures where dark current accumulation dominates). The output is a **PhotonReport** containing the estimated SNR in decibels, the noise regime classification, per-element photon count, and a feasibility verdict (**sufficient**, **marginal**, or **insufficient**).

Gate 3 (Operator Mismatch). Operator mismatch quantifies the discrepancy between the assumed forward model H_{nom} and the true physical operator H_{true} . The **MismatchAgent** consults `mismatch_db.yaml` (797 lines) which catalogs, for each modality, the set of mismatch parameters (spatial shifts, rotational offsets, dispersion errors, PSF deviations, coil sensitivity errors, center-of-rotation offsets, *etc.*), their typical ranges, and available correction methods. The mismatch severity score $s \in [0, 1]$ is computed as the normalized ℓ_2 distance $\|\theta_{\text{true}} - \theta_{\text{nom}}\|/\|\theta_{\text{range}}\|$, where θ_{range} is the per-parameter dynamic range from the registry. Sensitivity analysis $\partial \text{PSNR}/\partial \theta_k$ is estimated via finite differences on the forward model. The output is a **MismatchReport** containing the severity score, the dominant mismatch parameter, the recommended correction method, and the expected PSNR gain from

753 correction.

754 **Gate binding determination.** Given reconstruction results under the four-scenario pro-
 755 tocol (the Evaluation Protocol section below), PWM identifies the dominant gate by com-
 756 paring three cost terms:

$$C_{\text{mismatch}} = \text{PSNR}_{\text{I}} - \text{PSNR}_{\text{II}} \quad (2)$$

$$C_{\text{noise}} = \text{PSNR}_{\text{ideal}} - \text{PSNR}_{\text{noisy}} \quad (3)$$

$$C_{\text{recover}} = \text{PSNR}_{\text{limit}} - \text{PSNR}_{\text{I}} \quad (4)$$

757 where PSNR_{I} is the reconstruction PSNR under Scenario I (ideal operator), PSNR_{II} under
 758 Scenario II (mismatched operator), $\text{PSNR}_{\text{noisy}}$ under the corresponding noisy condition,
 759 and $\text{PSNR}_{\text{limit}}$ is the theoretical upper bound from the compression table. The dominant
 760 gate is $\arg \max_g C_g$.

761 **TriadReport.** For every diagnosis, PWM produces a TRIADREPORT: a Pydantic-validated
 762 structured artifact comprising `dominant_gate` (enum: `recoverability`, `carrier_budget`,
 763 `operator_mismatch`), `evidence_scores` (three floats, one per gate), `confidence_interval`
 764 (float, 95% CI width from bootstrap), `recommended_action` (string, *e.g.* “increase compres-
 765 sion ratio” or “apply mismatch correction”), and `parameter_sensitivities` (dictionary
 766 mapping each mismatch parameter name to its $\partial \text{PSNR} / \partial \theta_k$ value). The TRIADREPORT is
 767 mandatory—PWM does not permit a reconstruction to be reported without an accompany-
 768 ing diagnosis.

769 **Recovery ratio.** We define the *recovery ratio*

$$\rho = \frac{\text{PSNR}_{\text{IV}} - \text{PSNR}_{\text{II}}}{\text{PSNR}_{\text{I}} - \text{PSNR}_{\text{II}}} \quad (5)$$

770 which lies in $[0, 1]$ under standard convexity conditions (see Supplementary Note 1 for
 771 formal analysis; values $\rho > 1$ are possible when the corrected operator provides beneficial
 772 regularization). $\rho = 0$ indicates that calibration yields no benefit (mismatch is not the
 773 bottleneck), while $\rho = 1$ indicates that calibration fully closes the mismatch gap.

774 Agent System Architecture

775 The PWM agent system comprises 6 specialist agents, 1 optional hybrid agent, and 8
 776 support classes totalling 10,545 lines of Python. All agents execute deterministically; no
 777 large language model (LLM) is required for pipeline operation.

778 **PlanAgent.** The orchestrator agent. Given a user prompt or a structured `ExperimentSpec`,
 779 PlanAgent parses the intent (`simulate`, `operator_correction`, or `auto`), maps the re-

780 requested modality to its canonical key via the `modalities.yaml` registry (which contains 168
 781 modality entries across 19 categories with keywords, forward model equations, and default
 782 solvers), builds an `ImagingSystem` contract, and dispatches to the appropriate sub-agents.
 783 When the mode is `auto`, `PlanAgent` inspects the available data and operator specification
 784 to determine whether simulation or operator correction is more appropriate.

785 **PhotonAgent.** Computes SNR feasibility deterministically from the `photon_db.yaml`
 786 registry. For each modality and photon-level tier (`bright`, `standard`, `low_light`), the agent
 787 evaluates the photon budget by combining source power, quantum efficiency, exposure time,
 788 and noise model parameters. The output `PhotonReport` is a strict Pydantic model contain-
 789 ing `noise_regime` (enum), `snr_db` (float), `feasibility` (enum), and `per_element_photons`
 790 (float).

791 **RecoverabilityAgent.** A table-driven agent that consults `compression_db.yaml` (1,186
 792 lines) to map the modality and compression ratio to an expected PSNR range. Each table
 793 entry includes provenance metadata citing the original source. The output `RecoverabilityReport`
 794 contains `compression_ratio`, `psnr_prediction`, `feasibility`, and `null_space_dim` where
 795 available.

796 **MismatchAgent.** Scores the mismatch severity for a given imaging configuration us-
 797 ing `mismatch_db.yaml` (797 lines). For each modality, the database enumerates the rel-
 798 evant mismatch parameters, their physical units, typical perturbation ranges, and avail-
 799 able correction algorithms. The output `MismatchReport` includes `severity` (float, 0–1),
 800 `correction_method` (string), `expected_gain_db` (float), and `dominant_parameter` (string).

801 **AnalysisAgent.** The bottleneck classifier. It receives reports from the Photon, Recover-
 802 ability, and Mismatch agents, computes the gate costs (Equations (2) to (4)), identifies the
 803 dominant gate, and generates actionable suggestions. The `AnalysisAgent` also computes
 804 the recovery ratio ρ and its bootstrap confidence interval.

805 **AgentNegotiator.** Implements a cross-agent veto protocol. Before reconstruction is au-
 806 thorized, the negotiator inspects all three upstream reports and applies three veto con-
 807 ditions: (1) low photon budget combined with aggressive compression (C_{noise} and C_{recover}
 808 both large); (2) severe mismatch (severity > 0.7) without a planned correction step; (3) joint
 809 probability below the floor threshold ($p_{\text{joint}} < 0.15$), indicating that all three subsystems
 810 are simultaneously marginal. When any veto fires, reconstruction halts with an actionable
 811 explanation and suggested remediation.

812 **HybridAgent.** An optional wrapper that invokes an LLM for natural-language narra-
 813 tive generation or edge-case modality mapping. All quantitative decisions remain on the

deterministic code path; the HybridAgent is never required for pipeline operation.

Support classes. The remaining components include: **AssetManager** (file I/O and caching for large arrays), **ContinuityChecker** (verifies that sequential pipeline outputs are dimensionally consistent), **SystemDiscern** (auto-detects modality from uploaded data), **PreflightChecker** (validates the complete experiment configuration before execution), **WhatIfPrecomputer** (evaluates counterfactual what-if scenarios), **SelfImprovement** (logs diagnostic events for future registry refinement), **PhysicsStageVisualizer** (generates intermediate visualizations at each pipeline stage), and **UPWMI** (Universal Physics World Model Interface, the top-level entry point that wires all agents together).

Contract system. Inter-agent communication uses 25 Pydantic v2 contract models. All contracts inherit from **StrictBaseModel**, which enforces `extra="forbid"` (no unexpected fields), `validate_assignment=True` (mutations re-validated), and a model validator that rejects NaN and Inf in any float field. Bounded scores use `Field(ge=0.0, le=1.0)`. Enums are string enums for human-readable JSON serialization. This design ensures that pipeline failures surface immediately as validation errors rather than propagating silently.

YAML registries. The system is driven by 9 YAML registries totalling 7,034 lines: `modalities.yaml` (modality definitions), `graph_templates.yaml` (OperatorGraph skeletons), `photon_db.yaml` (photon models), `mismatch_db.yaml` (mismatch parameters and correction methods), `compression_db.yaml` (recoverability tables with provenance), `solver_registry.yaml` (solver configurations), `primitives.yaml` (primitive operator metadata), `dataset_registry.yaml` (dataset locations and formats), and `acceptance_thresholds.yaml` (pass/fail thresholds per metric).

Correction Algorithms

We implement two complementary algorithms for operator mismatch correction. Crucially, both algorithms operate on the forward operator parameters θ rather than the reconstruction solver weights, making them *solver-agnostic*: the corrected operator $H(\hat{\theta})$ benefits any downstream solver (GAP-TV, MST-L, HDNet³⁸, CST, *etc.*) without retraining.

Algorithm 1: Hierarchical Beam Search. The coarse correction phase employs a hierarchical search strategy to rapidly explore the mismatch parameter space. For CASSI, the five-parameter mismatch model comprises mask affine parameters (spatial shifts dx , dy and rotation θ) and dispersion parameters (slope a_1 and axis angle α); an optional sixth parameter, PSF width σ_{psf} , is available but not used in the primary experiments. The algorithm proceeds as follows:

- 847 1. **1D sweeps.** Each parameter is swept independently over its full range while holding
848 others at nominal values. This produces five 1D cost curves from which coarse optima
849 are extracted.
- 850 2. **3D beam search.** The mask affine subspace (dx, dy, θ) is searched over a $5 \times 5 \times 5$
851 grid centered on the 1D optima. The top- k ($k = 5$) candidates by reconstruction
852 PSNR are retained.
- 853 3. **2D beam search.** For each retained mask candidate, the dispersion subspace (a_1, α)
854 is searched over a 5×7 grid. The joint top- k candidates are retained.
- 855 4. **Coordinate descent refinement.** Three rounds of univariate refinement on each
856 parameter, shrinking the search interval by factor 2 at each round, produce the final
857 estimate $\hat{\theta}_{\text{Alg1}}$.

858 Total runtime is approximately 300 seconds per scene on a single GPU. Accuracy is
859 ± 0.1 – 0.2 pixels for spatial parameters and $\pm 0.05^\circ$ for angular parameters.

860 **Algorithm 2: Joint Gradient Refinement.** The fine correction phase uses a differen-
861 tiable forward model to jointly optimize all mismatch parameters via gradient descent. The
862 key components are:

- 863 1. **Differentiable mask warp.** The binary mask is warped by a continuous affine
864 transformation using bilinear interpolation, implemented as a custom PyTorch module
865 (`DifferentiableMaskWarpFixed`). The mask values are passed through a straight-
866 through estimator (STE) to maintain binary structure while permitting gradient flow.
- 867 2. **Differentiable forward model.** The CASSI forward model $y = \text{CASSI}(x; \theta)$ is
868 implemented as a differentiable PyTorch module (`DifferentiableCassiForwardSTE`)
869 that accepts mismatch parameters as differentiable inputs.
- 870 3. **GPU grid initialization.** A full-range 3D grid search over (dx, dy, θ) with $9 \times 9 \times 7 =$
871 567 points provides diverse starting candidates. The top 9 candidates seed multi-start
872 gradient refinement.
- 873 4. **Staged gradient refinement.** Each of the 9 candidates is refined using Adam
874 optimization (learning rate 10^{-2} , decaying to 10^{-3}) for 200 steps. For each candidate,
875 4 random restarts with jittered initialization guard against local minima. The loss
876 function is the negative PSNR computed via an unrolled K -iteration differentiable
877 GAP-TV solver (`DifferentiableGAPTV`, $K = 10$ unrolled iterations).

878 Total runtime for Algorithm 2 is approximately 3,200 seconds (200 steps $\times$ 4 restarts $\times$
879 9 candidates with early stopping). Accuracy improves to ± 0.05 – 0.1 pixels, a 3–5 $\times$ improve-
880 ment over Algorithm 1. The two algorithms are used sequentially in practice: Algorithm 1
881 provides a warm start, and Algorithm 2 refines to sub-pixel precision.

882 Evaluation Protocol

883 **Four-Scenario Protocol.** We evaluate every modality under four standardized scenarios
884 that isolate different sources of quality degradation:

885 **Scenario I (Ideal):** $\mathbf{y}_{\text{obs}} = H_{\text{true}} \mathbf{x}_{\text{gt}}$; reconstruct with H_{true} . In this scenario the system
886 is perfectly calibrated ($H_{\text{true}} = H_{\text{nom}}$), so the operator used for reconstruction matches
887 the one that generated the data. This yields the oracle upper bound on reconstruction
888 quality, limited only by the sensing geometry and solver convergence.

889 **Scenario II (Mismatch):** $\mathbf{y}_{\text{obs}} = H_{\text{true}} \mathbf{x}_{\text{gt}}$; reconstruct with H_{nom} ($H_{\text{nom}} \neq H_{\text{true}}$). This
890 is the standard operating condition in practice: the measurement is generated by the
891 true physics, but the reconstruction uses a nominal (potentially mismatched) forward
892 model.

893 **Scenario III (Oracle):** Same measurements as Scenario II ($\mathbf{y}_{\text{obs}} = H_{\text{true}} \mathbf{x}_{\text{gt}}$ with $H_{\text{true}} \neq$
894 H_{nom}); reconstruct with H_{true} instead of H_{nom} . Provides the correction ceiling: the
895 best reconstruction achievable when the true operator is known exactly, applied to
896 data that were sensed by the mismatched system. The gap between Scenario III and
897 Scenario I reveals the irreducible loss from the degraded sensing configuration itself
898 (e.g., a shifted mask pattern is suboptimal even when perfectly known).

899 **Scenario IV (Corrected):** $\mathbf{y}_{\text{obs}} = H_{\text{true}} \mathbf{x}_{\text{gt}}$; reconstruct with $\hat{H} = H(\hat{\theta})$ where $\hat{\theta}$ is
900 estimated by Algorithms 1 and 2. This quantifies the benefit of mismatch calibration.

901 **Metrics.** Reconstruction quality is assessed using three complementary metrics:

- 902 • **PSNR** (peak signal-to-noise ratio, in dB): the primary metric, computed per scene
903 and averaged. For signals normalized to $[0, 1]$, $\text{PSNR} = 10 \log_{10}(1/\text{MSE})$. For SPC
904 data normalized to $[0, 255]$, the peak value is 255.
- 905 • **SSIM** (structural similarity index): captures perceptual quality including luminance,
906 contrast, and structural components, computed with a Gaussian window of width 11
907 and standard deviation 1.5.
- 908 • **SAM** (spectral angle mapper): for hyperspectral modalities (CASSI), measures the
909 angle between predicted and true spectral vectors at each spatial location, reported
910 in degrees. Lower is better.

911 Datasets.

- 912 • **CASSI:** 10 scenes from the KAIST dataset³⁵, each a $256 \times 256 \times 28$ spectral cube
913 (28 spectral bands from 450 nm to 650 nm). Data range $[0, 1]$.

914 • **CACTI**: 6 benchmark videos, each $256 \times 256 \times 8$ (8 temporal frames encoded per
915 snapshot). Data range $[0, 1]$.

916 • **SPC**: 11 natural images from the Set11 benchmark, each 256×256 grayscale. Data
917 range $[0, 255]$.

918 All per-scene metrics are reported individually as well as averaged, and all reconstruction
919 arrays are saved as NumPy NPZ files.

920 Experimental Details

921 **Hardware.** All experiments are conducted on a single NVIDIA GPU. Algorithm 1 (beam
922 search) and all solver-based reconstructions use the GPU for matrix–vector products and
923 FFT operations. Algorithm 2 (gradient refinement) additionally uses PyTorch automatic
924 differentiation on the same GPU.

925 **CASSI configuration.** The coded aperture snapshot spectral imaging (CASSI) system
926 uses a TSA-Net binary mask of dimensions 256×256 , with 28 spectral bands dispersed along
927 the spatial dimension. The five-parameter mismatch model $\theta = (dx, dy, \theta, a_1, \alpha)$ describes:
928 mask spatial shift in x (dx , pixels), mask spatial shift in y (dy , pixels), mask rotation angle
929 (θ , degrees), dispersion slope (a_1 , pixels per band), and dispersion axis angle (α , degrees).
930 An optional sixth parameter, PSF blur width (σ_{psf} , pixels), is available but not used in the
931 primary experiments. These mismatch parameter values were determined through system-
932 atic characterization of realistic CASSI assembly errors (Supplementary Note 8). The true
933 mismatch parameters are $\theta_{\text{true}} = (dx = 0.5 \text{ px}, dy = 0.3 \text{ px}, \theta = 0.1^\circ, a_1 = 2.02, \alpha =$
934 $0.15^\circ)$. Solvers evaluated include TwIST³⁹, GAP-TV⁴⁰, DGSMP⁴¹, MST-L⁷, and CST-
935 L⁴², all of which receive the same operator and differ only in their reconstruction algorithm.
936 The supplementary per-scene analysis additionally includes DeSCI⁴³ and HDNet³⁸.

937 **CACTI configuration.** The coded aperture compressive temporal imaging system uses
938 binary temporal masks of dimensions 256×256 , encoding 8 video frames into a single
939 snapshot measurement. Mismatch is parameterized as a temporal mask timing offset (sub-
940 frame shift). The default solver is EfficientSCI²¹.

941 **SPC configuration.** The single-pixel camera uses random binary measurement patterns
942 at three compression ratios: 10%, 25%, and 50% ($r = m/n \in \{0.10, 0.25, 0.50\}$). Mismatch
943 is modeled as an exponential gain drift ($g_i = \exp(-\alpha \cdot i)$) on the measurement matrix. The
944 default solver is FISTA-TV with total-variation regularization.

945 **MRI configuration.** Cartesian k -space sampling with $4\times$ acceleration (25% of k -space
946 lines acquired). Mismatch is parameterized as a 5% multiplicative error in the coil sensitivity

947 maps used for parallel imaging reconstruction. The default solver is SENSE¹⁷ with ℓ_1 -
 948 wavelet regularization.

949 **ESPIRiT comparison protocol.** To quantitatively compare PWM with ESPIRiT¹⁰,
 950 we evaluate four conditions on the same multi-coil MRI configuration: (1) Scenario I (true
 951 maps), (2) Scenario II (5% mismatched maps), (3) ESPIRiT auto-calibrated maps esti-
 952 mated from the 24-line ACS region via eigenvalue decomposition of the calibration matrix,
 953 and (4) PWM beam-search corrected maps (grid search over per-coil amplitude and phase
 954 scaling). All four conditions use the same CG-SENSE solver ($\ell = 10^{-3}$, 30 iterations). The
 955 comparison isolates the effect of map quality on reconstruction, holding the solver constant.

956 **CT configuration.** Fan-beam geometry with 180 projections over 180° . Mismatch is
 957 modeled as a center-of-rotation (CoR) offset, which produces characteristic arc artifacts in
 958 the reconstruction. The default solver is filtered back-projection (FBP)¹⁵ with a Ram-Lak
 959 filter, supplemented by iterative SART for comparison.

960 **CASSI real-data configuration.** The TSA real hyperspectral dataset¹² consists of 5
 961 scenes at 660×660 spatial resolution with 28 spectral bands and mask-shift step 2. Four
 962 solvers are evaluated: GAP-TV (200 iterations), HDNet (pre-trained checkpoint, full spatial
 963 resolution), MST-S and MST-L (pre-trained checkpoints, centre-cropped to 256×256 due
 964 to hardcoded spatial assumptions in the model architecture). The coded aperture mask
 965 is perturbed by $dx = 0.5$ px, $dy = 0.3$ px to simulate assembly-induced mismatch. No
 966 ground truth is available; quality is assessed via the normalised measurement residual $r =$
 967 $\|\mathbf{y} - H\hat{\mathbf{x}}\|^2 / \|\mathbf{y}\|^2$.

968 **CACTI real-data configuration.** The EfficientSCI real temporal dataset²¹ consists of
 969 4 dynamic scenes (duomino, hand, pendulumBall, waterBalloon) at 512×512 with compres-
 970 sion ratio 10. The real mask is stored separately from the measurement data. Two solvers
 971 are evaluated: GAP-TV (50 iterations) and PnP-FFDNet (50 iterations with FFDNet de-
 972 noiser). Mismatch is induced by shifting the mask by $dx = 0.5$ px, $dy = 0.3$ px. Quality is
 973 assessed via the normalised measurement residual and total variation of the reconstruction.

974 **Controlled hardware experiment protocol.** The software-perturbation protocol above
 975 applies calibrated mask shifts to existing real measurements. A full hardware-in-the-loop
 976 validation requires physically displacing the coded aperture mask and re-acquiring data.
 977 The protocol proceeds as follows: (i) acquire a baseline dataset with the mask at its
 978 factory-calibrated position; (ii) physically translate the mask by a known displacement
 979 ($\Delta x \in \{0.25, 0.5, 1.0\}$ px equivalent, verified by micrometer stage) and re-acquire under
 980 identical illumination; (iii) reconstruct both datasets with the factory mask specification and
 981 compute the PSNR degradation and measurement residual; (iv) apply PWM autonomous

calibration and measure recovery. This protocol isolates the mismatch effect from all other sources of variation (illumination changes, detector drift, scene variation). Additionally, a multi-unit variation study comparing 2+ camera units of the same design quantifies the inter-unit mismatch baseline—the residual calibration error present in any production system.

Clinical CT phantom configuration. For clinical translation, PWM is evaluated on CT quality assurance using the ACR CT accreditation phantom (Gammex 464). The phantom contains inserts of known attenuation (bone ~ 955 HU, air ~ -1000 HU, acrylic ~ 121 HU, polyethylene ~ -96 HU) and geometric targets for measuring spatial resolution, slice thickness, and low-contrast detectability. Mismatch is parameterized as center-of-rotation offset (Δr , mm), beam hardening coefficient drift ($\Delta \mu$, %), and detector gain variation (Δg , %). Ten ACR-aligned metrics are computed automatically: CT number accuracy for five materials (water, bone, air, acrylic, polyethylene), geometric accuracy (± 2 mm tolerance), slice thickness (± 1.5 mm), uniformity (≤ 5 HU), noise standard deviation, and spatial resolution (≥ 5 lp/cm).

Clinical MRI validation configuration. For MRI clinical validation, PWM processes multi-coil k -space data from public datasets (fastMRI⁴⁴). Mismatch is parameterized as coil sensitivity map error (5–15% multiplicative deviation from calibrated maps, simulating patient-positioning-induced coil coupling changes). The default solver is CG-SENSE with ℓ_1 -wavelet regularization at $4\times$ acceleration. Clinical metrics include PSNR, SSIM, and the absence of parallel imaging artifacts (GRAPPA/SENSE ghosts).

Ultrasound configuration. Pulse-echo B-mode imaging is simulated with a 64-element linear array, 5 MHz center frequency, element pitch 0.3 mm, sampling rate 100 MHz with 1024 time samples, and speed of sound (SoS) $c = 1540$ m/s. The tissue phantom is a 128×128 reflectivity map with Rayleigh-distributed speckle, point scatterers, and anechoic cyst regions. Frequency-dependent attenuation follows $\alpha(f) = 0.5 \text{ dB cm}^{-1} \text{ MHz}^{-1}$. Mismatch is parameterized as SoS error $\Delta c \in \{50, 100, 200, 400\}$ m/s, perturbing time-of-flight calculations. The default solver is delay-and-sum (DAS) beamforming (operator adjoint). Calibration uses grid search over 21 SoS candidates in $[1440, 1640]$ m/s, selecting the value that maximizes reconstruction PSNR.

Cryo-EM configuration. Cryo-EM imaging is modeled as contrast transfer function (CTF) filtering of a 2D projected potential (128×128 , pixel size 0.1 nm). The CTF is parameterized by defocus $\Delta f = -2000$ nm, spherical aberration $C_s = 2.0$ mm, electron wavelength $\lambda = 0.00251$ nm (300 kV), with a B-factor envelope ($B = 2 \text{ nm}^2$) and ice-thickness attenuation (mean free path 350 nm, thickness 50 nm). Mismatch is parameterized as defocus error $\Delta(\Delta f) \in \{100, 200, 500, 1000\}$ nm. The default solver is Wiener filtering with $\text{SNR} = 50$.

1018 Calibration uses grid search over 51 defocus values in $[-3000, -500]$ nm, selecting the value
1019 that maximizes reconstruction PSNR.

1020 **CT detector offset configuration.** Parallel-beam CT is simulated with 360 projection
1021 angles over $[0, \pi)$ and 128 detector elements. The object is a 128×128 Shepp–Logan
1022 phantom. Detector offset mismatch is simulated by shifting the sinogram along the detector
1023 axis by $\Delta s \in \{2, 5, 10, 20\}$ pixels. The default solver is filtered backprojection (FBP) with
1024 Shepp–Logan filter. Calibration uses grid search over 21 offset candidates in $[-25, 25]$ px,
1025 selecting the value that maximizes reconstruction PSNR.

1026 **Compressive holography configuration.** Multi-depth inline holography encodes a $(4 \times$
1027 $64 \times 64)$ 3D object into a single (64×64) hologram via Fresnel propagation at four depth
1028 planes with spacing $200 \mu\text{m}$, wavelength $\lambda = 532 \text{ nm}$, and pixel pitch $5 \mu\text{m}$. Mismatch is
1029 parameterized as propagation distance error $\Delta z \in \{10, 50, 100, 200\} \mu\text{m}$ applied uniformly to
1030 all depth planes. The default solver is FISTA-TV (80 iterations, $\lambda_{\text{TV}} = 0.005$). Calibration
1031 uses hologram residual minimisation: grid search over distance error candidates with the
1032 search range clamped to ensure all propagation distances remain positive.

1033 **Fluorescence microscopy configuration.** Widefield fluorescence imaging uses a dual-
1034 PSF Stokes-shift model with Gaussian excitation PSF ($\sigma_{\text{ex}} = 1.5 \text{ px}$) and emission PSF
1035 ($\sigma_{\text{em}} = 2.0 \text{ px}$), quantum yield $\eta = 0.7$, and background level $b = 0.02$. The specimen
1036 is a 64×64 image with fluorescent puncta, diffuse organelles, and filamentary structures.
1037 Mismatch is parameterized as PSF sigma error $\Delta\sigma \in \{0.1, 0.3, 0.5, 1.0\} \text{ px}$ applied to both
1038 excitation and emission PSFs. The default solver is Richardson–Lucy (80 iterations). Cal-
1039 ibration uses a 9×9 grid search over $(\sigma_{\text{ex}}, \sigma_{\text{em}})$ centered at the nominal values, selecting
1040 the pair that minimizes the measurement residual.

1041 **Real experimental data.** Three of the five multi-phantom modalities are validated on
1042 real experimental data from established benchmarks: ultrasound uses PICMUS experimen-
1043 tal phantom recordings and in-vivo carotid scans²² plus one DeepUS CIRS-040GSE acqui-
1044 sition²³; cryo-EM uses five EMDB 3D density maps (EMD-5778, EMD-2984, EMD-6287,
1045 EMD-11103, EMD-21375) projected at random orientations; CT uses three LoDoPaB-CT
1046 patient slices²⁴, one FIPS walnut μCT scan²⁵, and one HTC 2022 sample²⁶. Compressive
1047 holography and fluorescence microscopy retain synthetic multi-phantom benchmarks as no
1048 public datasets with calibrated ground truth are available for these modalities.

1049 Statistical Analysis

1050 **Per-scene reporting.** All metrics are reported per scene, not merely as dataset averages.
1051 This enables identification of scene-dependent failure modes (*e.g.*, spectrally flat scenes that

are inherently harder for CASSI, or textureless regions that challenge SPC).

Summary statistics. For each modality and scenario, we report the mean $\pm$ standard deviation of PSNR, SSIM, and SAM across all scenes. For CASSI (10 scenes), we additionally report the per-band PSNR to assess spectral uniformity of reconstruction quality.

Recovery ratio confidence intervals. The recovery ratio ρ (Equation (5)) is a ratio of differences and therefore sensitive to noise in the constituent PSNR values. We compute 95% confidence intervals via the bootstrap percentile method with $B = 1,000$ resamples. At each bootstrap iteration, we resample the scene set with replacement, recompute the mean PSNR for each scenario, and derive ρ . The 2.5th and 97.5th percentiles of the bootstrap distribution define the 95% CI.

Parameter recovery accuracy. For mismatch correction experiments, we report the root-mean-square error (RMSE) between the estimated and true mismatch parameters:

$$\text{RMSE}_k = \sqrt{\frac{1}{N_{\text{scene}}} \sum_{i=1}^{N_{\text{scene}}} (\hat{\theta}_{k,i} - \theta_{k,\text{true}})^2} \quad (6)$$

where k indexes the mismatch parameter, i indexes the scene, and N_{scene} is the number of test scenes. Uncertainty in the RMSE is estimated via bootstrap ($B = 1,000$).

Ablation significance. Ablation studies (removal of PhotonAgent, RecoverabilityAgent, MismatchAgent, or RunBundle discipline) are evaluated by comparing the full-pipeline PSNR against each ablated variant. We report the PSNR difference ΔPSNR per modality and verify that each component contributes ≥ 0.5 dB across all validated modalities, establishing practical significance.

Code and Data Availability

Source code. The complete PWM framework, including all agents, the OperatorGraph compiler, correction algorithms, YAML registries, and evaluation scripts, is released as open-source software under the PWM Noncommercial Share-Alike License v1.0 at https://github.com/integritynoble/Physics_World_Model. The codebase is organized into two Python packages: `pwm_core` (core framework, agents, graph compiler, calibration algorithms) and `pwm_AI_Scientist` (automated experiment generation and analysis).

Reconstruction data. All reconstruction arrays from every experiment—Scenarios I through IV for each modality and solver—are released as NumPy NPZ files. Files are

1080 stored using Git LFS and require `allow_pickle=True` for loading. Data ranges are stan-
1081 dardized: CASSI and CACTI reconstructions are normalized to $[0, 1]$; SPC reconstructions
1082 are in $[0, 255]$.

1083 **Experiment manifests.** Every experiment is recorded in a RunBundle v0.3.0 manifest
1084 containing: the git commit hash at execution time, all random number generator seeds,
1085 platform information (Python version, GPU model, CUDA version), SHA-256 hashes of all
1086 input data and output artifacts, metric values, and wall-clock timestamps. These manifests
1087 enable exact reproduction of every reported result.

1088 **Registry data.** All 9 YAML registries (7,034 lines total) that drive the agent system—
1089 including modality definitions, graph templates, photon models, mismatch databases, com-
1090 pression tables, solver configurations, primitive specifications, dataset paths, and acceptance
1091 thresholds—are publicly available in the repository under `packages/pwm_core/contrib/`.
1092 The `ExperimentSpec` JSON schemas used for pipeline input validation are included along-
1093 side worked examples in `examples/`.

1094 References

- 1095 [1] Chengshuai Yang. The finite primitive basis theorem for computational imaging:
1096 Formal foundations of the OperatorGraph representation. *Under review at SIAM*
1097 *Journal on Imaging Sciences*, 2026. Manuscript available at [https://github.com/](https://github.com/integritynoble/Physics_World_Model)
1098 [integritynoble/Physics_World_Model](https://github.com/integritynoble/Physics_World_Model).
- 1099 [2] Emmanuel J. Candès and Michael B. Wakin. An introduction to compressive sampling.
1100 *IEEE Signal Processing Magazine*, 25(2):21–30, 2008. doi: 10.1109/MSP.2007.914731.
- 1101 [3] David L. Donoho. Compressed sensing. *IEEE Transactions on Information Theory*,
1102 52(4):1289–1306, 2006. doi: 10.1109/TIT.2006.871582.
- 1103 [4] Singanallur V. Venkatakrishnan, Charles A. Bouman, and Brendt Wohlberg. Plug-
1104 and-play priors for model based reconstruction. In *Proceedings of the IEEE Global*
1105 *Conference on Signal and Information Processing (GlobalSIP)*, pages 945–948, 2013.
1106 doi: 10.1109/GlobalSIP.2013.6737048.
- 1107 [5] Xin Yuan, Yang Liu, Jinli Suo, and Qionghai Dai. Plug-and-play algorithms for video
1108 snapshot compressive imaging. *IEEE Transactions on Pattern Analysis and Machine*
1109 *Intelligence*, 44(10):7093–7111, 2021.
- 1110 [6] Vishal Monga, Yuelong Li, and Yonina C. Eldar. Algorithm unrolling: Interpretable,
1111 efficient deep learning for signal and image processing. *IEEE Signal Processing Maga-*
1112 *zine*, 38(2):18–44, 2021. doi: 10.1109/MSP.2020.3016905.

- [7] Yuanhao Cai, Jing Lin, Xiaowan Hu, Haoqian Wang, Xin Yuan, Yulun Zhang, Radu Timofte, and Luc Van Gool. Mask-guided spectral-wise transformer for efficient hyperspectral image reconstruction. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 17502–17511, 2022.
- [8] Xin Yuan, David J. Brady, and Aggelos K. Katsaggelos. Snapshot compressive imaging: Theory, algorithms, and applications. *IEEE Signal Processing Magazine*, 38(2):65–88, 2021.
- [9] Vegard Antun, Francesco Renna, Clarice Poon, Ben Adcock, and Anders C. Hansen. On instabilities of deep learning in image reconstruction and the potential costs of AI. *Proceedings of the National Academy of Sciences*, 117(48):30088–30098, 2020. doi: 10.1073/pnas.1907377117.
- [10] Martin Uecker, Peng Lai, Mark J. Murphy, Patrick Virtue, Michael Elad, John M. Pauly, Shreyas S. Vasanawala, and Michael Lustig. ESPIRiT — an eigenvalue approach to autocalibrating parallel MRI: where SENSE meets GRAPPA. *Magnetic Resonance in Medicine*, 71(3):990–1001, 2014. doi: 10.1002/mrm.24751.
- [11] Andrew M. Maiden and John M. Rodenburg. An improved ptychographical phase retrieval algorithm for diffractive imaging. *Ultramicroscopy*, 109(10):1256–1262, 2009. doi: 10.1016/j.ultramic.2009.05.012.
- [12] Ashwin A. Wagadarikar, Renu John, Rebecca Willett, and David J. Brady. Single disperser design for coded aperture snapshot spectral imaging. *Applied Optics*, 47(10):B44–B51, 2008. doi: 10.1364/AO.47.000B44.
- [13] Patrick Llull, Xuejun Liao, Xin Yuan, Jianbo Yang, David Kittle, Lawrence Carin, Guillermo Sapiro, and David J. Brady. Coded aperture compressive temporal imaging. *Optics Express*, 21(9):10526–10545, 2013. doi: 10.1364/OE.21.010526.
- [14] Marco F. Duarte, Mark A. Davenport, Dharmpal Takhar, Jason N. Laska, Ting Sun, Kevin F. Kelly, and Richard G. Baraniuk. Single-pixel imaging via compressive sampling. *IEEE Signal Processing Magazine*, 25(2):83–91, 2008. doi: 10.1109/MSP.2007.914730.
- [15] L. A. Feldkamp, L. C. Davis, and J. W. Kress. Practical cone-beam algorithm. *Journal of the Optical Society of America A*, 1(6):612–619, 1984. doi: 10.1364/JOSAA.1.000612.
- [16] J. M. Rodenburg and H. M. L. Faulkner. A phase retrieval algorithm for shifting illumination. *Applied Physics Letters*, 85(20):4795–4797, 2004. doi: 10.1063/1.1823034.

- [17] Klaas P. Pruessmann, Markus Weiger, Markus B. Scheidegger, and Peter Boesiger. SENSE: Sensitivity encoding for fast MRI. *Magnetic Resonance in Medicine*, 42(5): 952–962, 1999. doi: 10.1002/(SICI)1522-2594(199911)42:5<952::AID-MRM16>3.0.CO; 2-S.
- [18] Michael Lustig, David Donoho, and John M. Pauly. Sparse MRI: The application of compressed sensing for rapid MR imaging. *Magnetic Resonance in Medicine*, 58(6): 1182–1195, 2007. doi: 10.1002/mrm.21391.
- [19] David J. Brady. *Optical Imaging and Spectroscopy*. John Wiley & Sons, 2009.
- [20] Chengshuai Yang. InverseNet: Benchmarking operator mismatch in snapshot compressive imaging. In *Under review at ECCV 2026*, 2026.
- [21] Lishun Wang, Miao Cao, and Xin Yuan. EfficientSCI: Densely connected network with space-time factorization for large-scale video snapshot compressive imaging. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 18477–18486, 2023.
- [22] Hervé Liebgott, Alfonso Rodriguez-Molares, Françoise Cervenansky, Jørgen Arendt Jensen, and Olivier Bernard. Plane-wave imaging challenge in medical ultrasound. In *Proc. IEEE International Ultrasonics Symposium (IUS)*, pages 1–4, 2016. doi: 10.1109/ULTSYM.2016.7728908.
- [23] Saif Khan, Jaeyoung Huh, Manish Bhatt, Md. Murad Hossain, and Arun Bhatt. A deep learning framework for ultrasound image formation and image quality assessment. *IEEE Transactions on Ultrasonics, Ferroelectrics, and Frequency Control*, 70(9):972–984, 2023. DeepUS dataset: <https://zenodo.org/doi/10.5281/zenodo.7986407>.
- [24] Johannes Leuschner, Maximilian Schmidt, Daniel Otero Baguer, and Peter Maaß. LoDoPaB-CT, a benchmark dataset for low-dose computed tomography reconstruction. *Scientific Data*, 8:109, 2021. doi: 10.1038/s41597-021-00893-z.
- [25] Keijo Hämäläinen, Lauri Harhanen, Andreas Hauptmann, Aki Kallonen, Esa Niemi, and Samuli Siltanen. Total variation regularization for large-scale X-ray tomography. *International Journal of Tomography and Simulation*, 28(1):1–25, 2015. FIPS walnut dataset: Zenodo 1254206.
- [26] Alexander Meaney, Andreas Hauptmann, and Samuli Siltanen. Helsinki tomography challenge 2022: description, methods, and results. *Applied Mathematics for Modern Challenges*, 1(1):1–28, 2024. doi: 10.3934/ammc.2024001.
- [27] David J. Brady, Kerkil Choi, Daniel L. Marks, Ryoichi Horisaki, and Sehoon Lim. Compressive holography. *Optics Express*, 17(15):13040–13049, 2009.

- [28] Hanchen Wang, Tianfan Fu, Yuanqi Du, Wenhao Gao, Kexin Huang, Ziming Liu, Payal Chandak, Shengchao Liu, Peter Van Katwyk, Andreea Deac, et al. Scientific discovery in the age of artificial intelligence. *Nature*, 620:47–60, 2023.
- [29] David J. Brady, Michael E. Gehm, Ronald A. Stack, Daniel L. Marks, David S. Kittle, Dathon R. Golish, E. M. Vera, and Steven D. Feller. Multiscale gigapixel photography. *Nature*, 486(7403):386–389, 2012.
- [30] Yuan Xu, Ti Bai, Hao Yan, Luo Ouyang, Arnold Pompos, Jing Wang, Linghong Zhou, Steve B. Jiang, and Xun Jia. A practical cone-beam CT scatter correction method with optimized Monte Carlo simulations for image-guided radiation therapy. *Physics in Medicine and Biology*, 60(9):3567–3587, 2015. doi: 10.1088/0031-9155/60/9/3567.
- [31] Xiao Liang, Liyuan Chen, Dan Nguyen, Zhiguo Zhou, Xuejun Gu, Ming Yang, Jing Wang, and Steve Jiang. Generating synthesized computed tomography (CT) from cone-beam computed tomography (CBCT) using CycleGAN for adaptive radiation therapy. *Physics in Medicine and Biology*, 64(12):125002, 2019. doi: 10.1088/1361-6560/ab22f9.
- [32] Chenyang Shen, Yesenia Gonzalez, Liyuan Chen, Steve B. Jiang, and Xun Jia. Intelligent parameter tuning in optimization-based iterative CT reconstruction via deep reinforcement learning. *IEEE Transactions on Medical Imaging*, 37(6):1430–1439, 2018. doi: 10.1109/TMI.2018.2823679.
- [33] Eric E. Klein, Joseph Hanley, John Bayouth, Fang-Fang Yin, William Simon, Sean Dresser, Christopher Serago, Francisco Aguirre, Lijun Ma, Bijan Arjomandy, Chihray Liu, Craig Sandin, and Timothy Holmes. Task Group 142 report: Quality assurance of medical accelerators. *Medical Physics*, 36(9):4197–4212, 2009.
- [34] Sasa Mutic, Jatinder R. Palta, Elizabeth K. Butker, Indra J. Das, M. Saiful Huq, Leh-Nien Dick Loo, Bill J. Salter, Cynthia H. McCollough, and Jacob Van Dyk. Quality assurance for computed-tomography simulators and the computed-tomography-simulation process: Report of the AAPM Radiation Therapy Committee Task Group No. 66. *Medical Physics*, 30(10):2762–2792, 2003.
- [35] Inchang Choi, Daniel S. Jeon, Giljoo Nam, Diego Gutierrez, and Min H. Kim. High-quality hyperspectral reconstruction using a spectral prior. *ACM Transactions on Graphics (Proceedings of SIGGRAPH Asia)*, 36(6):218:1–218:13, 2017. doi: 10.1145/3130800.3130810.
- [36] George Cybenko. Approximation by superpositions of a sigmoidal function. *Mathematics of Control, Signals and Systems*, 2(4):303–314, 1989.
- [37] Maziar Raissi, Paris Perdikaris, and George E. Karniadakis. Physics-informed neural networks: A deep learning framework for solving forward and inverse problems involv-

- ing nonlinear partial differential equations. *Journal of Computational Physics*, 378: 686–707, 2019.
- [38] Xiaowan Hu, Yuanhao Cai, Jing Lin, Haoqian Wang, Xin Yuan, Yulun Zhang, Radu Timofte, and Luc Van Gool. HDNet: High-resolution dual-domain learning for spectral compressive imaging. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 17542–17551, 2022.
- [39] José M. Bioucas-Dias and Mário A. T. Figueiredo. A new TwIST: Two-step iterative shrinkage/thresholding algorithms for image restoration. *IEEE Transactions on Image Processing*, 16(12):2992–3004, 2007. doi: 10.1109/TIP.2007.909319.
- [40] Xin Yuan. Generalized alternating projection based total variation minimization for compressive sensing. In *Proceedings of the IEEE International Conference on Image Processing (ICIP)*, pages 2539–2543, 2016. doi: 10.1109/ICIP.2016.7532817.
- [41] Tao Huang, Weisheng Dong, Xin Yuan, Jinjian Wu, and Guangming Shi. Deep gaussian scale mixture prior for spectral compressive imaging. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 16216–16225, 2021.
- [42] Yuanhao Cai, Jing Lin, Xiaowan Hu, Haoqian Wang, Xin Yuan, Yulun Zhang, Radu Timofte, and Luc Van Gool. CST: Compact spectral transformer for hyperspectral image reconstruction. In *Proceedings of the European Conference on Computer Vision (ECCV)*, 2022.
- [43] Yang Liu, Xin Yuan, Jinli Suo, David J. Brady, and Qionghai Dai. Rank minimization for snapshot compressive imaging. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 41(12):2990–3006, 2019.
- [44] Jure Zbontar, Florian Knoll, Anuroop Sriram, Tullie Murrell, Zhengnan Huang, Matthew J. Muckley, Aaron Defazio, Ruben Stern, Patricia Johnson, Mary Bruno, Marc Parente, Krzysztof J. Geras, Joe Katsnelson, Hersh Chandarana, Zizhao Zhang, Michal Drozdal, Adriana Romero, Michael Rabbat, Pascal Vincent, Nafissa Yakubova, James Pinkerton, Duo Wang, Erich Owens, C. Lawrence Zitnick, Michael P. Recht, Daniel K. Sodickson, and Yvonne W. Lui. fastMRI: An open dataset and benchmarks for accelerated MRI. *arXiv preprint arXiv:1811.08839*, 2018.




Figure 1: **PWM overview.** The Physics World Models pipeline. **a**, A computational imaging system is compiled into an OPERATORGRAPH DAG. **b**, The TRIAD DECOMPOSITION diagnostic agents evaluate each gate. **c**, The dominant gate is identified and a TRIADREPORT is produced. **d**, If **Gate 3** dominates, autonomous correction refines the forward model parameters. **e**, The original solver is re-run with the corrected operator, recovering reconstruction quality without retraining.




Figure 2: **OperatorGraph IR and Physics Fidelity Ladder.** **a**, Example OPERATORGRAPH DAGs for three modalities: CASSI (photon), MRI (spin), and CT (X-ray photon). Each node wraps a primitive operator; edges define data flow. **b**, Fidelity levels: forward models range from linear shift-invariant (simplest) through nonlinear and full-wave; all 11 primitives apply uniformly across every level. **c**, Summary statistics: 168 registered modality templates across 19 categories (12 with full correction validation), 5 physical carriers.




Figure 3: **Triad Decomposition structure and gate binding.** **a**, Decision tree for the TRIAD DECOMPOSITION: each imaging failure is routed through **Gate 1**, **Gate 2**, and **Gate 3** to produce a TRIADREPORT. **b**, Gate binding heatmap across 14 correction configurations (12 distinct modalities, 5 carrier families). Red indicates **Gate 3** dominance (all 14/14 configurations), blue indicates **Gate 1**, and amber indicates **Gate 2**. **c**, Recovery ratio ρ distribution across all 14 correction configurations, demonstrating that Gate 3 dominance holds universally across optical, X-ray, electron, spin, and acoustic carriers.




Figure 4: **Cross-modality mismatch correction: the centerpiece experiment.** **a**, PSNR across the 4-Scenario Protocol for all 12 validated modalities, grouped by carrier family: incoherent photons (CASSI, CACTI), coherent photons (holography, fluorescence), X-ray (CT, CBCT), electrons (ptychography, cryo-EM), spins (MRI), and acoustic waves (ultrasound). For each modality, four bars show: Scenario I (ideal, ceiling), Scenario II (mismatched, floor), Scenario III (oracle correction ceiling), and Scenario IV (PWM-corrected). Error bars show 95% bootstrap CIs over test scenes. The gap between Scenario I and II is the mismatch-induced degradation; the gap between Scenario II and IV is the autonomous recovery. In 10 of 12 modalities, a simple solver on the corrected operator (Scenario IV) outperforms the best available solver on the mismatched operator (Scenario II), demonstrating that calibration beats solver upgrades. **b**, Zero-shot generalization: hyperparameters tuned on photon-domain modalities (CASSI, CACTI, SPC) transfer without modification to coherent-photon, spin, X-ray, electron, and acoustic domains, with comparable correction gains (< 0.5 dB difference). Dark bars: modality-specific tuning; light bars: zero-shot transfer.




Figure 5: **Modality deep dives and visual comparison.** **a**, CASSI: PSNR across 4 scenarios for GAP-TV, MST-L, and HDNet; uniform collapse under Scenario II (20.83–21.88 dB) confirms operator-driven failure. **b**, CACTI: four methods across 4 scenarios, showing up to 20.58 dB mismatch degradation. **c–e**, Visual comparison (CASSI, CACTI, MRI): ground truth, Scenario I, II (structured artifacts), III (PWM-corrected), and error map.




Figure 6: **Hardware validation on real CASSI and CACTI instruments.** **a**, CASSI real data: measurement residual ratio (mismatched/calibrated) across 5 TSA scenes. GAP-TV shows $1.8\times$ mean ratio. **b**, CACTI real data: residual ratio across 4 scenes. GAP-TV shows $10.4\times$ mean ratio; PnP-FFDNet shows $2.0\times$. **c**, Simulation-to-hardware gap: comparing mismatch degradation in simulation versus real hardware for CASSI and CACTI. **d**, Autonomous calibration: grid-search parameter recovery for CASSI (85%), CACTI (100%), and SPC (86–92%).




Figure 7: **Basis-growth saturation and primitive decomposition.** **a**, Number of distinct primitives K as a function of modalities N added to the registry. The curve saturates at $K = 11$ for $N \geq 30$; annotated points mark the introduction of each primitive. **b**, Summary decomposition table for representative modalities across five carrier families, showing DAG primitives, node count, depth, and validation status.

Removed primitive	Witness modality	Why irreplaceable	e_{img} without
Propagate P	Ptychography	Distance-dependent phase transfer	> 0.05
Modulate M	CASSI	Arbitrary element-wise mask	> 0.10
Project Π	CT	Line-integral (Radon) geometry	> 0.15
Encode F	MRI	Non-uniform Fourier sampling	> 0.20
Convolve C	Lensless imaging	Arbitrary shift-invariant PSF kernel	> 0.08
Accumulate Σ	SPC	Dimensional summation (not scaling)	> 0.12
Detect D	All modalities	Carrier-to-measurement conversion	> 0.50
Sample S	MRI	Index-set restriction (k -space)	> 0.10
Disperse W	CASSI	λ -parameterized spatial shift	> 0.06
Scatter R	Compton imaging	Direction change + energy shift	0.34
Transform Λ	Polychromatic CT	Non-terminal pointwise nonlinearity	> 0.05

Table 1: **Primitive necessity ablation.** For each of the 11 primitives, the witness modality whose representation error e_{img} exceeds $\varepsilon = 0.01$ when that primitive is removed from $\mathcal{B}$. Eight primitives (P , M , Π , F , Σ , D , R , Λ) are strictly necessary; the remaining three (C , S , W) are necessary under the stated complexity bounds ($N_{\text{max}} = 20$, $D_{\text{max}} = 10$). Data from¹, Proposition 1.

1245 **Extended Data Table 1** | Oracle correction ceiling (Scenario III) across 14 validated
1246 configurations (12 modalities, 5 carrier families). Full table with per-scene metrics in Sup-
1247plementary Table S1; autonomous correction recovery (Scenario IV) in Supplementary Ta-
1248ble S9. Per-modality supplementary tables: ultrasound (S14), cryo-EM (S15), CBCT (S16),
1249compressive holography (S17), fluorescence microscopy (S18), SPC (S19), lensless imaging
1250(S20).

1251 **Extended Data Table 2** | 26-modality template registry spanning five physical carrier
1252 families. Full table in Supplementary Table S3.