

# Ten Primitives and Three Gates: The Universal Structure of Computational Imaging

Chengshuai Yang<sup>1,\*</sup>      Xin Yuan<sup>2</sup>

## Abstract

Computational imaging systems routinely underperform because the assumed forward model diverges from the true physics. Here we prove two results. First, the **Finite Primitive Basis Theorem**: every linear, shift-variant imaging forward model admits an  $\varepsilon$ -approximate representation as a typed directed acyclic graph over exactly 10 canonical primitives. Second, the **Triad Decomposition**: every reconstruction failure decomposes into three root causes—information deficiency, carrier noise, and operator mismatch—with mismatch dominant across all validated modalities. Together these results yield a modality-agnostic diagnostic and correction framework. Across seven modalities and five physical carriers, autonomous correction recovers +0.8 to +10.7 dB of mismatch-induced degradation without retraining the solver. Hardware validation on real instruments confirms mismatch dominance. A held-out closure test on 8 additional modalities confirms basis completeness with saturating growth.

## Introduction

Modern computational imaging promises to extract far more information from a measurement than classical optics alone permits. Coded aperture spectral cameras compress three-dimensional scenes into a single snapshot<sup>1</sup>; compressive temporal imagers freeze high-speed video into one exposure<sup>2</sup>; single-pixel cameras reconstruct images from far fewer measurements than pixels<sup>3</sup>; accelerated MRI scanners reconstruct diagnostic-quality images from a fraction of the acquired  $k$ -space data<sup>4</sup>. Over the past decade, the community has invested enormous effort in improving reconstruction algorithms—progressing from compressed sensing<sup>5,6</sup> and plug-and-play priors<sup>7</sup> to deep unrolling networks<sup>8</sup> and vision transformers<sup>9</sup>—yet these algorithms routinely fail when deployed on real instruments, because the forward model assumed by the solver does not match the physics that generated the data.

These failures persist because computational imaging lacks two theoretical foundations. First, a **representation theory**: no formal result guarantees that a finite set of primitive operators suffices to represent all imaging forward models. Without this, every new modality requires bespoke engineering and the field cannot reason about completeness. Second, a

---

<sup>\*1</sup>NextGen PlatformAI C Corp, USA. <sup>2</sup>School of Engineering, Westlake University, Hangzhou, China. <sup>\*</sup>Correspondence: [integrityyang@gmail.com](mailto:integrityyang@gmail.com)

**diagnostic theory:** no systematic framework decomposes reconstruction failures into root causes. Calibration errors, noise, and information deficiency are conflated, and failures are addressed *ad hoc*. Calibration methods exist for specific instruments<sup>10,11</sup>, but they do not generalise; robustness studies perturb individual systems<sup>12</sup>, but they lack a unifying formalism.

Here we establish both foundations within Physics World Models (PWM). We prove the **Finite Primitive Basis Theorem** (Theorem 3): every imaging forward model in the operator class  $\mathcal{C}_{\text{Tier2}}$ —encompassing all clinical, scientific, and industrial modalities at Tier-2 fidelity—admits an  $\varepsilon$ -approximate representation as a typed DAG over exactly 10 canonical primitives. We establish the **Triad Decomposition**: every reconstruction failure decomposes into three gates—information deficiency (**Gate 1**), carrier noise (**Gate 2**), and operator mismatch (**Gate 3**)—with **Gate 3** dominant across all validated modalities. These two results are complementary: the Finite Primitive Basis provides a universal representation for forward models; the Triad Decomposition provides a universal diagnostic law over that representation. Together they yield a practical framework that diagnoses and corrects imaging failures across modalities without modality-specific tuning.

We validate both contributions across seven modalities spanning five physical carriers (CASSI, CACTI, SPC, lensless, ptychography<sup>13</sup>, MRI<sup>14</sup>, CT<sup>15</sup>), with hardware validation on real DD-CASSI and CACTI instruments, confirming Triad predictions on physical measurements. A held-out closure test on 8 additional modalities—including quantum ghost imaging, THz time-domain spectroscopy, and Compton scatter imaging—confirms basis completeness. The proof of Theorem 3 is in Supplementary Note 10; formal primitive semantics are developed in<sup>16</sup>.

## The Finite Primitive Basis

A foundational question in computational imaging is whether the diversity of imaging forward models—from coded aperture cameras to MRI scanners to electron microscopes—can be captured by a small, fixed set of primitive operators. We answer affirmatively.

**The OperatorGraph representation.** Every imaging forward model is encoded as a typed directed acyclic graph (DAG) in which each node wraps a single primitive physical operator and edges define the data flow from source to detector. Every primitive implements both a `forward()` and an `adjoint()` method, with validated adjoint consistency ensuring  $\langle H\mathbf{x}, \mathbf{y} \rangle = \langle \mathbf{x}, H^\dagger \mathbf{y} \rangle$  to within numerical precision. Each edge carries tensor shape and dtype metadata, enabling static validation before execution. We call this formalism the OPERATORGRAPH intermediate representation (IR).

**Ten canonical primitives.** The OPERATORGRAPH IR defines a library of exactly 10 canonical primitives  $\mathcal{B} = \{P, M, \Pi, F, C, \Sigma, D, S, W, R\}$ :

| #  | Primitive  | Notation                    | Physical action                                      |
|----|------------|-----------------------------|------------------------------------------------------|
| 1  | Propagate  | $P(d, \lambda)$             | Free-space wave propagation                          |
| 2  | Modulate   | $M(\mathbf{m})$             | Element-wise multiplication (mask, coil, absorption) |
| 3  | Project    | $\Pi(\theta)$               | Radon line-integral projection                       |
| 4  | Encode     | $F(\mathbf{k})$             | Fourier-domain encoding ( $k$ -space)                |
| 5  | Convolve   | $C(\mathbf{h})$             | Spatial convolution (PSF)                            |
| 6  | Accumulate | $\Sigma$                    | Summation over spectral/temporal axis                |
| 7  | Detect     | $D(g, \eta)$                | Detector response (5 canonical families)             |
| 8  | Sample     | $S(\Omega)$                 | Sub-sampling on index set $\Omega$                   |
| 9  | Disperse   | $W(\alpha, a)$              | Wavelength-dependent spatial shift                   |
| 10 | Scatter    | $R(\sigma, \Delta\epsilon)$ | Direction change and/or energy shift                 |

68 The Detect nonlinearity  $\eta$  is restricted to five canonical families (linear-intensity, logarithmic,  
69 mic, sigmoid, Poisson-rate, coherent-field), each with at most 2 scalar parameters, preventing Detect from becoming a universal approximator (see<sup>16</sup> for formal definitions and  
70 non-universality proof).  
71

72 **Physics-stage mapping.** Each factor of a forward model is classified into one of four  
73 physics-stage families—propagation, interaction, encoding–projection, detection–readout—  
74 and mapped to primitives from  $\mathcal{B}$  (propagation  $\rightarrow \{P, C\}$ ; elastic interaction  $\rightarrow \{M\}$ ;  
75 scattering  $\rightarrow \{R\}$ ; encoding–projection  $\rightarrow \{\Pi, F\}$ ; detection–readout  $\rightarrow \{\Sigma, S, W, C, D\}$ ).

76 **Physics Fidelity Ladder.** The OPERATORGRAPH IR defines a four-tier Physics Fidelity  
77 Ladder: Tier 1 (linear, shift-invariant), Tier 2 (linear, shift-variant), Tier 3 (nonlinear,  
78 ray/wave-based), and Tier 4 (full-wave/Monte Carlo). The Finite Primitive Basis applies to  
79 Tiers 1–2; Tiers 3–4 are handled by refinement sub-DAGs that expand individual primitive  
80 nodes into higher-fidelity sub-graphs while preserving the top-level DAG structure.

81 **Five physical carriers.** The library spans five carrier families—photons, electrons, spins,  
82 acoustic waves, and particles—with 26 validated modality templates (Extended Data Table 2; Supplementary Table S3).  
83

84 **Theorem statement.**

85 **Definition 1** (Tier-2 Operator Class).  $\mathcal{C}_{\text{Tier2}}$  consists of all imaging forward models expressible as a finite sequential-parallel composition of linear, shift-variant stages with bounded operator norm ( $\|H_k\| \leq B$ ) and stage count  $K \leq N_{\max}$ .  
86  
87

88 **Definition 2** ( $\varepsilon$ -Approximate Representation). A typed DAG  $G$  with nodes  $V \subseteq \mathcal{B}$  is an  $\varepsilon$ -approximate representation of  $H \in \mathcal{C}_{\text{Tier2}}$  if  $\sup_{\|\mathbf{x}\| \leq 1} \|H(\mathbf{x}) - H_G(\mathbf{x})\| / \|H(\mathbf{x})\| \leq \varepsilon$  and  $|V| \leq N_{\max}$ ,  $\text{depth}(G) \leq D_{\max}$ .  
89  
90

91 **Theorem 3** (Finite Primitive Basis). For every  $H \in \mathcal{C}_{\text{Tier2}}$ , there exists a typed DAG  $G$   
92 with  $V \subseteq \mathcal{B}$  that is an  $\varepsilon$ -approximate representation of  $H$ .

93 *Proof sketch.* By Definition 1,  $H = H_K \circ \dots \circ H_1$  with  $K \leq N_{\max}$ . Each factor is classified  
 94 into one of four physics-stage families and realized by primitives: propagation factors  $\rightarrow P$   
 95 or  $C$  (wave-equation solutions); elastic interaction  $\rightarrow M$  (exact); scattering  $\rightarrow R$ ; encoding–  
 96 projection  $\rightarrow \Pi$  or  $F$  (exact); detection–readout  $\rightarrow$  finite composition from  $\{\Sigma, S, W, C, D\}$ .  
 97 The per-factor approximation errors compose via sub-multiplicativity:  $\|H - H_G\| \leq K \cdot$   
 98  $\max_k(\varepsilon_k) \cdot B^{K-1} \leq \varepsilon$  for  $\varepsilon_k$  bounded by the Tier-2 truncation. Full proof in Supplementary  
 99 Note 10; formal primitive semantics and typed DAG denotation in <sup>16</sup>.  $\square$

100 **Scope of  $\mathcal{C}_{\text{Tier2}}$ .** Covers: all clinical, scientific, and industrial modalities at Tier 1–2 (coded  
 aperture, interferometric, projection, Fourier-encoded, acoustic, scattering, THz). Excludes:  
 Tier 3–4 nonlinear models (handled by refinement sub-DAGs), quantum state tomography, rel-  
 ativistic regimes.

101 **Extension protocol.** A new primitive is warranted when no DAG over  $\mathcal{B}$  achieves  $e_{\text{Tier2}} \leq$   
 102  $\varepsilon$  within the complexity bounds. The extension requires: (1) validated forward/adjoint,  
 103 (2) demonstrated representation gap, (3) error reduction below  $\varepsilon$ , (4) need by  $\geq 2$  modal-  
 104 ities, (5) backward-compatible closure re-test. Scatter ( $R$ ) was added via this protocol when  
 105 Compton imaging produced  $e_{\text{Tier2}} = 0.34$  without it.

106 **Held-out closure test.** To validate Theorem 3 empirically, we conduct a closure test  
 107 under a frozen protocol: the primitive library  $\mathcal{B}$  (10 primitives), Detect families (5), fidelity  
 108 threshold ( $\varepsilon = 0.01$ ), and complexity bounds ( $N_{\max} = 20$ ,  $D_{\max} = 10$ ) are all frozen before  
 109 evaluation. We test 5 Tier-1 held-out modalities (OCT, photoacoustic, SIM, phase-contrast  
 110 X-ray, electron ptychography) and 3 exotic modalities (quantum ghost imaging, THz-TDS,  
 111 Compton scatter imaging):

| Modality             | $e_{\text{Tier2}}$ | #Nodes/Depth | Triad transfer | New prim.? |
|----------------------|--------------------|--------------|----------------|------------|
| OCT                  | < 0.01             | 5 / 3        | Y              | N          |
| Photoacoustic        | < 0.01             | 4 / 3        | Y              | N          |
| SIM                  | < 0.01             | 4 / 3        | Y              | N          |
| Phase-contrast X-ray | < 0.01             | 5 / 4        | Y              | N          |
| Electron ptycho      | < 0.01             | 4 / 3        | Y              | N          |
| Ghost imaging        | < 0.01             | 4 / 3        | Y              | N          |
| THz-TDS              | < 0.01             | 3 / 2        | Y              | N          |
| Compton scatter      | < 0.01*            | 4 / 3        | Y              | Y ( $R$ )  |

\*With  $R$  in library;  $e_{\text{Tier2}} > 0.01$  without  $R$ .

Max observed: 5 nodes / depth 4; median: 4 nodes / depth 3 (cf. bounds  $N_{\max} = 20$ ,  $D_{\max} = 10$ ).

112 Seven of eight modalities decompose with the frozen library. Quantum ghost imaging  
 113 is operator-equivalent to a single-pixel camera at the image-formation level—sharing the  
 114 canonical DAG Source  $\rightarrow M \rightarrow \Sigma \rightarrow D$  at Tier-2 abstraction despite fundamentally dif-  
 115 ferent source physics. THz-TDS requires only the coherent-field Detect family. Compton  
 116 scatter imaging triggered the extension protocol (Section “Extension protocol” above): its  
 117 representation gap ( $e_{\text{Tier2}} = 0.34 \gg \varepsilon$ ) met the falsifiable criterion for basis incompleteness,

and the disciplined resolution—adding Scatter ( $R$ )—covers 5+ additional modalities (Raman, fluorescence, DOT, Brillouin) while preserving all existing decompositions. The basis grows from 9 to 10 (11% increase) while modality coverage grows by 19%+.

**Basis-growth saturation.** Plotting the number of distinct primitives  $K$  against the number of registered modalities  $N$  reveals clear saturation (Figure 7): 8 of 10 primitives are introduced by the first 10 modalities, Disperse ( $W$ ) by CASSI-type spectral systems, and Scatter ( $R$ ) only when Compton/Raman-class modalities enter. The growth is sublinear and saturating:  $K = 10$  at  $N = 31$ , with no new primitive required for the most recent 19 modalities added. This saturation is consistent with Theorem 3: once all four physics-stage families are covered by primitives, new modalities compose existing primitives rather than requiring new ones.

**Theorem tightness.** The theoretical complexity bounds ( $N_{\max} = 20$ ,  $D_{\max} = 10$ ) are conservative: empirically, all 34 validated modalities require  $\leq 6$  nodes and depth  $\leq 5$  (held-out closure test table above), with the median modality using 4 nodes at depth 3. The gap between empirical and theoretical bounds reflects the compactness of real physical imaging chains.

## The Triad Decomposition

The TRIAD DECOMPOSITION asserts that every failure in computational image recovery within  $\mathcal{C}_{\text{Tier2}}$  can be attributed to one or more of exactly three root causes, which we term *gates*. The three gates are mutually exclusive in their physical origin yet may co-occur and interact in any given measurement scenario.

**Gate 1: Recoverability.** **Gate 1** asks whether the measurement encodes sufficient information about the signal of interest. Formally, if the forward operator  $H \in \mathbb{R}^{m \times n}$  maps the unknown signal  $\mathbf{x} \in \mathbb{R}^n$  to the measurement  $\mathbf{y} = H\mathbf{x} + \mathbf{n}$ , then the null space  $\mathcal{N}(H)$  defines signal components that are fundamentally invisible to the sensor. When  $\mathcal{N}(H)$  is large, no solver can recover the missing information. **Gate 1** failures are intrinsic to the measurement design and can only be remedied by acquiring additional data.

**Gate 2: Carrier Budget.** **Gate 2** asks whether the signal-to-noise ratio (SNR) is sufficient. Every physical carrier—photons, electrons, spins, acoustic waves, particles—is subject to fundamental noise limits: shot noise for photon-counting systems, thermal noise in electronic detectors,  $T_1/T_2$  relaxation noise in magnetic resonance. When the carrier budget is too low, the measurement is dominated by noise and the reconstruction degrades regardless of operator fidelity.

**Gate 3: Operator Mismatch.** **Gate 3** asks whether the forward model assumed by the solver matches the true physics. The solver operates with a nominal operator  $H_{\text{nom}}$ , but the data were generated by a true operator  $H_{\text{true}}$ . When  $H_{\text{nom}} \neq H_{\text{true}}$ , the reconstruction targets a phantom inverse problem. **Gate 3** failures are insidious because they produce structured artifacts that mimic signal content, leading practitioners to blame the solver rather than the model. Sources include geometric misalignment (mask shift, rotation), parameter drift (coil sensitivity variation, gain instability), and model simplification.

**Definition 4** (Triad Decomposition). *Let  $H \in \mathcal{C}_{\text{Tier2}}$  with measurement  $\mathbf{y} = H\mathbf{x} + \mathbf{n}$  and solver using nominal operator  $H_{\text{nom}}$ . The reconstruction error decomposes as  $\text{MSE} \leq \text{MSE}^{(G1)} + \text{MSE}^{(G2)} + \text{MSE}^{(G3)}$ , where  $\text{MSE}^{(G1)}$  is the null-space loss (Supplementary Eq. S1),  $\text{MSE}^{(G2)}$  is the noise-floor term (Supplementary Eq. S3), and  $\text{MSE}^{(G3)}$  is the mismatch-induced term (Supplementary Eq. S5).*

**Theorem 5** (Gate 3 Dominance). *For any  $H \in \mathcal{C}_{\text{Tier2}}$  operating above its Gate 1 floor ( $\gamma > \gamma_{\min}$ ) and Gate 2 floor ( $\text{SNR} > \text{SNR}_{\min}$ ), Gate 3 is the binding constraint whenever  $\|\delta\boldsymbol{\theta}\|_{\mathbf{J}^\top \mathbf{J}} > (\sigma_n / \|\mathbf{x}\|_\infty) \cdot \kappa(\mathbf{H})$ .*

*Proof.* See Supplementary Note 1, Proposition 2. □

**Key finding: Gate 3 dominates.** Across all 9 validated configurations (7 modalities), **Gate 3** is the dominant failure gate. In CASSI, a sub-pixel mask shift degrades MST-L by 13.98 dB; in MRI, a 5% coil sensitivity mismatch produces 6.94 dB degradation. The imaging community has been optimizing the wrong variable.

**Relationship to the Finite Primitive Basis.** The two contributions are complementary: the Finite Primitive Basis provides a universal representation (every forward model is a DAG over 10 primitives), and the Triad provides a universal diagnostic law over that representation. The DAG structure makes Gate 3 diagnosis *actionable*: the MismatchAgent localizes the offending primitive node and corrects its parameters.

## Consequences: Diagnosis and Correction

The two theoretical results directly imply a practical framework. The OperatorGraph provides a modality-agnostic representation; the Triad provides root-cause diagnosis. PWM implements this through three deterministic agents—no large language model, no learned parameters, no human intervention.

**Diagnostic agents.** Three deterministic agents evaluate each gate: the RecoverabilityAgent (Gate 1) estimates null-space dimension; the PhotonAgent (Gate 2) computes the carrier budget; and the MismatchAgent (Gate 3) detects and localizes operator mismatch to a specific OperatorGraph node. Details in Methods.

185 **Correction pipeline.** When **Gate 3** is dominant, PWM activates a two-stage correction  
 186 pipeline. **Algorithm 1 (Beam Search)** performs a coarse grid search over the declared  
 187 mismatch parameter family  $\theta = (\theta_1, \dots, \theta_k)$  associated with the offending node, scoring  
 188 candidates by reconstruction sharpness. **Algorithm 2 (Gradient Refinement)** takes  
 189 the top candidates and performs continuous optimization via backpropagation through the  
 190 DAG, combining data-fidelity and regularization losses. Correction operates exclusively on  
 191 the forward model, not on the solver: any existing solver—iterative, plug-and-play, or deep  
 192 unrolling—benefits without modification.

193 **4-Scenario Protocol.** To rigorously evaluate correction quality, PWM defines four canon-  
 194 ical scenarios. **Scenario I** (Ideal): the solver reconstructs using the true operator  $H_{\text{true}}$ .  
 195 **Scenario II** (Mismatch): the solver uses the nominal operator  $H_{\text{nom}}$  on data generated  
 196 by  $H_{\text{true}}$ . **Scenario III** (Corrected): the solver uses the PWM-corrected operator. **Sce-**  
 197 **nario IV** (Oracle): the true operator on mismatched data, providing the correction ceiling.  
 198 The recovery ratio  $\rho = (\text{PSNR}_{\text{III}} - \text{PSNR}_{\text{II}}) / (\text{PSNR}_{\text{I}} - \text{PSNR}_{\text{II}})$  quantifies how much of the  
 199 mismatch-induced degradation is recovered (see Methods, Equation (5)).

## 200 Empirical Validation

201 We validate both theoretical contributions in two stages: controlled simulation experiments  
 202 across seven modalities using the 4-Scenario Protocol, and hardware validation on real  
 203 CASSI and CACTI instruments.

### 204 Simulation experiments

205 **Correction results.** Extended Data Table 1 (Supplementary Table S1) summarizes the  
 206 correction performance across 9 configurations spanning 7 modalities and multiple carrier  
 207 families. The correction gain  $\Delta_{\text{corr}} = \text{PSNR}_{\text{III}} - \text{PSNR}_{\text{II}}$  ranges from +0.76 dB (CASSI)  
 208 to +10.68 dB (CT) across the six non-MRI modalities, with recovery ratios up to  $\rho =$   
 209 93.3%. The validated modalities span photon-domain (CASSI: +0.76/+6.50 dB; CACTI:  
 210 +10.21 dB; SPC: +10.38 dB; Lensless: +3.55 dB), coherent-photon (Ptychography: +7.09 dB),  
 211 and X-ray (CT: +10.68 dB). For MRI, the single-coil stress test yields +48.25 dB, but this  
 212 reflects a pathological scenario; under clinically realistic multi-coil conditions (8 coils, 4×  
 213 acceleration), the correction gain is +1.75 to +7.14 dB at 3–5% mismatch (Supplementary  
 214 Note S13)—clinically meaningful but commensurate with other modalities.

215 **Modality deep dives.** In CASSI, a 5-parameter mismatch<sup>17</sup> collapses all solvers to  
 216 ~21 dB regardless of ideal performance (24–35 dB); oracle recovery varies by solver ( $\rho = 0$ –  
 217 0.46), confirming that multi-parameter mismatches are harder to correct than isolated  
 218 shifts. In CACTI, EfficientSCI<sup>18</sup> drops by 20.58 dB, yet GAP-TV recovers  $\rho = 93.3\%$ —the

best ideal-condition solver suffers the largest degradation. SPC gain drift is corrected to  $\rho = 89.6\%$ . **Gate 3** is dominant in every case; Gates 1–2 impose irrecoverable information-theoretic limits only at extreme compression or photon starvation (Supplementary Tables S10–S11).

**Zero-shot generalization.** Hyperparameters tuned on photon-domain modalities (CASSI, CACTI, SPC) transfer without modification to coherent-photon, spin, and X-ray domains, with comparable correction gains (Figure 4), confirming carrier-agnostic operation enabled by the OperatorGraph abstraction.

## Hardware validation

**CASSI real data.** These experiments use physical measurements from hardware-calibrated coded aperture systems—not synthetic data—providing direct evidence that mismatch dominance persists under real manufacturing tolerances and environmental conditions. We reconstruct 5 real TSA scenes<sup>1</sup> under calibrated and perturbed masks. GAP-TV shows a mean residual ratio of  $1.8\times$  ( $0.00189 \rightarrow 0.00333$ ); HDNet shows  $1.0\times$  (mask-oblivious). MST-S/L show ratios near  $1.0\times$  on real data, in contrast to severe synthetic degradation—revealing that real hardware already contains manufacturing imperfections that absorb additional perturbations (Supplementary Table S5). These real-data results independently confirm the Triad predictions from simulation: the InverseNet benchmark<sup>17</sup> validates the same mismatch-dominance pattern on physical DD-CASSI and CACTI instruments.

**CACTI real data.** GAP-TV shows  $10.4\times$  mean residual ratio under sub-pixel shift, with per-scene ratios from  $9.4\times$  to  $11.0\times$ . The order-of-magnitude sensitivity arises because the temporal mask pattern replicates across all compressed frames, amplifying errors multiplicatively.

**Autonomous calibration.** Grid-search calibration recovers 85% (CASSI, 1,140 s), 100% (CACTI, 60 s), and 86–92% (SPC via TV objective) of the oracle correction. CACTI achieves full recovery because the mismatch manifold is low-dimensional and sensitivity is high.

**Simulation-to-hardware gap.** CASSI shows smaller real degradation ( $1.8\times$ ) than predicted, because pre-existing manufacturing errors absorb perturbations; CACTI shows larger real sensitivity ( $10.4\times$ ), because simpler optics lack this buffering.

## Discussion

The central finding is that computational imaging has a universal structure captured by two results: a finite primitive basis for forward models and a tripartite decomposition



for failures. Operator mismatch—not solver weakness, not information deficiency, not noise—is the dominant bottleneck across all validated modalities. The implication is direct: the research community should rebalance effort from solver-centric to operator-centric approaches. A single calibration step recovers more reconstruction quality than years of algorithmic innovation.

An instructive analogy is the periodic table of elements. Just as Mendeleev organized known elements by atomic number and predicted gaps, the OPERATORGRAPH organizes imaging modalities by their primitive composition. The analogy is pedagogical, not mathematical: imaging primitives do not have atomic numbers, and the DAG structure is richer than a two-dimensional table. Nevertheless, Theorem 3 implies that the space of imaging forward models, like the space of chemical elements, has a discoverable and finite structure—and for a related reason: the underlying physical interactions are finite.

Prior calibration frameworks—ESPIRiT<sup>10</sup> for MRI, ePIE<sup>11</sup> for ptychography, plug-and-play methods<sup>7</sup> at the solver level—operate within a single modality or abstraction layer. A direct comparison on multi-coil MRI (Supplementary Note S12) shows that ESPIRiT and PWM are complementary: under limited calibration data (24 ACS lines), ESPIRiT’s data-driven estimation degrades by  $-9.29$  dB while PWM’s model-based correction recovers  $+0.75$  dB; with abundant calibration data, ESPIRiT excels within MRI but does not transfer to other modalities. The OPERATORGRAPH is the first provably complete forward-model basis across carrier families, enabling a single diagnostic and correction pipeline to serve all modalities in  $\mathcal{C}_{\text{Tier2}}$ .

The coverage of the primitive basis is broader than it may initially appear. Quantum ghost imaging is operator-equivalent to a classical single-pixel camera at the image-formation level, sharing the same DAG despite fundamentally different source physics. THz time-domain imaging requires only the coherent-field Detect family. Of three exotic modalities analyzed, only Compton scatter imaging required a new primitive (Scatter), and that single primitive covers 5+ scattering and fluorescence modalities. This pattern—that exotic modalities decompose into existing primitives, and when they do not, a single new primitive covers a whole family—is the hallmark of a well-chosen basis.

The hardware validation reveals a nuanced picture. On real CASSI, degradation is smaller than predicted ( $1.8\times$  vs.  $3.38$  dB), because as-built masks absorb perturbations. On real CACTI, degradation is severe ( $10.4\times$ ), because temporal compression amplifies errors multiplicatively. This asymmetry is invisible without a unified framework. The simulation-to-hardware gap carries a methodological lesson: synthetic mismatch studies overestimate marginal impact while underestimating cumulative burden.

To demonstrate translational potential, we prototype a CT QC Copilot mapping the Triad gates to clinical failure modes. On a simulated 30-scanner fleet, it detects calibration drift with 100% sensitivity/specificity and reduces per-scanner QC time by 94% (Supplementary Note 7). Consistent with the research findings, Gate 3 dominates.

Several limitations should be noted. First, real-data experiments apply software-simulated perturbations rather than physical displacement; controlled hardware experiments are the natural next step (Supplementary Note 8). Second, Tier 1–2 models may not capture all failure modes visible at Tier 3–4. Third, the correction pipeline is limited to declared mismatch parameter families. Fourth, the CT QC validation uses a simulated fleet.

**Roadmap to physical validation.** The controlled hardware experiment protocol is fully specified (see Methods). For CASSI: physically translate the coded aperture mask by  $\Delta x \in \{0.25, 0.5, 1.0\}$  px-equivalent via micrometer stage, re-acquire under identical illumination, and compare against PWM-predicted degradation and autonomous recovery. For CACTI: apply equivalent temporal mask shifts. A multi-unit variation study comparing 2+ camera units of the same design will quantify inter-unit mismatch baselines—the residual calibration error present in any production system. For clinical MRI: process multi-coil  $k$ -space data under controlled coil repositioning to validate the +1.75 to +7.14 dB realistic correction range (Supplementary Note S13). These experiments require no methodological innovation, only instrument access and controlled acquisition; the PWM pipeline is ready to process the resulting data without modification.

Looking forward, we envision hardware-in-the-loop validation across additional instruments, real-time adaptive calibration during acquisition, prospective clinical deployment of the CT QC Copilot, and scaling the OperatorGraph library to compile a comprehensive atlas of imaging failure modes.

**Acknowledgements.** We thank the open-source computational imaging community for making reconstruction code and benchmark datasets publicly available. This work was supported by NextGen PlatformAI C Corp.

**Author Contributions.** C.Y. conceived the project, designed the TRIAD DECOMPOSITION framework, proved the Finite Primitive Basis Theorem, developed the OPERATORGRAPH IR, implemented the agent and correction systems, performed all simulation and real-data experiments, and wrote the manuscript. X.Y. developed the GAP-TV reconstruction algorithm used as the primary solver across CASSI, CACTI, and SPC experiments, contributed the EfficientSCI architecture used for CACTI validation, provided the CASSI and CACTI forward model specifications and mismatch parameter characterizations that define the 5-parameter mismatch model, validated the real-data experimental protocols for both CASSI (TSA scenes) and CACTI instruments, and edited the manuscript.

**Competing Interests.** C.Y. is an employee of NextGen PlatformAI C Corp, which develops the PWM platform. The authors declare no other competing interests.

325 **Ethics Declarations.** This study does not involve human participants, human tissue, or  
326 animals. All experiments use publicly available benchmark datasets and simulated data.

327 **Data Availability.** All synthetic measurement data can be regenerated using the OPER-  
328 ATORGRAPH templates and mismatch parameters in the Supplementary Information. The  
329 KAIST hyperspectral dataset<sup>19</sup> and TSA real-data scenes used for CASSI experiments are  
330 publicly available. CACTI real-data scenes are available from the EfficientSCI repository<sup>18</sup>.

331 **Code Availability.** The PWM codebase, including all OPERATORGRAPH templates, agent  
332 implementations, real-data validation scripts, and evaluation pipelines, is available at [https://github.com/integritynoble/Physics\\_World\\_Model](https://github.com/integritynoble/Physics_World_Model)  
333 under the PWM Noncommercial  
334 Share-Alike License v1.0 (see LICENSE in the repository).

335 **Correspondence.** Correspondence and requests for materials should be addressed to  
336 C.Y. (integrityyang@gmail.com).

## 337 Online Methods

### 338 OperatorGraph Specification

339 **Formal definition.** The OPERATORGRAPH intermediate representation encodes the for-  
340 ward physics of any computational imaging modality as a directed acyclic graph (DAG)  
341  $\mathcal{G} = (\mathcal{V}, \mathcal{E})$ . Each node  $v_i \in \mathcal{V}$  wraps a *primitive operator* and implements two entry points:  
342 `forward( $x$ )  $\rightarrow y$`  and `adjoint( $y$ )  $\rightarrow x$` , the latter defined only when the primitive is lin-  
343 ear. Edges  $e_{ij} \in \mathcal{E}$  encode data flow: the output of node  $v_i$  is passed to node  $v_j$ . Each  
344 node additionally exposes a set of learnable parameters  $\theta_i$  that may be perturbed during  
345 mismatch simulation or optimized during calibration, as well as read-only metadata flags  
346 (`is_linear`, `is_stochastic`, `is_differentiable`). The graph is stored as a declarative  
347 YAML specification (`OperatorGraphSpec`) and compiled to an executable `GraphOperator`  
348 object by the `GraphCompiler`.

349 **Node types.** Primitive operators fall into two categories:

- 350 • **Linear operators.** Convolution (`conv2d`), mask modulation (`mask_modulate`), sub-  
351 pixel shift (`subpixel_shift_2d`), Radon transform (`radon_fanbeam`), Fourier encod-  
352 ing (`fourier_encode`), spectral dispersion (`spectral_disperse`), Fresnel propaga-  
353 tion (`fresnel_propagate`), random projection (`random_project`), and structured il-  
354 lumination (`sim_modulate`). Each implements both `forward()` and `adjoint()`.
- 355 • **Nonlinear operators.** Squared magnitude (`magnitude_sq`), Poisson–Gaussian noise  
356 (`poisson_gaussian`), saturation clipping (`saturation_clip`), phase retrieval nonlin-  
357 earity (`phase_abs`), and detector quantization (`quantize`). These set `is_linear` =

358 `False` and raise `NotImplementedError` on `adjoint()`, except where a well-defined  
359 pseudo-adjoint exists (*e.g.*, the identity adjoint for magnitude-squared in Gerchberg–  
360 Saxton-type algorithms).

**Adjoint validation.** Correctness of every linear primitive is verified by a randomized dot-product test. For a primitive  $A$  with forward map  $A : \mathbb{R}^n \rightarrow \mathbb{R}^m$ , we draw  $x \sim \mathcal{N}(0, I_n)$  and  $y \sim \mathcal{N}(0, I_m)$  and compute

$$\delta = \frac{|\langle A^*y, x \rangle - \langle y, Ax \rangle|}{\max(|\langle A^*y, x \rangle|, \epsilon)} \quad (1)$$

361 where  $\epsilon = 10^{-12}$  guards against division by zero. The test is repeated  $n_{\text{trials}} = 5$  times  
362 with independent random draws; the primitive passes if  $\delta_{\text{max}} < 10^{-6}$ . At the graph level, a  
363 compiled `GraphOperator` composed entirely of linear nodes executes the same test over the  
364 composed forward–adjoint chain. A `GraphAdjointCheckReport` records  $n_{\text{trials}}$ ,  $\delta_{\text{max}}$ , and  $\bar{\delta}$   
365 for audit. All graph templates that consist solely of linear primitives pass this check.

366 **Graph compilation.** The compiler executes a four-stage pipeline:

- 367 1. **Validate.** Confirm acyclicity via topological sort (Kahn’s algorithm), verify that ev-  
368 ery `primitive_id` exists in the global `PRIMITIVE_REGISTRY`, reject duplicate `node_id`  
369 values, and optionally verify shape compatibility along edges when a `canonical_chain`  
370 metadata flag is set.
- 371 2. **Bind.** Instantiate each primitive with its parameter dictionary  $\theta_i$ .
- 372 3. **Plan forward.** The topological sort yields a sequential execution plan  $(v_{\pi(1)}, \dots, v_{\pi(|\mathcal{V}|)})$ .
- 373 4. **Plan adjoint.** For graphs where `all_linear = True`, the adjoint plan reverses the  
374 topological order and applies each node’s individual adjoint in sequence, implementing  
375 the chain rule  $A^* = A_1^* \circ \dots \circ A_{|\mathcal{V}|}^*$  for a composition  $A = A_{|\mathcal{V}|} \circ \dots \circ A_1$ . For  
376 graphs containing nonlinear nodes, the adjoint plan is not generated, and any call to  
377 `adjoint()` raises `NotImplementedError` at runtime.

378 The compiled `GraphOperator` is serializable to JSON and hashable via SHA-256 for prove-  
379 nance tracking in `RunBundle` manifests.

380 **Template library.** The `graph_templates.yaml` registry contains templates organized  
381 across 26 registered modalities (7 with full end-to-end correction validation, 1 with Scenario I  
382 baseline, 18 with template-level validation), grouped by physical carrier:

- 383 • **Photons (optical and X-ray):** CASSI, SPC, CACTI, structured illumination  
384 microscopy (SIM), confocal, light-sheet, holography, ptychography, Fourier ptycho-  
385 graphic microscopy (FPM), optical coherence tomography (OCT), lensless imaging,

light field, integral imaging, neural radiance fields (NeRF), Gaussian splatting, fluorescence lifetime imaging (FLIM), diffuse optical tomography (DOT), phase retrieval, X-ray computed tomography (CT), and cone-beam CT (CBCT).

- **Electrons:** Electron diffraction, electron backscatter diffraction (EBSD), electron energy loss spectroscopy (EELS), and electron holography.
- **Spins (MRI):** Functional MRI (fMRI), diffusion-weighted MRI (DW-MRI), and magnetic resonance spectroscopy (MRS).
- **Acoustic:** Ultrasound B-mode, Doppler ultrasound, shear-wave elastography, sonar, and photoacoustic tomography (combines optical excitation with acoustic detection).
- **Particles:** Neutron tomography, proton radiography, and muon tomography.

**Physics Fidelity Ladder.** Each template is parameterized by a fidelity tier that controls the degree of physical realism in the simulated forward model:

**Tier 1 (Linear, shift-invariant):** The forward model is a linear, spatially uniform operator—the simplest approximation, suitable for initial diagnostics and rapid prototyping.

**Tier 2 (Linear, shift-variant):** Spatially varying operator parameters (e.g. non-uniform illumination, position-dependent PSF, multi-coil sensitivity maps in MRI). Adds a modality-appropriate noise model (Poisson shot noise plus Gaussian read noise for photon-counting modalities, Rician noise for MRI, Poisson for CT).

**Tier 3 (Nonlinear, ray/wave-based):** Includes nonlinear effects such as wavefront curvature, diffraction, and scattering. Perturbation families and ranges are specified in `mismatch_db.yaml`.

**Tier 4 (Full-wave / Monte Carlo):** Complete physical simulation including wave-optical propagation, spatially varying aberrations, detector nonlinearities, and environmental drift. Currently implemented for holography and ptychography; other modalities degrade gracefully to Tier 3.

## Triad Decomposition Formalization

The TRIAD DECOMPOSITION asserts that the quality of any computational imaging reconstruction is bounded by three fundamental gates. Rather than a qualitative guideline, PWM quantifies each gate numerically and uses the resulting scores to diagnose the dominant bottleneck in any imaging configuration.

416 **Gate 1 (Recoverability).** Recoverability measures the information-theoretic capacity  
 417 of the sensing geometry. We quantify it via the *effective compression ratio*  $r = m/n$ , where  
 418  $m$  is the number of independent measurements and  $n$  the dimension of the signal. The  
 419 `compression_db.yaml` registry (1,186 lines) stores, for each modality, a lookup table map-  
 420 ping compression ratio to expected reconstruction PSNR under ideal conditions, obtained  
 421 from calibration experiments or published benchmarks. Each entry carries a **provenance**  
 422 field citing the source (paper DOI, internal experiment ID, or theoretical formula). Addi-  
 423 tional recoverability indicators include the effective rank of the measurement matrix (esti-  
 424 mated via randomized SVD for large operators), the dimension of the null space, and the  
 425 restricted isometry property (RIP) constant where analytically tractable (*e.g.*, for Gaussian  
 426 random projections in SPC).

427 **Gate 2 (Carrier Budget).** The carrier budget quantifies the signal-to-noise ratio (SNR)  
 428 of the measurement channel. The **PhotonAgent** consumes the `photon_db.yaml` registry  
 429 (624 lines) which stores, per modality, a deterministic photon model parameterized by  
 430 source power, quantum efficiency, exposure time, and detector characteristics. The agent  
 431 classifies the noise regime into one of three categories: *shot-limited* (Poisson-dominated,  
 432  $\text{SNR} \propto \sqrt{N_{\text{photon}}}$ ), *read-limited* (Gaussian read noise dominates,  $\text{SNR} \propto N_{\text{photon}}/\sigma_{\text{read}}$ ),  
 433 and *dark-current-limited* (long exposures where dark current accumulation dominates). The  
 434 output is a **PhotonReport** containing the estimated SNR in decibels, the noise regime  
 435 classification, per-element photon count, and a feasibility verdict (**sufficient**, **marginal**,  
 436 or **insufficient**).

437 **Gate 3 (Operator Mismatch).** Operator mismatch quantifies the discrepancy between  
 438 the assumed forward model  $H_{\text{nom}}$  and the true physical operator  $H_{\text{true}}$ . The **MismatchAgent**  
 439 consults `mismatch_db.yaml` (797 lines) which catalogs, for each modality, the set of mis-  
 440 match parameters (spatial shifts, rotational offsets, dispersion errors, PSF deviations, coil  
 441 sensitivity errors, center-of-rotation offsets, *etc.*), their typical ranges, and available cor-  
 442 rection methods. The mismatch severity score  $s \in [0, 1]$  is computed as the normalized  $\ell_2$   
 443 distance  $\|\theta_{\text{true}} - \theta_{\text{nom}}\|/\|\theta_{\text{range}}\|$ , where  $\theta_{\text{range}}$  is the per-parameter dynamic range from the  
 444 registry. Sensitivity analysis  $\partial\text{PSNR}/\partial\theta_k$  is estimated via finite differences on the forward  
 445 model. The output is a **MismatchReport** containing the severity score, the dominant mis-  
 446 match parameter, the recommended correction method, and the expected PSNR gain from  
 447 correction.

**Gate binding determination.** Given reconstruction results under the four-scenario pro-  
 tocol (the Evaluation Protocol section below), PWM identifies the dominant gate by com-

paring three cost terms:

$$C_{\text{mismatch}} = \text{PSNR}_{\text{I}} - \text{PSNR}_{\text{II}} \quad (2)$$

$$C_{\text{noise}} = \text{PSNR}_{\text{ideal}} - \text{PSNR}_{\text{noisy}} \quad (3)$$

$$C_{\text{recover}} = \text{PSNR}_{\text{limit}} - \text{PSNR}_{\text{I}} \quad (4)$$

where  $\text{PSNR}_{\text{I}}$  is the reconstruction PSNR under Scenario I (ideal operator),  $\text{PSNR}_{\text{II}}$  under Scenario II (mismatched operator),  $\text{PSNR}_{\text{noisy}}$  under the corresponding noisy condition, and  $\text{PSNR}_{\text{limit}}$  is the theoretical upper bound from the compression table. The dominant gate is  $\arg \max_g C_g$ .

**TriadReport.** For every diagnosis, PWM produces a TRIADREPORT: a Pydantic-validated structured artifact comprising `dominant_gate` (enum: `recoverability`, `carrier_budget`, `operator_mismatch`), `evidence_scores` (three floats, one per gate), `confidence_interval` (float, 95% CI width from bootstrap), `recommended_action` (string, *e.g.* “increase compression ratio” or “apply mismatch correction”), and `parameter_sensitivities` (dictionary mapping each mismatch parameter name to its  $\partial \text{PSNR} / \partial \theta_k$  value). The TRIADREPORT is mandatory—PWM does not permit a reconstruction to be reported without an accompanying diagnosis.

**Recovery ratio.** We define the *recovery ratio*

$$\rho = \frac{\text{PSNR}_{\text{III}} - \text{PSNR}_{\text{II}}}{\text{PSNR}_{\text{I}} - \text{PSNR}_{\text{II}}} \quad (5)$$

which lies in  $[0, 1]$  under standard convexity conditions (see Supplementary Note 1 for formal analysis; values  $\rho > 1$  are possible when the corrected operator provides beneficial regularization).  $\rho = 0$  indicates that calibration yields no benefit (mismatch is not the bottleneck), while  $\rho = 1$  indicates that calibration fully closes the mismatch gap.

## Agent System Architecture

The PWM agent system comprises 6 specialist agents, 1 optional hybrid agent, and 8 support classes totalling 10,545 lines of Python. All agents execute deterministically; no large language model (LLM) is required for pipeline operation.

**PlanAgent.** The orchestrator agent. Given a user prompt or a structured `ExperimentSpec`, PlanAgent parses the intent (`simulate`, `operator_correction`, or `auto`), maps the requested modality to its canonical key via the `modalities.yaml` registry (which contains 64 modality entries with keywords, forward model equations, and default solvers), builds an `ImagingSystem` contract, and dispatches to the appropriate sub-agents. When the mode is

473 **auto**, PlanAgent inspects the available data and operator specification to determine whether  
474 simulation or operator correction is more appropriate.

475 **PhotonAgent.** Computes SNR feasibility deterministically from the `photon_db.yaml`  
476 registry. For each modality and photon-level tier (`bright`, `standard`, `low_light`), the  
477 agent evaluates the photon budget by combining source power, quantum efficiency, ex-  
478 posure time, and noise model parameters. The output `PhotonReport` is a strict Pydan-  
479 tic model containing `noise_regime` (enum), `snr_db` (float), `feasibility` (enum), and  
480 `per_element_photons` (float).

481 **RecoverabilityAgent.** A table-driven agent that consults `compression_db.yaml` (1,186  
482 lines) to map the modality and compression ratio to an expected PSNR range. Each table  
483 entry includes provenance metadata citing the original source. The output `RecoverabilityReport`  
484 contains `compression_ratio`, `psnr_prediction`, `feasibility`, and `null_space_dim` where  
485 available.

486 **MismatchAgent.** Scores the mismatch severity for a given imaging configuration us-  
487 ing `mismatch_db.yaml` (797 lines). For each modality, the database enumerates the rel-  
488 evant mismatch parameters, their physical units, typical perturbation ranges, and avail-  
489 able correction algorithms. The output `MismatchReport` includes `severity` (float, 0–1),  
490 `correction_method` (string), `expected_gain_db` (float), and `dominant_parameter` (string).

491 **AnalysisAgent.** The bottleneck classifier. It receives reports from the Photon, Recover-  
492 ability, and Mismatch agents, computes the gate costs (Equations (2) to (4)), identifies the  
493 dominant gate, and generates actionable suggestions. The AnalysisAgent also computes  
494 the recovery ratio  $\rho$  and its bootstrap confidence interval.

495 **AgentNegotiator.** Implements a cross-agent veto protocol. Before reconstruction is au-  
496 thorized, the negotiator inspects all three upstream reports and applies three veto con-  
497 ditions: (1) low photon budget combined with aggressive compression ( $C_{\text{noise}}$  and  $C_{\text{recover}}$   
498 both large); (2) severe mismatch ( $\text{severity} > 0.7$ ) without a planned correction step; (3) joint  
499 probability below the floor threshold ( $p_{\text{joint}} < 0.15$ ), indicating that all three subsystems  
500 are simultaneously marginal. When any veto fires, reconstruction halts with an actionable  
501 explanation and suggested remediation.

502 **HybridAgent.** An optional wrapper that invokes an LLM for natural-language narra-  
503 tive generation or edge-case modality mapping. All quantitative decisions remain on the  
504 deterministic code path; the HybridAgent is never required for pipeline operation.



**Support classes.** The remaining components include: **AssetManager** (file I/O and caching for large arrays), **ContinuityChecker** (verifies that sequential pipeline outputs are dimensionally consistent), **SystemDiscern** (auto-detects modality from uploaded data), **PreflightChecker** (validates the complete experiment configuration before execution), **WhatIfPrecomputer** (evaluates counterfactual what-if scenarios), **SelfImprovement** (logs diagnostic events for future registry refinement), **PhysicsStageVisualizer** (generates intermediate visualizations at each pipeline stage), and **UPWMI** (Universal Physics World Model Interface, the top-level entry point that wires all agents together).

**Contract system.** Inter-agent communication uses 25 Pydantic v2 contract models. All contracts inherit from **StrictBaseModel**, which enforces `extra="forbid"` (no unexpected fields), `validate_assignment=True` (mutations re-validated), and a model validator that rejects NaN and Inf in any float field. Bounded scores use `Field(ge=0.0, le=1.0)`. Enums are string enums for human-readable JSON serialization. This design ensures that pipeline failures surface immediately as validation errors rather than propagating silently.

**YAML registries.** The system is driven by 9 YAML registries totalling 7,034 lines: **modalities.yaml** (modality definitions), **graph\_templates.yaml** (OperatorGraph skeletons), **photon\_db.yaml** (photon models), **mismatch\_db.yaml** (mismatch parameters and correction methods), **compression\_db.yaml** (recoverability tables with provenance), **solver\_registry.yaml** (solver configurations), **primitives.yaml** (primitive operator metadata), **dataset\_registry.yaml** (dataset locations and formats), and **acceptance\_thresholds.yaml** (pass/fail thresholds per metric).

## Correction Algorithms

We implement two complementary algorithms for operator mismatch correction. Crucially, both algorithms operate on the forward operator parameters  $\theta$  rather than the reconstruction solver weights, making them *solver-agnostic*: the corrected operator  $H(\hat{\theta})$  benefits any downstream solver (GAP-TV, MST-L, HDNet<sup>20</sup>, CST, *etc.*) without retraining.

**Algorithm 1: Hierarchical Beam Search.** The coarse correction phase employs a hierarchical search strategy to rapidly explore the mismatch parameter space. For CASSI, the five-parameter mismatch model comprises mask affine parameters (spatial shifts  $dx$ ,  $dy$  and rotation  $\theta$ ) and dispersion parameters (slope  $a_1$  and axis angle  $\alpha$ ); an optional sixth parameter, PSF width  $\sigma_{\text{psf}}$ , is available but not used in the primary experiments. The algorithm proceeds as follows:

1. **1D sweeps.** Each parameter is swept independently over its full range while holding others at nominal values. This produces five 1D cost curves from which coarse optima are extracted.

- 540 2. **3D beam search.** The mask affine subspace  $(dx, dy, \theta)$  is searched over a  $5 \times 5 \times 5$   
541 grid centered on the 1D optima. The top- $k$  ( $k = 5$ ) candidates by reconstruction  
542 PSNR are retained.
- 543 3. **2D beam search.** For each retained mask candidate, the dispersion subspace  $(a_1, \alpha)$   
544 is searched over a  $5 \times 7$  grid. The joint top- $k$  candidates are retained.
- 545 4. **Coordinate descent refinement.** Three rounds of univariate refinement on each  
546 parameter, shrinking the search interval by factor 2 at each round, produce the final  
547 estimate  $\hat{\theta}_{\text{Alg1}}$ .

548 Total runtime is approximately 300 seconds per scene on a single GPU. Accuracy is  
549  $\pm 0.1$ – $0.2$  pixels for spatial parameters and  $\pm 0.05^\circ$  for angular parameters.

550 **Algorithm 2: Joint Gradient Refinement.** The fine correction phase uses a differ-  
551 entiable forward model to jointly optimize all mismatch parameters via gradient descent.  
552 The key components are:

- 553 1. **Differentiable mask warp.** The binary mask is warped by a continuous affine  
554 transformation using bilinear interpolation, implemented as a custom PyTorch module  
555 (`DifferentiableMaskWarpFixed`). The mask values are passed through a straight-  
556 through estimator (STE) to maintain binary structure while permitting gradient flow.
- 557 2. **Differentiable forward model.** The CASSI forward model  $y = \text{CASSI}(x; \theta)$  is  
558 implemented as a differentiable PyTorch module (`DifferentiableCassiForwardSTE`)  
559 that accepts mismatch parameters as differentiable inputs.
- 560 3. **GPU grid initialization.** A full-range 3D grid search over  $(dx, dy, \theta)$  with  $9 \times 9 \times 7 =$   
561 567 points provides diverse starting candidates. The top 9 candidates seed multi-start  
562 gradient refinement.
- 563 4. **Staged gradient refinement.** Each of the 9 candidates is refined using Adam  
564 optimization (learning rate  $10^{-2}$ , decaying to  $10^{-3}$ ) for 200 steps. For each candidate,  
565 4 random restarts with jittered initialization guard against local minima. The loss  
566 function is the negative PSNR computed via an unrolled  $K$ -iteration differentiable  
567 GAP-TV solver (`DifferentiableGAPTV`,  $K = 10$  unrolled iterations).

568 Total runtime for Algorithm 2 is approximately 3,200 seconds (200 steps  $\times$  4 restarts  $\times$   
569 9 candidates with early stopping). Accuracy improves to  $\pm 0.05$ – $0.1$  pixels, a 3–5 $\times$  improve-  
570 ment over Algorithm 1. The two algorithms are used sequentially in practice: Algorithm 1  
571 provides a warm start, and Algorithm 2 refines to sub-pixel precision.

## Evaluation Protocol

**Four-Scenario Protocol.** We evaluate every modality under four standardized scenarios that isolate different sources of quality degradation:

**Scenario I (Ideal):**  $\mathbf{y}_{\text{obs}} = H_{\text{true}} \mathbf{x}_{\text{gt}}$ ; reconstruct with  $H_{\text{true}}$ . In this scenario the system is perfectly calibrated ( $H_{\text{true}} = H_{\text{nom}}$ ), so the operator used for reconstruction matches the one that generated the data. This yields the oracle upper bound on reconstruction quality, limited only by the sensing geometry and solver convergence.

**Scenario II (Mismatch):**  $\mathbf{y}_{\text{obs}} = H_{\text{true}} \mathbf{x}_{\text{gt}}$ ; reconstruct with  $H_{\text{nom}}$  ( $H_{\text{nom}} \neq H_{\text{true}}$ ). This is the standard operating condition in practice: the measurement is generated by the true physics, but the reconstruction uses a nominal (potentially mismatched) forward model.

**Scenario III (Corrected):**  $\mathbf{y}_{\text{obs}} = H_{\text{true}} \mathbf{x}_{\text{gt}}$ ; reconstruct with  $\hat{H} = H(\hat{\theta})$  where  $\hat{\theta}$  is estimated by Algorithms 1 and 2. This quantifies the benefit of mismatch calibration.

**Scenario IV (Oracle Mask):** Same measurements as Scenario II ( $\mathbf{y}_{\text{obs}} = H_{\text{true}} \mathbf{x}_{\text{gt}}$  with  $H_{\text{true}} \neq H_{\text{nom}}$ ); reconstruct with  $H_{\text{true}}$  instead of  $H_{\text{nom}}$ . Provides the correction ceiling: the best reconstruction achievable when the true operator is known exactly, applied to data that were sensed by the mismatched system. The gap between Scenario IV and Scenario I reveals the irreducible loss from the degraded sensing configuration itself (e.g., a shifted mask pattern is suboptimal even when perfectly known).

**Metrics.** Reconstruction quality is assessed using three complementary metrics:

- **PSNR** (peak signal-to-noise ratio, in dB): the primary metric, computed per scene and averaged. For signals normalized to  $[0, 1]$ ,  $\text{PSNR} = 10 \log_{10}(1/\text{MSE})$ . For SPC data normalized to  $[0, 255]$ , the peak value is 255.
- **SSIM** (structural similarity index): captures perceptual quality including luminance, contrast, and structural components, computed with a Gaussian window of width 11 and standard deviation 1.5.
- **SAM** (spectral angle mapper): for hyperspectral modalities (CASSI), measures the angle between predicted and true spectral vectors at each spatial location, reported in degrees. Lower is better.

## Datasets.

- **CASSI:** 10 scenes from the KAIST dataset<sup>19</sup>, each a  $256 \times 256 \times 28$  spectral cube (28 spectral bands from 450 nm to 650 nm). Data range  $[0, 1]$ .

604 • **CACTI**: 6 benchmark videos, each  $256 \times 256 \times 8$  (8 temporal frames encoded per  
605 snapshot). Data range  $[0, 1]$ .

606 • **SPC**: 11 natural images from the Set11 benchmark, each  $256 \times 256$  grayscale. Data  
607 range  $[0, 255]$ .

608 All per-scene metrics are reported individually as well as averaged, and all reconstruction  
609 arrays are saved as NumPy NPZ files.

## 610 Experimental Details

611 **Hardware.** All experiments are conducted on a single NVIDIA GPU. Algorithm 1 (beam  
612 search) and all solver-based reconstructions use the GPU for matrix-vector products and  
613 FFT operations. Algorithm 2 (gradient refinement) additionally uses PyTorch automatic  
614 differentiation on the same GPU.

615 **CASSI configuration.** The coded aperture snapshot spectral imaging (CASSI) system  
616 uses a TSA-Net binary mask of dimensions  $256 \times 256$ , with 28 spectral bands dispersed along  
617 the spatial dimension. The five-parameter mismatch model  $\theta = (dx, dy, \theta, a_1, \alpha)$  describes:  
618 mask spatial shift in  $x$  ( $dx$ , pixels), mask spatial shift in  $y$  ( $dy$ , pixels), mask rotation angle  
619 ( $\theta$ , degrees), dispersion slope ( $a_1$ , pixels per band), and dispersion axis angle ( $\alpha$ , degrees).  
620 An optional sixth parameter, PSF blur width ( $\sigma_{\text{psf}}$ , pixels), is available but not used in the  
621 primary experiments. These mismatch parameter values were determined through system-  
622 atic characterization of realistic CASSI assembly errors (Supplementary Note 9). The true  
623 mismatch parameters are  $\theta_{\text{true}} = (dx = 0.5 \text{ px}, dy = 0.3 \text{ px}, \theta = 0.1^\circ, a_1 = 2.02, \alpha =$   
624  $0.15^\circ)$ . Solvers evaluated include TwIST<sup>21</sup>, GAP-TV<sup>22</sup>, DGSMP<sup>23</sup>, MST-L<sup>9</sup>, and CST-L<sup>24</sup>,  
625 all of which receive the same operator and differ only in their reconstruction algorithm. The  
626 supplementary per-scene analysis additionally includes DeSCI<sup>25</sup> and HDNet<sup>20</sup>.

627 **CACTI configuration.** The coded aperture compressive temporal imaging system uses  
628 binary temporal masks of dimensions  $256 \times 256$ , encoding 8 video frames into a single  
629 snapshot measurement. Mismatch is parameterized as a temporal mask timing offset (sub-  
630 frame shift). The default solver is EfficientSCI<sup>18</sup>.

631 **SPC configuration.** The single-pixel camera uses random binary measurement patterns  
632 at three compression ratios: 10%, 25%, and 50% ( $r = m/n \in \{0.10, 0.25, 0.50\}$ ). Mismatch  
633 is modeled as an exponential gain drift ( $g_i = \exp(-\alpha \cdot i)$ ) on the measurement matrix. The  
634 default solver is FISTA-TV with total-variation regularization.

635 **MRI configuration.** Cartesian  $k$ -space sampling with  $4\times$  acceleration (25% of  $k$ -space  
636 lines acquired). Mismatch is parameterized as a 5% multiplicative error in the coil sensitivity

maps used for parallel imaging reconstruction. The default solver is SENSE<sup>14</sup> with  $\ell_1$ -wavelet regularization.

**ESPIRiT comparison protocol.** To quantitatively compare PWM with ESPIRiT<sup>10</sup>, we evaluate four conditions on the same multi-coil MRI configuration: (1) Scenario I (true maps), (2) Scenario II (5% mismatched maps), (3) ESPIRiT auto-calibrated maps estimated from the 24-line ACS region via eigenvalue decomposition of the calibration matrix, and (4) PWM beam-search corrected maps (grid search over per-coil amplitude and phase scaling). All four conditions use the same CG-SENSE solver ( $\ell = 10^{-3}$ , 30 iterations). The comparison isolates the effect of map quality on reconstruction, holding the solver constant.

**CT configuration.** Fan-beam geometry with 180 projections over  $180^\circ$ . Mismatch is modeled as a center-of-rotation (CoR) offset, which produces characteristic arc artifacts in the reconstruction. The default solver is filtered back-projection (FBP)<sup>15</sup> with a Ram-Lak filter, supplemented by iterative SART for comparison.

**CASSI real-data configuration.** The TSA real hyperspectral dataset<sup>1</sup> consists of 5 scenes at  $660 \times 660$  spatial resolution with 28 spectral bands and mask-shift step 2. Four solvers are evaluated: GAP-TV (200 iterations), HDNet (pre-trained checkpoint, full spatial resolution), MST-S and MST-L (pre-trained checkpoints, centre-cropped to  $256 \times 256$  due to hardcoded spatial assumptions in the model architecture). The coded aperture mask is perturbed by  $dx = 0.5$  px,  $dy = 0.3$  px to simulate assembly-induced mismatch. No ground truth is available; quality is assessed via the normalised measurement residual  $r = \|\mathbf{y} - H\hat{\mathbf{x}}\|^2 / \|\mathbf{y}\|^2$ .

**CACTI real-data configuration.** The EfficientSCI real temporal dataset<sup>18</sup> consists of 4 dynamic scenes (duomino, hand, pendulumBall, waterBalloon) at  $512 \times 512$  with compression ratio 10. The real mask is stored separately from the measurement data. Two solvers are evaluated: GAP-TV (50 iterations) and PnP-FFDNet (50 iterations with FFDNet denoiser). Mismatch is induced by shifting the mask by  $dx = 0.5$  px,  $dy = 0.3$  px. Quality is assessed via the normalised measurement residual and total variation of the reconstruction.

**Controlled hardware experiment protocol.** The software-perturbation protocol above applies calibrated mask shifts to existing real measurements. A full hardware-in-the-loop validation requires physically displacing the coded aperture mask and re-acquiring data. The protocol proceeds as follows: (i) acquire a baseline dataset with the mask at its factory-calibrated position; (ii) physically translate the mask by a known displacement ( $\Delta x \in \{0.25, 0.5, 1.0\}$  px equivalent, verified by micrometer stage) and re-acquire under identical illumination; (iii) reconstruct both datasets with the factory mask specification and compute the PSNR degradation and measurement residual; (iv) apply PWM autonomous

calibration and measure recovery. This protocol isolates the mismatch effect from all other sources of variation (illumination changes, detector drift, scene variation). Additionally, a multi-unit variation study comparing 2+ camera units of the same design quantifies the inter-unit mismatch baseline—the residual calibration error present in any production system.

**Clinical CT phantom configuration.** For clinical translation, PWM is evaluated on CT quality assurance using the ACR CT accreditation phantom (Gammex 464)<sup>26</sup>. The phantom contains inserts of known attenuation (bone  $\sim 955$  HU, air  $\sim -1000$  HU, acrylic  $\sim 121$  HU, polyethylene  $\sim -96$  HU) and geometric targets for measuring spatial resolution, slice thickness, and low-contrast detectability. Mismatch is parameterized as center-of-rotation offset ( $\Delta r$ , mm), beam hardening coefficient drift ( $\Delta\mu$ , %), and detector gain variation ( $\Delta g$ , %). Nine ACR-aligned metrics are computed automatically: CT number accuracy (5 materials), geometric accuracy ( $\pm 2$  mm tolerance), slice thickness ( $\pm 1.5$  mm), uniformity ( $\leq 5$  HU), noise standard deviation, spatial resolution ( $\geq 5$  lp/cm), low-contrast detectability ( $\geq 4$  targets), and artifact score.

**Clinical MRI validation configuration.** For MRI clinical validation, PWM processes multi-coil  $k$ -space data from public datasets (fastMRI<sup>27</sup>). Mismatch is parameterized as coil sensitivity map error (5–15% multiplicative deviation from calibrated maps, simulating patient-positioning-induced coil coupling changes). The default solver is CG-SENSE with  $\ell_1$ -wavelet regularization at  $4\times$  acceleration. Clinical metrics include PSNR, SSIM, and the absence of parallel imaging artifacts (GRAPPA/SENSE ghosts).

## Statistical Analysis

**Per-scene reporting.** All metrics are reported per scene, not merely as dataset averages. This enables identification of scene-dependent failure modes (*e.g.*, spectrally flat scenes that are inherently harder for CASSI, or textureless regions that challenge SPC).

**Summary statistics.** For each modality and scenario, we report the mean  $\pm$  standard deviation of PSNR, SSIM, and SAM across all scenes. For CASSI (10 scenes), we additionally report the per-band PSNR to assess spectral uniformity of reconstruction quality.

**Recovery ratio confidence intervals.** The recovery ratio  $\rho$  (Equation (5)) is a ratio of differences and therefore sensitive to noise in the constituent PSNR values. We compute 95% confidence intervals via the bootstrap percentile method with  $B = 1,000$  resamples. At each bootstrap iteration, we resample the scene set with replacement, recompute the mean PSNR for each scenario, and derive  $\rho$ . The 2.5th and 97.5th percentiles of the bootstrap distribution define the 95% CI.

**Parameter recovery accuracy.** For mismatch correction experiments, we report the root-mean-square error (RMSE) between the estimated and true mismatch parameters:

$$\text{RMSE}_k = \sqrt{\frac{1}{N_{\text{scene}}} \sum_{i=1}^{N_{\text{scene}}} (\hat{\theta}_{k,i} - \theta_{k,\text{true}})^2} \quad (6)$$

where  $k$  indexes the mismatch parameter,  $i$  indexes the scene, and  $N_{\text{scene}}$  is the number of test scenes. Uncertainty in the RMSE is estimated via bootstrap ( $B = 1,000$ ).

**Ablation significance.** Ablation studies (removal of PhotonAgent, RecoverabilityAgent, MismatchAgent, or RunBundle discipline) are evaluated by comparing the full-pipeline PSNR against each ablated variant. We report the PSNR difference  $\Delta\text{PSNR}$  per modality and verify that each component contributes  $\geq 0.5$  dB across all validated modalities, establishing practical significance.

## Code and Data Availability

**Source code.** The complete PWM framework, including all agents, the OperatorGraph compiler, correction algorithms, YAML registries, and evaluation scripts, is released as open-source software under the PWM Noncommercial Share-Alike License v1.0 at [https://github.com/integritynoble/Physics\\_World\\_Model](https://github.com/integritynoble/Physics_World_Model). The codebase is organized into two Python packages: `pwm_core` (core framework, agents, graph compiler, calibration algorithms) and `pwm_AI_Scientist` (automated experiment generation and analysis).

**Reconstruction data.** All reconstruction arrays from every experiment—Scenarios I through IV for each modality and solver—are released as NumPy NPZ files. Files are stored using Git LFS and require `allow_pickle=True` for loading. Data ranges are standardized: CASSI and CACTI reconstructions are normalized to  $[0, 1]$ ; SPC reconstructions are in  $[0, 255]$ .

**Experiment manifests.** Every experiment is recorded in a RunBundle v0.3.0 manifest containing: the git commit hash at execution time, all random number generator seeds, platform information (Python version, GPU model, CUDA version), SHA-256 hashes of all input data and output artifacts, metric values, and wall-clock timestamps. These manifests enable exact reproduction of every reported result.

**Registry data.** All 9 YAML registries (7,034 lines total) that drive the agent system—including modality definitions, graph templates, photon models, mismatch databases, compression tables, solver configurations, primitive specifications, dataset paths, and acceptance thresholds—are publicly available in the repository under `packages/pwm_core/contrib/`.

734 The `ExperimentSpec` JSON schemas used for pipeline input validation are included along-  
735 side worked examples in `examples/`.

## 736 References

- 737 [1] Ashwin A. Wagadarikar, Renu John, Rebecca Willett, and David J. Brady. Single  
738 disperser design for coded aperture snapshot spectral imaging. *Applied Optics*, 47(10):  
739 B44–B51, 2008. doi: 10.1364/AO.47.000B44.
- 740 [2] Patrick Llull, Xuejun Liao, Xin Yuan, Jianbo Yang, David Kittle, Lawrence Carin,  
741 Guillermo Sapiro, and David J. Brady. Coded aperture compressive temporal imaging.  
742 *Optics Express*, 21(9):10526–10545, 2013. doi: 10.1364/OE.21.010526.
- 743 [3] Marco F. Duarte, Mark A. Davenport, Dharmpal Takhar, Jason N. Laska, Ting  
744 Sun, Kevin F. Kelly, and Richard G. Baraniuk. Single-pixel imaging via com-  
745 pressive sampling. *IEEE Signal Processing Magazine*, 25(2):83–91, 2008. doi:  
746 10.1109/MSP.2007.914730.
- 747 [4] Michael Lustig, David Donoho, and John M. Pauly. Sparse MRI: The application of  
748 compressed sensing for rapid MR imaging. *Magnetic Resonance in Medicine*, 58(6):  
749 1182–1195, 2007. doi: 10.1002/mrm.21391.
- 750 [5] Emmanuel J. Candès and Michael B. Wakin. An introduction to compressive sampling.  
751 *IEEE Signal Processing Magazine*, 25(2):21–30, 2008. doi: 10.1109/MSP.2007.914731.
- 752 [6] David L. Donoho. Compressed sensing. *IEEE Transactions on Information Theory*,  
753 52(4):1289–1306, 2006. doi: 10.1109/TIT.2006.871582.
- 754 [7] Singanallur V. Venkatakrishnan, Charles A. Bouman, and Brendt Wohlberg. Plug-  
755 and-play priors for model based reconstruction. In *Proceedings of the IEEE Global  
756 Conference on Signal and Information Processing (GlobalSIP)*, pages 945–948, 2013.  
757 doi: 10.1109/GlobalSIP.2013.6737048.
- 758 [8] Vishal Monga, Yuelong Li, and Yonina C. Eldar. Algorithm unrolling: Interpretable,  
759 efficient deep learning for signal and image processing. *IEEE Signal Processing Maga-  
760 zine*, 38(2):18–44, 2021. doi: 10.1109/MSP.2020.3016905.
- 761 [9] Yuanhao Cai, Jing Lin, Xiaowan Hu, Haoqian Wang, Xin Yuan, Yulun Zhang, Radu  
762 Timofte, and Luc Van Gool. Mask-guided spectral-wise transformer for efficient hy-  
763 perspectral image reconstruction. In *Proceedings of the IEEE/CVF Conference on  
764 Computer Vision and Pattern Recognition (CVPR)*, pages 17502–17511, 2022.
- 765 [10] Martin Uecker, Peng Lai, Mark J. Murphy, Patrick Virtue, Michael Elad, John M.  
766 Pauly, Shreyas S. Vasanawala, and Michael Lustig. ESPIRiT — an eigenvalue approach



- 767 to autocalibrating parallel MRI: where SENSE meets GRAPPA. *Magnetic Resonance*  
768 *in Medicine*, 71(3):990–1001, 2014. doi: 10.1002/mrm.24751.
- 769 [11] Andrew M. Maiden and John M. Rodenburg. An improved ptychographical phase  
770 retrieval algorithm for diffractive imaging. *Ultramicroscopy*, 109(10):1256–1262, 2009.  
771 doi: 10.1016/j.ultramic.2009.05.012.
- 772 [12] Vegard Antun, Francesco Renna, Clarice Poon, Ben Adcock, and Anders C. Hansen.  
773 On instabilities of deep learning in image reconstruction and the potential costs of  
774 AI. *Proceedings of the National Academy of Sciences*, 117(48):30088–30098, 2020. doi:  
775 10.1073/pnas.1907377117.
- 776 [13] J. M. Rodenburg and H. M. L. Faulkner. A phase retrieval algorithm for shifting  
777 illumination. *Applied Physics Letters*, 85(20):4795–4797, 2004. doi: 10.1063/1.1823034.
- 778 [14] Klaas P. Pruessmann, Markus Weiger, Markus B. Scheidegger, and Peter Boesiger.  
779 SENSE: Sensitivity encoding for fast MRI. *Magnetic Resonance in Medicine*, 42(5):  
780 952–962, 1999. doi: 10.1002/(SICI)1522-2594(199911)42:5<952::AID-MRM16>3.0.  
781 CO;2-S.
- 782 [15] L. A. Feldkamp, L. C. Davis, and J. W. Kress. Practical cone-beam algorithm. *Journal*  
783 *of the Optical Society of America A*, 1(6):612–619, 1984. doi: 10.1364/JOSAA.1.000612.
- 784 [16] Chengshuai Yang. The finite primitive basis theorem for computational imaging: For-  
785 mal foundations of the OperatorGraph representation. *arXiv preprint*, 2026. In prepa-  
786 ration. Full formal definitions, primitive semantics, typed DAG denotation, and con-  
787 structive proof.
- 788 [17] Chengshuai Yang. InverseNet: Benchmarking operator mismatch in snapshot compres-  
789 sive imaging. In *Under review at ECCV 2026*, 2026.
- 790 [18] Lishun Wang, Miao Cao, and Xin Yuan. EfficientSCI: Densely connected network  
791 with space-time factorization for large-scale video snapshot compressive imaging. In  
792 *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*  
793 *(CVPR)*, pages 18477–18486, 2023.
- 794 [19] Inchang Choi, Daniel S. Jeon, Giljoo Nam, Diego Gutierrez, and Min H. Kim. High-  
795 quality hyperspectral reconstruction using a spectral prior. *ACM Transactions on*  
796 *Graphics (Proceedings of SIGGRAPH Asia)*, 36(6):218:1–218:13, 2017. doi: 10.1145/  
797 3130800.3130810.
- 798 [20] Xiaowan Hu, Yuanhao Cai, Jing Lin, Haoqian Wang, Xin Yuan, Yulun Zhang, Radu  
799 Timofte, and Luc Van Gool. HDNet: High-resolution dual-domain learning for spectral  
800 compressive imaging. In *Proceedings of the IEEE/CVF Conference on Computer Vision*  
801 *and Pattern Recognition (CVPR)*, pages 17542–17551, 2022.

- [21] José M. Bioucas-Dias and Mário A. T. Figueiredo. A new TwIST: Two-step iterative shrinkage/thresholding algorithms for image restoration. *IEEE Transactions on Image Processing*, 16(12):2992–3004, 2007. doi: 10.1109/TIP.2007.909319.
- [22] Xin Yuan. Generalized alternating projection based total variation minimization for compressive sensing. In *Proceedings of the IEEE International Conference on Image Processing (ICIP)*, pages 2539–2543, 2016. doi: 10.1109/ICIP.2016.7532817.
- [23] Tao Huang, Weisheng Dong, Xin Yuan, Jinjian Wu, and Guangming Shi. Deep gaussian scale mixture prior for spectral compressive imaging. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 16216–16225, 2021.
- [24] Yuanhao Cai, Jing Lin, Xiaowan Hu, Haoqian Wang, Xin Yuan, Yulun Zhang, Radu Timofte, and Luc Van Gool. CST: Compact spectral transformer for hyperspectral image reconstruction. In *Proceedings of the European Conference on Computer Vision (ECCV)*, 2022.
- [25] Yang Liu, Xin Yuan, Jinli Suo, David J. Brady, and Qionghai Dai. Rank minimization for snapshot compressive imaging. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 41(12):2990–3006, 2019.
- [26] Hu Chen, Yi Zhang, Mannudeep K. Kalra, Feng Lin, Yang Chen, Peixi Liao, Jiliu Zhou, and Ge Wang. Low-dose CT with a residual encoder-decoder convolutional neural network. *IEEE Transactions on Medical Imaging*, 36(12):2524–2535, 2017. doi: 10.1109/TMI.2017.2715284.
- [27] Jure Zbontar, Florian Knoll, Anuroop Sriram, Tullie Murrell, Zhengnan Huang, Matthew J. Muckley, Aaron Defazio, Ruben Stern, Patricia Johnson, Mary Bruno, Marc Parente, Krzysztof J. Geras, Joe Katsnelson, Hersh Chandarana, Zizhao Zhang, Michal Drozdal, Adriana Romero, Michael Rabbat, Pascal Vincent, Nafissa Yakubova, James Pinkerton, Duo Wang, Erich Owens, C. Lawrence Zitnick, Michael P. Recht, Daniel K. Sodickson, and Yvonne W. Lui. fastMRI: An open dataset and benchmarks for accelerated MRI. *arXiv preprint arXiv:1811.08839*, 2018.




Figure 1: **PWM overview.** The Physics World Models pipeline. **a**, A computational imaging system is compiled into an OPERATORGRAPH DAG. **b**, The TRIAD DECOMPOSITION diagnostic agents evaluate each gate. **c**, The dominant gate is identified and a TRIADREPORT is produced. **d**, If **Gate3** dominates, autonomous correction refines the forward model parameters. **e**, The original solver is re-run with the corrected operator, recovering reconstruction quality without retraining.




Figure 2: **OperatorGraph IR and Physics Fidelity Ladder.** **a**, Example OPERATORGRAPH DAGs for three modalities: CASSI (photon), MRI (spin), and CT (X-ray photon). Each node wraps a primitive operator; edges define data flow. **b**, The Physics Fidelity Ladder. Tier 1: linear shift-invariant. Tier 2: linear shift-variant. Tier 3: nonlinear ray/wave-based. Tier 4: full-wave/Monte Carlo. **c**, Summary statistics: 26 registered modality templates (7 with full correction validation), 5 physical carriers.




Figure 3: **Triad Decomposition structure and gate binding.** **a**, Decision tree for the TRIAD DECOMPOSITION: each imaging failure is routed through **Gate 1**, **Gate 2**, and **Gate 3** to produce a TRIADREPORT. **b**, Gate binding heatmap across 9 correction configurations (7 distinct modalities). Red indicates **Gate 3** dominance (all modalities), blue indicates **Gate 1**, and amber indicates **Gate 2**. **c**, Recovery ratio  $\rho$  distribution across all 9 correction configurations.





Figure 4: **Correction results and zero-shot generalization.** **a**, Correction gain  $\Delta_{\text{corr}}$  (dB) across 9 validated configurations (7 distinct modalities), grouped by carrier family. **b**, Zero-shot generalization: hyperparameters tuned on photon-domain modalities transfer without modification to coherent-photon, spin, and X-ray domains. Dark bars: modality-specific tuning; light bars: zero-shot transfer.



**Visual reconstruction comparison across three modalities**




Figure 5: **Modality deep dives and visual comparison.** **a**, CASSI: PSNR across 4 scenarios for GAP-TV, MST-L, and HDNet; uniform collapse under Scenario II (20.83–21.88 dB) confirms operator-driven failure. **b**, CACTI: four methods across 4 scenarios, showing up to 20.58 dB mismatch degradation. **c–e**, Visual reconstruction comparison (CASSI, CACTI, MRI): ground truth, Scenario I, Scenario II (structured artifacts), Scenario III (PWM-corrected), and error map. Mismatch artifacts are qualitatively distinct from missing data artifacts.




Figure 6: **Hardware validation on real CASSI and CACTI instruments.** **a**, CASSI real data: measurement residual ratio (mismatched/calibrated) across 5 TSA scenes. GAP-TV shows 1.8 $\times$  mean ratio. **b**, CACTI real data: residual ratio across 4 scenes. GAP-TV shows 10.4 $\times$  mean ratio; PnP-FFDNet shows 2.0 $\times$ . **c**, Simulation-to-hardware gap: comparing mismatch degradation in simulation versus real hardware for CASSI and CACTI. **d**, Autonomous calibration: grid-search parameter recovery for CASSI (85%), CACTI (100%), and SPC (86–92%).




Figure 7: **Basis-growth saturation and primitive decomposition.** **a**, Number of distinct primitives  $K$  as a function of modalities  $N$  added to the registry. The curve saturates at  $K = 10$  for  $N \geq 27$ ; annotated points mark the introduction of each primitive. **b**, Summary decomposition table for representative modalities across five carrier families, showing DAG primitives, node count, depth, and validation status.



830 **Extended Data Table 1** | Master correction results across 9 validated configurations (7  
831 modalities). Full table with per-scene metrics in Supplementary Table S1.

832 **Extended Data Table 2** | 26-modality template registry spanning five physical carrier  
833 families. Full table in Supplementary Table S3.