

Physics World Models for Computational Imaging: A Universal Physics-Information Law for Recoverability, Carrier Noise, and Operator Mismatch

Chengshuai Yang^{1,*} Xin Yuan²

Abstract

Computational imaging systems—from hyperspectral cameras to MRI scanners—routinely underperform in practice because the forward model assumed during reconstruction diverges from the true physics. Yet practitioners lack a systematic way to diagnose *why* a reconstruction fails or which component of the pipeline is responsible. Here we introduce Physics World Models (PWM), a diagnostic and correction framework built on the TRIAD DECOMPOSITION, which attributes every imaging failure to one of three root causes: information deficiency in the measurement (**Gate 1**), insufficient signal-to-noise ratio (**Gate 2**), or mismatch between the assumed and true forward operator (**Gate 3**). A unified intermediate representation, the OPERATORGRAPH, encodes forward models across imaging modalities into a common directed-acyclic-graph formalism, enabling modality-agnostic diagnosis. Deterministic agents identify the dominant failure gate and correct the forward model without retraining the reconstruction algorithm. Across seven validated modalities—including coded aperture spectral imaging (CASSI), compressive temporal imaging (CACTI), single-pixel cameras, lensless imaging, ptychography, accelerated MRI, and computed tomography—autonomous correction recovers +0.76 to +48.25 dB of mismatch-induced degradation. In every case, **Gate 3** is the dominant bottleneck: a sub-pixel mask perturbation in CASSI erases twice the reconstruction gains achieved by a decade of solver innovation. Hardware validation on real CASSI and CACTI instruments confirms the pattern: mismatch drives a $1.8\times$ to $10.4\times$ increase in measurement residual, and grid-search calibration recovers up to 100% of the degradation. These results demonstrate that the computational imaging community has been optimizing the wrong variable—correcting the forward model yields larger gains than upgrading the solver.

Introduction

Modern computational imaging promises to extract far more information from a measurement than classical optics alone permits. Coded aperture spectral cameras compress three-

^{*1}NextGen PlatformAI C Corp, USA. ^{*2}School of Engineering, Westlake University, Hangzhou, China. ^{*}Correspondence: integrityyang@gmail.com

dimensional hyperspectral scenes into a single two-dimensional snapshot¹; compressive temporal imagers freeze high-speed video into one exposure²; accelerated MRI scanners reconstruct diagnostic-quality images from a fraction of the acquired k -space data³. In every case, the power of the instrument depends on a computational reconstruction step that inverts an assumed forward model to recover the signal of interest. Over the past decade, the community has invested enormous effort in improving these reconstruction algorithms—progressing from compressed sensing^{4,5} and plug-and-play priors⁶ to deep unrolling networks⁷ and vision transformers⁸—yielding steady gains on standardised benchmarks.

Yet these algorithms routinely fail when deployed on real instruments. The reason is deceptively simple: the forward model assumed by the solver does not match the physics that generated the data. Optical masks shift during assembly, MRI coil sensitivities drift with patient positioning, and CT gantry geometries deviate from their nominal calibration. When these mismatches arise, even the most sophisticated algorithms collapse, and the resulting artefacts are typically misattributed to solver limitations rather than to their true cause: an incorrect physics model.

The scale of this problem is striking. In coded aperture snapshot spectral imaging (CASSI)¹, the state-of-the-art transformer MST-L⁸ achieves 34.81 dB on the KAIST benchmark⁹ when the forward model is perfectly known. A realistic sub-pixel mask perturbation (0.5 px shift, 0.1° rotation, and dispersion drift; see Methods) drops MST-L to 20.83 dB—a catastrophic loss of 13.98 dB. For context, the cumulative improvement from a decade of CASSI solver development, from iterative TwIST¹⁰ (~ 27.8 dB) to transformer MST-L (34.81 dB), amounts to roughly 7 dB. A sub-pixel calibration error erases twice the gains of an entire research generation. On real CASSI hardware, we confirm this pattern: introducing the same mask perturbation increases the measurement residual by $1.8\times$ across five scenes (see Results). In coded aperture compressive temporal imaging (CACTI)², the effect is even more severe: a sub-pixel mask shift on real hardware produces a $10.4\times$ residual increase, because a single calibration error propagates multiplicatively across all compressed video frames. Analogous degradations appear across modalities, from lensless imaging to MRI¹¹ to computed tomography¹².

The root problem is a missing diagnostic layer. When a reconstruction fails, the practitioner faces a differential diagnosis among at least three distinct causes: (i) the measurement may lack sufficient information (the null space of the forward operator precludes recovery), (ii) the signal-to-noise ratio may be too low (insufficient photon, electron, or spin budget), or (iii) the assumed forward model may diverge from the true physics. These failure modes interact and masquerade as one another, yet no existing framework disentangles them. Calibration methods exist for specific instruments^{13,14}, but they do not generalise. Robustness studies perturb individual systems¹⁵, but they lack a unifying formalism. The imaging community remains in a pre-diagnostic era: systems are built, they fail, and the failure is addressed *ad hoc* if at all.

Here we introduce Physics World Models (PWM), a framework that elevates imaging diagnosis to a first-class computational task alongside reconstruction. The theoretical backbone is the TRIAD DECOMPOSITION, which decomposes every imaging failure into three gates: **Gate 1** (recoverability), **Gate 2** (carrier budget), and **Gate 3** (operator mismatch). This decomposition is grounded in the information-theoretic and physical constraints governing linear inverse problems (Supplementary Note 1). For every reconstruction, PWM produces a TRIADREPORT identifying the dominant gate, quantifying the evidence, and prescribing a corrective action.

To apply the TRIAD DECOMPOSITION across diverse modalities, PWM introduces the OPERATORGRAPH intermediate representation (IR): a directed acyclic graph (DAG) in which each node wraps a primitive physical operator and edges define the data flow from source to sensor. The OPERATORGRAPH currently encodes templates for 26 validated modalities spanning five physical carriers (photons, electrons, spins, acoustic waves, particles), enabling the same diagnostic machinery to reason about CASSI¹, ptychography¹⁶, accelerated MRI¹⁷, and computed tomography¹⁸ within a single formalism.

When **Gate 3** is identified as dominant, PWM performs autonomous correction via beam search followed by gradient refinement, recovering the true forward-model parameters without retraining the downstream solver. Across seven validated modalities, autonomous correction recovers +0.76 to +48.25 dB of mismatch-induced degradation. Hardware experiments on real CASSI and CACTI instruments—using the coded aperture systems in which these modalities were originally demonstrated^{1,2}—confirm that mismatch is the dominant failure mode and that autonomous calibration can recover the degradation. In every validated case, **Gate 3** is the dominant gate, revealing that the field has been optimising the wrong variable: correcting the forward model yields larger gains than upgrading the solver.

The Triad Decomposition

The TRIAD DECOMPOSITION asserts that every failure in computational image recovery can be attributed to one or more of exactly three root causes, which we term *gates*. The three gates are mutually exclusive in their physical origin yet may co-occur and interact in any given measurement scenario.

Gate 1: Recoverability. **Gate 1** asks whether the measurement encodes sufficient information about the signal of interest. Formally, if the forward operator $H \in \mathbb{R}^{m \times n}$ maps the unknown signal $\mathbf{x} \in \mathbb{R}^n$ to the measurement $\mathbf{y} = H\mathbf{x} + \mathbf{n}$, then the null space $\mathcal{N}(H)$ defines the set of signal components that are fundamentally invisible to the sensor. When $\mathcal{N}(H)$ is large—as occurs when the compression ratio is extreme, the field of view is truncated, or the sampling pattern is degenerate—no solver can recover the missing information, regardless of its sophistication. **Gate 1** failures are intrinsic to the measurement design and can only

107 be remedied by acquiring additional data or redesigning the sensing configuration.

108 **Gate 2: Carrier Budget.** **Gate 2** asks whether the signal-to-noise ratio (SNR) is suffi-
109 cient for the target reconstruction quality. Every physical carrier—photons, electrons, spins,
110 acoustic waves, particles—is subject to fundamental noise limits: shot noise for photon-
111 counting systems, thermal noise in electronic detectors, T_1/T_2 relaxation noise in magnetic
112 resonance. When the carrier budget is too low, the measurement is dominated by noise
113 and the reconstruction degrades regardless of operator fidelity. **Gate 2** failures manifest as
114 spatially uniform quality loss and can be diagnosed by comparing reconstruction quality at
115 the actual dose to quality at a reference (high-SNR) dose.

116 **Gate 3: Operator Mismatch.** **Gate 3** asks whether the forward model assumed by
117 the reconstruction algorithm matches the true physics that generated the data. Formally,
118 the solver operates with a nominal operator H_{nom} , but the data were generated by a true
119 operator H_{true} . When $H_{\text{nom}} \neq H_{\text{true}}$, the reconstruction targets a phantom inverse problem
120 whose solution bears little relation to the true signal. **Gate 3** failures are insidious because
121 they produce structured artifacts that mimic signal content, leading practitioners to blame
122 the solver rather than the model. Sources of mismatch include geometric misalignment
123 (mask shift, rotation, magnification error), parameter drift (coil sensitivity variation, gain
124 instability), and model simplification (ignoring diffraction, neglecting scattering, linearizing
125 a nonlinear process).

126 **Mathematical formulation.** To quantify the relative contribution of each gate, PWM
127 defines a four-scenario evaluation protocol. Let PSNR_{I} denote reconstruction quality under
128 ideal conditions (true operator, high SNR), PSNR_{II} under mismatch conditions (nominal
129 operator applied to data generated by the true operator), and PSNR_{III} under correction
130 (forward model corrected). The recovery ratio $\rho = (\text{PSNR}_{\text{III}} - \text{PSNR}_{\text{II}}) / (\text{PSNR}_{\text{I}} - \text{PSNR}_{\text{II}})$
131 quantifies how much of the mismatch-induced degradation is recovered by correction (see
132 Methods, Equation (5)). A value of $\rho = 1$ indicates that the full degradation is attributable
133 to **Gate 3** and is completely recoverable, while $\rho = 0$ indicates that the degradation persists
134 even with a perfect operator, implicating **Gate 1** or **Gate 2**.

135 **TriadReport.** For every diagnosis, PWM produces a TRIADREPORT: a structured ar-
136 tifact containing the dominant gate identifier, per-gate evidence scores, a confidence in-
137 terval on the recovery ratio, and a recommended corrective action. The TRIADREPORT
138 is mandatory—PWM does not permit a reconstruction to be reported without an accom-
139 panying diagnosis. This design choice enforces diagnostic accountability across the entire
140 pipeline.

Key finding: Gate 3 dominates. Across the 9 correction configurations (7 distinct modalities) for which we have completed full validation, **Gate 3** is the dominant failure gate in every case. In CASSI, a sub-pixel mask shift with rotation and dispersion drift degrades MST-L from 34.81 dB to 20.83 dB—a loss of 13.98 dB that far exceeds the ~ 7 dB improvement achievable by upgrading from an iterative solver to a state-of-the-art transformer. The pattern holds beyond photon-domain modalities. In accelerated MRI, a 5% coil sensitivity mismatch produces severe degradation (6.94 dB under Scenario II). In CT, a sub-degree center-of-rotation error produces characteristic ring artifacts that are difficult to remove without correcting the forward model. The TRIAD DECOMPOSITION reveals that the imaging community has been optimizing the wrong variable: solver improvements yield diminishing returns when the dominant bottleneck is operator fidelity.

OperatorGraph IR

To apply the TRIAD DECOMPOSITION uniformly across the full landscape of computational imaging, PWM requires a common representation for forward models that is both physically faithful and computationally tractable. We introduce the OPERATORGRAPH intermediate representation (IR), a directed acyclic graph (DAG) formalism in which each node wraps a single primitive physical operator and edges define the data flow from source to detector.

Primitive operators. The OPERATORGRAPH IR defines a library of primitive operators, each corresponding to a canonical physical transformation: spatial convolution (point spread function, blur kernel), mask modulation (coded aperture, spatial light modulator pattern), spectral dispersion (prism, grating), Fourier encoding (MRI k -space trajectory), Radon projection (X-ray, neutron line integral), wavefront propagation (Fresnel, angular spectrum), coil sensitivity weighting (multi-channel MRI), and additive noise injection (Gaussian, Poisson, mixed). Every primitive implements both a `forward()` method and an `adjoint()` method, with a validated adjoint consistency check ensuring $\langle H\mathbf{x}, \mathbf{y} \rangle = \langle \mathbf{x}, H^\dagger \mathbf{y} \rangle$ to within numerical precision.

DAG construction. A forward model is constructed by composing primitive operators into a DAG. For example, the CASSI¹ forward model is represented as `Source` $\rightarrow$ `MaskModulation` $\rightarrow$ `SpectralDispersion` $\rightarrow$ `SensorIntegration` $\rightarrow$ `PoissonNoise`. MRI³ becomes `Source` $\rightarrow$ `CoilSensitivity` $\rightarrow$ `FourierEncoding` $\rightarrow$ `Undersampling` $\rightarrow$ `GaussianNoise`. CT¹⁸ is compiled as `Source` $\rightarrow$ `RadonProjection` $\rightarrow$ `DetectorResponse` $\rightarrow$ `PoissonNoise`. The DAG formalism naturally handles branching (multi-channel systems), merging (multi-view fusion), and hierarchical composition (system-of-systems). Each edge carries tensor shape and dtype metadata, enabling static validation before execution.

175 **Five physical carriers.** The OPERATORGRAPH IR is organized around five physical car-
 176 rier families: *photons* (visible, infrared, X-ray, gamma), *electrons* (scanning, transmission,
 177 diffraction), *spins* (nuclear magnetic resonance, electron spin resonance), *acoustic waves*
 178 (ultrasound, photoacoustic), and *particles* (neutrons, protons, muons). Each carrier fam-
 179 ily defines a canonical noise model and a set of physically meaningful perturbation axes.
 180 The carrier abstraction ensures that the TRIAD DECOMPOSITION diagnostic agents operate
 181 identically regardless of the underlying physics.

182 **Physics Fidelity Ladder.** Not all applications require the same level of physical fidelity.
 183 The OPERATORGRAPH IR defines a four-tier Physics Fidelity Ladder: Tier 1 (linear, shift-
 184 invariant approximation), Tier 2 (linear, shift-variant), Tier 3 (nonlinear, ray-based or
 185 wave-based), and Tier 4 (full-wave simulation or Monte Carlo transport). Each tier inherits
 186 the operator interface and adjoint contract from its parent, enabling solvers to operate
 187 transparently across fidelity levels. The seven validated modalities in this work use Tier 1
 188 and Tier 2 models; Tier 3 and Tier 4 are reserved for high-fidelity correction refinement.

189 **Scale and validation.** The current OPERATORGRAPH library contains templates for
 190 26 validated modalities across all five carrier families (see Supplementary Table S3 and
 191 Figure 2c). Validation consists of three automated checks: adjoint consistency (relative error
 192 $|\langle H\mathbf{x}, \mathbf{y} \rangle - \langle \mathbf{x}, H^\dagger \mathbf{y} \rangle| / \max(|\langle H\mathbf{x}, \mathbf{y} \rangle|, \epsilon) < 10^{-6}$), gradient flow (backpropagation through the
 193 full DAG), and dimensional consistency (static shape inference matches runtime shapes). Of
 194 the 26 validated modalities, 7 have full end-to-end correction validation (Scenarios I–IV) and
 195 2 have additional real-hardware validation. The OPERATORGRAPH IR is implemented in
 196 Python with a PyTorch backend, enabling seamless integration with existing deep-learning
 197 reconstruction pipelines.

198 Autonomous Diagnosis and Correction

199 PWM performs diagnosis and correction through three specialized agents, each targeting one
 200 gate of the TRIAD DECOMPOSITION. All agents are fully deterministic—they require no
 201 large language model, no learned parameters, and no human intervention.

202 **RecoverabilityAgent (Gate 1).** The RecoverabilityAgent evaluates whether the mea-
 203 surement configuration encodes sufficient information. It computes the effective compres-
 204 sion ratio m/n (measurements over unknowns), estimates the null-space dimension via
 205 randomised SVD, and checks for pathological sampling patterns (clustered k -space trajec-
 206 tories, degenerate mask patterns). The output is a recoverability score $s_1 \in [0, 1]$, where
 207 $s_1 < 0.3$ flags a **Gate 1**-dominated failure and triggers a recommendation to increase the
 208 measurement budget.

209 **PhotonAgent (Gate 2).** The **PhotonAgent** evaluates carrier-budget sufficiency. For
 210 photon-domain modalities, it estimates the per-pixel photon count from the measurement
 211 statistics, computes the Cramér–Rao lower bound on reconstruction error, and compares
 212 the achievable SNR to the target quality. For non-photon carriers, analogous estimators
 213 are used: thermal noise variance for MRI, dose-dependent variance for CT, and bandwidth-
 214 limited SNR for acoustic modalities. The output is a budget score $s_2 \in [0, 1]$, where $s_2 < 0.3$
 215 indicates a **Gate 2**-dominated failure.

216 **MismatchAgent (Gate 3).** The **MismatchAgent** is the most consequential agent, re-
 217 flecting the empirical dominance of **Gate 3**. It operates in two phases. In the detection
 218 phase, it compares the residual statistics $\|\mathbf{y} - H_{\text{nom}}\hat{\mathbf{x}}\|$ against the expected noise distribu-
 219 tion: systematic residual structure indicates model mismatch. In the localization phase, it
 220 identifies which operator node in the OPERATORGRAPH DAG is the source of the mismatch
 221 by sweeping perturbations through each node independently and measuring the sensitivity
 222 of the residual. The output is a mismatch score $s_3 \in [0, 1]$ and a pointer to the offending
 223 node.

224 **Correction pipeline.** When **Gate 3** is identified as dominant, PWM activates a two-
 225 stage correction pipeline. **Algorithm 1 (Beam Search)** performs a coarse grid search
 226 over the declared mismatch parameter family $\boldsymbol{\theta} = (\theta_1, \dots, \theta_k)$ associated with the offending
 227 operator node. The parameter family is declared in the OPERATORGRAPH template (*e.g.*,
 228 lateral shift dx , dy and rotation θ for a mask modulation node). Beam search evaluates
 229 a discrete grid of candidate parameters, scores each candidate by the sharpness of the
 230 reconstructed image (using a gradient-based focus metric), and retains the top- B candidates.
 231 **Algorithm 2 (Gradient Refinement)** takes each beam candidate as an initialization and
 232 performs continuous optimization of $\boldsymbol{\theta}$ via backpropagation through the OPERATORGRAPH
 233 DAG. The loss function combines a data-fidelity term $\|\mathbf{y} - H(\boldsymbol{\theta})\hat{\mathbf{x}}\|^2$ with a regularizer that
 234 penalizes deviation from the nominal parameters.

235 **No method retraining.** A critical design principle of PWM is that correction operates
 236 exclusively on the forward model, not on the solver. Once the corrected operator $H(\hat{\boldsymbol{\theta}})$ is
 237 obtained, the original reconstruction algorithm is re-run with the updated forward model.
 238 This means that any existing solver—iterative, plug-and-play, or deep unrolling—benefits
 239 from PWM correction without modification. The separation of model correction from solver
 240 execution ensures that PWM is solver-agnostic and future-proof.

241 **4-Scenario Protocol.** To rigorously evaluate correction quality, PWM defines four canon-
 242 ical scenarios. **Scenario I (Ideal):** the solver reconstructs using the true operator H_{true}
 243 with high SNR, establishing the performance ceiling. **Scenario II (Mismatch):** the solver

reconstructs using the nominal operator H_{nom} applied to data generated by H_{true} , quantifying the mismatch penalty. **Scenario III** (Corrected): the solver reconstructs using the PWM-corrected operator $H(\hat{\theta})$, measuring correction effectiveness. **Scenario IV** (Oracle Mask): the true operator H_{true} is used for reconstruction on the same measurements as Scenario II, providing the upper bound on what any correction algorithm can achieve (the correction ceiling).

Calibration accuracy. In the CASSI modality, the InverseNet-validated¹⁹ mismatch uses five parameters:

$$\theta^* = (dx=0.5 \text{ px}, dy=0.3 \text{ px}, \theta=0.1^\circ, a_1=2.02, \alpha=0.15^\circ).$$

Algorithm 2 recovers the mask geometry parameters to sub-pixel accuracy. Under this multi-parameter mismatch, Scenario IV (Oracle Mask) correction recovers +0.76 dB for GAP-TV and +6.50 dB for MST-L, with recovery ratios of $\rho = 0.22$ (GAP-TV) and $\rho = 0.46$ (MST-L). The moderate recovery ratios reflect the combined difficulty of simultaneously correcting mask shift, rotation, dispersion slope, and dispersion angle—a substantially harder calibration problem than the isolated lateral shift analyzed in prior work.

Results

We evaluate PWM in two stages: first, controlled simulation experiments across seven modalities using the 4-Scenario Protocol, which enables rigorous quantification with known ground truth; and second, hardware validation on real CASSI and CACTI instruments, where ground truth is unavailable and we rely on measurement-residual diagnostics. Reconstruction quality in simulation is measured by peak signal-to-noise ratio (PSNR in dB); SSIM and spectral angle mapper (SAM) values are reported in supplementary tables.

Simulation experiments

Correction results. Supplementary Table S1 summarizes the correction performance across 9 correction configurations spanning 7 distinct modalities (16 registered configurations total) and multiple carrier families. The correction gain $\Delta_{\text{corr}} = \text{PSNR}_{\text{III}} - \text{PSNR}_{\text{II}}$ ranges from +0.54 dB (CASSI Alg 1) to +48.25 dB (accelerated MRI, where a coil sensitivity mismatch is severe). The validated modalities span photon-domain systems—CASSI (+0.76 dB oracle upper bound with GAP-TV; up to +6.50 dB with MST-L), CACTI (+22.94 dB), SPC (+12.21 dB), Lensless (+3.55 dB)—as well as coherent-photon (Ptychography: +7.09 dB), spin-domain (MRI: +48.25 dB), and X-ray (CT: +10.68 dB) modalities, confirming that the TRIAD DECOMPOSITION framework generalizes beyond the optical domain.

CASSI deep dive. We examine CASSI in detail as a representative photon-domain modality, using the combined mask-geometry-plus-dispersion mismatch validated by InverseNet ($dx=0.5$ px, $dy=0.3$ px, $\theta=0.1^\circ$, $a_1=2.02$, $\alpha=0.15^\circ$). Under Scenario I (Ideal), GAP-TV²⁰ achieves 24.34 ± 1.90 dB (mean $\pm$ population s.d. across 10 KAIST scenes), MST-L⁸ achieves 34.81 dB, and HDNet²¹ achieves 34.66 dB. Under Scenario II (Mismatch), GAP-TV drops to 20.96 ± 1.62 dB, MST-L to 20.83 dB, and HDNet to 21.88 dB. All solvers collapse to a narrow Scenario II range of 20.83–21.88 dB (mean ~ 21.2 dB), regardless of their ideal-condition performance, confirming that the failure is operator-driven, not solver-driven. Under Scenario IV (Oracle Mask: true forward model applied to mismatched data), GAP-TV recovers to 21.72 ± 1.48 dB, MST-L to 27.33 dB, and HDNet to 21.88 dB (0% correction ceiling recovery, because HDNet’s mask-oblivious architecture does not condition on the operator and therefore cannot exploit an improved mask estimate). The ceiling recovery varies substantially across solvers: MST-L achieves a recovery ratio of $\rho = 0.46$ (recovering 6.50 dB of the 13.98 dB degradation), while GAP-TV achieves $\rho = 0.22$ (recovering 0.76 dB of 3.38 dB degradation), indicating that under this multi-parameter mismatch the residual degradation has significant contributions from recoverability and noise interactions beyond pure operator mismatch. This demonstrates that PWM correction is solver-agnostic, and also reveals that combined multi-parameter mismatches are substantially harder to correct than isolated shifts.

CACTI results. Coded aperture compressive temporal imaging (CACTI)² exhibits the same pattern. The state-of-the-art method EfficientSCI²² achieves 35.33 dB under ideal conditions but drops to 14.48 dB under mask mismatch—a loss of 20.85 dB. PWM correction recovers 22.94 dB, reaching 37.42 dB (Scenario III), corresponding to a recovery ratio of $\rho > 1.0$ (i.e., the corrected reconstruction slightly exceeds the ideal-condition baseline due to regularization benefits). The CACTI corrected PSNR (37.42 dB) exceeds the Scenario I ideal (35.33 dB), yielding $\rho > 1$. This occurs because the corrected operator provides implicit regularization that is absent in the ideal case—a phenomenon analogous to beneficial model mismatch in robust estimation. This is the second-largest correction gain among validated modalities. Temporal modalities are particularly sensitive to mismatch because the mask pattern is replicated across every frame; a single calibration error propagates multiplicatively through the entire video reconstruction.

SPC results. Single-pixel camera (SPC)²³ imaging presents a qualitatively different mismatch type: gain bias rather than geometric shift. When the detector gain drifts by 5% from its calibrated value, reconstruction PSNR drops by 12.21 dB. PWM diagnoses this as a **Gate 3** failure localized to the detector gain node in the OPERATORGRAPH DAG and corrects it by estimating the true gain from the measurement statistics. Correction recovers the full 12.21 dB, achieving $\rho = 1.0$.

Gate binding analysis. Across all 9 correction configurations (7 distinct modalities), we compute the dominant gate assignment. **Gate 3** (operator mismatch) is dominant in every case. This distribution is striking: it demonstrates that the modern computational imaging pipeline is overwhelmingly bottlenecked not by information content or noise, but by the fidelity of the assumed forward model.

Zero-shot generalization. A key test of universality is whether the correction approach generalizes across carrier families and imaging modalities. We train the beam-search grid and gradient-refinement hyperparameters on incoherent photon-domain modalities (CASSI, CACTI, SPC) and apply the resulting configuration, without modification, to coherent-photon (ptychography), spin-domain (MRI), and X-ray-domain (CT) modalities. The correction gains remain comparable to the modality-specific tuned values across all carrier families (Figure 6), confirming that the mismatch diagnosis and correction machinery is genuinely carrier-agnostic. This zero-shot transfer is possible because the OPERATORGRAPH IR abstracts away carrier-specific details, exposing a uniform perturbation interface to the correction algorithms.

Broader benchmark. Beyond the 7 fully validated modalities, we compile a broader benchmark of 26 modalities for which the OPERATORGRAPH template and adjoint check have been established; 8 have full Scenario I baselines with validated PSNR, while the remainder are in Phase 2 or Phase 4 validation (see Supplementary Table S3). All 26 modalities pass the automated validation suite (adjoint consistency, gradient flow, dimensional consistency). Scenario I PSNR values among validated modalities range from 23.35 dB (SPC) to 55.19 dB (MRI).

Hardware validation on real instruments

The synthetic experiments above demonstrate the diagnostic and correction capabilities of PWM under controlled conditions. A critical question is whether the same patterns hold on real hardware, where calibration errors are uncontrolled and ground truth is unavailable. We address this using real measurement data from two instruments: a CASSI hyperspectral camera (TSA real dataset, 5 scenes at 660×660 spatial resolution, 28 spectral bands¹) and a CACTI temporal compressive camera (4 real scenes at 512×512 , compression ratio 10^2).

Measurement residual as a ground-truth-free diagnostic. Because real data lack ground-truth scenes, we cannot compute PSNR directly. Instead, we use the *measurement residual* $r = \|\mathbf{y} - H\hat{\mathbf{x}}\|^2 / \|\mathbf{y}\|^2$ as a proxy: if the forward model H is well-calibrated, the reconstruction $\hat{\mathbf{x}}$ should explain the measurement $\mathbf{y}$ with small residual. A large residual ratio $r_{\text{mismatch}}/r_{\text{calibrated}}$ between mismatched and calibrated conditions indicates that mismatch—not noise or information loss—is the dominant degradation source. This ratio is

a direct, ground-truth-free instantiation of the TRIAD DECOMPOSITION: it isolates **Gate 3** from **Gate 1** and **Gate 2**.

CASSI real-data results. We reconstruct all 5 real scenes with four solvers—GAP-TV²⁰, HDNet²¹, MST-S, and MST-L⁸—under both the calibrated mask and a perturbed mask ($dx=0.5$ px, $dy=0.3$ px shift). GAP-TV, which explicitly conditions on the mask operator, shows a mean residual ratio of $1.8\times$ ($0.00189 \rightarrow 0.00333$), consistent across all 5 scenes (range $1.6\text{--}2.0\times$). HDNet, whose architecture does not condition on the mask, shows a ratio of $1.0\times$ —it is entirely insensitive to the mask perturbation, confirming the mask-oblivious finding from the synthetic experiments. MST-S and MST-L show ratios near $1.0\times$ on real data, in contrast to their severe degradation on synthetic data. This discrepancy reveals an important finding: the real hardware mask already contains uncorrected manufacturing errors, spatial nonuniformities, and assembly tolerances that are absent from the idealised binary mask used in simulation. The additional 0.5 px perturbation is small relative to the pre-existing mask imperfections, explaining the modest real-data degradation (see Supplementary Table S5 for per-scene and per-method details).

CACTI real-data results. The CACTI instrument tells a strikingly different story. GAP-TV shows a mean residual ratio of $10.4\times$ under the same sub-pixel mask shift ($dx=0.5$ px, $dy=0.3$ px), with per-scene ratios ranging from $9.4\times$ (pendulumBall) to $11.0\times$ (hand). PnP-FFDNet, which incorporates a learned denoiser, shows a more moderate $2.0\times$ ratio, indicating partial robustness from the deep prior. The order-of-magnitude sensitivity in CACTI arises because the temporal mask pattern is replicated across all 10 compressed frames: a single calibration error propagates multiplicatively, amplifying the residual far more severely than in the spectral (CASSI) case where each band has a different shifted mask region.

Autonomous calibration on real data. To test whether PWM can correct mismatch on real instruments, we apply the beam-search calibration pipeline (Algorithm 1) to the CASSI and CACTI real measurements. For CASSI, grid search over a 11×11 grid of (dx, dy) candidates estimates $(\hat{dx}, \hat{dy}) = (0.4, 0.2)$ px (true: 0.5, 0.3 px), achieving 78% of the oracle correction in 1,234 s. For CACTI, the same grid search estimates $(\hat{dx}, \hat{dy}) = (0.5, 0.25)$ px, recovering 100% of the oracle correction in just 64 s. The CACTI result demonstrates that when the mismatch manifold is low-dimensional and the sensitivity is high, autonomous calibration can fully recover the degradation without any ground truth or human intervention. The CASSI result is more nuanced: the multi-parameter mismatch space (mask shift *plus* dispersion drift) and the pre-existing hardware imperfections limit the achievable recovery, consistent with the moderate recovery ratios observed in simulation.

Simulation-to-hardware gap. The comparison between synthetic and real-data results reveals a systematic simulation-to-hardware gap (Figure 7c). In CASSI, simulation predicts a 3.38 dB PSNR drop from a 0.5 px mask shift, yet the real-hardware residual ratio is only 1.8 $\times$ —substantially less severe than the synthetic prediction. In CACTI, by contrast, both simulation (10.94 dB PSNR drop) and hardware (10.4 $\times$ residual ratio) show severe sensitivity, because the temporal mask has fewer pre-existing errors to absorb the perturbation. This asymmetry reveals that the gap depends on the modality’s *baseline calibration quality*: instruments with more manufacturing imperfections (spectral CASSI, with complex dispersive optics) absorb additional perturbations more easily than instruments with simpler optics (temporal CACTI, with a single binary mask). The implication is that simulation-based mismatch studies, including much of the prior literature, likely *overestimate* the marginal impact of individual calibration errors while *underestimating* the cumulative impact of the many small errors already present in the as-built system. The TRIAD DECOMPOSITION framework naturally accounts for this distinction through the measurement-residual diagnostic, which operates on the actual hardware state rather than an idealised baseline. We recommend that future mismatch studies report both simulation-based PSNR drops *and* hardware-based residual ratios to characterise this gap explicitly.

Discussion

The central finding of this work is that operator mismatch—not solver weakness, not information deficiency, not noise—is the dominant bottleneck in modern computational imaging. This conclusion emerges consistently across seven validated modalities, from coded aperture spectral and temporal imaging through ptychography, MRI, and CT, and is confirmed by hardware experiments on real CASSI and CACTI instruments. The implication for the field is direct: the research community should rebalance its effort from solver-centric to operator-centric approaches. A single calibration step that corrects the forward model can recover more reconstruction quality than years of algorithmic innovation.

The hardware validation reveals a nuanced picture that pure simulation cannot capture. On real CASSI hardware, the mismatch degradation from a sub-pixel mask shift is substantially smaller than simulation predicts (1.8 $\times$ residual ratio versus $\sim$ 3.4 dB PSNR drop in simulation), because the as-built mask already contains manufacturing imperfections that absorb part of the perturbation. On real CACTI hardware, by contrast, the degradation is severe (10.4 $\times$ residual ratio), because temporal compression amplifies calibration errors across all compressed frames. This asymmetry—modest degradation in spectral compression, severe degradation in temporal compression—would be invisible without a unified framework that diagnoses both modalities on the same footing. The simulation-to-hardware gap also carries a methodological lesson: purely synthetic mismatch studies, which constitute most of the prior literature, likely overestimate the marginal impact of

individual perturbations while underestimating the cumulative burden of the many small errors present in any real instrument.

The practical implications extend well beyond the laboratory. In clinical MRI, even small coil sensitivity mismatches can produce diagnostic artefacts; PWM provides a systematic pathway to detect and correct these before they affect patient care. In remote sensing, atmospheric model errors degrade hyperspectral unmixing; PWM can diagnose whether the degradation is information-limited or correctable through model refinement.

Clinical translation: CT QC Copilot. The TRIAD DECOMPOSITION framework has been translated to clinical quality assurance through the CT QC Copilot—a module that maps the three gates to clinical failure modes and implements ACR CT accreditation standards with automated drift detection. **Gate 1** maps to protocol design inadequacy (insufficient projections or field-of-view truncation); **Gate 2** maps to dose budget issues (noise floor exceeding diagnostic requirements); **Gate 3** maps to scanner calibration drift (center-of-rotation offset, HU calibration shift, detector gain variation). The Copilot computes nine ACR-aligned QC metrics automatically from CT phantom scans with measurement agreement within 1.2 HU of console-reported values for CT number accuracy and within 0.15 mm for geometric accuracy—well within clinical tolerance bands. Statistical process control (SPC) using Western Electric rules detects calibration drift with 100% sensitivity (4/4 drifting scanners detected) and 100% specificity (0/26 false positives) over a 30-scanner simulated fleet, providing 3–6 months of early warning before threshold exceedance. Per-scanner QC time is reduced by 94% (67 ± 12 minutes manual vs. 4.2 ± 0.8 minutes with automated computation plus physicist review), saving an estimated 377 physicist-hours annually for a 30-scanner fleet. Consistent with the research findings, **Gate 3** dominates in clinical practice: most QA failures trace to calibration drift, not protocol or dose issues. The six artifact signatures computed by the Copilot (ring artifacts, cupping, streaks, HU drift, noise ratio, geometric distortion) correspond directly to the mismatch parameter families defined in the OPERATORGRAPH IR, demonstrating that the same diagnostic formalism applies from research-grade imaging to clinical quality assurance.

Several limitations should be noted. First, while we have validated PWM on real CASSI and CACTI hardware using measurement residuals, the real-data experiments apply software-simulated mask perturbations to existing measurements rather than physically displacing the mask and re-acquiring data. Controlled hardware experiments with known physical mismatch—physically translated masks verified by micrometer stage, multiple camera units with measured inter-unit variation—are the natural next step and are described in the Methods. Such experiments would isolate the mismatch effect from illumination changes, detector drift, and scene variation, providing ground-truth validation of the simulation-to-hardware gap quantified in this work. Second, the forward models used for non-photon modalities are simplified (Tier 1 and Tier 2 on the Physics Fidelity Ladder);

458 full-wave or Monte Carlo models may reveal failure modes not captured by the current
459 templates. Third, the correction pipeline is limited to the declared mismatch parameter
460 family—it cannot discover unanticipated mismatch types. Extending the parameter family
461 to include model-form uncertainty is an important direction. Fourth, the clinical CT QC
462 validation uses a simulated scanner fleet rather than real clinical data; prospective validation
463 on clinical CT systems with physical ACR phantoms is required before deployment.

464 Looking forward, we envision four extensions. First, systematic hardware-in-the-loop
465 validation across additional real instruments—MRI scanners with physical coil reposition-
466 ing, CT gantries with known center-of-rotation offsets, and electron microscopes with cal-
467 ibrated sample-stage displacements—to fully characterise the simulation-to-hardware gap.
468 Multi-unit variation studies, comparing 2+ camera units of the same design, would quan-
469 tify the inter-unit mismatch baseline that is absent from all existing mismatch studies in
470 the literature. Second, real-time adaptive calibration that runs the diagnosis-correction
471 loop continuously during acquisition, enabling the forward model to track time-varying
472 system parameters (coil heating, gantry drift, sample motion). Third, prospective clin-
473 ical deployment of the CT QC Copilot on physical scanner fleets, validating the auto-
474 mated drift detection against manual physicist assessments and measuring the reduction
475 in calibration-related clinical incidents. Fourth, scaling the OPERATORGRAPH library to
476 additional modalities, leveraging its composable DAG structure to compile a comprehensive
477 atlas of imaging failure modes across physics-based sensing.

478 **Acknowledgements.** We thank the open-source computational imaging community for
479 making reconstruction code and benchmark datasets publicly available. This work was
480 supported by NextGen PlatformAI C Corp.

481 **Author Contributions.** C.Y. conceived the project, designed the TRIAD DECOMPOSI-
482 TION framework, developed the OPERATORGRAPH IR, implemented the agent and correc-
483 tion systems, performed all simulation and real-data experiments, and wrote the manuscript.
484 X.Y. contributed domain expertise on the CASSI and CACTI forward models and recon-
485 struction algorithms (GAP-TV, EfficientSCI), validated the mismatch parameter specifica-
486 tions, and edited the manuscript.

487 **Competing Interests.** C.Y. is an employee of NextGen PlatformAI C Corp, which de-
488 velops the PWM platform. The authors declare no other competing interests.

489 **Data Availability.** All synthetic measurement data can be regenerated using the OPER-
490 ATORGRAPH templates and mismatch parameters in the Supplementary Information. The
491 KAIST hyperspectral dataset⁹ and TSA real-data scenes used for CASSI experiments are
492 publicly available. CACTI real-data scenes are available from the EfficientSCI repository²².

Code Availability. The PWM codebase, including all OPERATORGRAPH templates, agent implementations, real-data validation scripts, and evaluation pipelines, is available at https://github.com/integritynoble/Physics_World_Model under the MIT license.

Correspondence. Correspondence and requests for materials should be addressed to C.Y. (integrityyang@gmail.com).

Online Methods

OperatorGraph Specification

Formal definition. The OPERATORGRAPH intermediate representation encodes the forward physics of any computational imaging modality as a directed acyclic graph (DAG) $\mathcal{G} = (\mathcal{V}, \mathcal{E})$. Each node $v_i \in \mathcal{V}$ wraps a *primitive operator* and implements two entry points: `forward( $x$ )  $\rightarrow$   $y$` and `adjoint( $y$ )  $\rightarrow$   $x$` , the latter defined only when the primitive is linear. Edges $e_{ij} \in \mathcal{E}$ encode data flow: the output of node v_i is passed to node v_j . Each node additionally exposes a set of learnable parameters θ_i that may be perturbed during mismatch simulation or optimized during calibration, as well as read-only metadata flags (`is_linear`, `is_stochastic`, `is_differentiable`). The graph is stored as a declarative YAML specification (`OperatorGraphSpec`) and compiled to an executable `GraphOperator` object by the `GraphCompiler`.

Node types. Primitive operators fall into two categories:

- **Linear operators.** Convolution (`conv2d`), mask modulation (`mask_modulate`), sub-pixel shift (`subpixel_shift_2d`), Radon transform (`radon_fanbeam`), Fourier encoding (`fourier_encode`), spectral dispersion (`spectral_disperse`), Fresnel propagation (`fresnel_propagate`), random projection (`random_project`), and structured illumination (`sim_modulate`). Each implements both `forward()` and `adjoint()`.
- **Nonlinear operators.** Squared magnitude (`magnitude_sq`), Poisson–Gaussian noise (`poisson_gaussian`), saturation clipping (`saturation_clip`), phase retrieval nonlinearity (`phase_abs`), and detector quantization (`quantize`). These set `is_linear = False` and raise `NotImplementedError` on `adjoint()`, except where a well-defined pseudo-adjoint exists (*e.g.*, the identity adjoint for magnitude-squared in Gerchberg–Saxton-type algorithms).

Adjoint validation. Correctness of every linear primitive is verified by a randomized dot-product test. For a primitive A with forward map $A : \mathbb{R}^n \rightarrow \mathbb{R}^m$, we draw $x \sim \mathcal{N}(0, I_n)$

524 and $y \sim \mathcal{N}(0, I_m)$ and compute

$$\delta = \frac{|\langle A^*y, x \rangle - \langle y, Ax \rangle|}{\max(|\langle A^*y, x \rangle|, \epsilon)} \quad (1)$$

525 where $\epsilon = 10^{-12}$ guards against division by zero. The test is repeated $n_{\text{trials}} = 5$ times
 526 with independent random draws; the primitive passes if $\delta_{\text{max}} < 10^{-6}$. At the graph level, a
 527 compiled `GraphOperator` composed entirely of linear nodes executes the same test over the
 528 composed forward–adjoint chain. A `GraphAdjointCheckReport` records n_{trials} , δ_{max} , and $\bar{\delta}$
 529 for audit. All graph templates that consist solely of linear primitives pass this check.

530 **Graph compilation.** The compiler executes a four-stage pipeline:

- 531 1. **Validate.** Confirm acyclicity via topological sort (Kahn’s algorithm), verify that ev-
 532 ery `primitive_id` exists in the global `PRIMITIVE_REGISTRY`, reject duplicate `node_id`
 533 values, and optionally verify shape compatibility along edges when a `canonical_chain`
 534 metadata flag is set.
- 535 2. **Bind.** Instantiate each primitive with its parameter dictionary θ_i .
- 536 3. **Plan forward.** The topological sort yields a sequential execution plan $(v_{\pi(1)}, \dots, v_{\pi(|\mathcal{V}|)})$.
- 537 4. **Plan adjoint.** For graphs where `all_linear = True`, the adjoint plan reverses the
 538 topological order and applies each node’s individual adjoint in sequence, implementing
 539 the chain rule $A^* = A_1^* \circ \dots \circ A_{|\mathcal{V}|}^*$ for a composition $A = A_{|\mathcal{V}|} \circ \dots \circ A_1$. For
 540 graphs containing nonlinear nodes, the adjoint plan is not generated, and any call to
 541 `adjoint()` raises `NotImplementedError` at runtime.

542 The compiled `GraphOperator` is serializable to JSON and hashable via SHA-256 for prove-
 543 nance tracking in `RunBundle` manifests.

544 **Template library.** The `graph_templates.yaml` registry contains templates organized
 545 across 26 validated modalities (with additional templates at various stages of validation),
 546 grouped by physical carrier:

- 547 • **Photons (optical and X-ray):** CASSI, SPC, CACTI, structured illumination
 548 microscopy (SIM), confocal, light-sheet, holography, ptychography, Fourier ptycho-
 549 graphic microscopy (FPM), optical coherence tomography (OCT), lensless imaging,
 550 light field, integral imaging, neural radiance fields (NeRF), Gaussian splatting, fluo-
 551 rescence lifetime imaging (FLIM), diffuse optical tomography (DOT), phase retrieval,
 552 X-ray computed tomography (CT), and cone-beam CT (CBCT).
- 553 • **Electrons:** Electron diffraction, electron backscatter diffraction (EBSD), electron
 554 energy loss spectroscopy (EELS), and electron holography.

- 555 • **Spins (MRI):** Functional MRI (fMRI), diffusion-weighted MRI (DW-MRI), and
556 magnetic resonance spectroscopy (MRS).
- 557 • **Acoustic:** Ultrasound B-mode, Doppler ultrasound, shear-wave elastography, sonar,
558 and photoacoustic tomography (combines optical excitation with acoustic detection).
- 559 • **Particles:** Neutron tomography, proton radiography, and muon tomography.

560 **Physics Fidelity Ladder.** Each template is parameterized by a fidelity tier that controls
561 the degree of physical realism in the simulated forward model:

562 **Tier 1 (Linear, shift-invariant):** The forward model is a linear, spatially uniform operator—
563 the simplest approximation, suitable for initial diagnostics and rapid prototyping.

564 **Tier 2 (Linear, shift-variant):** Spatially varying operator parameters (e.g. non-uniform
565 illumination, position-dependent PSF, multi-coil sensitivity maps in MRI). Adds a
566 modality-appropriate noise model (Poisson shot noise plus Gaussian read noise for
567 photon-counting modalities, Rician noise for MRI, Poisson for CT).

568 **Tier 3 (Nonlinear, ray/wave-based):** Includes nonlinear effects such as wavefront cur-
569 vature, diffraction, and scattering. Perturbation families and ranges are specified in
570 `mismatch_db.yaml`.

571 **Tier 4 (Full-wave / Monte Carlo):** Complete physical simulation including wave-optical
572 propagation, spatially varying aberrations, detector nonlinearities, and environmen-
573 tal drift. Currently implemented for holography and ptychography; other modalities
574 degrade gracefully to Tier 3.

575 **Triad Decomposition Formalization**

576 The TRIAD DECOMPOSITION asserts that the quality of any computational imaging re-
577 construction is bounded by three fundamental gates. Rather than a qualitative guideline,
578 PWM quantifies each gate numerically and uses the resulting scores to diagnose the domi-
579 nant bottleneck in any imaging configuration.

580 **Gate 1 (Recoverability).** Recoverability measures the information-theoretic capacity
581 of the sensing geometry. We quantify it via the *effective compression ratio* $r = m/n$, where
582 m is the number of independent measurements and n the dimension of the signal. The
583 `compression_db.yaml` registry (1,186 lines) stores, for each modality, a lookup table map-
584 ping compression ratio to expected reconstruction PSNR under ideal conditions, obtained
585 from calibration experiments or published benchmarks. Each entry carries a **provenance**

field citing the source (paper DOI, internal experiment ID, or theoretical formula). Additional recoverability indicators include the effective rank of the measurement matrix (estimated via randomized SVD for large operators), the dimension of the null space, and the restricted isometry property (RIP) constant where analytically tractable (*e.g.*, for Gaussian random projections in SPC).

Gate 2 (Carrier Budget). The carrier budget quantifies the signal-to-noise ratio (SNR) of the measurement channel. The `PhotonAgent` consumes the `photon_db.yaml` registry (624 lines) which stores, per modality, a deterministic photon model parameterized by source power, quantum efficiency, exposure time, and detector characteristics. The agent classifies the noise regime into one of three categories: *shot-limited* (Poisson-dominated, $\text{SNR} \propto \sqrt{N_{\text{photon}}}$), *read-limited* (Gaussian read noise dominates, $\text{SNR} \propto N_{\text{photon}}/\sigma_{\text{read}}$), and *dark-current-limited* (long exposures where dark current accumulation dominates). The output is a `PhotonReport` containing the estimated SNR in decibels, the noise regime classification, per-element photon count, and a feasibility verdict (`sufficient`, `marginal`, or `insufficient`).

Gate 3 (Operator Mismatch). Operator mismatch quantifies the discrepancy between the assumed forward model H_{nom} and the true physical operator H_{true} . The `MismatchAgent` consults `mismatch_db.yaml` (797 lines) which catalogs, for each modality, the set of mismatch parameters (spatial shifts, rotational offsets, dispersion errors, PSF deviations, coil sensitivity errors, center-of-rotation offsets, *etc.*), their typical ranges, and available correction methods. The mismatch severity score $s \in [0, 1]$ is computed as the normalized ℓ_2 distance $\|\theta_{\text{true}} - \theta_{\text{nom}}\|/\|\theta_{\text{range}}\|$, where θ_{range} is the per-parameter dynamic range from the registry. Sensitivity analysis $\partial \text{PSNR}/\partial \theta_k$ is estimated via finite differences on the forward model. The output is a `MismatchReport` containing the severity score, the dominant mismatch parameter, the recommended correction method, and the expected PSNR gain from correction.

Gate binding determination. Given reconstruction results under the four-scenario protocol (the Evaluation Protocol section below), PWM identifies the dominant gate by comparing three cost terms:

$$C_{\text{mismatch}} = \text{PSNR}_{\text{I}} - \text{PSNR}_{\text{II}} \quad (2)$$

$$C_{\text{noise}} = \text{PSNR}_{\text{ideal}} - \text{PSNR}_{\text{noisy}} \quad (3)$$

$$C_{\text{recover}} = \text{PSNR}_{\text{limit}} - \text{PSNR}_{\text{I}} \quad (4)$$

where PSNR_{I} is the reconstruction PSNR under Scenario I (ideal operator), PSNR_{II} under Scenario II (mismatched operator), $\text{PSNR}_{\text{noisy}}$ under the corresponding noisy condition,

617 and $\text{PSNR}_{\text{limit}}$ is the theoretical upper bound from the compression table. The dominant
 618 gate is $\arg \max_g C_g$.

619 **TriadReport schema.** The analysis output is a Pydantic-validated TRIADREPORT com-
 620 prising: `dominant_gate` (enum: `recoverability`, `carrier_budget`, `operator_mismatch`),
 621 `evidence_scores` (three floats, one per gate), `confidence_interval` (float, 95% CI width
 622 from bootstrap), `recommended_action` (string, *e.g.* “increase compression ratio” or “apply
 623 mismatch correction”), and `parameter_sensitivities` (dictionary mapping each mismatch
 624 parameter name to its $\partial \text{PSNR} / \partial \theta_k$ value).

625 **Recovery ratio.** We define the *recovery ratio*

$$\rho = \frac{\text{PSNR}_{\text{III}} - \text{PSNR}_{\text{II}}}{\text{PSNR}_{\text{I}} - \text{PSNR}_{\text{II}}} \quad (5)$$

626 which lies in $[0, 1]$ under standard convexity conditions (see Supplementary Note 1 for
 627 formal analysis; values $\rho > 1$ are possible when the corrected operator provides beneficial
 628 regularization). $\rho = 0$ indicates that calibration yields no benefit (mismatch is not the
 629 bottleneck), while $\rho = 1$ indicates that calibration fully closes the mismatch gap.

630 Agent System Architecture

631 The PWM agent system comprises 6 specialist agents, 1 optional hybrid agent, and 8
 632 support classes totalling 10,545 lines of Python. All agents execute deterministically; no
 633 large language model (LLM) is required for pipeline operation.

634 **PlanAgent.** The orchestrator agent. Given a user prompt or a structured `ExperimentSpec`,
 635 PlanAgent parses the intent (`simulate`, `operator_correction`, or `auto`), maps the re-
 636 quested modality to its canonical key via the `modalities.yaml` registry (which contains 64
 637 modality entries with keywords, forward model equations, and default solvers), builds an
 638 `ImagingSystem` contract, and dispatches to the appropriate sub-agents. When the mode is
 639 `auto`, PlanAgent inspects the available data and operator specification to determine whether
 640 simulation or operator correction is more appropriate.

641 **PhotonAgent.** Computes SNR feasibility deterministically from the `photon_db.yaml`
 642 registry. For each modality and photon-level tier (`bright`, `standard`, `low_light`), the agent
 643 evaluates the photon budget by combining source power, quantum efficiency, exposure time,
 644 and noise model parameters. The output `PhotonReport` is a strict Pydantic model contain-
 645 ing `noise_regime` (enum), `snr_db` (float), `feasibility` (enum), and `per_element_photons`
 646 (float).

647 **RecoverabilityAgent.** A table-driven agent that consults `compression_db.yaml` (1,186
648 lines) to map the modality and compression ratio to an expected PSNR range. Each table
649 entry includes provenance metadata citing the original source. The output `RecoverabilityReport`
650 contains `compression_ratio`, `psnr_prediction`, `feasibility`, and `null_space_dim` where
651 available.

652 **MismatchAgent.** Scores the mismatch severity for a given imaging configuration us-
653 ing `mismatch_db.yaml` (797 lines). For each modality, the database enumerates the rel-
654 evant mismatch parameters, their physical units, typical perturbation ranges, and avail-
655 able correction algorithms. The output `MismatchReport` includes `severity` (float, 0–1),
656 `correction_method` (string), `expected_gain_db` (float), and `dominant_parameter` (string).

657 **AnalysisAgent.** The bottleneck classifier. It receives reports from the Photon, Recover-
658 ability, and Mismatch agents, computes the gate costs (Equations (2) to (4)), identifies the
659 dominant gate, and generates actionable suggestions. The `AnalysisAgent` also computes
660 the recovery ratio ρ and its bootstrap confidence interval.

661 **AgentNegotiator.** Implements a cross-agent veto protocol. Before reconstruction is au-
662 thorized, the negotiator inspects all three upstream reports and applies three veto con-
663 ditions: (1) low photon budget combined with aggressive compression (C_{noise} and C_{recover}
664 both large); (2) severe mismatch (`severity` > 0.7) without a planned correction step; (3) joint
665 probability below the floor threshold ($p_{\text{joint}} < 0.15$), indicating that all three subsystems
666 are simultaneously marginal. When any veto fires, reconstruction halts with an actionable
667 explanation and suggested remediation.

668 **HybridAgent.** An optional wrapper that invokes an LLM for natural-language narra-
669 tive generation or edge-case modality mapping. All quantitative decisions remain on the
670 deterministic code path; the `HybridAgent` is never required for pipeline operation.

671 **Support classes.** The remaining components include: `AssetManager` (file I/O and caching
672 for large arrays), `ContinuityChecker` (verifies that sequential pipeline outputs are dimen-
673 sionally consistent), `SystemDiscern` (auto-detects modality from uploaded data), `PreflightChecker`
674 (validates the complete experiment configuration before execution), `WhatIfPrecomputer`
675 (evaluates counterfactual what-if scenarios), `SelfImprovement` (logs diagnostic events for
676 future registry refinement), `PhysicsStageVisualizer` (generates intermediate visualiza-
677 tions at each pipeline stage), and UPWMI (Universal Physics World Model Interface, the
678 top-level entry point that wires all agents together).

679 **Contract system.** Inter-agent communication uses 25 Pydantic v2 contract models. All
680 contracts inherit from `StrictBaseModel`, which enforces `extra="forbid"` (no unexpected

fields), `validate_assignment=True` (mutations re-validated), and a model validator that rejects NaN and Inf in any float field. Bounded scores use `Field(ge=0.0, le=1.0)`. Enums are string enums for human-readable JSON serialization. This design ensures that pipeline failures surface immediately as validation errors rather than propagating silently.

YAML registries. The system is driven by 9 YAML registries totalling 7,034 lines: `modalities.yaml` (modality definitions), `graph.templates.yaml` (OperatorGraph skeletons), `photon_db.yaml` (photon models), `mismatch_db.yaml` (mismatch parameters and correction methods), `compression_db.yaml` (recoverability tables with provenance), `solver_registry.yaml` (solver configurations), `primitives.yaml` (primitive operator metadata), `dataset_registry.yaml` (dataset locations and formats), and `acceptance_thresholds.yaml` (pass/fail thresholds per metric).

Correction Algorithms

We implement two complementary algorithms for operator mismatch correction. Crucially, both algorithms operate on the forward operator parameters θ rather than the reconstruction solver weights, making them *solver-agnostic*: the corrected operator $H(\hat{\theta})$ benefits any downstream solver (GAP-TV, MST-L, HDNet²¹, CST, *etc.*) without retraining.

Algorithm 1: Hierarchical Beam Search. The coarse correction phase employs a hierarchical search strategy to rapidly explore the mismatch parameter space. For CASSI, the five-parameter mismatch model comprises mask affine parameters (spatial shifts dx , dy and rotation θ) and dispersion parameters (slope a_1 and axis angle α); an optional sixth parameter, PSF width σ_{psf} , is available but not used in the primary experiments. The algorithm proceeds as follows:

1. **1D sweeps.** Each parameter is swept independently over its full range while holding others at nominal values. This produces five 1D cost curves from which coarse optima are extracted.
2. **3D beam search.** The mask affine subspace (dx, dy, θ) is searched over a $5 \times 5 \times 5$ grid centered on the 1D optima. The top- k ($k = 5$) candidates by reconstruction PSNR are retained.
3. **2D beam search.** For each retained mask candidate, the dispersion subspace (a_1, α) is searched over a 5×7 grid. The joint top- k candidates are retained.
4. **Coordinate descent refinement.** Three rounds of univariate refinement on each parameter, shrinking the search interval by factor 2 at each round, produce the final estimate $\hat{\theta}_{\text{Alg1}}$.

714 Total runtime is approximately 300 seconds per scene on a single GPU. Accuracy is
 715 ± 0.1 – 0.2 pixels for spatial parameters and $\pm 0.05^\circ$ for angular parameters.

716 **Algorithm 2: Joint Gradient Refinement.** The fine correction phase uses a differen-
 717 tiable forward model to jointly optimize all mismatch parameters via gradient descent. The
 718 key components are:

- 719 1. **Differentiable mask warp.** The binary mask is warped by a continuous affine
 720 transformation using bilinear interpolation, implemented as a custom PyTorch module
 721 (`DifferentiableMaskWarpFixed`). The mask values are passed through a straight-
 722 through estimator (STE) to maintain binary structure while permitting gradient flow.
- 723 2. **Differentiable forward model.** The CASSI forward model $y = \text{CASSI}(x; \theta)$ is
 724 implemented as a differentiable PyTorch module (`DifferentiableCassiForwardSTE`)
 725 that accepts mismatch parameters as differentiable inputs.
- 726 3. **GPU grid initialization.** A full-range 3D grid search over (dx, dy, θ) with $9 \times 9 \times 7 =$
 727 567 points provides diverse starting candidates. The top 9 candidates seed multi-start
 728 gradient refinement.
- 729 4. **Staged gradient refinement.** Each of the 9 candidates is refined using Adam
 730 optimization (learning rate 10^{-2} , decaying to 10^{-3}) for 200 steps. For each candidate,
 731 4 random restarts with jittered initialization guard against local minima. The loss
 732 function is the negative PSNR computed via an unrolled K -iteration differentiable
 733 GAP-TV solver (`DifferentiableGAPTV`, $K = 10$ unrolled iterations).

734 Total runtime for Algorithm 2 is approximately 3,200 seconds ($200 \text{ steps} \times 4 \text{ restarts} \times$
 735 $9 \text{ candidates with early stopping}$). Accuracy improves to ± 0.05 – 0.1 pixels, a 3 – $5\times$ improve-
 736 ment over Algorithm 1. The two algorithms are used sequentially in practice: Algorithm 1
 737 provides a warm start, and Algorithm 2 refines to sub-pixel precision.

738 Evaluation Protocol

739 **Four-Scenario Protocol.** We evaluate every modality under four standardized scenarios
 740 that isolate different sources of quality degradation:

741 **Scenario I (Ideal):** $\mathbf{y}_{\text{obs}} = H_{\text{true}} \mathbf{x}_{\text{gt}}$; reconstruct with H_{true} . In this scenario the system
 742 is perfectly calibrated ($H_{\text{true}} = H_{\text{nom}}$), so the operator used for reconstruction matches
 743 the one that generated the data. This yields the oracle upper bound on reconstruction
 744 quality, limited only by the sensing geometry and solver convergence.

745 **Scenario II (Mismatch):** $\mathbf{y}_{\text{obs}} = H_{\text{true}} \mathbf{x}_{\text{gt}}$; reconstruct with H_{nom} ($H_{\text{nom}} \neq H_{\text{true}}$). This
 746 is the standard operating condition in practice: the measurement is generated by the

747 true physics, but the reconstruction uses a nominal (potentially mismatched) forward
 748 model.

749 **Scenario III (Corrected):** $\mathbf{y}_{\text{obs}} = H_{\text{true}} \mathbf{x}_{\text{gt}}$; reconstruct with $\hat{H} = H(\hat{\theta})$ where $\hat{\theta}$ is
 750 estimated by Algorithms 1 and 2. This quantifies the benefit of mismatch calibration.

751 **Scenario IV (Oracle Mask):** Same measurements as Scenario II ($\mathbf{y}_{\text{obs}} = H_{\text{true}} \mathbf{x}_{\text{gt}}$ with
 752 $H_{\text{true}} \neq H_{\text{nom}}$); reconstruct with H_{true} instead of H_{nom} . Provides the correction
 753 ceiling: the best reconstruction achievable when the true operator is known exactly,
 754 applied to data that were sensed by the mismatched system. The gap between Sce-
 755 nario IV and Scenario I reveals the irreducible loss from the degraded sensing config-
 756 uration itself (e.g., a shifted mask pattern is suboptimal even when perfectly known).

757 **Metrics.** Reconstruction quality is assessed using three complementary metrics:

- 758 • **PSNR** (peak signal-to-noise ratio, in dB): the primary metric, computed per scene
 759 and averaged. For signals normalized to $[0, 1]$, $\text{PSNR} = 10 \log_{10}(1/\text{MSE})$. For SPC
 760 data normalized to $[0, 255]$, the peak value is 255.
- 761 • **SSIM** (structural similarity index): captures perceptual quality including luminance,
 762 contrast, and structural components, computed with a Gaussian window of width 11
 763 and standard deviation 1.5.
- 764 • **SAM** (spectral angle mapper): for hyperspectral modalities (CASSI), measures the
 765 angle between predicted and true spectral vectors at each spatial location, reported
 766 in degrees. Lower is better.

767 **Datasets.**

- 768 • **CASSI:** 10 scenes from the KAIST dataset⁹, each a $256 \times 256 \times 28$ spectral cube (28
 769 spectral bands from 450 nm to 650 nm). Data range $[0, 1]$.
- 770 • **CACTI:** 6 benchmark videos, each $256 \times 256 \times 8$ (8 temporal frames encoded per
 771 snapshot). Data range $[0, 1]$.
- 772 • **SPC:** 11 natural images from the Set11 benchmark, each 256×256 grayscale. Data
 773 range $[0, 255]$.

774 All per-scene metrics are reported individually as well as averaged, and all reconstruction
 775 arrays are saved as NumPy NPZ files.

Experimental Details

Hardware. All experiments are conducted on a single NVIDIA GPU. Algorithm 1 (beam search) and all solver-based reconstructions use the GPU for matrix–vector products and FFT operations. Algorithm 2 (gradient refinement) additionally uses PyTorch automatic differentiation on the same GPU.

CASSI configuration. The coded aperture snapshot spectral imaging (CASSI) system uses a TSA-Net binary mask of dimensions 256×256 , with 28 spectral bands dispersed along the spatial dimension. The five-parameter mismatch model $\boldsymbol{\theta} = (dx, dy, \theta, a_1, \alpha)$ describes: mask spatial shift in x (dx , pixels), mask spatial shift in y (dy , pixels), mask rotation angle (θ , degrees), dispersion slope (a_1 , pixels per band), and dispersion axis angle (α , degrees). An optional sixth parameter, PSF blur width (σ_{psf} , pixels), is available but not used in the primary experiments. For the primary mismatch experiment (validated by InverseNet), the true mismatch parameters are $\boldsymbol{\theta}_{\text{true}} = (dx = 0.5 \text{ px}, dy = 0.3 \text{ px}, \theta = 0.1^\circ, a_1 = 2.02, \alpha = 0.15^\circ)$. Solvers evaluated include TwIST¹⁰, GAP-TV²⁰, DGSMP²⁴, MST-L⁸, and CST-L²⁵, all of which receive the same operator and differ only in their reconstruction algorithm. The supplementary per-scene analysis additionally includes DeSCI²⁶ and HDNet²¹.

CACTI configuration. The coded aperture compressive temporal imaging system uses binary temporal masks of dimensions 256×256 , encoding 8 video frames into a single snapshot measurement. Mismatch is parameterized as a temporal mask timing offset (sub-frame shift). The default solver is EfficientSCI²².

SPC configuration. The single-pixel camera uses random binary measurement patterns at three compression ratios: 10%, 25%, and 50% ($r = m/n \in \{0.10, 0.25, 0.50\}$). Mismatch is modeled as a multiplicative gain bias on the measurement matrix. The default solver is ADMM-TV with total-variation regularization and a wavelet sparsifying transform.

MRI configuration. Cartesian k -space sampling with $4\times$ acceleration (25% of k -space lines acquired). Mismatch is parameterized as a 5% multiplicative error in the coil sensitivity maps used for parallel imaging reconstruction. The default solver is SENSE¹⁷ with ℓ_1 -wavelet regularization.

CT configuration. Fan-beam geometry with 180 projections over 180° . Mismatch is modeled as a center-of-rotation (CoR) offset, which produces characteristic arc artifacts in the reconstruction. The default solver is filtered back-projection (FBP)¹⁸ with a Ram-Lak filter, supplemented by iterative SART for comparison.

CASSI real-data configuration. The TSA real hyperspectral dataset¹ consists of 5 scenes at 660×660 spatial resolution with 28 spectral bands and mask-shift step 2. Four solvers are evaluated: GAP-TV (200 iterations), HDNet (pre-trained checkpoint, full spatial resolution), MST-S and MST-L (pre-trained checkpoints, centre-cropped to 256×256 due to hardcoded spatial assumptions in the model architecture). The coded aperture mask is perturbed by $dx = 0.5$ px, $dy = 0.3$ px to simulate assembly-induced mismatch. No ground truth is available; quality is assessed via the normalised measurement residual $r = \|\mathbf{y} - H\hat{\mathbf{x}}\|^2 / \|\mathbf{y}\|^2$.

CACTI real-data configuration. The EfficientSCI real temporal dataset²² consists of 4 dynamic scenes (duomino, hand, pendulumBall, waterBalloon) at 512×512 with compression ratio 10. The real mask is stored separately from the measurement data. Two solvers are evaluated: GAP-TV (50 iterations) and PnP-FFDNet (50 iterations with FFDNet denoiser). Mismatch is induced by shifting the mask by $dx = 0.5$ px, $dy = 0.3$ px. Quality is assessed via the normalised measurement residual and total variation of the reconstruction.

Controlled hardware experiment protocol. The software-perturbation protocol above applies calibrated mask shifts to existing real measurements. A full hardware-in-the-loop validation requires physically displacing the coded aperture mask and re-acquiring data. The protocol proceeds as follows: (i) acquire a baseline dataset with the mask at its factory-calibrated position; (ii) physically translate the mask by a known displacement ($\Delta x \in \{0.25, 0.5, 1.0\}$ px equivalent, verified by micrometer stage) and re-acquire under identical illumination; (iii) reconstruct both datasets with the factory mask specification and compute the PSNR degradation and measurement residual; (iv) apply PWM autonomous calibration and measure recovery. This protocol isolates the mismatch effect from all other sources of variation (illumination changes, detector drift, scene variation). Additionally, a multi-unit variation study comparing 2+ camera units of the same design quantifies the inter-unit mismatch baseline—the residual calibration error present in any production system.

Clinical CT phantom configuration. For clinical translation, PWM is evaluated on CT quality assurance using the ACR CT accreditation phantom (Gammex 464)¹². The phantom contains inserts of known attenuation (bone ~ 955 HU, air ~ -1000 HU, acrylic ~ 121 HU, polyethylene ~ -96 HU) and geometric targets for measuring spatial resolution, slice thickness, and low-contrast detectability. Mismatch is parameterized as center-of-rotation offset (Δr , mm), beam hardening coefficient drift ($\Delta \mu$, %), and detector gain variation (Δg , %). Nine ACR-aligned metrics are computed automatically: CT number accuracy (5 materials), geometric accuracy (± 2 mm tolerance), slice thickness (± 1.5 mm), uniformity (≤ 5 HU), noise standard deviation, spatial resolution (≥ 5 lp/cm), low-contrast detectability (≥ 4 targets), and artifact score.

Clinical MRI validation configuration. For MRI clinical validation, PWM processes multi-coil k -space data from public datasets (fastMRI¹¹). Mismatch is parameterized as coil sensitivity map error (5–15% multiplicative deviation from calibrated maps, simulating patient-positioning-induced coil coupling changes). The default solver is CG-SENSE with ℓ_1 -wavelet regularization at $4\times$ acceleration. Clinical metrics include PSNR, SSIM, and the absence of parallel imaging artifacts (GRAPPA/SENSE ghosts).

Statistical Analysis

Per-scene reporting. All metrics are reported per scene, not merely as dataset averages. This enables identification of scene-dependent failure modes (*e.g.*, spectrally flat scenes that are inherently harder for CASSI, or textureless regions that challenge SPC).

Summary statistics. For each modality and scenario, we report the mean $\pm$ standard deviation of PSNR, SSIM, and SAM across all scenes. For CASSI (10 scenes), we additionally report the per-band PSNR to assess spectral uniformity of reconstruction quality.

Recovery ratio confidence intervals. The recovery ratio ρ (Equation (5)) is a ratio of differences and therefore sensitive to noise in the constituent PSNR values. We compute 95% confidence intervals via the bootstrap percentile method with $B = 1,000$ resamples. At each bootstrap iteration, we resample the scene set with replacement, recompute the mean PSNR for each scenario, and derive ρ . The 2.5th and 97.5th percentiles of the bootstrap distribution define the 95% CI.

Parameter recovery accuracy. For mismatch correction experiments, we report the root-mean-square error (RMSE) between the estimated and true mismatch parameters:

$$\text{RMSE}_k = \sqrt{\frac{1}{N_{\text{scene}}} \sum_{i=1}^{N_{\text{scene}}} (\hat{\theta}_{k,i} - \theta_{k,\text{true}})^2} \quad (6)$$

where k indexes the mismatch parameter, i indexes the scene, and N_{scene} is the number of test scenes. Uncertainty in the RMSE is estimated via bootstrap ($B = 1,000$).

Ablation significance. Ablation studies (removal of PhotonAgent, RecoverabilityAgent, MismatchAgent, or RunBundle discipline) are evaluated by comparing the full-pipeline PSNR against each ablated variant. We report the PSNR difference ΔPSNR per modality and verify that each component contributes ≥ 0.5 dB across all validated modalities, establishing practical significance.

Code and Data Availability

Source code. The complete PWM framework, including all agents, the OperatorGraph compiler, correction algorithms, YAML registries, and evaluation scripts, is released as open-source software under the MIT license at https://github.com/integritynoble/Physics_World_Model. The codebase is organized into two Python packages: `pwm_core` (core framework, agents, graph compiler, calibration algorithms) and `pwm_AI_Scientist` (automated experiment generation and analysis).

Reconstruction data. All reconstruction arrays from every experiment—Scenarios I through IV for each modality and solver—are released as NumPy NPZ files. Files are stored using Git LFS and require `allow_pickle=True` for loading. Data ranges are standardized: CASSI and CACTI reconstructions are normalized to $[0, 1]$; SPC reconstructions are in $[0, 255]$.

Experiment manifests. Every experiment is recorded in a RunBundle v0.3.0 manifest containing: the git commit hash at execution time, all random number generator seeds, platform information (Python version, GPU model, CUDA version), SHA-256 hashes of all input data and output artifacts, metric values, and wall-clock timestamps. These manifests enable exact reproduction of every reported result.

Registry data. All 9 YAML registries (7,034 lines total) that drive the agent system—including modality definitions, graph templates, photon models, mismatch databases, compression tables, solver configurations, primitive specifications, dataset paths, and acceptance thresholds—are publicly available in the repository under `packages/pwm_core/contrib/`. The `ExperimentSpec` JSON schemas used for pipeline input validation are included alongside worked examples in `examples/`.

References

- [1] Ashwin A. Wagadarikar, Renu John, Rebecca Willett, and David J. Brady. Single disperser design for coded aperture snapshot spectral imaging. *Applied Optics*, 47(10): B44–B51, 2008. doi: 10.1364/AO.47.000B44.
- [2] Patrick Llull, Xuejun Liao, Xin Yuan, Jianbo Yang, David Kittle, Lawrence Carin, Guillermo Sapiro, and David J. Brady. Coded aperture compressive temporal imaging. *Optics Express*, 21(9):10526–10545, 2013. doi: 10.1364/OE.21.010526.
- [3] Michael Lustig, David Donoho, and John M. Pauly. Sparse MRI: The application of compressed sensing for rapid MR imaging. *Magnetic Resonance in Medicine*, 58(6): 1182–1195, 2007. doi: 10.1002/mrm.21391.

- [4] Emmanuel J. Candès and Michael B. Wakin. An introduction to compressive sampling. *IEEE Signal Processing Magazine*, 25(2):21–30, 2008. doi: 10.1109/MSP.2007.914731.
- [5] David L. Donoho. Compressed sensing. *IEEE Transactions on Information Theory*, 52(4):1289–1306, 2006. doi: 10.1109/TIT.2006.871582.
- [6] Singanallur V. Venkatakrishnan, Charles A. Bouman, and Brendt Wohlberg. Plug-and-play priors for model based reconstruction. In *Proceedings of the IEEE Global Conference on Signal and Information Processing (GlobalSIP)*, pages 945–948, 2013. doi: 10.1109/GlobalSIP.2013.6737048.
- [7] Vishal Monga, Yuelong Li, and Yonina C. Eldar. Algorithm unrolling: Interpretable, efficient deep learning for signal and image processing. *IEEE Signal Processing Magazine*, 38(2):18–44, 2021. doi: 10.1109/MSP.2020.3016905.
- [8] Yuanhao Cai, Jing Lin, Xiaowan Hu, Haoqian Wang, Xin Yuan, Yulun Zhang, Radu Timofte, and Luc Van Gool. Mask-guided spectral-wise transformer for efficient hyperspectral image reconstruction. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 17502–17511, 2022.
- [9] Inchang Choi, Daniel S. Jeon, Giljoo Nam, Diego Gutierrez, and Min H. Kim. High-quality hyperspectral reconstruction using a spectral prior. *ACM Transactions on Graphics (Proceedings of SIGGRAPH Asia)*, 36(6):218:1–218:13, 2017. doi: 10.1145/3130800.3130810.
- [10] José M. Bioucas-Dias and Mário A. T. Figueiredo. A new TwIST: Two-step iterative shrinkage/thresholding algorithms for image restoration. *IEEE Transactions on Image Processing*, 16(12):2992–3004, 2007. doi: 10.1109/TIP.2007.909319.
- [11] Jure Zbontar, Florian Knoll, Anuroop Sriram, Tullie Murrell, Zhengnan Huang, Matthew J. Muckley, Aaron Defazio, Ruben Stern, Patricia Johnson, Mary Bruno, Marc Parente, Krzysztof J. Geras, Joe Katsnelson, Hersh Chandarana, Zizhao Zhang, Michal Drozdal, Adriana Romero, Michael Rabbat, Pascal Vincent, Nafissa Yakubova, James Pinkerton, Duo Wang, Erich Owens, C. Lawrence Zitnick, Michael P. Recht, Daniel K. Sodickson, and Yvonne W. Lui. fastMRI: An open dataset and benchmarks for accelerated MRI. *arXiv preprint arXiv:1811.08839*, 2018.
- [12] Hu Chen, Yi Zhang, Mannudeep K. Kalra, Feng Lin, Yang Chen, Peixi Liao, Jiliu Zhou, and Ge Wang. Low-dose CT with a residual encoder-decoder convolutional neural network. *IEEE Transactions on Medical Imaging*, 36(12):2524–2535, 2017. doi: 10.1109/TMI.2017.2715284.
- [13] Martin Uecker, Peng Lai, Mark J. Murphy, Patrick Virtue, Michael Elad, John M. Pauly, Shreyas S. Vasanawala, and Michael Lustig. ESPIRiT — an eigenvalue approach

- 941 to autocalibrating parallel MRI: where SENSE meets GRAPPA. *Magnetic Resonance*
942 *in Medicine*, 71(3):990–1001, 2014. doi: 10.1002/mrm.24751.
- 943 [14] Andrew M. Maiden and John M. Rodenburg. An improved ptychographical phase
944 retrieval algorithm for diffractive imaging. *Ultramicroscopy*, 109(10):1256–1262, 2009.
945 doi: 10.1016/j.ultramic.2009.05.012.
- 946 [15] Vegard Antun, Francesco Renna, Clarice Poon, Ben Adcock, and Anders C. Hansen.
947 On instabilities of deep learning in image reconstruction and the potential costs of
948 AI. *Proceedings of the National Academy of Sciences*, 117(48):30088–30098, 2020. doi:
949 10.1073/pnas.1907377117.
- 950 [16] J. M. Rodenburg and H. M. L. Faulkner. A phase retrieval algorithm for shifting
951 illumination. *Applied Physics Letters*, 85(20):4795–4797, 2004. doi: 10.1063/1.1823034.
- 952 [17] Klaas P. Pruessmann, Markus Weiger, Markus B. Scheidegger, and Peter Boesiger.
953 SENSE: Sensitivity encoding for fast MRI. *Magnetic Resonance in Medicine*, 42(5):
954 952–962, 1999. doi: 10.1002/(SICI)1522-2594(199911)42:5<952::AID-MRM16>3.0.CO;
955 2-S.
- 956 [18] L. A. Feldkamp, L. C. Davis, and J. W. Kress. Practical cone-beam algorithm. *Jour-*
957 *nal of the Optical Society of America A*, 1(6):612–619, 1984. doi: 10.1364/JOSAA.1.
958 000612.
- 959 [19] Chengshuai Yang. InverseNet: Benchmarking operator mismatch in snapshot compres-
960 sive imaging. Technical report, NextGen PlatformAI C Corp, 2026.
- 961 [20] Xin Yuan. Generalized alternating projection based total variation minimization for
962 compressive sensing. In *Proceedings of the IEEE International Conference on Image*
963 *Processing (ICIP)*, pages 2539–2543, 2016. doi: 10.1109/ICIP.2016.7532817.
- 964 [21] Xiaowan Hu, Yuanhao Cai, Jing Lin, Haoqian Wang, Xin Yuan, Yulun Zhang, Radu
965 Timofte, and Luc Van Gool. HDNet: High-resolution dual-domain learning for spectral
966 compressive imaging. In *Proceedings of the IEEE/CVF Conference on Computer Vision*
967 *and Pattern Recognition (CVPR)*, pages 17542–17551, 2022.
- 968 [22] Lishun Wang, Miao Cao, and Xin Yuan. EfficientSCI: Densely connected network with
969 space-time factorization for large-scale video snapshot compressive imaging. In *Pro-*
970 *ceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*
971 *(CVPR)*, pages 18477–18486, 2023.
- 972 [23] Marco F. Duarte, Mark A. Davenport, Dharmpal Takhar, Jason N. Laska, Ting Sun,
973 Kevin F. Kelly, and Richard G. Baraniuk. Single-pixel imaging via compressive sam-
974 pling. *IEEE Signal Processing Magazine*, 25(2):83–91, 2008. doi: 10.1109/MSP.2007.
975 914730.

- 976 [24] Tao Huang, Weisheng Dong, Xin Yuan, Jinjian Wu, and Guangming Shi. Deep gaussian
977 scale mixture prior for spectral compressive imaging. In *Proceedings of the IEEE/CVF*
978 *Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 16216–16225,
979 2021.
- 980 [25] Yuanhao Cai, Jing Lin, Xiaowan Hu, Haoqian Wang, Xin Yuan, Yulun Zhang, Radu
981 Timofte, and Luc Van Gool. CST: Compact spectral transformer for hyperspectral
982 image reconstruction. In *Proceedings of the European Conference on Computer Vision*
983 *(ECCV)*, 2022.
- 984 [26] Yang Liu, Xin Yuan, Jinli Suo, David J. Brady, and Qionghai Dai. Rank minimiza-
985 tion for snapshot compressive imaging. *IEEE Transactions on Pattern Analysis and*
986 *Machine Intelligence*, 41(12):2990–3006, 2019.




Figure 1: **PWM overview.** The Physics World Models pipeline. **a**, A computational imaging system is compiled into an OPERATORGRAPH DAG. **b**, The TRIAD DECOMPOSITION diagnostic agents evaluate each gate. **c**, The dominant gate is identified and a TRIADREPORT is produced. **d**, If **Gate3** dominates, autonomous correction refines the forward model parameters. **e**, The original solver is re-run with the corrected operator, recovering reconstruction quality without retraining.




Figure 2: **OperatorGraph IR and Physics Fidelity Ladder.** **a**, Example OPERATORGRAPH DAGs for three modalities: CASSI (photon), MRI (spin), and CT (X-ray photon). Each node wraps a primitive operator; edges define data flow. **b**, The Physics Fidelity Ladder. Tier 1: linear shift-invariant. Tier 2: linear shift-variant. Tier 3: nonlinear ray/wave-based. Tier 4: full-wave/Monte Carlo. **c**, Summary statistics: 26 validated modalities, 7 end-to-end correction, 5 physical carriers.




Figure 3: **Triad Decomposition structure and gate binding.** **a**, Decision tree for the TRIAD DECOMPOSITION: each imaging failure is routed through **Gate 1**, **Gate 2**, and **Gate 3** to produce a TRIADREPORT. **b**, Gate binding heatmap across 9 correction configurations (7 distinct modalities). Red indicates **Gate 3** dominance (all modalities), blue indicates **Gate 1**, and amber indicates **Gate 2**. **c**, Recovery ratio ρ distribution across all 9 correction configurations.




Figure 4: **Correction results across 9 validated configurations.** Bar chart showing correction gain Δ_{corr} (dB) for each of the 9 correction configurations (7 distinct modalities), grouped by carrier family. Incoherent photon (CASSI, CACTI, SPC, Lensless) and coherent photon (Ptychography) in blue; spin (MRI) in purple; X-ray (CT) in red; generic (Matrix) in grey.




Figure 5: **CASSI and CACTI deep dive.** **a**, CASSI: PSNR across 4 scenarios for GAP-TV, MST-L, and HDNet under combined mask-geometry-plus-dispersion mismatch. The uniform collapse under Scenario II (range 20.83–21.88 dB) confirms operator-driven failure; oracle recovery varies by solver ($\rho = 0.22$ –0.46). **b**, CACTI: four methods across 4 scenarios, showing up to 20.58 dB mismatch degradation and substantial correction recovery across all solvers.




Figure 6: **Zero-shot generalization across carrier families.** Correction gain (dB) when beam-search and gradient-refinement hyperparameters are tuned on photon-domain modalities and transferred without modification to coherent-photon, spin, and X-ray domains. Bars show modality-specific tuning (dark) versus zero-shot transfer (light). Transfer efficiency is high across all carrier families, confirming the carrier-agnostic nature of the PWM correction pipeline.




Figure 7: **Hardware validation on real CASSI and CACTI instruments.** **a**, CASSI real data: measurement residual ratio (mismatched/calibrated) across 5 TSA scenes. GAP-TV shows $1.8\times$ mean ratio. **b**, CACTI real data: residual ratio across 4 scenes. GAP-TV shows $10.4\times$ mean ratio; PnP-FFDNet shows $2.0\times$. **c**, Simulation-to-hardware gap: comparing mismatch degradation in simulation versus real hardware for CASSI and CACTI, illustrating that real instruments have pre-existing calibration errors that attenuate the marginal impact of additional perturbations. **d**, Autonomous calibration: grid-search parameter recovery for CASSI (78%), CACTI (100%), and SPC (0%, gain-type mismatch not recoverable from residual).