

Ten Primitives and Three Gates: The Universal Structure of Computational Imaging

Chengshuai Yang^{1,*} Xin Yuan²

Abstract

Computational imaging systems routinely underperform because the assumed forward model diverges from the true physics. Here we prove two results. First, the **Finite Primitive Basis Theorem**: every linear, shift-variant imaging forward model admits an ε -approximate representation as a typed directed acyclic graph over exactly 10 canonical primitives. Second, the **Triad Decomposition**: every reconstruction failure decomposes into three root causes—information deficiency, carrier noise, and operator mismatch—with mismatch dominant across all validated modalities. Together these results yield a modality-agnostic diagnostic and correction framework. Across seven modalities spanning three carrier families (optical photons, X-ray photons, and nuclear spins), autonomous correction recovers +0.8 to +10.7 dB of mismatch-induced degradation without retraining the solver. Hardware validation on real instruments confirms mismatch dominance. A held-out closure test on 8 additional modalities confirms basis completeness with saturating growth.

Introduction

Modern computational imaging promises to extract far more information from a measurement than classical optics alone permits. Coded aperture spectral cameras compress three-dimensional scenes into a single snapshot¹; compressive temporal imagers freeze high-speed video into one exposure²; single-pixel cameras reconstruct images from far fewer measurements than pixels³; accelerated MRI scanners reconstruct diagnostic-quality images from a fraction of the acquired k -space data⁴. Over the past decade, the community has invested enormous effort in improving reconstruction algorithms—progressing from compressed sensing^{5,6} and plug-and-play priors⁷ to deep unrolling networks⁸ and vision transformers⁹—yet these algorithms routinely fail when deployed on real instruments, because the forward model assumed by the solver does not match the physics that generated the data.

These failures persist because computational imaging lacks two theoretical foundations. First, a **representation theory**: no formal result guarantees that a finite set of primitive operators suffices to represent all imaging forward models. Without this, every new modality

^{*1}NextGen PlatformAI C Corp, USA. ²School of Engineering, Westlake University, Hangzhou, China. ^{*}Correspondence: integrityyang@gmail.com

requires bespoke engineering and the field cannot reason about completeness. Second, a **diagnostic theory**: no systematic framework decomposes reconstruction failures into root causes. Calibration errors, noise, and information deficiency are conflated, and failures are addressed *ad hoc*. Calibration methods exist for specific instruments^{10,11}, but they do not generalise; robustness studies perturb individual systems¹², but they lack a unifying formalism.

Here we establish both foundations within Physics World Models (PWM). We prove the **Finite Primitive Basis Theorem** (Theorem 3): every imaging forward model in the operator class $\mathcal{C}_{\text{Tier2}}$ —encompassing all clinical, scientific, and industrial modalities at Tier-2 fidelity—admits an ε -approximate representation as a typed DAG over exactly 10 canonical primitives. We establish the **Triad Decomposition**: every reconstruction failure decomposes into three gates—information deficiency (**Gate 1**), carrier noise (**Gate 2**), and operator mismatch (**Gate 3**)—with **Gate 3** dominant across all validated modalities. These two results are complementary: the Finite Primitive Basis provides a universal representation for forward models; the Triad Decomposition provides a universal diagnostic law over that representation. Together they yield a practical framework that diagnoses and corrects imaging failures across modalities without modality-specific tuning.

We validate both contributions across seven modalities spanning three carrier families—optical photons (CASSI, CACTI, SPC, lensless), coherent photons (ptychography¹³), X-ray photons (CT¹⁴), and nuclear spins (MRI¹⁵)—with hardware validation on real DD-CASSI and CACTI instruments, confirming Triad predictions on physical measurements. The OperatorGraph library additionally registers templates for electron, acoustic, and particle modalities (Extended Data Table 2), pending future validation. A held-out closure test on 8 additional modalities—including quantum ghost imaging, THz time-domain spectroscopy, and Compton scatter imaging—confirms basis completeness. The proof of Theorem 3 is in Supplementary Note 10; formal primitive semantics are developed in¹⁶.

The Finite Primitive Basis

A foundational question in computational imaging is whether the diversity of imaging forward models—from coded aperture cameras to MRI scanners to electron microscopes—can be captured by a small, fixed set of primitive operators. We answer affirmatively.

The OperatorGraph representation. Every imaging forward model is encoded as a typed directed acyclic graph (DAG) in which each node wraps a single primitive physical operator and edges define the data flow from source to detector. Every primitive implements both a `forward()` and an `adjoint()` method, with validated adjoint consistency ensuring $\langle H\mathbf{x}, \mathbf{y} \rangle = \langle \mathbf{x}, H^\dagger \mathbf{y} \rangle$ to within numerical precision. Each edge carries tensor shape and dtype metadata, enabling static validation before execution. We call this formalism the

68 OPERATORGRAPH intermediate representation (IR).

69 **Ten canonical primitives.** The OPERATORGRAPH IR defines a library of exactly 10
70 canonical primitives $\mathcal{B} = \{P, M, \Pi, F, C, \Sigma, D, S, W, R\}$:

#	Primitive	Notation	Physical action
1	Propagate	$P(d, \lambda)$	Free-space wave propagation
2	Modulate	$M(\mathbf{m})$	Element-wise multiplication (mask, coil, absorption)
3	Project	$\Pi(\theta)$	Radon line-integral projection
4	Encode	$F(\mathbf{k})$	Fourier-domain encoding (k -space)
5	Convolve	$C(\mathbf{h})$	Spatial convolution (PSF)
6	Accumulate	Σ	Summation over spectral/temporal axis
7	Detect	$D(g, \eta)$	Detector response (5 canonical families)
8	Sample	$S(\Omega)$	Sub-sampling on index set Ω
9	Disperse	$W(\alpha, a)$	Wavelength-dependent spatial shift
10	Scatter	$R(\sigma, \Delta\varepsilon)$	Direction change and/or energy shift

71 The Detect nonlinearity η is restricted to five canonical families (linear-field: $\eta(x) = gx$;
72 logarithmic; sigmoid; intensity-square-law: $\eta(x) = g|x|^2$; coherent-field), each with at most
73 2 scalar parameters, preventing Detect from becoming a universal approximator (see¹⁶ for
74 formal definitions and non-universality proof).

75 **Physics-stage mapping.** Each factor of a forward model is classified into one of five
76 physics-stage families and mapped to primitives from $\mathcal{B}$: propagation $\rightarrow \{P, C\}$; elastic in-
77 teraction $\rightarrow \{M\}$; inelastic interaction (scattering) $\rightarrow \{R\}$; encoding–projection $\rightarrow \{\Pi, F\}$;
78 detection–readout $\rightarrow \{\Sigma, S, W, C, D\}$.

79 **Physics Fidelity Ladder.** The OPERATORGRAPH IR defines a four-tier Physics Fidelity
80 Ladder: Tier 1 (linear, shift-invariant), Tier 2 (linear, shift-variant), Tier 3 (nonlinear,
81 ray/wave-based), and Tier 4 (full-wave/Monte Carlo). The Finite Primitive Basis applies to
82 Tiers 1–2; Tiers 3–4 are handled by refinement sub-DAGs that expand individual primitive
83 nodes into higher-fidelity sub-graphs while preserving the top-level DAG structure.

84 **Five physical carriers.** The library spans five carrier families—photons, electrons, spins,
85 acoustic waves, and particles—with 26 registered modality templates (7 with full end-to-end
86 correction validation; Extended Data Table 2; Supplementary Table S3).

87 **Theorem statement.**

88 **Definition 1** (Tier-2 Operator Class). $\mathcal{C}_{\text{Tier2}}$ consists of all imaging forward models express-
89 ible as a finite sequential-parallel composition of linear, shift-variant stages with bounded
90 operator norm ($\|H_k\| \leq B$) and stage count $K \leq N_{\text{max}}$.

91 **Definition 2** (ε -Approximate Representation). A typed DAG G with nodes $V \subseteq \mathcal{B}$ is an
 92 ε -approximate representation of $H \in \mathcal{C}_{\text{Tier2}}$ if $\sup_{\|\mathbf{x}\| \leq 1} \|H(\mathbf{x}) - H_G(\mathbf{x})\| / (\|H(\mathbf{x})\| + \delta) \leq \varepsilon$,
 93 where $\delta > 0$ is a regularization constant, and $|V| \leq N_{\max}$, $\text{depth}(G) \leq D_{\max}$.

94 **Theorem 3** (Finite Primitive Basis). For every $H \in \mathcal{C}_{\text{Tier2}}$, there exists a typed DAG G
 95 with $V \subseteq \mathcal{B}$ that is an ε -approximate representation of H .

96 *Proof sketch.* By Definition 1, $H = H_K \circ \dots \circ H_1$ with $K \leq N_{\max}$. Each factor is clas-
 97 sified into one of five physics-stage families and realized by primitives: propagation $\rightarrow P$
 98 or C (wave-equation solutions); elastic interaction $\rightarrow M$ (exact); inelastic interaction $\rightarrow$
 99 R ; encoding–projection $\rightarrow \Pi$ or F (exact); detection–readout $\rightarrow$ finite composition from
 100 $\{\Sigma, S, W, C, D\}$. The per-factor approximation errors compose via sub-multiplicativity:
 101 $\|H - H_G\| \leq K \cdot \max_k(\varepsilon_k) \cdot B^{K-1} \leq \varepsilon$ for ε_k bounded by the Tier-2 truncation. Full
 102 proof in Supplementary Note 10; formal primitive semantics and typed DAG denotation
 103 in¹⁶. □

104 **Scope of $\mathcal{C}_{\text{Tier2}}$.** Covers: all clinical, scientific, and industrial modalities at Tier 1–2 (coded
 aperture, interferometric, projection, Fourier-encoded, acoustic, scattering, THz). Excludes:
 Tier 3–4 nonlinear models (handled by refinement sub-DAGs), quantum state tomography, rel-
 ativistic regimes.

105 **Extension protocol.** A new primitive is warranted when no DAG over $\mathcal{B}$ achieves $e_{\text{Tier2}} \leq$
 106 ε within the complexity bounds. The extension requires: (1) validated forward/adjoint,
 107 (2) demonstrated representation gap, (3) error reduction below ε , (4) need by ≥ 2 modali-
 108 ties, (5) backward-compatible closure re-test. Scatter (R) was added via this protocol when
 109 Compton imaging produced $e_{\text{Tier2}} = 0.34$ without it.

110 **Held-out closure test.** To validate Theorem 3 empirically, we conduct a closure test
 111 under a frozen protocol: the primitive library $\mathcal{B}$ (10 primitives), Detect families (5), fidelity
 112 threshold ($\varepsilon = 0.01$), and complexity bounds ($N_{\max} = 20$, $D_{\max} = 10$) are all frozen before
 113 evaluation. We test 5 Tier-1 held-out modalities (OCT, photoacoustic, SIM, phase-contrast
 114 X-ray, electron ptychography) and 3 exotic modalities (quantum ghost imaging, THz-TDS,
 115 Compton scatter imaging):

116 Seven of eight modalities decompose with the frozen library. Quantum ghost imaging
 117 is operator-equivalent to a single-pixel camera at the image-formation level—sharing the
 118 canonical DAG $\text{Source} \rightarrow M \rightarrow \Sigma \rightarrow D$ at Tier-2 abstraction despite fundamentally dif-
 119 ferent source physics. THz-TDS requires only the coherent-field Detect family. Compton
 120 scatter imaging triggered the extension protocol (Section “Extension protocol” above): its
 121 representation gap ($e_{\text{Tier2}} = 0.34 \gg \varepsilon$) met the falsifiable criterion for basis incompleteness,
 122 and the disciplined resolution—adding Scatter (R)—covers 5+ additional modalities (Ra-
 123 man, fluorescence, DOT, Brillouin) while preserving all existing decompositions. The basis
 124 grows from 9 to 10 (11% increase) while modality coverage grows by 19%+.

Modality	e_{Tier2}	#Nodes/Depth	Triad transfer	New prim.?
OCT	< 0.01	5 / 3	Y	N
Photoacoustic	< 0.01	4 / 3	Y	N
SIM	< 0.01	4 / 3	Y	N
Phase-contrast X-ray	< 0.01	5 / 4	Y	N
Electron ptycho	< 0.01	4 / 3	Y	N
Ghost imaging	< 0.01	4 / 3	Y	N
THz-TDS	< 0.01	3 / 2	Y	N
Compton scatter	< 0.01*	4 / 3	Y	Y (R)

*With R in library; $e_{\text{Tier2}} > 0.01$ without R .

Max observed: 5 nodes / depth 4; median: 4 nodes / depth 3 (cf. bounds $N_{\text{max}} = 20$, $D_{\text{max}} = 10$).

Basis-growth saturation. Plotting the number of distinct primitives K against the number of registered modalities N reveals clear saturation (Figure 7): 8 of 10 primitives are introduced by the first 10 modalities, Disperse (W) by CASSI-type spectral systems, and Scatter (R) only when Compton/Raman-class modalities enter. The growth is sublinear and saturating: $K = 10$ at $N = 31$, with no new primitive required for the most recent 19 modalities added. This saturation is consistent with Theorem 3: once all five physics-stage families are covered by primitives, new modalities compose existing primitives rather than requiring new ones.

Theorem tightness. The theoretical complexity bounds ($N_{\text{max}} = 20$, $D_{\text{max}} = 10$) are conservative: empirically, all 26 registered and 8 held-out modalities (~ 31 unique, accounting for overlaps) require ≤ 6 nodes and depth ≤ 5 (held-out closure test table above), with the median modality using 4 nodes at depth 3. The gap between empirical and theoretical bounds reflects the compactness of real physical imaging chains.

The Triad Decomposition

The TRIAD DECOMPOSITION asserts that every failure in computational image recovery within $\mathcal{C}_{\text{Tier2}}$ can be attributed to one or more of exactly three root causes, which we term *gates*. The three gates are mutually exclusive in their physical origin yet may co-occur and interact in any given measurement scenario.

Gate 1: Recoverability. **Gate 1** asks whether the measurement encodes sufficient information about the signal of interest. Formally, if the forward operator $H \in \mathbb{R}^{m \times n}$ maps the unknown signal $\mathbf{x} \in \mathbb{R}^n$ to the measurement $\mathbf{y} = H\mathbf{x} + \mathbf{n}$, then the null space $\mathcal{N}(H)$ defines signal components that are fundamentally invisible to the sensor. When $\mathcal{N}(H)$ is large, no solver can recover the missing information. **Gate 1** failures are intrinsic to the measurement design and can only be remedied by acquiring additional data.

149 **Gate 2: Carrier Budget.** **Gate 2** asks whether the signal-to-noise ratio (SNR) is suffi-
 150 cient. Every physical carrier—photons, electrons, spins, acoustic waves, particles—is sub-
 151 ject to fundamental noise limits: shot noise for photon-counting systems, thermal noise
 152 in electronic detectors, T_1/T_2 relaxation noise in magnetic resonance. When the carrier
 153 budget is too low, the measurement is dominated by noise and the reconstruction degrades
 154 regardless of operator fidelity.

155 **Gate 3: Operator Mismatch.** **Gate 3** asks whether the forward model assumed by the
 156 solver matches the true physics. The solver operates with a nominal operator H_{nom} , but
 157 the data were generated by a true operator H_{true} . When $H_{\text{nom}} \neq H_{\text{true}}$, the reconstruction
 158 targets a phantom inverse problem. **Gate 3** failures are insidious because they produce
 159 structured artifacts that mimic signal content, leading practitioners to blame the solver
 160 rather than the model. Sources include geometric misalignment (mask shift, rotation),
 161 parameter drift (coil sensitivity variation, gain instability), and model simplification.

162 **Definition 4** (Triad Decomposition). *Let $H \in \mathcal{C}_{\text{Tier2}}$ with measurement $\mathbf{y} = H\mathbf{x} + \mathbf{n}$ and*
 163 *solver using nominal operator H_{nom} . The reconstruction error decomposes as $\text{MSE} \leq$*
 164 *$\text{MSE}^{(G1)} + \text{MSE}^{(G2)} + \text{MSE}^{(G3)}$, where $\text{MSE}^{(G1)}$ is the null-space loss (Supplementary*
 165 *Eq. S1), $\text{MSE}^{(G2)}$ is the noise-floor term (Supplementary Eq. S3), and $\text{MSE}^{(G3)}$ is the*
 166 *mismatch-induced term (Supplementary Eq. S5).*

167 **Theorem 5** (Gate 3 Dominance). *For any $H \in \mathcal{C}_{\text{Tier2}}$ operating above its Gate 1 floor*
 168 *($\gamma > \gamma_{\text{min}}$) and Gate 2 floor ($\text{SNR} > \text{SNR}_{\text{min}}$), Gate 3 is the binding constraint whenever*
 169 *$\|\delta\boldsymbol{\theta}\|_{\mathbf{J}^\top \mathbf{J}} > (\sigma_n / \|\mathbf{x}\|_\infty) \cdot \kappa(\mathbf{H})$.*

170 *Proof.* See Supplementary Note 1, Proposition 2. □

171 **Key finding: Gate 3 dominates.** Across all 9 validated configurations (7 modalities),
 172 **Gate 3** is the dominant failure gate. In CASSI, a sub-pixel mask shift degrades MST-L by
 173 13.98 dB; in MRI, a 5% coil sensitivity mismatch collapses the single-coil reconstruction to
 174 6.94 dB (a ~ 48 dB degradation; under clinically realistic multi-coil conditions the degrada-
 175 tion is 11.2 dB, see Supplementary Note S13). The imaging community has been optimizing
 176 the wrong variable.

177 **Relationship to the Finite Primitive Basis.** The two contributions are complemen-
 178 tary: the Finite Primitive Basis provides a universal representation (every forward model
 179 is a DAG over 10 primitives), and the Triad provides a universal diagnostic law over that
 180 representation. The DAG structure makes Gate 3 diagnosis *actionable*: the MismatchAgent
 181 localizes the offending primitive node and corrects its parameters.

Consequences: Diagnosis and Correction

The two theoretical results directly imply a practical framework. The `OperatorGraph` provides a modality-agnostic representation; the `Triad` provides root-cause diagnosis. PWM implements this through three deterministic agents—no large language model, no learned parameters, no human intervention.

Diagnostic agents. Three deterministic agents evaluate each gate: the `RecoverabilityAgent` (Gate 1) estimates null-space dimension; the `PhotonAgent` (Gate 2) computes the carrier budget; and the `MismatchAgent` (Gate 3) detects and localizes operator mismatch to a specific `OperatorGraph` node. Details in Methods.

Correction pipeline. When **Gate 3** is dominant, PWM activates a two-stage correction pipeline. **Algorithm 1 (Beam Search)** performs a coarse grid search over the declared mismatch parameter family $\theta = (\theta_1, \dots, \theta_k)$ associated with the offending node, scoring candidates by reconstruction sharpness. **Algorithm 2 (Gradient Refinement)** takes the top candidates and performs continuous optimization via backpropagation through the DAG, combining data-fidelity and regularization losses. Correction operates exclusively on the forward model, not on the solver: any existing solver—iterative, plug-and-play, or deep unrolling—benefits without modification.

4-Scenario Protocol. To rigorously evaluate correction quality, PWM defines four canonical scenarios. **Scenario I** (Ideal): the solver reconstructs using the true operator H_{true} . **Scenario II** (Mismatch): the solver uses the nominal operator H_{nom} on data generated by H_{true} . **Scenario III** (Corrected): the solver uses the PWM-corrected operator. **Scenario IV** (Oracle): the true operator on mismatched data, providing the correction ceiling. The recovery ratio $\rho = (\text{PSNR}_{\text{III}} - \text{PSNR}_{\text{II}}) / (\text{PSNR}_{\text{I}} - \text{PSNR}_{\text{II}})$ quantifies how much of the mismatch-induced degradation is recovered (see Methods, Equation (5)).

Empirical Validation

We validate both theoretical contributions in two stages: controlled simulation experiments across seven modalities using the 4-Scenario Protocol, and hardware validation on real CASSI and CACTI instruments.

Simulation experiments

Correction results. Extended Data Table 1 (Supplementary Table S1) summarizes the oracle correction ceiling across 9 configurations spanning 7 modalities. The oracle gain $\Delta_{\text{oracle}} = \text{PSNR}_{\text{IV}} - \text{PSNR}_{\text{II}}$ (Scenario IV: true operator applied to mismatched data) ranges from +0.76 dB (CASSI) to +10.68 dB (CT) across the six non-MRI modalities.

Autonomous grid-search calibration achieves 85–100% of the oracle ceiling (Supplementary Table S7). The validated modalities span photon-domain (CASSI: +0.76/+6.50 dB; CACTI: +10.21 dB; SPC: +10.38 dB; Lensless: +3.55 dB), coherent-photon (Ptychography: +7.09 dB), and X-ray (CT: +10.68 dB). For MRI, the single-coil stress test yields +48.25 dB, but this reflects a pathological scenario; under clinically realistic multi-coil conditions (8 coils, 4× acceleration), the correction gain is +1.75 to +7.14 dB at 3–5% mismatch (Supplementary Note S13)—clinically meaningful but commensurate with other modalities.

Modality deep dives. In CASSI, a 5-parameter mismatch¹⁷ collapses all solvers to ~21 dB regardless of ideal performance (24–35 dB); oracle recovery varies by solver ($\rho_{IV} = 0$ –0.46; Supplementary Table S2), confirming that multi-parameter mismatches are harder to correct than isolated shifts. In CACTI, EfficientSCI¹⁸ drops by 20.58 dB, yet GAP-TV achieves $\rho_{IV} = 93.3\%$ oracle recovery—the best ideal-condition solver suffers the largest degradation. SPC gain drift is corrected to $\rho = 89.6\%$ (Table S7). **Gate3** is dominant in every case; Gates 1–2 impose irrecoverable information-theoretic limits only at extreme compression or photon starvation (Supplementary Tables S10–S11).

Zero-shot generalization. Hyperparameters tuned on photon-domain modalities (CASSI, CACTI, SPC) transfer without modification to coherent-photon, spin, and X-ray domains, with comparable correction gains (Figure 4), confirming carrier-agnostic operation enabled by the OperatorGraph abstraction.

Hardware validation

CASSI real data. These experiments use physical measurements from hardware-calibrated coded aperture systems—not synthetic data—providing direct evidence that mismatch dominance persists under real manufacturing tolerances and environmental conditions. We reconstruct 5 real TSA scenes¹ under calibrated and perturbed masks. GAP-TV shows a mean residual ratio of $1.8\times$ ($0.00189 \rightarrow 0.00333$); HDNet shows $1.0\times$ (mask-oblivious). MST-S/L show ratios near $1.0\times$ on real data, in contrast to severe synthetic degradation—revealing that real hardware already contains manufacturing imperfections that absorb additional perturbations (Supplementary Table S5). These real-data results independently confirm the Triad predictions from simulation: the InverseNet benchmark¹⁷ validates the same mismatch-dominance pattern on physical DD-CASSI and CACTI instruments.

CACTI real data. GAP-TV shows $10.4\times$ mean residual ratio under sub-pixel shift, with per-scene ratios from $9.4\times$ to $11.0\times$. The order-of-magnitude sensitivity arises because the temporal mask pattern replicates across all compressed frames, amplifying errors multiplicatively.

Autonomous calibration. Grid-search calibration recovers 85% (CASSI, 1,140 s), 100% (CACTI, 60 s), and 86–92% (SPC via TV objective) of the oracle correction. CACTI achieves full recovery because the mismatch manifold is low-dimensional and sensitivity is high.

Simulation-to-hardware gap. CASSI shows smaller real degradation ($1.8\times$) than predicted, because pre-existing manufacturing errors absorb perturbations; CACTI shows larger real sensitivity ($10.4\times$), because simpler optics lack this buffering.

Discussion

The central finding is that computational imaging has a universal structure captured by two results: a finite primitive basis for forward models and a tripartite decomposition for failures. Operator mismatch—not solver weakness, not information deficiency, not noise—is the dominant bottleneck across all validated modalities. The implication is direct: the research community should rebalance effort from solver-centric to operator-centric approaches. A single calibration step recovers more reconstruction quality than years of algorithmic innovation.

An instructive analogy is the periodic table of elements. Just as Mendeleev organized known elements by atomic number and predicted gaps, the OPERATORGRAPH organizes imaging modalities by their primitive composition. The analogy is pedagogical, not mathematical: imaging primitives do not have atomic numbers, and the DAG structure is richer than a two-dimensional table. Nevertheless, Theorem 3 implies that the space of imaging forward models, like the space of chemical elements, has a discoverable and finite structure—and for a related reason: the underlying physical interactions are finite.

Prior calibration frameworks—ESPIRiT¹⁰ for MRI, ePIE¹¹ for ptychography, plug-and-play methods⁷ at the solver level—operate within a single modality or abstraction layer. A direct comparison on multi-coil MRI (Supplementary Note S12) shows that ESPIRiT and PWM are complementary: under limited calibration data (24 ACS lines), ESPIRiT’s data-driven estimation degrades by -9.29 dB while PWM’s model-based correction recovers $+0.75$ dB; with abundant calibration data, ESPIRiT excels within MRI but does not transfer to other modalities. The OPERATORGRAPH is the first provably complete forward-model basis across carrier families, enabling a single diagnostic and correction pipeline to serve all modalities in $\mathcal{C}_{\text{Tier2}}$.

The coverage of the primitive basis is broader than it may initially appear. Quantum ghost imaging is operator-equivalent to a classical single-pixel camera at the image-formation level, sharing the same DAG despite fundamentally different source physics. THz time-domain imaging requires only the coherent-field Detect family. Of three exotic modalities analyzed, only Compton scatter imaging required a new primitive (Scatter), and that

single primitive covers 5+ scattering and fluorescence modalities. This pattern—that exotic modalities decompose into existing primitives, and when they do not, a single new primitive covers a whole family—is the hallmark of a well-chosen basis.

The hardware validation reveals a nuanced picture. On real CASSI, degradation is smaller than predicted ($1.8\times$ vs. 3.38 dB), because as-built masks absorb perturbations. On real CACTI, degradation is severe ($10.4\times$), because temporal compression amplifies errors multiplicatively. This asymmetry is invisible without a unified framework. The simulation-to-hardware gap carries a methodological lesson: synthetic mismatch studies overestimate marginal impact while underestimating cumulative burden.

To demonstrate translational potential, we prototype a CT QC Copilot mapping the Triad gates to clinical failure modes. On a simulated 30-scanner fleet, it detects calibration drift with 100% sensitivity/specificity and reduces per-scanner QC time by 94% (Supplementary Note 7). Consistent with the research findings, Gate 3 dominates.

Several limitations should be noted. First, real-data experiments apply software-simulated perturbations rather than physical displacement; controlled hardware experiments are the natural next step (Supplementary Note 8). Second, Tier 1–2 models may not capture all failure modes visible at Tier 3–4. Third, the correction pipeline is limited to declared mismatch parameter families. Fourth, the CT QC validation uses a simulated fleet.

Roadmap to physical validation. The controlled hardware experiment protocol is fully specified (see Methods). For CASSI: physically translate the coded aperture mask by $\Delta x \in \{0.25, 0.5, 1.0\}$ px-equivalent via micrometer stage, re-acquire under identical illumination, and compare against PWM-predicted degradation and autonomous recovery. For CACTI: apply equivalent temporal mask shifts. A multi-unit variation study comparing 2+ camera units of the same design will quantify inter-unit mismatch baselines—the residual calibration error present in any production system. For clinical MRI: process multi-coil k -space data under controlled coil repositioning to validate the $+1.75$ to $+7.14\text{ dB}$ realistic correction range (Supplementary Note S13). These experiments require no methodological innovation, only instrument access and controlled acquisition; the PWM pipeline is ready to process the resulting data without modification.

Looking forward, we envision hardware-in-the-loop validation across additional instruments, real-time adaptive calibration during acquisition, prospective clinical deployment of the CT QC Copilot, and scaling the OperatorGraph library to compile a comprehensive atlas of imaging failure modes.

Acknowledgements. We thank the open-source computational imaging community for making reconstruction code and benchmark datasets publicly available. This work was supported by NextGen PlatformAI C Corp.

Author Contributions. C.Y. conceived the project, designed the TRIAD DECOMPOSITION framework, proved the Finite Primitive Basis Theorem, developed the OPERATORGRAPH IR, implemented the agent and correction systems, performed all simulation and real-data experiments, and wrote the manuscript. X.Y. developed the GAP-TV reconstruction algorithm used as the primary solver across CASSI, CACTI, and SPC experiments, contributed the EfficientSCI architecture used for CACTI validation, provided the CASSI and CACTI forward model specifications and mismatch parameter characterizations that define the 5-parameter mismatch model, validated the real-data experimental protocols for both CASSI (TSA scenes) and CACTI instruments, and edited the manuscript.

Competing Interests. C.Y. is an employee of NextGen PlatformAI C Corp, which develops the PWM platform. The authors declare no other competing interests.

Ethics Declarations. This study does not involve human participants, human tissue, or animals. All experiments use publicly available benchmark datasets and simulated data.

Data Availability. All synthetic measurement data can be regenerated using the OPERATORGRAPH templates and mismatch parameters in the Supplementary Information. The KAIST hyperspectral dataset¹⁹ and TSA real-data scenes used for CASSI experiments are publicly available. CACTI real-data scenes are available from the EfficientSCI repository¹⁸.

Code Availability. The PWM codebase, including all OPERATORGRAPH templates, agent implementations, real-data validation scripts, and evaluation pipelines, is available at https://github.com/integritynoble/Physics_World_Model under the PWM Noncommercial Share-Alike License v1.0 (see LICENSE in the repository).

Correspondence. Correspondence and requests for materials should be addressed to C.Y. (integrityyang@gmail.com).

Online Methods

OperatorGraph Specification

Formal definition. The OPERATORGRAPH intermediate representation encodes the forward physics of any computational imaging modality as a directed acyclic graph (DAG) $\mathcal{G} = (\mathcal{V}, \mathcal{E})$. Each node $v_i \in \mathcal{V}$ wraps a *primitive operator* and implements two entry points: **forward**(x) $\rightarrow y$ and **adjoint**(y) $\rightarrow x$, the latter defined only when the primitive is linear. Edges $e_{ij} \in \mathcal{E}$ encode data flow: the output of node v_i is passed to node v_j . Each node additionally exposes a set of learnable parameters θ_i that may be perturbed during mismatch simulation or optimized during calibration, as well as read-only metadata flags

(`is_linear`, `is_stochastic`, `is_differentiable`). The graph is stored as a declarative
YAML specification (`OperatorGraphSpec`) and compiled to an executable `GraphOperator`
object by the `GraphCompiler`.

Node types. Primitive operators fall into two categories:

- **Linear operators.** Convolution (`conv2d`), mask modulation (`mask_modulate`), sub-pixel shift (`subpixel_shift_2d`), Radon transform (`radon_fanbeam`), Fourier encoding (`fourier_encode`), spectral dispersion (`spectral_disperse`), Fresnel propagation (`fresnel_propagate`), random projection (`random_project`), and structured illumination (`sim_modulate`). Each implements both `forward()` and `adjoint()`.
- **Nonlinear operators.** Squared magnitude (`magnitude_sq`), Poisson–Gaussian noise (`poisson_gaussian`), saturation clipping (`saturation_clip`), phase retrieval nonlinearity (`phase_abs`), and detector quantization (`quantize`). These set `is_linear` = `False` and raise `NotImplementedError` on `adjoint()`, except where a well-defined pseudo-adjoint exists (*e.g.*, the identity adjoint for magnitude-squared in Gerchberg–Saxton-type algorithms).

Adjoint validation. Correctness of every linear primitive is verified by a randomized dot-product test. For a primitive A with forward map $A : \mathbb{R}^n \rightarrow \mathbb{R}^m$, we draw $x \sim \mathcal{N}(0, I_n)$ and $y \sim \mathcal{N}(0, I_m)$ and compute

$$\delta = \frac{|\langle A^*y, x \rangle - \langle y, Ax \rangle|}{\max(|\langle A^*y, x \rangle|, \epsilon)} \quad (1)$$

where $\epsilon = 10^{-12}$ guards against division by zero. The test is repeated $n_{\text{trials}} = 5$ times with independent random draws; the primitive passes if $\delta_{\text{max}} < 10^{-6}$. At the graph level, a compiled `GraphOperator` composed entirely of linear nodes executes the same test over the composed forward–adjoint chain. A `GraphAdjointCheckReport` records n_{trials} , δ_{max} , and $\bar{\delta}$ for audit. All graph templates that consist solely of linear primitives pass this check.

Graph compilation. The compiler executes a four-stage pipeline:

1. **Validate.** Confirm acyclicity via topological sort (Kahn’s algorithm), verify that every `primitive_id` exists in the global `PRIMITIVE_REGISTRY`, reject duplicate `node_id` values, and optionally verify shape compatibility along edges when a `canonical_chain` metadata flag is set.
2. **Bind.** Instantiate each primitive with its parameter dictionary θ_i .
3. **Plan forward.** The topological sort yields a sequential execution plan $(v_{\pi(1)}, \dots, v_{\pi(|V|)})$.

380 **4. Plan adjoint.** For graphs where `all_linear = True`, the adjoint plan reverses the
 381 topological order and applies each node’s individual adjoint in sequence, implementing
 382 the chain rule $A^* = A_1^* \circ \dots \circ A_{|\mathcal{V}|}^*$ for a composition $A = A_{|\mathcal{V}|} \circ \dots \circ A_1$. For
 383 graphs containing nonlinear nodes, the adjoint plan is not generated, and any call to
 384 `adjoint()` raises `NotImplementedError` at runtime.

385 The compiled `GraphOperator` is serializable to JSON and hashable via SHA-256 for prove-
 386 nance tracking in `RunBundle` manifests.

387 **Template library.** The `graph_templates.yaml` registry contains templates organized
 388 across 26 registered modalities (7 with full end-to-end correction validation, 1 with Scenario I
 389 baseline, 18 with template-level validation), grouped by physical carrier:

- 390 • **Photons (optical and X-ray):** CASSI, SPC, CACTI, structured illumination
 391 microscopy (SIM), confocal, light-sheet, holography, ptychography, Fourier ptycho-
 392 graphic microscopy (FPM), optical coherence tomography (OCT), lensless imaging,
 393 light field, integral imaging, neural radiance fields (NeRF), Gaussian splatting, fluo-
 394 rescence lifetime imaging (FLIM), diffuse optical tomography (DOT), phase retrieval,
 395 X-ray computed tomography (CT), and cone-beam CT (CBCT).
- 396 • **Electrons:** Electron diffraction, electron backscatter diffraction (EBSD), electron
 397 energy loss spectroscopy (EELS), and electron holography.
- 398 • **Spins (MRI):** Functional MRI (fMRI), diffusion-weighted MRI (DW-MRI), and
 399 magnetic resonance spectroscopy (MRS).
- 400 • **Acoustic:** Ultrasound B-mode, Doppler ultrasound, shear-wave elastography, sonar,
 401 and photoacoustic tomography (combines optical excitation with acoustic detection).
- 402 • **Particles:** Neutron tomography, proton radiography, and muon tomography.

403 **Physics Fidelity Ladder.** Each template is parameterized by a fidelity tier that controls
 404 the degree of physical realism in the simulated forward model:

405 **Tier 1 (Linear, shift-invariant):** The forward model is a linear, spatially uniform operator—
 406 the simplest approximation, suitable for initial diagnostics and rapid prototyping.

407 **Tier 2 (Linear, shift-variant):** Spatially varying operator parameters (e.g. non-uniform
 408 illumination, position-dependent PSF, multi-coil sensitivity maps in MRI). Adds a
 409 modality-appropriate noise model (Poisson shot noise plus Gaussian read noise for
 410 photon-counting modalities, Rician noise for MRI, Poisson for CT).

411 **Tier 3 (Nonlinear, ray/wave-based):** Includes nonlinear effects such as wavefront cur-
 412 vature, diffraction, and scattering. Perturbation families and ranges are specified in
 413 `mismatch_db.yaml`.

414 **Tier 4 (Full-wave / Monte Carlo):** Complete physical simulation including wave-optical
 415 propagation, spatially varying aberrations, detector nonlinearities, and environmen-
 416 tal drift. Currently implemented for holography and ptychography; other modalities
 417 degrade gracefully to Tier 3.

418 **Triad Decomposition Formalization**

419 The TRIAD DECOMPOSITION asserts that the quality of any computational imaging re-
 420 construction is bounded by three fundamental gates. Rather than a qualitative guideline,
 421 PWM quantifies each gate numerically and uses the resulting scores to diagnose the domi-
 422 nant bottleneck in any imaging configuration.

423 **Gate 1 (Recoverability).** Recoverability measures the information-theoretic capacity
 424 of the sensing geometry. We quantify it via the *effective compression ratio* $r = m/n$, where
 425 m is the number of independent measurements and n the dimension of the signal. The
 426 `compression_db.yaml` registry (1,186 lines) stores, for each modality, a lookup table map-
 427 ping compression ratio to expected reconstruction PSNR under ideal conditions, obtained
 428 from calibration experiments or published benchmarks. Each entry carries a **provenance**
 429 field citing the source (paper DOI, internal experiment ID, or theoretical formula). Addi-
 430 tional recoverability indicators include the effective rank of the measurement matrix (esti-
 431 mated via randomized SVD for large operators), the dimension of the null space, and the
 432 restricted isometry property (RIP) constant where analytically tractable (*e.g.*, for Gaussian
 433 random projections in SPC).

434 **Gate 2 (Carrier Budget).** The carrier budget quantifies the signal-to-noise ratio (SNR)
 435 of the measurement channel. The **PhotonAgent** consumes the `photon_db.yaml` registry
 436 (624 lines) which stores, per modality, a deterministic photon model parameterized by
 437 source power, quantum efficiency, exposure time, and detector characteristics. The agent
 438 classifies the noise regime into one of three categories: *shot-limited* (Poisson-dominated,
 439 $\text{SNR} \propto \sqrt{N_{\text{photon}}}$), *read-limited* (Gaussian read noise dominates, $\text{SNR} \propto N_{\text{photon}}/\sigma_{\text{read}}$),
 440 and *dark-current-limited* (long exposures where dark current accumulation dominates). The
 441 output is a **PhotonReport** containing the estimated SNR in decibels, the noise regime
 442 classification, per-element photon count, and a feasibility verdict (**sufficient**, **marginal**,
 443 or **insufficient**).

444 **Gate 3 (Operator Mismatch).** Operator mismatch quantifies the discrepancy between
 445 the assumed forward model H_{nom} and the true physical operator H_{true} . The **MismatchAgent**
 446 consults `mismatch_db.yaml` (797 lines) which catalogs, for each modality, the set of mis-
 447 match parameters (spatial shifts, rotational offsets, dispersion errors, PSF deviations, coil
 448 sensitivity errors, center-of-rotation offsets, *etc.*), their typical ranges, and available cor-

449 rection methods. The mismatch severity score $s \in [0, 1]$ is computed as the normalized ℓ_2
 450 distance $\|\boldsymbol{\theta}_{\text{true}} - \boldsymbol{\theta}_{\text{nom}}\| / \|\boldsymbol{\theta}_{\text{range}}\|$, where $\boldsymbol{\theta}_{\text{range}}$ is the per-parameter dynamic range from the
 451 registry. Sensitivity analysis $\partial \text{PSNR} / \partial \theta_k$ is estimated via finite differences on the forward
 452 model. The output is a `MismatchReport` containing the severity score, the dominant mis-
 453 match parameter, the recommended correction method, and the expected PSNR gain from
 454 correction.

Gate binding determination. Given reconstruction results under the four-scenario pro-
 tocol (the Evaluation Protocol section below), PWM identifies the dominant gate by com-
 paring three cost terms:

$$C_{\text{mismatch}} = \text{PSNR}_{\text{I}} - \text{PSNR}_{\text{II}} \quad (2)$$

$$C_{\text{noise}} = \text{PSNR}_{\text{ideal}} - \text{PSNR}_{\text{noisy}} \quad (3)$$

$$C_{\text{recover}} = \text{PSNR}_{\text{limit}} - \text{PSNR}_{\text{I}} \quad (4)$$

455 where PSNR_{I} is the reconstruction PSNR under Scenario I (ideal operator), PSNR_{II} under
 456 Scenario II (mismatched operator), $\text{PSNR}_{\text{noisy}}$ under the corresponding noisy condition,
 457 and $\text{PSNR}_{\text{limit}}$ is the theoretical upper bound from the compression table. The dominant
 458 gate is $\arg \max_g C_g$.

459 **TriadReport.** For every diagnosis, PWM produces a `TRIADREPORT`: a Pydantic-validated
 460 structured artifact comprising `dominant_gate` (enum: `recoverability`, `carrier_budget`,
 461 `operator_mismatch`), `evidence_scores` (three floats, one per gate), `confidence_interval`
 462 (float, 95% CI width from bootstrap), `recommended_action` (string, *e.g.* “increase compres-
 463 sion ratio” or “apply mismatch correction”), and `parameter_sensitivities` (dictionary
 464 mapping each mismatch parameter name to its $\partial \text{PSNR} / \partial \theta_k$ value). The `TRIADREPORT` is
 465 mandatory—PWM does not permit a reconstruction to be reported without an accompany-
 466 ing diagnosis.

Recovery ratio. We define the *recovery ratio*

$$\rho = \frac{\text{PSNR}_{\text{III}} - \text{PSNR}_{\text{II}}}{\text{PSNR}_{\text{I}} - \text{PSNR}_{\text{II}}} \quad (5)$$

467 which lies in $[0, 1]$ under standard convexity conditions (see Supplementary Note 1 for
 468 formal analysis; values $\rho > 1$ are possible when the corrected operator provides beneficial
 469 regularization). $\rho = 0$ indicates that calibration yields no benefit (mismatch is not the
 470 bottleneck), while $\rho = 1$ indicates that calibration fully closes the mismatch gap.

471 Agent System Architecture

472 The PWM agent system comprises 6 specialist agents, 1 optional hybrid agent, and 8
473 support classes totalling 10,545 lines of Python. All agents execute deterministically; no
474 large language model (LLM) is required for pipeline operation.

475 **PlanAgent.** The orchestrator agent. Given a user prompt or a structured `ExperimentSpec`,
476 `PlanAgent` parses the intent (`simulate`, `operator_correction`, or `auto`), maps the re-
477 quested modality to its canonical key via the `modalities.yaml` registry (which contains 64
478 modality entries with keywords, forward model equations, and default solvers), builds an
479 `ImagingSystem` contract, and dispatches to the appropriate sub-agents. When the mode is
480 `auto`, `PlanAgent` inspects the available data and operator specification to determine whether
481 simulation or operator correction is more appropriate.

482 **PhotonAgent.** Computes SNR feasibility deterministically from the `photon_db.yaml`
483 registry. For each modality and photon-level tier (`bright`, `standard`, `low_light`), the
484 agent evaluates the photon budget by combining source power, quantum efficiency, ex-
485 posure time, and noise model parameters. The output `PhotonReport` is a strict Pydan-
486 tic model containing `noise_regime` (enum), `snr_db` (float), `feasibility` (enum), and
487 `per_element_photons` (float).

488 **RecoverabilityAgent.** A table-driven agent that consults `compression_db.yaml` (1,186
489 lines) to map the modality and compression ratio to an expected PSNR range. Each table
490 entry includes provenance metadata citing the original source. The output `RecoverabilityReport`
491 contains `compression_ratio`, `psnr_prediction`, `feasibility`, and `null_space_dim` where
492 available.

493 **MismatchAgent.** Scores the mismatch severity for a given imaging configuration us-
494 ing `mismatch_db.yaml` (797 lines). For each modality, the database enumerates the rel-
495 evant mismatch parameters, their physical units, typical perturbation ranges, and avail-
496 able correction algorithms. The output `MismatchReport` includes `severity` (float, 0–1),
497 `correction_method` (string), `expected_gain_db` (float), and `dominant_parameter` (string).

498 **AnalysisAgent.** The bottleneck classifier. It receives reports from the Photon, Recover-
499 ability, and Mismatch agents, computes the gate costs (Equations (2) to (4)), identifies the
500 dominant gate, and generates actionable suggestions. The `AnalysisAgent` also computes
501 the recovery ratio ρ and its bootstrap confidence interval.

502 **AgentNegotiator.** Implements a cross-agent veto protocol. Before reconstruction is au-
503 thorized, the negotiator inspects all three upstream reports and applies three veto con-
504 ditions: (1) low photon budget combined with aggressive compression (C_{noise} and C_{recover}

505 both large); (2) severe mismatch (severity > 0.7) without a planned correction step; (3) joint
506 probability below the floor threshold ($p_{\text{joint}} < 0.15$), indicating that all three subsystems
507 are simultaneously marginal. When any veto fires, reconstruction halts with an actionable
508 explanation and suggested remediation.

509 **HybridAgent.** An optional wrapper that invokes an LLM for natural-language narra-
510 tive generation or edge-case modality mapping. All quantitative decisions remain on the
511 deterministic code path; the HybridAgent is never required for pipeline operation.

512 **Support classes.** The remaining components include: **AssetManager** (file I/O and caching
513 for large arrays), **ContinuityChecker** (verifies that sequential pipeline outputs are dimen-
514 sionally consistent), **SystemDiscern** (auto-detects modality from uploaded data), **PreflightChecker**
515 (validates the complete experiment configuration before execution), **WhatIfPrecomputer**
516 (evaluates counterfactual what-if scenarios), **SelfImprovement** (logs diagnostic events for
517 future registry refinement), **PhysicsStageVisualizer** (generates intermediate visualiza-
518 tions at each pipeline stage), and **UPWMI** (Universal Physics World Model Interface, the
519 top-level entry point that wires all agents together).

520 **Contract system.** Inter-agent communication uses 25 Pydantic v2 contract models. All
521 contracts inherit from **StrictBaseModel**, which enforces `extra="forbid"` (no unexpected
522 fields), `validate_assignment=True` (mutations re-validated), and a model validator that
523 rejects NaN and Inf in any float field. Bounded scores use `Field(ge=0.0, le=1.0)`. Enums
524 are string enums for human-readable JSON serialization. This design ensures that pipeline
525 failures surface immediately as validation errors rather than propagating silently.

526 **YAML registries.** The system is driven by 9 YAML registries totalling 7,034 lines:
527 `modalities.yaml` (modality definitions), `graph_templates.yaml` (OperatorGraph skele-
528 tons), `photon_db.yaml` (photon models), `mismatch_db.yaml` (mismatch parameters and
529 correction methods), `compression_db.yaml` (recoverability tables with provenance), `solver_registry.yaml`
530 (solver configurations), `primitives.yaml` (primitive operator metadata), `dataset_registry.yaml`
531 (dataset locations and formats), and `acceptance_thresholds.yaml` (pass/fail thresholds
532 per metric).

533 Correction Algorithms

534 We implement two complementary algorithms for operator mismatch correction. Crucially,
535 both algorithms operate on the forward operator parameters θ rather than the reconstruc-
536 tion solver weights, making them *solver-agnostic*: the corrected operator $H(\hat{\theta})$ benefits any
537 downstream solver (GAP-TV, MST-L, HDNet²⁰, CST, *etc.*) without retraining.

538 **Algorithm 1: Hierarchical Beam Search.** The coarse correction phase employs a
 539 hierarchical search strategy to rapidly explore the mismatch parameter space. For CASSI,
 540 the five-parameter mismatch model comprises mask affine parameters (spatial shifts dx , dy
 541 and rotation θ) and dispersion parameters (slope a_1 and axis angle α); an optional sixth
 542 parameter, PSF width σ_{psf} , is available but not used in the primary experiments. The
 543 algorithm proceeds as follows:

- 544 1. **1D sweeps.** Each parameter is swept independently over its full range while holding
 545 others at nominal values. This produces five 1D cost curves from which coarse optima
 546 are extracted.
- 547 2. **3D beam search.** The mask affine subspace (dx, dy, θ) is searched over a $5 \times 5 \times 5$
 548 grid centered on the 1D optima. The top- k ($k = 5$) candidates by reconstruction
 549 PSNR are retained.
- 550 3. **2D beam search.** For each retained mask candidate, the dispersion subspace (a_1, α)
 551 is searched over a 5×7 grid. The joint top- k candidates are retained.
- 552 4. **Coordinate descent refinement.** Three rounds of univariate refinement on each
 553 parameter, shrinking the search interval by factor 2 at each round, produce the final
 554 estimate $\hat{\theta}_{\text{Alg1}}$.

555 Total runtime is approximately 300 seconds per scene on a single GPU. Accuracy is
 556 ± 0.1 – 0.2 pixels for spatial parameters and $\pm 0.05^\circ$ for angular parameters.

557 **Algorithm 2: Joint Gradient Refinement.** The fine correction phase uses a differ-
 558 entiable forward model to jointly optimize all mismatch parameters via gradient descent.
 559 The key components are:

- 560 1. **Differentiable mask warp.** The binary mask is warped by a continuous affine
 561 transformation using bilinear interpolation, implemented as a custom PyTorch module
 562 (`DifferentiableMaskWarpFixed`). The mask values are passed through a straight-
 563 through estimator (STE) to maintain binary structure while permitting gradient flow.
- 564 2. **Differentiable forward model.** The CASSI forward model $y = \text{CASSI}(x; \theta)$ is
 565 implemented as a differentiable PyTorch module (`DifferentiableCassiForwardSTE`)
 566 that accepts mismatch parameters as differentiable inputs.
- 567 3. **GPU grid initialization.** A full-range 3D grid search over (dx, dy, θ) with $9 \times 9 \times 7 =$
 568 567 points provides diverse starting candidates. The top 9 candidates seed multi-start
 569 gradient refinement.

570 **4. Staged gradient refinement.** Each of the 9 candidates is refined using Adam
 571 optimization (learning rate 10^{-2} , decaying to 10^{-3}) for 200 steps. For each candidate,
 572 4 random restarts with jittered initialization guard against local minima. The loss
 573 function is the negative PSNR computed via an unrolled K -iteration differentiable
 574 GAP-TV solver (`DifferentiableGAPTV`, $K = 10$ unrolled iterations).

575 Total runtime for Algorithm 2 is approximately 3,200 seconds ($200 \text{ steps} \times 4 \text{ restarts} \times$
 576 $9 \text{ candidates with early stopping}$). Accuracy improves to ± 0.05 – 0.1 pixels, a 3 – $5\times$ improve-
 577 ment over Algorithm 1. The two algorithms are used sequentially in practice: Algorithm 1
 578 provides a warm start, and Algorithm 2 refines to sub-pixel precision.

579 Evaluation Protocol

580 **Four-Scenario Protocol.** We evaluate every modality under four standardized scenarios
 581 that isolate different sources of quality degradation:

582 **Scenario I (Ideal):** $\mathbf{y}_{\text{obs}} = H_{\text{true}} \mathbf{x}_{\text{gt}}$; reconstruct with H_{true} . In this scenario the system
 583 is perfectly calibrated ($H_{\text{true}} = H_{\text{nom}}$), so the operator used for reconstruction matches
 584 the one that generated the data. This yields the oracle upper bound on reconstruction
 585 quality, limited only by the sensing geometry and solver convergence.

586 **Scenario II (Mismatch):** $\mathbf{y}_{\text{obs}} = H_{\text{true}} \mathbf{x}_{\text{gt}}$; reconstruct with H_{nom} ($H_{\text{nom}} \neq H_{\text{true}}$). This
 587 is the standard operating condition in practice: the measurement is generated by the
 588 true physics, but the reconstruction uses a nominal (potentially mismatched) forward
 589 model.

590 **Scenario III (Corrected):** $\mathbf{y}_{\text{obs}} = H_{\text{true}} \mathbf{x}_{\text{gt}}$; reconstruct with $\hat{H} = H(\hat{\theta})$ where $\hat{\theta}$ is
 591 estimated by Algorithms 1 and 2. This quantifies the benefit of mismatch calibration.

592 **Scenario IV (Oracle Mask):** Same measurements as Scenario II ($\mathbf{y}_{\text{obs}} = H_{\text{true}} \mathbf{x}_{\text{gt}}$ with
 593 $H_{\text{true}} \neq H_{\text{nom}}$); reconstruct with H_{true} instead of H_{nom} . Provides the correction
 594 ceiling: the best reconstruction achievable when the true operator is known exactly,
 595 applied to data that were sensed by the mismatched system. The gap between Sce-
 596 nario IV and Scenario I reveals the irreducible loss from the degraded sensing config-
 597 uration itself (e.g., a shifted mask pattern is suboptimal even when perfectly known).

598 **Metrics.** Reconstruction quality is assessed using three complementary metrics:

- 599 • **PSNR** (peak signal-to-noise ratio, in dB): the primary metric, computed per scene
 600 and averaged. For signals normalized to $[0, 1]$, $\text{PSNR} = 10 \log_{10}(1/\text{MSE})$. For SPC
 601 data normalized to $[0, 255]$, the peak value is 255.

- **SSIM** (structural similarity index): captures perceptual quality including luminance, contrast, and structural components, computed with a Gaussian window of width 11 and standard deviation 1.5.
- **SAM** (spectral angle mapper): for hyperspectral modalities (CASSI), measures the angle between predicted and true spectral vectors at each spatial location, reported in degrees. Lower is better.

Datasets.

- **CASSI**: 10 scenes from the KAIST dataset¹⁹, each a $256 \times 256 \times 28$ spectral cube (28 spectral bands from 450 nm to 650 nm). Data range $[0, 1]$.
- **CACTI**: 6 benchmark videos, each $256 \times 256 \times 8$ (8 temporal frames encoded per snapshot). Data range $[0, 1]$.
- **SPC**: 11 natural images from the Set11 benchmark, each 256×256 grayscale. Data range $[0, 255]$.

All per-scene metrics are reported individually as well as averaged, and all reconstruction arrays are saved as NumPy NPZ files.

Experimental Details

Hardware. All experiments are conducted on a single NVIDIA GPU. Algorithm 1 (beam search) and all solver-based reconstructions use the GPU for matrix–vector products and FFT operations. Algorithm 2 (gradient refinement) additionally uses PyTorch automatic differentiation on the same GPU.

CASSI configuration. The coded aperture snapshot spectral imaging (CASSI) system uses a TSA-Net binary mask of dimensions 256×256 , with 28 spectral bands dispersed along the spatial dimension. The five-parameter mismatch model $\theta = (dx, dy, \theta, a_1, \alpha)$ describes: mask spatial shift in x (dx , pixels), mask spatial shift in y (dy , pixels), mask rotation angle (θ , degrees), dispersion slope (a_1 , pixels per band), and dispersion axis angle (α , degrees). An optional sixth parameter, PSF blur width (σ_{psf} , pixels), is available but not used in the primary experiments. These mismatch parameter values were determined through systematic characterization of realistic CASSI assembly errors (Supplementary Note 9). The true mismatch parameters are $\theta_{\text{true}} = (dx = 0.5 \text{ px}, dy = 0.3 \text{ px}, \theta = 0.1^\circ, a_1 = 2.02, \alpha = 0.15^\circ)$. Solvers evaluated include TwIST²¹, GAP-TV²², DGSMP²³, MST-L⁹, and CST-L²⁴, all of which receive the same operator and differ only in their reconstruction algorithm. The supplementary per-scene analysis additionally includes DeSCI²⁵ and HDNet²⁰.

634 **CACTI configuration.** The coded aperture compressive temporal imaging system uses
 635 binary temporal masks of dimensions 256×256 , encoding 8 video frames into a single
 636 snapshot measurement. Mismatch is parameterized as a temporal mask timing offset (sub-
 637 frame shift). The default solver is EfficientSCI¹⁸.

638 **SPC configuration.** The single-pixel camera uses random binary measurement patterns
 639 at three compression ratios: 10%, 25%, and 50% ($r = m/n \in \{0.10, 0.25, 0.50\}$). Mismatch
 640 is modeled as an exponential gain drift ($g_i = \exp(-\alpha \cdot i)$) on the measurement matrix. The
 641 default solver is FISTA-TV with total-variation regularization.

642 **MRI configuration.** Cartesian k -space sampling with $4\times$ acceleration (25% of k -space
 643 lines acquired). Mismatch is parameterized as a 5% multiplicative error in the coil sensitivity
 644 maps used for parallel imaging reconstruction. The default solver is SENSE¹⁵ with ℓ_1 -
 645 wavelet regularization.

646 **ESPIRiT comparison protocol.** To quantitatively compare PWM with ESPIRiT¹⁰,
 647 we evaluate four conditions on the same multi-coil MRI configuration: (1) Scenario I (true
 648 maps), (2) Scenario II (5% mismatched maps), (3) ESPIRiT auto-calibrated maps esti-
 649 mated from the 24-line ACS region via eigenvalue decomposition of the calibration matrix,
 650 and (4) PWM beam-search corrected maps (grid search over per-coil amplitude and phase
 651 scaling). All four conditions use the same CG-SENSE solver ($\ell = 10^{-3}$, 30 iterations). The
 652 comparison isolates the effect of map quality on reconstruction, holding the solver constant.

653 **CT configuration.** Fan-beam geometry with 180 projections over 180° . Mismatch is
 654 modeled as a center-of-rotation (CoR) offset, which produces characteristic arc artifacts in
 655 the reconstruction. The default solver is filtered back-projection (FBP)¹⁴ with a Ram-Lak
 656 filter, supplemented by iterative SART for comparison.

657 **CASSI real-data configuration.** The TSA real hyperspectral dataset¹ consists of 5
 658 scenes at 660×660 spatial resolution with 28 spectral bands and mask-shift step 2. Four
 659 solvers are evaluated: GAP-TV (200 iterations), HDNet (pre-trained checkpoint, full spatial
 660 resolution), MST-S and MST-L (pre-trained checkpoints, centre-cropped to 256×256 due
 661 to hardcoded spatial assumptions in the model architecture). The coded aperture mask
 662 is perturbed by $dx = 0.5$ px, $dy = 0.3$ px to simulate assembly-induced mismatch. No
 663 ground truth is available; quality is assessed via the normalised measurement residual $r =$
 664 $\|\mathbf{y} - H\hat{\mathbf{x}}\|^2 / \|\mathbf{y}\|^2$.

665 **CACTI real-data configuration.** The EfficientSCI real temporal dataset¹⁸ consists of
 666 4 dynamic scenes (duomino, hand, pendulumBall, waterBalloon) at 512×512 with compres-
 667 sion ratio 10. The real mask is stored separately from the measurement data. Two solvers

are evaluated: GAP-TV (50 iterations) and PnP-FFDNet (50 iterations with FFDNet de-noiser). Mismatch is induced by shifting the mask by $dx = 0.5$ px, $dy = 0.3$ px. Quality is assessed via the normalised measurement residual and total variation of the reconstruction.

Controlled hardware experiment protocol. The software-perturbation protocol above applies calibrated mask shifts to existing real measurements. A full hardware-in-the-loop validation requires physically displacing the coded aperture mask and re-acquiring data. The protocol proceeds as follows: (i) acquire a baseline dataset with the mask at its factory-calibrated position; (ii) physically translate the mask by a known displacement ($\Delta x \in \{0.25, 0.5, 1.0\}$ px equivalent, verified by micrometer stage) and re-acquire under identical illumination; (iii) reconstruct both datasets with the factory mask specification and compute the PSNR degradation and measurement residual; (iv) apply PWM autonomous calibration and measure recovery. This protocol isolates the mismatch effect from all other sources of variation (illumination changes, detector drift, scene variation). Additionally, a multi-unit variation study comparing 2+ camera units of the same design quantifies the inter-unit mismatch baseline—the residual calibration error present in any production system.

Clinical CT phantom configuration. For clinical translation, PWM is evaluated on CT quality assurance using the ACR CT accreditation phantom (Gammex 464)²⁶. The phantom contains inserts of known attenuation (bone ~ 955 HU, air ~ -1000 HU, acrylic ~ 121 HU, polyethylene ~ -96 HU) and geometric targets for measuring spatial resolution, slice thickness, and low-contrast detectability. Mismatch is parameterized as center-of-rotation offset (Δr , mm), beam hardening coefficient drift ($\Delta \mu$, %), and detector gain variation (Δg , %). Nine ACR-aligned metrics are computed automatically: CT number accuracy (5 materials), geometric accuracy (± 2 mm tolerance), slice thickness (± 1.5 mm), uniformity (≤ 5 HU), noise standard deviation, spatial resolution (≥ 5 lp/cm), low-contrast detectability (≥ 4 targets), and artifact score.

Clinical MRI validation configuration. For MRI clinical validation, PWM processes multi-coil k -space data from public datasets (fastMRI²⁷). Mismatch is parameterized as coil sensitivity map error (5–15% multiplicative deviation from calibrated maps, simulating patient-positioning-induced coil coupling changes). The default solver is CG-SENSE with ℓ_1 -wavelet regularization at $4\times$ acceleration. Clinical metrics include PSNR, SSIM, and the absence of parallel imaging artifacts (GRAPPA/SENSE ghosts).

Statistical Analysis

Per-scene reporting. All metrics are reported per scene, not merely as dataset averages. This enables identification of scene-dependent failure modes (*e.g.*, spectrally flat scenes that

are inherently harder for CASSI, or textureless regions that challenge SPC).

Summary statistics. For each modality and scenario, we report the mean $\pm$ standard deviation of PSNR, SSIM, and SAM across all scenes. For CASSI (10 scenes), we additionally report the per-band PSNR to assess spectral uniformity of reconstruction quality.

Recovery ratio confidence intervals. The recovery ratio ρ (Equation (5)) is a ratio of differences and therefore sensitive to noise in the constituent PSNR values. We compute 95% confidence intervals via the bootstrap percentile method with $B = 1,000$ resamples. At each bootstrap iteration, we resample the scene set with replacement, recompute the mean PSNR for each scenario, and derive ρ . The 2.5th and 97.5th percentiles of the bootstrap distribution define the 95% CI.

Parameter recovery accuracy. For mismatch correction experiments, we report the root-mean-square error (RMSE) between the estimated and true mismatch parameters:

$$\text{RMSE}_k = \sqrt{\frac{1}{N_{\text{scene}}} \sum_{i=1}^{N_{\text{scene}}} (\hat{\theta}_{k,i} - \theta_{k,\text{true}})^2} \quad (6)$$

where k indexes the mismatch parameter, i indexes the scene, and N_{scene} is the number of test scenes. Uncertainty in the RMSE is estimated via bootstrap ($B = 1,000$).

Ablation significance. Ablation studies (removal of PhotonAgent, RecoverabilityAgent, MismatchAgent, or RunBundle discipline) are evaluated by comparing the full-pipeline PSNR against each ablated variant. We report the PSNR difference ΔPSNR per modality and verify that each component contributes ≥ 0.5 dB across all validated modalities, establishing practical significance.

Code and Data Availability

Source code. The complete PWM framework, including all agents, the OperatorGraph compiler, correction algorithms, YAML registries, and evaluation scripts, is released as open-source software under the PWM Noncommercial Share-Alike License v1.0 at https://github.com/integritynoble/Physics_World_Model. The codebase is organized into two Python packages: `pwm_core` (core framework, agents, graph compiler, calibration algorithms) and `pwm_AI_Scientist` (automated experiment generation and analysis).

Reconstruction data. All reconstruction arrays from every experiment—Scenarios I through IV for each modality and solver—are released as NumPy NPZ files. Files are

729 stored using Git LFS and require `allow_pickle=True` for loading. Data ranges are stan-
730 dardized: CASSI and CACTI reconstructions are normalized to $[0, 1]$; SPC reconstructions
731 are in $[0, 255]$.

732 **Experiment manifests.** Every experiment is recorded in a RunBundle v0.3.0 manifest
733 containing: the git commit hash at execution time, all random number generator seeds,
734 platform information (Python version, GPU model, CUDA version), SHA-256 hashes of all
735 input data and output artifacts, metric values, and wall-clock timestamps. These manifests
736 enable exact reproduction of every reported result.

737 **Registry data.** All 9 YAML registries (7,034 lines total) that drive the agent system—
738 including modality definitions, graph templates, photon models, mismatch databases, com-
739 pression tables, solver configurations, primitive specifications, dataset paths, and acceptance
740 thresholds—are publicly available in the repository under `packages/pwm_core/contrib/`.
741 The `ExperimentSpec` JSON schemas used for pipeline input validation are included along-
742 side worked examples in `examples/`.

743 References

- 744 [1] Ashwin A. Wagadarikar, Renu John, Rebecca Willett, and David J. Brady. Single
745 disperser design for coded aperture snapshot spectral imaging. *Applied Optics*, 47(10):
746 B44–B51, 2008. doi: 10.1364/AO.47.000B44.
- 747 [2] Patrick Llull, Xuejun Liao, Xin Yuan, Jianbo Yang, David Kittle, Lawrence Carin,
748 Guillermo Sapiro, and David J. Brady. Coded aperture compressive temporal imaging.
749 *Optics Express*, 21(9):10526–10545, 2013. doi: 10.1364/OE.21.010526.
- 750 [3] Marco F. Duarte, Mark A. Davenport, Dharmpal Takhar, Jason N. Laska, Ting
751 Sun, Kevin F. Kelly, and Richard G. Baraniuk. Single-pixel imaging via com-
752 pressive sampling. *IEEE Signal Processing Magazine*, 25(2):83–91, 2008. doi:
753 10.1109/MSP.2007.914730.
- 754 [4] Michael Lustig, David Donoho, and John M. Pauly. Sparse MRI: The application of
755 compressed sensing for rapid MR imaging. *Magnetic Resonance in Medicine*, 58(6):
756 1182–1195, 2007. doi: 10.1002/mrm.21391.
- 757 [5] Emmanuel J. Candès and Michael B. Wakin. An introduction to compressive sampling.
758 *IEEE Signal Processing Magazine*, 25(2):21–30, 2008. doi: 10.1109/MSP.2007.914731.
- 759 [6] David L. Donoho. Compressed sensing. *IEEE Transactions on Information Theory*,
760 52(4):1289–1306, 2006. doi: 10.1109/TIT.2006.871582.

- [7] Singanallur V. Venkatakrishnan, Charles A. Bouman, and Brendt Wohlberg. Plug-and-play priors for model based reconstruction. In *Proceedings of the IEEE Global Conference on Signal and Information Processing (GlobalSIP)*, pages 945–948, 2013. doi: 10.1109/GlobalSIP.2013.6737048.
- [8] Vishal Monga, Yuelong Li, and Yonina C. Eldar. Algorithm unrolling: Interpretable, efficient deep learning for signal and image processing. *IEEE Signal Processing Magazine*, 38(2):18–44, 2021. doi: 10.1109/MSP.2020.3016905.
- [9] Yuanhao Cai, Jing Lin, Xiaowan Hu, Haoqian Wang, Xin Yuan, Yulun Zhang, Radu Timofte, and Luc Van Gool. Mask-guided spectral-wise transformer for efficient hyperspectral image reconstruction. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 17502–17511, 2022.
- [10] Martin Uecker, Peng Lai, Mark J. Murphy, Patrick Virtue, Michael Elad, John M. Pauly, Shreyas S. Vasanawala, and Michael Lustig. ESPIRiT — an eigenvalue approach to autocalibrating parallel MRI: where SENSE meets GRAPPA. *Magnetic Resonance in Medicine*, 71(3):990–1001, 2014. doi: 10.1002/mrm.24751.
- [11] Andrew M. Maiden and John M. Rodenburg. An improved ptychographical phase retrieval algorithm for diffractive imaging. *Ultramicroscopy*, 109(10):1256–1262, 2009. doi: 10.1016/j.ultramic.2009.05.012.
- [12] Vegard Antun, Francesco Renna, Clarice Poon, Ben Adcock, and Anders C. Hansen. On instabilities of deep learning in image reconstruction and the potential costs of AI. *Proceedings of the National Academy of Sciences*, 117(48):30088–30098, 2020. doi: 10.1073/pnas.1907377117.
- [13] J. M. Rodenburg and H. M. L. Faulkner. A phase retrieval algorithm for shifting illumination. *Applied Physics Letters*, 85(20):4795–4797, 2004. doi: 10.1063/1.1823034.
- [14] L. A. Feldkamp, L. C. Davis, and J. W. Kress. Practical cone-beam algorithm. *Journal of the Optical Society of America A*, 1(6):612–619, 1984. doi: 10.1364/JOSAA.1.000612.
- [15] Klaas P. Pruessmann, Markus Weiger, Markus B. Scheidegger, and Peter Boesiger. SENSE: Sensitivity encoding for fast MRI. *Magnetic Resonance in Medicine*, 42(5):952–962, 1999. doi: 10.1002/(SICI)1522-2594(199911)42:5<952::AID-MRM16>3.0.CO;2-S.
- [16] Chengshuai Yang. The finite primitive basis theorem for computational imaging: Formal foundations of the OperatorGraph representation. *Under review at SIAM Journal on Imaging Sciences*, 2026. Manuscript available at https://github.com/integritynoble/Physics_World_Model.

- [17] Chengshuai Yang. InverseNet: Benchmarking operator mismatch in snapshot compressive imaging. In *Under review at ECCV 2026*, 2026.
- [18] Lishun Wang, Miao Cao, and Xin Yuan. EfficientSCI: Densely connected network with space-time factorization for large-scale video snapshot compressive imaging. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 18477–18486, 2023.
- [19] Inchang Choi, Daniel S. Jeon, Giljoo Nam, Diego Gutierrez, and Min H. Kim. High-quality hyperspectral reconstruction using a spectral prior. *ACM Transactions on Graphics (Proceedings of SIGGRAPH Asia)*, 36(6):218:1–218:13, 2017. doi: 10.1145/3130800.3130810.
- [20] Xiaowan Hu, Yuanhao Cai, Jing Lin, Haoqian Wang, Xin Yuan, Yulun Zhang, Radu Timofte, and Luc Van Gool. HDNet: High-resolution dual-domain learning for spectral compressive imaging. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 17542–17551, 2022.
- [21] José M. Bioucas-Dias and Mário A. T. Figueiredo. A new TwIST: Two-step iterative shrinkage/thresholding algorithms for image restoration. *IEEE Transactions on Image Processing*, 16(12):2992–3004, 2007. doi: 10.1109/TIP.2007.909319.
- [22] Xin Yuan. Generalized alternating projection based total variation minimization for compressive sensing. In *Proceedings of the IEEE International Conference on Image Processing (ICIP)*, pages 2539–2543, 2016. doi: 10.1109/ICIP.2016.7532817.
- [23] Tao Huang, Weisheng Dong, Xin Yuan, Jinjian Wu, and Guangming Shi. Deep gaussian scale mixture prior for spectral compressive imaging. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 16216–16225, 2021.
- [24] Yuanhao Cai, Jing Lin, Xiaowan Hu, Haoqian Wang, Xin Yuan, Yulun Zhang, Radu Timofte, and Luc Van Gool. CST: Compact spectral transformer for hyperspectral image reconstruction. In *Proceedings of the European Conference on Computer Vision (ECCV)*, 2022.
- [25] Yang Liu, Xin Yuan, Jinli Suo, David J. Brady, and Qionghai Dai. Rank minimization for snapshot compressive imaging. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 41(12):2990–3006, 2019.
- [26] Hu Chen, Yi Zhang, Mannudeep K. Kalra, Feng Lin, Yang Chen, Peixi Liao, Jiliu Zhou, and Ge Wang. Low-dose CT with a residual encoder-decoder convolutional neural network. *IEEE Transactions on Medical Imaging*, 36(12):2524–2535, 2017. doi: 10.1109/TMI.2017.2715284.

830 [27] Jure Zbontar, Florian Knoll, Anuroop Sriram, Tullie Murrell, Zhengnan Huang,
831 Matthew J. Muckley, Aaron Defazio, Ruben Stern, Patricia Johnson, Mary Bruno,
832 Marc Parente, Krzysztof J. Geras, Joe Katsnelson, Hersh Chandarana, Zizhao Zhang,
833 Michal Drozdal, Adriana Romero, Michael Rabbat, Pascal Vincent, Nafissa Yakubova,
834 James Pinkerton, Duo Wang, Erich Owens, C. Lawrence Zitnick, Michael P. Recht,
835 Daniel K. Sodickson, and Yvonne W. Lui. fastMRI: An open dataset and benchmarks
836 for accelerated MRI. *arXiv preprint arXiv:1811.08839*, 2018.




Figure 1: **PWM overview.** The Physics World Models pipeline. **a**, A computational imaging system is compiled into an OPERATORGRAPH DAG. **b**, The TRIAD DECOMPOSITION diagnostic agents evaluate each gate. **c**, The dominant gate is identified and a TRIADREPORT is produced. **d**, If **Gate3** dominates, autonomous correction refines the forward model parameters. **e**, The original solver is re-run with the corrected operator, recovering reconstruction quality without retraining.




Figure 2: **OperatorGraph IR and Physics Fidelity Ladder.** **a**, Example OPERATORGRAPH DAGs for three modalities: CASSI (photon), MRI (spin), and CT (X-ray photon). Each node wraps a primitive operator; edges define data flow. **b**, The Physics Fidelity Ladder. Tier 1: linear shift-invariant. Tier 2: linear shift-variant. Tier 3: nonlinear ray/wave-based. Tier 4: full-wave/Monte Carlo. **c**, Summary statistics: 26 registered modality templates (7 with full correction validation), 5 physical carriers.




Figure 3: **Triad Decomposition structure and gate binding.** **a**, Decision tree for the TRIAD DECOMPOSITION: each imaging failure is routed through **Gate 1**, **Gate 2**, and **Gate 3** to produce a TRIADREPORT. **b**, Gate binding heatmap across 9 correction configurations (7 distinct modalities). Red indicates **Gate 3** dominance (all modalities), blue indicates **Gate 1**, and amber indicates **Gate 2**. **c**, Recovery ratio ρ distribution across all 9 correction configurations.





Figure 4: **Correction results and zero-shot generalization.** **a**, Correction gain Δ_{corr} (dB) across 9 validated configurations (7 distinct modalities), grouped by carrier family. **b**, Zero-shot generalization: hyperparameters tuned on photon-domain modalities transfer without modification to coherent-photon, spin, and X-ray domains. Dark bars: modality-specific tuning; light bars: zero-shot transfer.



Visual reconstruction comparison across three modalities




Figure 5: **Modality deep dives and visual comparison.** **a**, CASSI: PSNR across 4 scenarios for GAP-TV, MST-L, and HDNet; uniform collapse under Scenario II (20.83–21.88 dB) confirms operator-driven failure. **b**, CACTI: four methods across 4 scenarios, showing up to 20.58 dB mismatch degradation. **c–e**, Visual reconstruction comparison (CASSI, CACTI, MRI): ground truth, Scenario I, Scenario II (structured artifacts), Scenario III (PWM-corrected), and error map. Mismatch artifacts are qualitatively distinct from missing data artifacts.




Figure 6: **Hardware validation on real CASSI and CACTI instruments.** **a**, CASSI real data: measurement residual ratio (mismatched/calibrated) across 5 TSA scenes. GAP-TV shows 1.8 $\times$ mean ratio. **b**, CACTI real data: residual ratio across 4 scenes. GAP-TV shows 10.4 $\times$ mean ratio; PnP-FFDNet shows 2.0 $\times$. **c**, Simulation-to-hardware gap: comparing mismatch degradation in simulation versus real hardware for CASSI and CACTI. **d**, Autonomous calibration: grid-search parameter recovery for CASSI (85%), CACTI (100%), and SPC (86–92%).




Figure 7: **Basis-growth saturation and primitive decomposition.** **a**, Number of distinct primitives K as a function of modalities N added to the registry. The curve saturates at $K = 10$ for $N \geq 27$; annotated points mark the introduction of each primitive. **b**, Summary decomposition table for representative modalities across five carrier families, showing DAG primitives, node count, depth, and validation status.

837 **Extended Data Table 1** | Oracle correction ceiling (Scenario IV) across 9 validated
838 configurations (7 modalities). Full table with per-scene metrics in Supplementary Table S1;
839 autonomous correction recovery (Scenario III) in Supplementary Table S7.

840 **Extended Data Table 2** | 26-modality template registry spanning five physical carrier
841 families. Full table in Supplementary Table S3.