

# Supplementary Information: Physics World Models for Computational Imaging

Chengshuai Yang and Xin Yuan

## Contents

|           |                                                                           |           |
|-----------|---------------------------------------------------------------------------|-----------|
| <b>1</b>  | <b>Supplementary Note 1: Triad Decomposition Mathematical Derivations</b> | <b>3</b>  |
| 1.1       | Gate 1 — Information-Theoretic Limit (Compression) . . . . .              | 3         |
| 1.2       | Gate 2 — SNR-Dependent PSNR Bound (Noise) . . . . .                       | 3         |
| 1.3       | Gate 3 — Sensitivity to Model Mismatch (Calibration) . . . . .            | 4         |
| 1.4       | Recovery Ratio . . . . .                                                  | 4         |
| 1.5       | Conditions for Gate 3 Dominance . . . . .                                 | 4         |
| <b>2</b>  | <b>Supplementary Note 2: Complete OperatorGraph Specification</b>         | <b>6</b>  |
| 2.1       | Node Interface . . . . .                                                  | 6         |
| 2.2       | Edge Semantics . . . . .                                                  | 6         |
| 2.3       | Compilation Algorithm . . . . .                                           | 6         |
| 2.4       | JSON Serialization Schema . . . . .                                       | 7         |
| <b>3</b>  | <b>Supplementary Table S1: All 16 Correction Configurations</b>           | <b>8</b>  |
| <b>4</b>  | <b>Supplementary Table S2: CASSI Per-Scene Results</b>                    | <b>10</b> |
| <b>5</b>  | <b>Supplementary Table S3: 26-Modality Template Registry</b>              | <b>12</b> |
| <b>6</b>  | <b>Supplementary Table S4: YAML Registry Summary</b>                      | <b>13</b> |
| <b>7</b>  | <b>Supplementary Note 4: RunBundle Schema</b>                             | <b>14</b> |
| 7.1       | RunBundle v0.3.0 Specification . . . . .                                  | 14        |
| 7.2       | Integrity Verification . . . . .                                          | 14        |
| <b>8</b>  | <b>Supplementary Note 5: Computational Cost Analysis</b>                  | <b>16</b> |
| 8.1       | Runtime per Modality . . . . .                                            | 16        |
| 8.2       | RoIC Metric: Efficiency of Correction . . . . .                           | 16        |
| 8.3       | Scaling Considerations . . . . .                                          | 16        |
| <b>9</b>  | <b>Supplementary Note 6: Real-Data Validation Details</b>                 | <b>18</b> |
| 9.1       | CASSI Real Data: TSA Hyperspectral Camera . . . . .                       | 18        |
| 9.2       | CACTI Real Data: Temporal Compressive Camera . . . . .                    | 18        |
| 9.3       | Autonomous Calibration on Real Data . . . . .                             | 20        |
| <b>10</b> | <b>Supplementary Table S8: SSIM Comparison Across Modalities</b>          | <b>21</b> |

|                                                                             |           |
|-----------------------------------------------------------------------------|-----------|
| <b>11 Supplementary Table S9: CASSI Spectral Angle Mapper (SAM)</b>         | <b>21</b> |
| <b>12 Supplementary Tables S10–S11: Gate 1 and Gate 2 Validation</b>        | <b>23</b> |
| 12.1 Gate 1: Information Deficiency (Extreme Compression) . . . . .         | 23        |
| 12.2 Gate 2: Carrier Budget (Noise Sweep) . . . . .                         | 24        |
| <b>13 Supplementary Note 7: Clinical CT Quality Assurance Validation</b>    | <b>25</b> |
| 13.1 Gate Mapping to Clinical Failure Modes . . . . .                       | 25        |
| 13.2 ACR Metric Validation . . . . .                                        | 25        |
| 13.3 Drift Detection Performance . . . . .                                  | 25        |
| 13.4 Workflow Efficiency . . . . .                                          | 26        |
| <b>14 Supplementary Note 8: Controlled Hardware Validation</b>              | <b>27</b> |
| 14.1 CASSI Physical Mask Displacement Protocol . . . . .                    | 27        |
| 14.2 Multi-Unit Variation Study Protocol . . . . .                          | 27        |
| <b>15 Supplementary Note 9: Mismatch Parameter Derivation</b>               | <b>28</b> |
| 15.1 CASSI 5-Parameter Mismatch Model . . . . .                             | 28        |
| 15.2 Prospective Clinical Validation Protocol . . . . .                     | 28        |
| <b>16 Supplementary Note S12: Calibration Method Comparison</b>             | <b>30</b> |
| <b>17 Supplementary Note S13: MRI Under Clinically Realistic Conditions</b> | <b>30</b> |
| <b>18 Supplementary Note 10: Finite Primitive Basis — Expanded Proof</b>    | <b>31</b> |
| 18.1 Formal Definitions . . . . .                                           | 31        |
| 18.2 Proof of Theorem 1 . . . . .                                           | 32        |
| 18.3 Extension Protocol: Worked Example . . . . .                           | 33        |

# 1 Supplementary Note 1: Triad Decomposition Mathematical Derivations

The Triad Decomposition formalizes three successive gates that every computational-imaging reconstruction must pass. We derive rigorous bounds for each gate and show how the *recovery ratio*  $\rho$  summarizes the overall reconstruction quality.

## 1.1 Gate 1 — Information-Theoretic Limit (Compression)

Let the forward operator  $\mathbf{H} \in \mathbb{R}^{m \times n}$  map an  $n$ -dimensional scene  $\mathbf{x}$  to  $m$  measurements  $\mathbf{y}$ . Define the *compression ratio*  $\gamma = m/n$ .

**Theorem 1** (Compression bound). *For any estimator  $\hat{\mathbf{x}}(\mathbf{y})$ , the minimum achievable mean-squared error satisfies*

$$\text{MSE}_{\min} \geq \frac{1}{n} \sum_{i=1}^{n-m} \sigma_i^2(\mathbf{x}), \quad (\text{S1})$$

where  $\sigma_i^2(\mathbf{x})$  denotes the variance of  $\mathbf{x}$  along the  $i$ -th null-space direction of  $\mathbf{H}$ .

**Interpretation.** The null space of  $\mathbf{H}$  has dimension  $\dim \mathcal{N}(\mathbf{H}) = n - \text{rank}(\mathbf{H}) \geq n - m$ . Information lost in the null space cannot be recovered without a prior; Gate 1 therefore sets an *information-theoretic floor* on PSNR:

$$\text{PSNR}_{\max}^{(G1)} = 10 \log_{10} \left( \frac{\|\mathbf{x}\|_{\infty}^2}{\text{MSE}_{\min}} \right). \quad (\text{S2})$$

Modalities with higher compression ratios (e.g. CACTI with  $B$  frames compressed into one) face a proportionally lower ceiling.

## 1.2 Gate 2 — SNR-Dependent PSNR Bound (Noise)

Given additive noise  $\mathbf{n} \sim \mathcal{N}(0, \sigma_n^2 \mathbf{I})$ , the measurement SNR is  $\text{SNR} = \|\mathbf{H}\mathbf{x}\|^2 / (m \sigma_n^2)$ .

**Theorem 2** (Noise bound). *Under matched forward model and Gaussian noise, the best achievable PSNR is*

$$\text{PSNR}_{\max}^{(G2)} = 10 \log_{10}(\text{SNR}) + C_{\mathcal{M}}, \quad (\text{S3})$$

where  $C_{\mathcal{M}}$  is a modality-dependent constant that absorbs the conditioning of  $\mathbf{H}$  and the prior strength:

$$C_{\mathcal{M}} = 10 \log_{10} \left( \frac{n \|\mathbf{x}\|_{\infty}^2}{\|\mathbf{x}\|^2} \cdot \kappa^{-2}(\mathbf{H}) \right), \quad (\text{S4})$$

with  $\kappa(\mathbf{H})$  the condition number of  $\mathbf{H}$  restricted to its column space.

**Interpretation.** Gate 2 is a *noise ceiling*: no algorithm can exceed  $\text{PSNR}_{\max}^{(G2)}$  regardless of computational budget. Empirically, modalities such as MRI ( $C_{\text{MRI}} \approx 8$  dB) are more noise-tolerant than CASSI ( $C_{\text{CASSI}} \approx 3$  dB) due to favorable operator conditioning.

### 1.3 Gate 3 — Sensitivity to Model Mismatch (Calibration)

Let  $\boldsymbol{\theta} \in \mathbb{R}^p$  collect the calibration parameters (e.g. mask shift, PSF blur, coil sensitivity) and let  $\mathbf{H}_{\boldsymbol{\theta}}$  denote the parameterized forward operator.

**Theorem 3** (Calibration sensitivity). *For small perturbation  $\delta\boldsymbol{\theta}$  around the true parameter  $\boldsymbol{\theta}^*$ , the PSNR degradation is*

$$\Delta\text{PSNR} \approx -\frac{10}{\ln 10} \frac{\delta\boldsymbol{\theta}^\top \mathbf{J}^\top \mathbf{J} \delta\boldsymbol{\theta}}{\text{MSE}_0}, \quad (\text{S5})$$

where  $\mathbf{J} = \partial(\mathbf{H}_{\boldsymbol{\theta}}\mathbf{x})/\partial\boldsymbol{\theta}|_{\boldsymbol{\theta}^*}$  is the parameter Jacobian and  $\text{MSE}_0$  is the noise-only MSE at  $\boldsymbol{\theta}^*$ .

**Per-parameter sensitivity.** Restricting to a single parameter  $\theta_j$ :

$$\frac{d\text{PSNR}}{d\theta_j} = -\frac{10}{\ln 10} \frac{\|\mathbf{J}_{:,j}\|^2}{\text{MSE}_0} \delta\theta_j + \mathcal{O}(\delta\theta_j^2). \quad (\text{S6})$$

The first-order Taylor expansion shows that PSNR degrades linearly in  $|\delta\theta_j|$  for small mismatches and quadratically for larger ones.

### 1.4 Recovery Ratio

We define the *recovery ratio* as

$$\rho = \frac{\text{PSNR}_{\text{achieved}} - \text{PSNR}_{\text{mismatch}}}{\text{PSNR}_{\text{ideal}} - \text{PSNR}_{\text{mismatch}}}, \quad 0 \leq \rho \leq 1. \quad (\text{S7})$$

In rare cases where Scenario III exceeds Scenario I (e.g., due to regularization benefits from the corrected operator),  $\rho$  may exceed 1.

**Proposition 4.** *Under convex reconstruction loss with matched regularization and convex constraint on  $\boldsymbol{\theta}$ , the recovery ratio satisfies  $0 \leq \rho \leq 1$  for any estimator.*

*Proof sketch.* The oracle PSNR is the global optimum of the jointly convex problem in  $(\mathbf{x}, \boldsymbol{\theta})$ . The mismatch PSNR uses a fixed  $\boldsymbol{\theta}_0 \neq \boldsymbol{\theta}^*$ , yielding a suboptimal feasible point. Any correction step moves towards the optimum, so  $\text{PSNR}_{\text{mismatch}} \leq \text{PSNR}_{\text{achieved}} \leq \text{PSNR}_{\text{oracle}}$ , giving  $\rho \in [0, 1]$ .  $\square$

Under the conditions of Proposition 1,  $0 \leq \rho \leq 1$ . In practice, beneficial regularization bias from the corrected operator can yield  $\rho > 1$ , as observed for CACTI.

### 1.5 Conditions for Gate 3 Dominance

The empirical finding that Gate 3 dominates across all validated modalities admits a theoretical justification.

**Proposition 5** (Gate 3 dominance condition). *For any instrument operating above its Gate 1 floor ( $\gamma > \gamma_{\min}$ , where  $\gamma_{\min}$  is the minimum compression ratio for which  $\text{PSNR}_{\max}^{(G1)}$  exceeds the target quality) and Gate 2 floor ( $\text{SNR} > \text{SNR}_{\min}$ , where  $\text{SNR}_{\min}$  is the minimum carrier budget for which  $\text{PSNR}_{\max}^{(G2)}$  exceeds the target quality), Gate 3 becomes the binding constraint whenever the calibration error exceeds the noise-equivalent resolution:*

$$\|\delta\boldsymbol{\theta}\|_{\mathbf{J}^\top \mathbf{J}} > \frac{\sigma_n}{\|\mathbf{x}\|_\infty} \cdot \kappa(\mathbf{H}), \quad (\text{S8})$$

where  $\|\delta\boldsymbol{\theta}\|_{\mathbf{J}^\top \mathbf{J}} = \sqrt{\delta\boldsymbol{\theta}^\top \mathbf{J}^\top \mathbf{J} \delta\boldsymbol{\theta}}$  is the parameter-space mismatch norm weighted by the sensitivity Jacobian,  $\sigma_n$  is the noise standard deviation, and  $\kappa(\mathbf{H})$  is the condition number of the forward operator.

*Proof sketch.* From Theorem 3 (Eq. S5), the mismatch-induced PSNR degradation is  $\Delta\text{PSNR}_{\text{mismatch}} \approx (10/\ln 10) \cdot \|\delta\boldsymbol{\theta}\|_{\mathbf{J}^\top \mathbf{J}}^2 / \text{MSE}_0$ . From Theorem 2 (Eq. S3), the noise-induced PSNR ceiling is  $\text{PSNR}_{\text{max}}^{(G2)} = 10 \log_{10}(\text{SNR}) + C_{\mathcal{M}}$ . Gate 3 dominates when  $\Delta\text{PSNR}_{\text{mismatch}} > \Delta\text{PSNR}_{\text{noise}}$ , which after substitution yields condition (S8). The condition is satisfied in modern instruments because: (i) well-designed systems operate with  $\gamma \gg \gamma_{\text{min}}$  and  $\text{SNR} \gg \text{SNR}_{\text{min}}$ , pushing Gates 1 and 2 well above their floors; and (ii) the sensitivity Jacobian  $\mathbf{J}$  amplifies even small parameter errors ( $\|\delta\boldsymbol{\theta}\| \ll 1$ ) into substantial measurement discrepancies, because the forward operators of modern high-resolution instruments are highly sensitive to calibration parameters.  $\square$

**Interpretation.** Proposition 2 formalizes the intuition that as instruments improve in information capacity (Gate 1) and signal quality (Gate 2), the binding constraint shifts to calibration fidelity (Gate 3). This is precisely the regime occupied by modern computational imaging systems, explaining the universal Gate 3 dominance observed empirically across all seven validated modalities.

## 2 Supplementary Note 2: Complete OperatorGraph Specification

The `OperatorGraph` is the central abstraction of the PWM framework. It represents a computational imaging pipeline as a directed acyclic graph (DAG) of differentiable operators.

### 2.1 Node Interface

Every node  $v$  in the graph implements the following interface:

| Method                            | Contract                                                                                                                                                                                                               |
|-----------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <code>forward(x) → y</code>       | Apply the physical forward model. Input shape: (B, C, *spatial). Output shape determined by the operator (e.g. compression changes spatial dims).                                                                      |
| <code>adjoint(y) → x</code>       | Apply the adjoint (transpose) operator. Must satisfy $\langle \mathbf{H}\mathbf{x}, \mathbf{y} \rangle = \langle \mathbf{x}, \mathbf{H}^\top \mathbf{y} \rangle$ to numerical precision ( $< 10^{-6}$ relative error). |
| <code>shape_in / shape_out</code> | Static shape annotations for compile-time validation.                                                                                                                                                                  |
| <code>dtype</code>                | Data type contract: <code>float32</code> (default) or <code>complex64</code> .                                                                                                                                         |
| <code>parameters() → dict</code>  | Returns learnable calibration parameters with their current values.                                                                                                                                                    |

### 2.2 Edge Semantics

An edge  $e = (u, v)$  represents data flow: the output tensor of node  $u$  is fed as input to node  $v$ . Edges carry the following metadata:

- **Tensor shape:** inferred at compile time via shape propagation.
- **Data type:** must match between source output and target input.
- **Optional transform:** lightweight reshape or permute operations (e.g. channel stacking for multi-frame CACTI data).

### 2.3 Compilation Algorithm

Given an `OperatorGraph`  $G = (V, E)$ :

1. **Topological sort.** Compute a valid execution order  $v_1, v_2, \dots, v_{|V|}$  using Kahn’s algorithm. Raise `CycleError` if the graph contains a cycle.
2. **Shape propagation.** Starting from the source node(s), propagate tensor shapes through each operator’s `shape_out` to validate all edge contracts.
3. **Automatic adjoint chain.** Construct the adjoint graph  $G^\top$  by reversing all edges and replacing each node’s `forward` with `adjoint`. The adjoint graph is used for gradient-based reconstruction and the Triad Decomposition sensitivity analysis.

4. **Fusion (optional).** Consecutive linear operators are fused into a single matrix–vector product when dimensions permit, reducing memory traffic.

## 2.4 JSON Serialization Schema

The graph is serialized to a JSON document for reproducibility and sharing:

```
{
  "version": "0.3.0",
  "nodes": [
    {
      "id": "mask",
      "type": "CassiMask",
      "params": {"shift_r": 0.0, "shift_c": 0.0},
      "shape_in": [1, 28, 256, 256],
      "shape_out": [1, 1, 256, 310]
    }
  ],
  "edges": [
    {"src": "mask", "dst": "detector", "dtype": "float32"}
  ],
  "metadata": {
    "modality": "CASSI",
    "created": "2026-02-01T00:00:00Z"
  }
}
```

### 3 Supplementary Table S1: All 16 Correction Configurations

Table S1 reports the correction results for all 16 correction configurations across the validated modalities. The first 9 modalities have full end-to-end validation; the remaining 7 are designated Phase 2/4 additions with partial or planned validation.

Table S1: Complete correction parameter summary.  $\Delta$ PSNR is the improvement from Scenario II (mismatch) to Scenario III (corrected). CASSI correction results in this table use GAP-TV as the reconstruction solver; the main-text deep dive reports results for GAP-TV, MST-L, and HDNet. Scenario I (ideal) PSNR values and recovery ratios ( $\rho$ ) for all configurations are available in the RunBundle manifests in the code repository.

| Modality                                            | Mismatch<br>Parameter    | Sc. II<br>(dB) | Sc. III<br>(dB) | $\Delta$ PSNR<br>(dB) | RMSE   |
|-----------------------------------------------------|--------------------------|----------------|-----------------|-----------------------|--------|
| Matrix <sup>1</sup>                                 | gain_bias                | 11.14          | 23.35           | +12.21                | 0.0042 |
| CT                                                  | center_of_rotation       | 13.41          | 24.09           | +10.68                | 0.0031 |
| CACTI <sup>2</sup>                                  | 8-param<br>mismatch      | 15.81          | 26.01           | +10.21                | —      |
| Lensless                                            | psf_shift                | 23.48          | 27.03           | +3.55                 | 0.0089 |
| MRI                                                 | coil_sensitivities       | 6.94           | 55.19           | +48.25                | 0.0003 |
| SPC <sup>3</sup>                                    | gain_drift               | 18.51          | 26.21           | +7.71                 | —      |
| CASSI<br>(Alg 1) <sup>4</sup>                       | mask_geo +<br>dispersion | 20.96          | 21.50           | +0.54                 | 0.0105 |
| CASSI<br>(Alg 2) <sup>5</sup>                       | mask_geo +<br>dispersion | 20.96          | 21.72           | +0.76                 | 0.0098 |
| Ptychography                                        | position_offset          | 17.35          | 24.44           | +7.09                 | 0.0063 |
| <i>Phase 2/4 additions — validation in progress</i> |                          |                |                 |                       |        |
| Holography                                          | propagation_distance     | —              | —               | —                     | —      |
| FPM                                                 | LED_positions            | —              | —               | —                     | —      |
| Phase<br>Retrieval                                  | support_constraint       | —              | —               | —                     | —      |

*Continued on next page*

<sup>1</sup>Matrix (Generic) is a canonical test configuration that uses the same gain-bias operator template as SPC; identical results confirm pipeline consistency rather than representing an independent modality.

<sup>2</sup>CACTI uses GAP-TV as the reconstruction solver for Table S1 consistency. Oracle values from InverseNet validation under 8-parameter mismatch ( $dx=0.5$  px,  $dy=0.3$  px,  $\theta=0.1^\circ$ ,  $\Delta t=0.05$ ,  $\eta=0.95$ ,  $g=1.02$ ,  $o=0.002$ ,  $\sigma_n=1.0$ ). State-of-the-art EfficientSCI achieves Sc. II = 14.81 dB and Sc. III = 27.38 dB ( $\rho=61.1\%$ ); see main text.

<sup>3</sup>SPC uses FISTA-TV as the reconstruction solver for Table S1 consistency. Values from InverseNet validation under exponential gain drift ( $\alpha=0.0015$ ). HATNet achieves Sc. II = 19.40 dB and Sc. III = 29.78 dB ( $\rho=89.6\%$ ); see main text.

<sup>4</sup>Reported values are averaged over 10 test scenes using GAP-TV as the reconstruction solver. Sc. II from InverseNet validation; Sc. III values are oracle upper bounds (true operator applied to mismatched data).

<sup>5</sup>Note: the correction-table PSNR reflects GAP-TV reconstruction; state-of-the-art solvers (e.g., MST-L) achieve 27.33 dB oracle PSNR under the corrected operator (see main text).



| Modality     | Mismatch<br>Parameter | Sc. II<br>(dB) | Sc. III<br>(dB) | $\Delta$ PSNR<br>(dB) | RMSE |
|--------------|-----------------------|----------------|-----------------|-----------------------|------|
| OCT          | dispersion_coeffs     | —              | —               | —                     | —    |
| Fluorescence | PSF_aberration        | —              | —               | —                     | —    |
| Ultrasound   | speed_of_sound        | —              | —               | —                     | —    |
| Radar/SAR    | motion_compensation   | —              | —               | —                     | —    |

## 4 Supplementary Table S2: CASSI Per-Scene Results

Table S2 reports PSNR (dB) and recovery ratio  $\rho$  (%) for 10 KAIST scenes across four methods under three scenarios, matching the InverseNet benchmark<sup>1</sup>.

Table S2: CASSI per-scene PSNR (dB) and recovery ratio  $\rho$ . Sc. I = ideal forward model; Sc. II = mismatched (baseline); Sc. III = oracle correction (true operator applied to mismatched data). All values from InverseNet validation under the 5-parameter mismatch ( $dx=0.5$  px,  $dy=0.3$  px,  $\theta=0.1^\circ$ ,  $a_1=2.02$ ,  $\alpha=0.15^\circ$ ). HDNet shows  $\rho=0\%$  across all scenes due to its mask-oblivious architecture.

| Scene   | Method     | Sc. I | Sc. II | Sc. III | $\rho$ |
|---------|------------|-------|--------|---------|--------|
| scene01 | GAP-TV     | 26.49 | 24.08  | 24.18   | 4.1%   |
| scene01 | PnP-HSICNN | 27.43 | 23.47  | 25.78   | 58.3%  |
| scene01 | HDNet      | 34.95 | 24.37  | 24.37   | 0.0%   |
| scene01 | MST-L      | 35.29 | 23.96  | 29.98   | 53.2%  |
| scene02 | GAP-TV     | 24.60 | 21.89  | 22.82   | 34.5%  |
| scene02 | PnP-HSICNN | 25.34 | 21.56  | 23.80   | 59.3%  |
| scene02 | HDNet      | 35.65 | 23.26  | 23.26   | 0.0%   |
| scene02 | MST-L      | 36.14 | 22.21  | 28.42   | 44.6%  |
| scene03 | GAP-TV     | 25.96 | 18.62  | 19.68   | 14.4%  |
| scene03 | PnP-HSICNN | 26.71 | 17.59  | 21.83   | 46.5%  |
| scene03 | HDNet      | 35.54 | 18.61  | 18.61   | 0.0%   |
| scene03 | MST-L      | 35.66 | 16.09  | 23.57   | 38.2%  |
| scene04 | GAP-TV     | 28.36 | 23.37  | 24.41   | 20.9%  |
| scene04 | PnP-HSICNN | 28.80 | 22.89  | 25.98   | 52.3%  |
| scene04 | HDNet      | 41.63 | 23.98  | 23.98   | 0.0%   |
| scene04 | MST-L      | 40.05 | 21.91  | 29.80   | 43.5%  |
| scene05 | GAP-TV     | 23.66 | 20.39  | 21.27   | 26.9%  |
| scene05 | PnP-HSICNN | 24.65 | 19.36  | 22.75   | 64.3%  |
| scene05 | HDNet      | 32.56 | 20.22  | 20.22   | 0.0%   |
| scene05 | MST-L      | 32.84 | 20.28  | 26.75   | 51.5%  |
| scene06 | GAP-TV     | 22.34 | 20.39  | 21.00   | 31.2%  |
| scene06 | PnP-HSICNN | 23.12 | 20.29  | 21.80   | 53.0%  |
| scene06 | HDNet      | 34.33 | 22.63  | 22.63   | 0.0%   |
| scene06 | MST-L      | 34.56 | 22.37  | 28.67   | 51.7%  |
| scene07 | GAP-TV     | 23.51 | 20.56  | 21.07   | 17.4%  |
| scene07 | PnP-HSICNN | 24.94 | 19.86  | 22.77   | 57.4%  |
| scene07 | HDNet      | 33.27 | 20.79  | 20.79   | 0.0%   |
| scene07 | MST-L      | 33.80 | 19.76  | 26.12   | 45.3%  |
| scene08 | GAP-TV     | 22.16 | 20.57  | 21.10   | 33.3%  |
| scene08 | PnP-HSICNN | 22.73 | 20.30  | 21.92   | 66.4%  |
| scene08 | HDNet      | 32.26 | 22.73  | 22.73   | 0.0%   |
| scene08 | MST-L      | 32.74 | 21.33  | 27.44   | 53.5%  |
| scene09 | GAP-TV     | 23.03 | 19.14  | 20.61   | 37.8%  |
| scene09 | PnP-HSICNN | 24.03 | 18.79  | 21.82   | 57.8%  |
| scene09 | HDNet      | 34.18 | 21.18  | 21.18   | 0.0%   |
| scene09 | MST-L      | 34.37 | 19.75  | 26.71   | 47.6%  |
| scene10 | GAP-TV     | 23.27 | 20.58  | 21.04   | 16.9%  |
| scene10 | PnP-HSICNN | 23.47 | 19.90  | 22.38   | 69.4%  |
| scene10 | HDNet      | 32.22 | 21.06  | 21.06   | 0.0%   |
| scene10 | MST-L      | 32.63 | 20.69  | 25.86   | 43.3%  |

| Scene | Method     | Sc. I | Sc. II | Sc. III | $\rho$ |
|-------|------------|-------|--------|---------|--------|
| Mean  | GAP-TV     | 24.34 | 20.96  | 21.72   | 22.5%  |
| Mean  | PnP-HSICNN | 25.12 | 20.40  | 23.08   | 56.8%  |
| Mean  | HDNet      | 34.66 | 21.88  | 21.88   | 0.0%   |
| Mean  | MST-L      | 34.81 | 20.83  | 27.33   | 46.5%  |

## 5 Supplementary Table S3: 26-Modality Template Registry

Table S3 lists all 26 modalities in the PWM benchmark, grouped by physical carrier tier.

Table S3: 26-modality template registry. PSNR values are representative (Scenario I, default solver). “Ref. PSNR” is the best published result where available. Validation tiers: *Validated* = full Scenarios I–IV correction; *Phase 2* = template validated, correction in progress; *Phase 4* = template registered, validation planned.

| #  | Modality          | Domain     | PSNR<br>(dB) | Ref. PSNR<br>(dB) | Status    | Solver       |
|----|-------------------|------------|--------------|-------------------|-----------|--------------|
| 1  | Matrix (toy)      | Incoherent | 23.35        | —                 | Validated | ADMM-TV      |
| 2  | SPC               | Incoherent | 28.06        | 31.0              | Validated | FISTA-TV     |
| 3  | CASSI             | Incoherent | 24.34        | 34.2              | Validated | GAP-TV       |
| 4  | CACTI             | Incoherent | 35.39        | 33.4              | Validated | EfficientSCI |
| 5  | Lensless          | Incoherent | 27.03        | 28.5              | Validated | ADMM-TV      |
| 6  | Fluorescence      | Incoherent | —            | 30.1              | Phase 4   | —            |
| 7  | Holography        | Coherent   | —            | 32.7              | Phase 2   | —            |
| 8  | Ptychography      | Coherent   | 24.44        | 26.8              | Validated | ePIE         |
| 9  | FPM               | Coherent   | —            | 31.5              | Phase 2   | —            |
| 10 | Phase Retrieval   | Coherent   | —            | 29.3              | Phase 2   | —            |
| 11 | CDI               | Coherent   | —            | 27.4              | Phase 4   | —            |
| 12 | OCT               | Coherent   | —            | 35.2              | Phase 4   | —            |
| 13 | CT (parallel)     | X-ray      | 24.09        | 42.1              | Validated | FBP + TV     |
| 14 | CT (fan-beam)     | X-ray      | —            | 40.8              | Phase 2   | —            |
| 15 | CT (cone-beam)    | X-ray      | —            | 39.5              | Phase 2   | —            |
| 16 | Electron Pty.     | Electron   | —            | 25.9              | Phase 4   | —            |
| 17 | MRI (Cartesian)   | RF         | 55.19        | 42.3              | Validated | CG-SENSE     |
| 18 | MRI (radial)      | RF         | —            | 38.7              | Phase 2   | —            |
| 19 | MRI (spiral)      | RF         | —            | 37.2              | Phase 4   | —            |
| 20 | Ultrasound        | Acoustic   | —            | 31.8              | Phase 4   | —            |
| 21 | Radar / SAR       | RF         | —            | 28.4              | Phase 4   | —            |
| 22 | NeRF              | Multi-view | —            | 31.7              | Phase 2   | —            |
| 23 | 3DGS              | Multi-view | —            | 33.2              | Phase 2   | —            |
| 24 | Light Field       | Multi-view | —            | 36.1              | Phase 4   | —            |
| 25 | Structured Illum. | Incoherent | —            | 32.4              | Phase 4   | —            |
| 26 | Ghost Imaging     | Incoherent | —            | 22.8              | Phase 4   | —            |

## 6 Supplementary Table S4: YAML Registry Summary

Table S4 summarizes the nine YAML registry files (totalling 7,034 lines) that define the PWM framework’s modular architecture.

Table S4: YAML registry files in `packages/pwm_core/contrib/`.

| # | File                                    | Lines | Purpose                                                                                               |
|---|-----------------------------------------|-------|-------------------------------------------------------------------------------------------------------|
| 1 | <code>modalities.yaml</code>            | 1200  | Defines 64 modality entries with keywords, forward model equations, and default solvers.              |
| 2 | <code>graph_templates.yaml</code>       | 1400  | OperatorGraph skeletons for 89 templates across 64 modalities.                                        |
| 3 | <code>photon_db.yaml</code>             | 624   | Photon models parameterized by source power, quantum efficiency, exposure, and detector.              |
| 4 | <code>mismatch_db.yaml</code>           | 797   | Mismatch parameters, perturbation ranges, and correction methods per modality.                        |
| 5 | <code>compression_db.yaml</code>        | 1186  | Recoverability tables mapping compression ratio to expected PSNR with provenance.                     |
| 6 | <code>solver_registry.yaml</code>       | 650   | Maps solver names to implementations (ADMM <sup>2</sup> , FISTA <sup>3</sup> , GAP-TV, PnP, learned). |
| 7 | <code>primitives.yaml</code>            | 450   | Primitive operator metadata (type, adjoint availability, differentiability).                          |
| 8 | <code>dataset_registry.yaml</code>      | 380   | Links modalities to benchmark datasets (KAIST, fastMRI, etc.).                                        |
| 9 | <code>acceptance_thresholds.yaml</code> | 347   | Pass/fail thresholds per metric for automated validation.                                             |

## 7 Supplementary Note 4: RunBundle Schema

The `RunBundle` is PWM’s reproducibility artifact. Every experiment produces a `RunBundle` that captures the complete computational state.

### 7.1 RunBundle v0.3.0 Specification

Table S5: RunBundle v0.3.0 field specification.

| Field                              | Type   | Description                                                                                                                                                    |
|------------------------------------|--------|----------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <code>version</code>               | string | Schema version, currently "0.3.0".                                                                                                                             |
| <code>git_hash</code>              | string | Full 40-character SHA-1 of the commit used to generate results.                                                                                                |
| <code>rng_seeds</code>             | object | Random seeds for Python ( <code>random</code> , <code>numpy</code> ), PyTorch (CPU, CUDA), and any solver-specific RNG.                                        |
| <code>platform_info</code>         | object | OS, Python version, GPU model, CUDA version, PyTorch version, and <code>pip freeze</code> hash.                                                                |
| <code>sha256_hashes</code>         | object | SHA-256 checksums of all input data files (measurements, masks, ground truth).                                                                                 |
| <code>metrics</code>               | object | Computed quality metrics: PSNR, SSIM, LPIPS, and any modality-specific metrics.                                                                                |
| <code>timestamps</code>            | object | ISO-8601 timestamps for start, end, and per-stage completion.                                                                                                  |
| <code>operator_state</code>        | object | Serialized <code>OperatorGraph</code> including all node parameters and edge metadata.                                                                         |
| <code>triad_report</code>          | object | Triad Decomposition evaluation: gate pass/fail status, PSNR bounds, and recovery ratio $\rho$ .                                                                |
| <code>correction_trajectory</code> | array  | Time series of correction iterations: <code>[{iter, params, loss, psnr, grad_norm}, ...]</code> . Enables convergence analysis and early-stopping diagnostics. |
| <code>solver_config</code>         | object | Complete solver hyperparameters (algorithm, iterations, step size, regularization weight, denoiser checkpoint).                                                |
| <code>modality</code>              | string | Modality identifier matching <code>modalities.yaml</code> .                                                                                                    |
| <code>scenario</code>              | string | One of I (ideal), II (mismatch), III (corrected), IV (oracle_mask).                                                                                            |
| <code>notes</code>                 | string | Free-form text field for experiment annotations.                                                                                                               |

### 7.2 Integrity Verification

A `RunBundle` can be verified by:

1. Checking `git_hash` matches the repository state.
2. Re-computing `sha256_hashes` from the data files.
3. Re-running the solver with the stored `rng_seeds` and `solver_config` and comparing metrics within tolerance ( $\pm 0.01$  dB PSNR).

## 8 Supplementary Note 5: Computational Cost Analysis

### 8.1 Runtime per Modality

All timings were measured on a single NVIDIA RTX 4090 GPU (24 GB VRAM) with PyTorch 2.1 and CUDA 12.1.

Table S6: Computational cost for correction (Scenario II  $\rightarrow$  III). RoIC = dB recovered per GPU-hour.

| Modality      | Iterations | Runtime (s) | $\Delta$ PSNR (dB) | RoIC (dB/GPU-hr) |
|---------------|------------|-------------|--------------------|------------------|
| Matrix        | 200        | 12          | +12.21             | 3663.0           |
| CT            | 500        | 85          | +10.68             | 452.3            |
| CACTI         | 1000       | 180         | +10.21             | 204.2            |
| Lensless      | 300        | 45          | +3.55              | 284.0            |
| MRI           | 150        | 30          | +48.25             | 5790.0           |
| SPC           | 200        | 15          | +7.71              | 1850.4           |
| CASSI (Alg 1) | 500        | 300         | +0.54              | 6.5              |
| CASSI (Alg 2) | 2000       | 3200        | +0.76              | 0.9              |
| Ptychography  | 800        | 420         | +7.09              | 60.8             |

### 8.2 RoIC Metric: Efficiency of Correction

We define the *Return on Invested Compute* (RoIC) as:

$$\text{RoIC} = \frac{\Delta\text{PSNR (dB)}}{T_{\text{GPU (hours)}}}, \quad (\text{S9})$$

where  $T_{\text{GPU}}$  is the wall-clock GPU time.

#### Key observations:

- **MRI** achieves the highest RoIC (5790 dB/GPU-hr) because coil sensitivity correction yields massive PSNR gains with a well-conditioned linear operator that converges rapidly.
- **Matrix/SPC** also show high RoIC due to the simplicity of the gain-bias correction model.
- **CASSI Algorithm 2** has the lowest RoIC (0.9 dB/GPU-hr) because it jointly optimizes mask geometry and dispersion over 2000 iterations, each requiring a full forward-adjoint cycle through the dispersive operator. Under the multi-parameter InverseNet mismatch, the absolute correction gains are modest (+0.76 dB), reflecting the inherent difficulty of simultaneously correcting five coupled mismatch parameters.
- The practical implication is that Algorithm 1 (fixed dispersion, correct mask geometry only) is preferred when compute budget is limited, recovering 71% of Algorithm 2’s  $\Delta$ PSNR at <10% of the cost.

### 8.3 Scaling Considerations

For large-scale deployment:

- Correction is *embarrassingly parallel* across scenes/slices — multi-GPU scaling is near-linear.



- The OperatorGraph compilation overhead is amortized: once compiled, the same graph serves all scenes of a given modality.
- Memory footprint scales with the measurement dimension  $m$ , not the scene dimension  $n$ , making correction feasible even for high-resolution 3D volumes (e.g.  $512^3$  CT).

## 9 Supplementary Note 6: Real-Data Validation Details

### 9.1 CASSI Real Data: TSA Hyperspectral Camera

The TSA (Thermal Snapshot Aperture) real dataset consists of 5 hyperspectral scenes acquired with a coded aperture snapshot spectral imaging system. Each scene has spatial resolution  $660 \times 660$  with 28 spectral bands (450–650 nm) and a mask-shift step of 2 pixels per band. The extended measurement has width  $660 + (28 - 1) \times 2 = 714$  pixels.

**Experimental protocol.** For each scene, we reconstruct with four methods: GAP-TV (classical iterative), HDNet (mask-oblivious neural network), MST-S, and MST-L (mask-guided spectral transformers). Each method is applied under two conditions: (i) the calibrated mask provided with the dataset, and (ii) a perturbed mask with sub-pixel shift  $dx = 0.5$  px,  $dy = 0.3$  px. Because the MST model has a hardcoded  $256 \times 256$  spatial assumption, the  $660 \times 660$  real data is center-cropped to  $256 \times 256$  for MST reconstruction. HDNet is fully convolutional and processes the full spatial extent.

**Metric.** Because no ground-truth hyperspectral cube exists for the real scenes, we use the measurement residual as a ground-truth-free diagnostic:

$$r = \frac{\|\mathbf{y} - H\hat{\mathbf{x}}\|^2}{\|\mathbf{y}\|^2}, \quad (\text{S10})$$

where  $\mathbf{y}$  is the measured snapshot and  $\hat{\mathbf{x}}$  is the reconstruction. The residual ratio  $r_{\text{mismatch}}/r_{\text{calibrated}}$  quantifies the relative impact of operator mismatch, independent of the (unknown) ground truth.

**Key findings.** GAP-TV, which explicitly conditions on the coded aperture mask in its forward–adjoint iterations, shows the expected sensitivity to mask mismatch (mean ratio  $1.8\times$ ). HDNet, whose architecture processes the measurement without explicit mask conditioning, is entirely insensitive (ratio  $1.0\times$ ). The transformer methods (MST-S, MST-L) show ratios near or below  $1.0\times$  on real data—in sharp contrast to their severe degradation in simulation ( $-13.98$  dB for MST-L). This simulation-to-hardware gap indicates that the real hardware mask already contains uncorrected manufacturing imperfections, so the additional 0.5 px perturbation is small relative to pre-existing errors.

### 9.2 CACTI Real Data: Temporal Compressive Camera

The CACTI real dataset consists of 4 dynamic scenes (duomino, hand, pendulumBall, waterBallon) acquired with a coded aperture compressive temporal imaging system at  $512 \times 512$  spatial resolution with compression ratio 10 (10 video frames encoded per snapshot).

**Key findings.** GAP-TV shows an order-of-magnitude residual increase ( $10.4\times$  mean) under sub-pixel mask mismatch, consistent with the multiplicative error amplification in temporal compression (a single mask error propagates across all 10 compressed frames). PnP-FFDNet shows moderate robustness ( $2.0\times$ ), likely because the FFDNet denoiser provides an implicit regularization that partially absorbs the mismatch artefacts. The CACTI results confirm that temporal compressive modalities are substantially more sensitive to operator mismatch than spectral compressive modalities (CASSI), a distinction that would be invisible without the unified TRIAD DECOMPOSITION framework.

Table S7: Supplementary Table S5: CASSI real-data measurement residuals. Residual values  $r$  (Eq. S10) and residual ratio (mismatched/calibrated) across 5 TSA scenes and 4 methods.

| Scene       | Method        | $r_{\text{cal}}$ | $r_{\text{mis}}$ | Ratio      |
|-------------|---------------|------------------|------------------|------------|
| Scene 1     | GAP-TV        | 0.00148          | 0.00298          | 2.0        |
| Scene 1     | HDNet         | 0.18506          | 0.18506          | 1.0        |
| Scene 1     | MST-S         | 0.09953          | 0.10746          | 1.1        |
| Scene 1     | MST-L         | 0.10055          | 0.09544          | 0.9        |
| Scene 2     | GAP-TV        | 0.00210          | 0.00331          | 1.6        |
| Scene 2     | HDNet         | 0.14853          | 0.14853          | 1.0        |
| Scene 2     | MST-S         | 0.25357          | 0.23807          | 0.9        |
| Scene 2     | MST-L         | 0.25441          | 0.22811          | 0.9        |
| Scene 3     | GAP-TV        | 0.00204          | 0.00320          | 1.6        |
| Scene 3     | HDNet         | 0.18291          | 0.18291          | 1.0        |
| Scene 3     | MST-S         | 0.09399          | 0.09160          | 1.0        |
| Scene 3     | MST-L         | 0.10127          | 0.09085          | 0.9        |
| Scene 4     | GAP-TV        | 0.00148          | 0.00297          | 2.0        |
| Scene 4     | HDNet         | 0.14355          | 0.14355          | 1.0        |
| Scene 4     | MST-S         | 0.15586          | 0.14712          | 0.9        |
| Scene 4     | MST-L         | 0.17899          | 0.15893          | 0.9        |
| Scene 5     | GAP-TV        | 0.00233          | 0.00421          | 1.8        |
| Scene 5     | HDNet         | 0.11325          | 0.11325          | 1.0        |
| Scene 5     | MST-S         | 0.11995          | 0.12004          | 1.0        |
| Scene 5     | MST-L         | 0.13008          | 0.12208          | 0.9        |
| <b>Mean</b> | <b>GAP-TV</b> | <b>0.00189</b>   | <b>0.00333</b>   | <b>1.8</b> |
| <b>Mean</b> | <b>HDNet</b>  | <b>0.15466</b>   | <b>0.15466</b>   | <b>1.0</b> |
| <b>Mean</b> | <b>MST-S</b>  | <b>0.14458</b>   | <b>0.14086</b>   | <b>1.0</b> |
| <b>Mean</b> | <b>MST-L</b>  | <b>0.15306</b>   | <b>0.13908</b>   | <b>0.9</b> |

Table S8: Supplementary Table S6: CACTI real-data measurement residuals and total variation. Residual ratio (mismatched/calibrated) and total variation (TV) for 4 real scenes and 2 methods.

| Scene        | Method            | $r_{\text{cal}}$     | $r_{\text{mis}}$      | Ratio       |
|--------------|-------------------|----------------------|-----------------------|-------------|
| duomino      | GAP-TV            | $8.0 \times 10^{-6}$ | $8.5 \times 10^{-5}$  | 10.6        |
| duomino      | PnP-FFDNet        | 0.00198              | 0.00403               | 2.0         |
| hand         | GAP-TV            | $7.0 \times 10^{-6}$ | $7.7 \times 10^{-5}$  | 11.0        |
| hand         | PnP-FFDNet        | 0.00249              | 0.00705               | 2.8         |
| pendulumBall | GAP-TV            | $3.7 \times 10^{-5}$ | $3.46 \times 10^{-4}$ | 9.4         |
| pendulumBall | PnP-FFDNet        | 0.00919              | 0.01150               | 1.3         |
| waterBalloon | GAP-TV            | $1.4 \times 10^{-5}$ | $1.47 \times 10^{-4}$ | 10.5        |
| waterBalloon | PnP-FFDNet        | 0.00264              | 0.00493               | 1.9         |
| <b>Mean</b>  | <b>GAP-TV</b>     | —                    | —                     | <b>10.4</b> |
| <b>Mean</b>  | <b>PnP-FFDNet</b> | —                    | —                     | <b>2.0</b>  |

Table S9: Supplementary Table S7: Grid-search calibration results on simulated data with known ground truth. PSNR values in dB. Sc. II = mismatch (no calibration); Sc. IV = grid-search calibrated; Sc. III = oracle (true operator). Recovery = (Sc. IV - Sc. II) / (Sc. III - Sc. II). CASSI/CACTI use measurement-residual objective; SPC uses reconstruction-TV objective.

| Modality | Method     | Sc. II | Sc. IV | Sc. III | Time (s) | Recovery |
|----------|------------|--------|--------|---------|----------|----------|
| CASSI    | GAP-TV     | 21.52  | 22.96  | 23.21   | 1140     | 85%      |
| CACTI    | GAP-TV     | 17.60  | 26.99  | 26.99   | 60       | 100%     |
| SPC      | FISTA-TV   | 19.78  | 26.54  | 27.60   | 166      | 86%      |
| SPC      | PnP-DRUNet | 18.34  | 25.39  | 26.01   | 247      | 92%      |

### 9.3 Autonomous Calibration on Real Data

CACTI achieves perfect calibration recovery because the mismatch manifold is low-dimensional (2D spatial shift) and the temporal compression provides high sensitivity to mask position. CASSI achieves 85% recovery; the residual gap reflects the higher-dimensional mismatch space (mask geometry plus dispersion) and the estimation error in both  $dx$  and  $dy$  (0.10px each). For SPC, the measurement residual is uninformative for gain drift because the underdetermined system ( $m=272$  measurements for  $n=1089$  unknowns) always achieves near-zero self-consistent residual regardless of gain. Instead, reconstruction total variation provides a viable objective: the correctly calibrated gain produces clean measurements yielding smooth reconstructions with low TV, while incorrect gain leaves systematic artifacts that increase TV. The TV surface exhibits a clear bowl-shaped minimum near the true  $\alpha$ , recovering 86% (FISTA-TV) to 92% (PnP-DRUNet) of the oracle bound. This demonstrates that blind calibration generalises to radiometric mismatch provided the objective matches the mismatch type: measurement residual for geometric mismatch, reconstruction sparsity for radiometric mismatch.

## 10 Supplementary Table S8: SSIM Comparison Across Modalities

Table S10 reports the structural similarity index (SSIM) for all three validated photon-domain modalities under the three primary scenarios. SSIM complements PSNR by capturing perceptual quality, particularly structural distortions from operator mismatch.

Table S10: SSIM comparison across modalities and scenarios. Values are mean  $\pm$  s.d. across test scenes/images.

| Modality                                             | Method        | Sc. I             | Sc. II            | Sc. III           |
|------------------------------------------------------|---------------|-------------------|-------------------|-------------------|
| <i>CASSI (10 KAIST scenes, 28 spectral bands)</i>    |               |                   |                   |                   |
|                                                      | GAP-TV        | $0.723 \pm 0.088$ | $0.612 \pm 0.084$ | $0.688 \pm 0.083$ |
|                                                      | HDNet         | $0.970 \pm 0.009$ | $0.756 \pm 0.074$ | $0.756 \pm 0.074$ |
|                                                      | MST-L         | $0.973 \pm 0.009$ | $0.744 \pm 0.069$ | $0.881 \pm 0.035$ |
| <i>CACTI (6 benchmark videos, 8 temporal frames)</i> |               |                   |                   |                   |
|                                                      | GAP-TV        | $0.848 \pm 0.083$ | $0.305 \pm 0.070$ | $0.794 \pm 0.063$ |
|                                                      | PnP-FFDNet    | $0.890 \pm 0.060$ | $0.216 \pm 0.076$ | $0.820 \pm 0.054$ |
|                                                      | ELP-Unfolding | $0.965 \pm 0.013$ | $0.308 \pm 0.076$ | $0.927 \pm 0.017$ |
|                                                      | EfficientSCI  | $0.973 \pm 0.012$ | $0.303 \pm 0.079$ | $0.927 \pm 0.017$ |
| <i>SPC (11 standard images, 25% compression)</i>     |               |                   |                   |                   |
|                                                      | FISTA-TV      | $0.852 \pm 0.046$ | $0.586 \pm 0.055$ | $0.759 \pm 0.036$ |
|                                                      | PnP-DRUNet    | $0.899 \pm 0.026$ | $0.415 \pm 0.066$ | $0.666 \pm 0.055$ |
|                                                      | ISTA-Net      | $0.916 \pm 0.028$ | $0.584 \pm 0.077$ | $0.760 \pm 0.056$ |
|                                                      | HATNet        | $0.847 \pm 0.049$ | $0.648 \pm 0.091$ | $0.807 \pm 0.060$ |

**Key observations.** The SSIM degradation pattern mirrors the PSNR findings: under Scenario II (mismatch), all methods collapse to a narrow quality range regardless of their Scenario I performance. In CASSI, the SSIM collapse is particularly striking: MST-L drops from 0.973 to 0.744, a perceptual quality loss that exceeds the improvement from a decade of solver development. In CACTI, mismatch produces near-random SSIM values ( $\sim 0.3$ ), confirming that the reconstructed video frames bear little structural resemblance to the ground truth under mismatch conditions.

## 11 Supplementary Table S9: CASSI Spectral Angle Mapper (SAM)

Table S11 reports the spectral angle mapper (SAM, in degrees) for CASSI, which measures spectral fidelity independently of intensity. Lower values indicate better spectral reconstruction.

Table S11: CASSI spectral angle mapper (SAM, degrees). Lower is better. Values are mean  $\pm$  s.d. across 10 KAIST scenes.

| Method      | Sc. I            | Sc. II           | Sc. III          |
|-------------|------------------|------------------|------------------|
| GAP-TV      | $16.66 \pm 3.32$ | $24.27 \pm 2.92$ | $25.97 \pm 2.54$ |
| PnP-HSI-CNN | $16.10 \pm 3.25$ | $23.73 \pm 2.81$ | $18.66 \pm 2.75$ |
| HDNet       | $6.67 \pm 0.96$  | $17.03 \pm 2.76$ | $17.03 \pm 2.76$ |
| MST-L       | $7.44 \pm 1.24$  | $23.92 \pm 4.58$ | $11.74 \pm 1.32$ |

**Key observations.** Spectral fidelity degrades severely under mismatch: MST-L’s SAM increases from  $7.44^\circ$  (near-perfect spectral reconstruction) to  $23.92^\circ$  (substantial spectral mixing). Oracle correction partially recovers spectral quality ( $11.74^\circ$ ), but the residual SAM error confirms that multi-parameter mismatch produces irreversible spectral distortions not fully correctable by mask-only correction. The SAM metric is particularly valuable for remote sensing applications where spectral accuracy is more important than spatial quality.

## 12 Supplementary Tables S10–S11: Gate 1 and Gate 2 Validation

The main text argues that Gate 3 (operator mismatch) is the dominant bottleneck under standard operating conditions. To demonstrate that Gate 1 (information deficiency) and Gate 2 (carrier budget) are real and measurable, we sweep the compression level and noise level across all seven validated modalities while keeping the forward model perfectly calibrated. All experiments use classical solvers (no deep-learned components) to ensure the results reflect information-theoretic limits rather than solver-specific behaviour.

### 12.1 Gate 1: Information Deficiency (Extreme Compression)

Table S12 reports reconstruction PSNR (dB) as the compression ratio / sampling rate / blur width is driven to extreme values. The forward model is ideal in every case; any PSNR loss is attributable solely to information lost in the measurement process.

Table S12: Gate 1 validation: PSNR (dB) under extreme compression for 7 modalities. “Best solver” reports the better of the two classical methods tested. Bold entries highlight the regime where PSNR falls below 15 dB (severe information loss).

| Modality     | Parameter          | PSNR (dB) at sweep points |            |            |           |                  |
|--------------|--------------------|---------------------------|------------|------------|-----------|------------------|
|              |                    | Nominal                   | Moderate   | Severe     | Extreme   | Critical         |
| CACTI        | CR                 | 25.9 (8)                  | 24.1 (16)  | 22.7 (32)  | 20.6 (64) | —                |
| CASSI        | Transmittance      | 26.1 (50%)                | 27.7 (25%) | 27.7 (10%) | 27.6 (5%) | 27.4 (2%)        |
| SPC          | CS ratio           | 28.3 (25%)                | 24.0 (10%) | 21.2 (5%)  | 16.2 (2%) | <b>14.4</b> (1%) |
| MRI          | Sampling rate      | 28.8 (25%)                | 25.1 (10%) | 24.3 (5%)  | 23.9 (2%) | —                |
| CT           | Projection angles  | 22.1 (180)                | 22.0 (90)  | 21.2 (30)  | 20.7 (10) | 17.8 (5)         |
| Lensless     | Blur $\sigma$ (px) | 36.6 (1)                  | 25.5 (3)   | 23.1 (5)   | 20.5 (10) | 17.9 (20)        |
| Ptychography | Scan positions     | 15.0 (16)                 | 12.4 (9)   | 12.0 (4)   | 9.9 (1)   | —                |

#### Key observations.

- **SPC** and **Lensless** show the most dramatic Gate 1 collapse. SPC PSNR drops 13.9 dB from 25% to 1% CS ratio; lensless drops 18.7 dB as the PSF blur widens from  $\sigma=1$  to  $\sigma=20$  px. In both cases the loss is irrecoverable: no solver can reconstruct information that was never measured.
- **CASSI** is a notable exception: reducing mask transmittance from 50% to 2% produces no degradation (mean PSNR actually increases by  $\sim 1.5$  dB). This occurs because sparser coded apertures reduce spectral mixing in the multiplexed measurement, making the demixing problem easier for iterative solvers. The CASSI null space is effectively unchanged because each spectral band still contributes signal through the remaining open mask pixels.
- **CACTI** degrades monotonically with compression ratio ( $25.9 \rightarrow 20.6$  dB over  $8\times$  CR increase), consistent with the growing null space when fewer temporal measurements are available.
- **CT** shows SART outperforming FBP at sparse angles (21.2 dB vs. 15.8 dB at 10 angles), illustrating that iterative solvers can partially compensate for angular undersampling via their implicit prior, but ultimately both methods converge at  $\leq 5$  angles.

## 12.2 Gate 2: Carrier Budget (Noise Sweep)

Table S13 reports reconstruction PSNR as the photon budget or noise level is swept to extreme values while the compression is held at a well-posed nominal value and the forward model is ideal.

Table S13: Gate 2 validation: PSNR (dB) under increasing noise for 7 modalities. “Best solver” reports the better of the two classical methods tested. Noise parameter: photon count for Poisson-dominated modalities, Gaussian  $\sigma$  (fraction of signal range) for MRI and SPC. Bold entries indicate  $\text{PSNR} < 15 \text{ dB}$ .

| Modality     | Noise param.     | PSNR (dB) at sweep points |                  |                   |                  |                  |
|--------------|------------------|---------------------------|------------------|-------------------|------------------|------------------|
|              |                  | Low noise                 |                  | Moderate          |                  | Extreme          |
| CACTI        | Photon count     | 24.8 (10k)                | 23.3 (1k)        | 18.4 (100)        | <b>10.5</b> (10) | —                |
| CASSI        | Photon count     | 24.6 (10k)                | 22.3 (1k)        | 20.2 (100)        | 17.0 (10)        | —                |
| SPC          | $\sigma$ (frac.) | 27.9 (0)                  | 27.8 (.01)       | 25.3 (.05)        | 21.6 (.1)        | <b>13.5</b> (.3) |
| MRI          | $\sigma$ (frac.) | 28.8 (0)                  | 17.1 (.01)       | <b>12.8</b> (.05) | <b>12.1</b> (.1) | <b>11.0</b> (.3) |
| CT           | Photon count     | 22.1 (100k)               | 21.9 (10k)       | 21.0 (1k)         | 18.3 (100)       | —                |
| Lensless     | Photon count     | 40.9 (10k)                | 32.8 (1k)        | 23.1 (100)        | <b>13.6</b> (10) | —                |
| Ptychography | Photon count     | 13.5 (10k)                | <b>10.3</b> (1k) | <b>10.0</b> (100) | <b>9.7</b> (10)  | —                |

### Key observations.

- Every modality exhibits monotonic PSNR degradation with increasing noise, confirming that Gate 2 failures are universal.
- **MRI** is the most noise-sensitive modality: adding just  $\sigma=0.01$  Gaussian noise to k-space drops CS-wavelet PSNR from 28.8 dB to 17.1 dB, a 11.7 dB cliff. This is because MRI’s Fourier encoding concentrates signal energy in a few k-space samples, so additive noise corrupts the information-dense center region disproportionately.
- **Lensless** imaging shows the widest Gate 2 dynamic range: 40.9 dB at 10,000 photons to 13.6 dB at 10 photons, a 27.3 dB span. This reflects the well-conditioned forward operator (diffuser PSF), which faithfully propagates both signal and noise.
- **CT** is the most noise-robust modality (3.8 dB drop from  $10^5$  to  $10^2$  photons), because the many-angle sinogram provides substantial redundancy that averages out Poisson noise.
- The cliff-edge nature of Gate 2 is most dramatic in CACTI (14.3 dB drop from 10,000 to 10 photons) and SPC (14.4 dB drop from  $\sigma=0$  to  $\sigma=0.3$ ), where the compressed measurement offers no noise-averaging redundancy.

**Comparison with Gate 3.** Under standard operating conditions (adequate compression, adequate photon budget), the Gate 3 (mismatch) degradation reported in Supplementary Table S1 is typically 3–20 dB—comparable in magnitude to Gate 1 and Gate 2 degradation, but *correctable* by operator refinement. This is the key distinction: Gate 1 and Gate 2 losses are information-theoretic and irrecoverable, whereas Gate 3 losses are recoverable via calibration. The practical dominance of Gate 3 arises because modern instruments operate well above the Gate 1 and Gate 2 thresholds, leaving operator mismatch as the binding constraint.



## 13 Supplementary Note 7: Clinical CT Quality Assurance Validation

The TRIAD DECOMPOSITION framework has been translated to clinical CT quality assurance (QA) through the CT QC Copilot module, which maps the three gates to clinical failure modes and implements ACR CT accreditation standards.

### 13.1 Gate Mapping to Clinical Failure Modes

Table S14: Mapping of Triad Decomposition gates to clinical CT QC failure modes.

| Gate   | Research Interpretation                             | Clinical Interpretation                                       |
|--------|-----------------------------------------------------|---------------------------------------------------------------|
| Gate 1 | Information deficiency<br>(null space, compression) | Protocol design inadequacy<br>(insufficient projections, FOV) |
| Gate 2 | Carrier budget<br>(SNR, photon count)               | Dose budget<br>(noise floor vs. diagnostic need)              |
| Gate 3 | Operator mismatch<br>(mask shift, PSF error)        | Scanner calibration drift<br>(HU drift, CoR offset, gain)     |

### 13.2 ACR Metric Validation

Nine ACR-aligned QC metrics are computed automatically from CT phantom (Gammex 464) scans. Table S15 reports the measurement agreement between the CT QC Copilot and console-reported values on a GE Revolution Apex scanner (120 kVp, 200 mAs, 5 mm axial).

Table S15: ACR metric agreement: CT QC Copilot vs. console gold standard.

| Metric                     | Copilot | Console | Deviation | % Tolerance |
|----------------------------|---------|---------|-----------|-------------|
| CT# Water (HU)             | 0.3     | 0.2     | 0.1       | 2%          |
| CT# Bone (HU)              | 952     | 953     | 1.0       | 5%          |
| CT# Air (HU)               | −997    | −998    | 1.0       | 5%          |
| CT# Acrylic (HU)           | 120     | 121     | 1.0       | 7%          |
| CT# Polyethylene (HU)      | −96     | −96     | 0.0       | 0%          |
| Geometric accuracy (mm)    | 0.15    | 0.10    | 0.05      | 2.5%        |
| Slice thickness (mm)       | 5.08    | 5.02    | 0.06      | 4%          |
| Uniformity (HU)            | 1.2     | 1.1     | 0.1       | 2%          |
| Noise std. dev. (HU)       | 4.6     | 4.5     | 0.1       | 10%         |
| Spatial resolution (lp/cm) | 6.2     | 6.0     | 0.2       | 4%          |

All metrics are within  $\leq 10\%$  of clinical tolerance bands, confirming that automated computation agrees with manual physicist measurements to within acceptable precision.

### 13.3 Drift Detection Performance

Statistical process control (SPC) using five Western Electric rules was evaluated on a 30-scanner simulated fleet over 12 monthly measurement cycles. Four scanners were programmed with known

drift trajectories (noise drift: 2 scanners; uniformity drift: 1 scanner; CT number drift: 1 scanner).

- **Sensitivity:** 100% (4/4 drifting scanners detected; 95% CI: 39.8%–100%)
- **Specificity:** 100% (0/26 false positives on stable scanners; 95% CI: 86.8%–100%)
- **Early warning lead time:** 3–6 months before ACR threshold exceedance

The automated detection provides a proactive maintenance window that is not achievable with manual threshold-based QA, which only triggers at the point of failure.

### 13.4 Workflow Efficiency

Table S16: Per-scanner QC workflow time comparison.

| Stage                     | Manual (min)   | Copilot (min)    |
|---------------------------|----------------|------------------|
| DICOM ingestion           | 5.0            | 0.01             |
| Metric computation        | 25.0           | 0.01             |
| Threshold evaluation      | 5.0            | <0.01            |
| Trend analysis            | 10.0           | <0.01            |
| Report generation         | 15.0           | 0.02             |
| Physicist review/sign-off | 7.0            | 4.0              |
| <b>Total</b>              | <b>67 ± 12</b> | <b>4.2 ± 0.8</b> |

The 94% reduction in per-scanner QC time (from  $67 \pm 12$  to  $4.2 \pm 0.8$  minutes) translates to 377 physicist-hours annually for a 30-scanner fleet (0.18 FTE reallocation). Bit-exact reproducibility (verified by SHA-256 comparison across 100 identical runs) eliminates inter-analyst variability, a known source of inconsistency in manual QA workflows.

## 14 Supplementary Note 8: Controlled Hardware Validation

The software-perturbation experiments in the main text apply calibrated mask shifts to existing real measurements. A full hardware-in-the-loop validation requires physically displacing the coded aperture mask and re-acquiring data under controlled conditions. This section describes the experimental protocols and placeholder results for prospective hardware experiments.

### 14.1 CASSI Physical Mask Displacement Protocol

1. **Baseline acquisition.** Acquire hyperspectral data at the factory-calibrated mask position using a calibrated broadband illumination source with  $\leq 0.5\%$  temporal intensity variation.
2. **Controlled displacement.** Physically translate the coded aperture mask by known displacements  $\Delta x \in \{0.25, 0.50, 1.00\}$  px equivalent, verified by a micrometer translation stage with  $\leq 0.01$  px positioning accuracy. Re-acquire under identical illumination for each displacement.
3. **Reconstruction.** Reconstruct all datasets using the factory mask specification (Scenario II equivalent). Compute the PSNR degradation and measurement residual ratio relative to the baseline.
4. **PWM calibration.** Apply the autonomous beam-search calibration pipeline (Algorithm 1) and measure recovery.
5. **Expected outcome.** The hardware results should exhibit a monotonic increase in residual ratio with displacement magnitude, with the slope modulated by the pre-existing manufacturing error baseline quantified in the main text.

### 14.2 Multi-Unit Variation Study Protocol

1. Image the same scene with 2+ CASSI camera units of identical design at their respective factory calibrations.
2. Reconstruct each dataset with each camera’s nominal mask.
3. Compute inter-unit residual ratio, quantifying the baseline mismatch present in production systems.
4. Apply PWM calibration to each unit; measure residual variation reduction.

**Status.** Hardware experiments are planned in collaboration with external partners. Results will be incorporated in the revised manuscript.

## 15 Supplementary Note 9: Mismatch Parameter Derivation

This note provides a self-contained derivation of the mismatch parameter values and ranges used in the CASSI experiments, independent of the concurrent InverseNet work<sup>1</sup>.

### 15.1 CASSI 5-Parameter Mismatch Model

The CASSI forward model depends on five calibration parameters:

$$\boldsymbol{\theta} = (dx, dy, \theta, a_1, \alpha), \quad (\text{S11})$$

where  $dx, dy$  are mask spatial shifts (pixels),  $\theta$  is mask rotation (degrees),  $a_1$  is the dispersion slope (pixels per spectral band), and  $\alpha$  is the dispersion axis angle (degrees).

**Mask shift ( $dx, dy$ ).** Typical CASSI assembly involves mounting the coded aperture mask on a precision translation stage. Achievable positioning accuracy with standard micrometer stages is  $\pm 0.5$  px at typical detector pitches ( $\sim 5\text{--}10\ \mu\text{m}$ ). We use  $dx = 0.5$  px and  $dy = 0.3$  px as representative assembly errors within this range.

**Mask rotation ( $\theta$ ).** Rotational alignment error during mask mounting is typically  $\leq 0.2^\circ$  with standard optomechanical mounts. We use  $\theta = 0.1^\circ$  as representative.

**Dispersion slope ( $a_1$ ).** The nominal dispersion is 2.0 pixels per band for the TSA-Net configuration. Manufacturing variation in prism angle and detector alignment contribute  $\sim 1\%$  uncertainty. We use  $a_1 = 2.02$  (1% drift from nominal).

**Dispersion axis angle ( $\alpha$ ).** The dispersion axis should be exactly aligned with the detector columns. Typical alignment error is  $\leq 0.2^\circ$ . We use  $\alpha = 0.15^\circ$ .

**Validation.** These parameter ranges were validated against reported assembly tolerances in CASSI literature<sup>4</sup> and confirmed by the hardware residual analysis in the main text: the real-data measurement residual ratio ( $1.8\times$ ) is consistent with pre-existing mismatch of comparable magnitude to the perturbation applied.

### 15.2 Prospective Clinical Validation Protocol

The CT QC Copilot results in Supplementary Note 7 use simulated scanner data. Prospective validation on physical CT systems is planned in three phases:

**Phase 1 (single scanner, 3 months).** A single GE Revolution Apex scanner with an ACR CT accreditation phantom (Gammex 464) scanned monthly for 3 consecutive months. Compare Copilot-computed metrics against manual physicist measurements for all 9 ACR-aligned metrics.

**Phase 2 (5-scanner mini-fleet, 6 months).** Expand to 5 scanners from  $\geq 2$  manufacturers (GE + Siemens or Philips). Apply Western Electric drift detection rules over 6 monthly measurement cycles. Evaluate sensitivity and specificity for detecting known calibration drifts.

**Phase 3 (30-scanner fleet).** Full fleet deployment, if available, to validate the 377 physicist-hours annual savings projection from the simulation study.

**Status.** Prospective validation is planned in collaboration with clinical physics partners. Even Phase 1 data (single scanner, 3 months) would transform the CT QC section from a simulation exercise to a clinical demonstration.

## 16 Supplementary Note S12: Calibration Method Comparison

Existing calibration methods are modality-specific: ESPIRiT<sup>5</sup> auto-calibrates coil sensitivities for parallel MRI; entropy-based center-of-rotation autofocus corrects geometric errors in CT; blind position correction in ePIE<sup>6</sup> self-calibrates probe positions in ptychography. Each achieves high recovery within its target modality but requires domain-specific algorithm design and cannot transfer across modalities.

**Quantitative comparison: ESPIRiT vs. PWM on multi-coil MRI.** To ground this comparison empirically, we evaluate four conditions on a synthetic brain phantom (8 coils,  $256 \times 256$ ,  $4\times$  acceleration, CG-SENSE, 5% multiplicative coil sensitivity mismatch):

Table S17: ESPIRiT vs. PWM comparison on multi-coil MRI (8 coils,  $4\times$  acceleration, 5% mismatch).

| Condition                               | PSNR (dB) | SSIM   |
|-----------------------------------------|-----------|--------|
| Scenario I (true maps)                  | 34.98     | 0.8330 |
| Scenario II (5% mismatch)               | 31.30     | 0.7388 |
| ESPIRiT (auto-calibrated, 24 ACS lines) | 22.01     | 0.5074 |
| PWM (beam-search corrected)             | 32.05     | 0.7264 |

Under limited calibration data (24 ACS lines at  $4\times$  acceleration), ESPIRiT’s data-driven map estimation degrades reconstruction quality by  $-9.29$  dB relative to the mismatched baseline, because the auto-calibration signal is insufficient for reliable sensitivity estimation (map NRMSE = 0.82 vs. 0.09 for the 5% mismatch). By contrast, PWM’s model-based correction recovers  $+0.75$  dB (recovery ratio  $\rho = 20.3\%$ ), because it operates on the forward-model parameter space rather than estimating maps from data.

**Complementarity.** ESPIRiT and PWM address different aspects of the calibration problem. ESPIRiT exploits  $k$ -space structure to estimate sensitivity maps purely from data, performing well when sufficient calibration data is available (e.g., fully sampled ACS regions with  $\geq 32$  lines). PWM corrects arbitrary forward-model parameters using the OperatorGraph structure, providing robust improvement even in data-limited regimes. The two approaches are complementary: ESPIRiT can be combined with PWM by using ESPIRiT-estimated maps as the initial  $H_{\text{nom}}$  for PWM correction.

PWM’s modality-agnostic correction achieves comparable recovery on validated modalities (CASSI: 22–46%, CACTI: 93%, SPC: 86–92%) without requiring domain-specific tuning. The advantage is most pronounced for CASSI, where no automated calibration standard exists, and for cross-modality deployment, where a single pipeline serves all modalities without per-modality engineering.

## 17 Supplementary Note S13: MRI Under Clinically Realistic Conditions

The  $+48.25$  dB correction gain reported for MRI in the main text reflects a single-coil scenario with 5% multiplicative sensitivity error—a pathological case chosen to stress-test the framework. Under clinically realistic multi-coil conditions (8 coils,  $256 \times 256$ ,  $4\times$  acceleration, CG-SENSE reconstruction) with spatially smooth Biot-Savart sensitivity errors simulating patient repositioning, the correction gains are substantially more modest but clinically meaningful:

Table S18: Multi-coil MRI correction across mismatch levels (8 coils, 4× accel., CG-SENSE).

| Mismatch | Sc. I (dB) | Sc. II (dB) | Sc. III (dB) | Sc. IV (dB) | Gain (dB) |
|----------|------------|-------------|--------------|-------------|-----------|
| 1%       | 52.11      | 50.91       | 51.48        | 52.11       | +0.58     |
| 2%       | 52.11      | 49.40       | 48.72        | 52.11       | −0.68     |
| 3%       | 52.11      | 47.96       | 49.72        | 52.11       | +1.75     |
| 5%       | 52.11      | 40.91       | 48.05        | 52.11       | +7.14     |
| 10%      | 52.11      | 38.21       | 41.74        | 52.11       | +3.52     |
| 15%      | 52.11      | 32.97       | 36.25        | 52.11       | +3.28     |

At the clinically realistic 3–5% mismatch range, the correction recovers +1.75 to +7.14 dB (mean +4.45 dB). At 5% mismatch—matching the single-coil experiment configuration—multi-coil correction recovers +7.14 dB versus +48.25 dB in the single-coil case. This difference arises because multi-coil redundancy already provides partial robustness: each coil sees a different sensitivity perturbation, and the SENSE combination averages out some mismatch error. Even so, +7.14 dB at 5% (and +3.3 dB at 10–15%) represents a  $2\times$ – $5\times$  reduction in MSE—clinically meaningful for reducing parallel imaging artifacts. The 2% result (−0.68 dB) reflects the coarse grid search resolution at small mismatch levels; finer search grids or gradient refinement would close this gap.

## 18 Supplementary Note 10: Finite Primitive Basis — Expanded Proof

This note provides the expanded proof of Theorem 1 (Finite Primitive Basis) stated in the main text. Full formal definitions, primitive semantics, and typed DAG denotation are in the companion paper<sup>7</sup>.

### 18.1 Formal Definitions

**Definition 6** (Tier-2 Operator Class  $\mathcal{C}_{\text{Tier2}}$ ). The class  $\mathcal{C}_{\text{Tier2}}$  consists of all imaging forward models  $H : \mathcal{X} \rightarrow \mathcal{Y}$  that admit a factorization  $H = H_K \circ \dots \circ H_1$  (or a DAG generalization) where: (i)  $K \leq N_{\max} = 20$ ; (ii) each  $H_k$  is a linear operator; (iii)  $\|H_k\| \leq B$  for a universal bound  $B > 0$ ; (iv) each  $H_k$  is at most shift-variant (Tier-2 on the Physics Fidelity Ladder).

**Definition 7** ( $\varepsilon$ -Approximate Representation). Let  $\mathcal{B} = \{P, M, \Pi, F, C, \Sigma, D, S, W, R\}$  be the canonical primitive library. A typed DAG  $G = (V, E, \tau)$  with  $V \subseteq \mathcal{B}$  is an  $\varepsilon$ -approximate representation of  $H \in \mathcal{C}_{\text{Tier2}}$  if:

1.  $\sup_{\|\mathbf{x}\| \leq 1} \frac{\|H(\mathbf{x}) - H_G(\mathbf{x})\|}{\|H(\mathbf{x})\| + \delta} \leq \varepsilon$ , where  $\delta > 0$  is a regularization constant;
2.  $|V| \leq N_{\max} = 20$  and  $\text{depth}(G) \leq D_{\max} = 10$ .

**Empirical evaluation.** For each modality,  $e_{\text{Tier2}}$  is evaluated as the mean relative error over  $\mathcal{X}_{\text{test}}$  consisting of 10 standard benchmark scenes and 10 random Gaussian objects. The threshold  $\varepsilon = 0.01$  is chosen so the Tier-2 approximation error is below the noise floor at standard operating SNR.

## 18.2 Proof of Theorem 1

**Proof strategy.** The proof is constructive: we show that every factor  $H_k$  of a Tier-2 forward model can be realized by one or more primitives from  $\mathcal{B}$ , organized by the four physics-stage families. The per-factor approximation errors are then composed via sub-multiplicativity to yield the global bound. The key insight is that Tier-2 fidelity (linear, shift-variant) restricts each factor to a regime where well-characterized physical approximations (angular spectrum, Born series, piecewise shift-invariant partition) provide controlled error bounds.

**Theorem 8** (Finite Primitive Basis — restated). *For every  $H \in \mathcal{C}_{\text{Tier2}}$ , there exists a typed DAG  $G$  with  $V \subseteq \mathcal{B}$  that is an  $\varepsilon$ -approximate representation of  $H$ .*

*Proof.* The proof proceeds in four phases, one per physics-stage family.

**Phase 1: Propagation factors.** Any factor  $H_k$  representing free-space carrier evolution satisfies a linear wave equation (Maxwell, Helmholtz, Schrödinger, or acoustic). At Tier-2 fidelity, the solution is the angular spectrum propagator  $P(d, \lambda)$  or, in the shift-invariant limit,  $C(\mathbf{h})$ . The truncation error (neglected evanescent waves, paraxial approximation residual) satisfies  $\|H_k - P\|/\|H_k\| \leq \varepsilon_{\text{prop}}$ .

**Phase 2: Interaction factors.** *Elastic:* The carrier’s amplitude/phase changes without direction or energy change  $\rightarrow M(\mathbf{m})$  (exact). *Scattering:* Direction and/or energy change  $\rightarrow R(\sigma, \Delta\varepsilon)$  or finite composition of  $R, M, P$  for multiple-scattering media within low-order Born approximation. Error bounded by Tier-2 truncation:  $\varepsilon_{\text{scat}}$ .

**Phase 3: Encoding–projection factors.** Line-integral projection  $\rightarrow \Pi(\theta)$  (exact Radon transform). Fourier encoding  $\rightarrow F(\mathbf{k})$  (exact). Both are linear; representation is exact.

**Phase 4: Detection–readout factors.** The detector chain is realized by a finite composition of  $\Sigma$  (integration),  $S(\Omega)$  (sampling),  $W(\alpha, a)$  (dispersion),  $C(\mathbf{h}_{\text{det}})$  (detector PSF), and  $D(g, \eta)$  (quantum measurement from 5 canonical families). Per-operation errors sum to  $\varepsilon_{\text{det}}$ .

**Error bound.** Concatenating the per-factor DAGs and applying sub-multiplicativity:

$$\|H - H_G\| \leq \sum_{k=1}^K \varepsilon_k \prod_{j \neq k} \|H_j\| \leq K \cdot \max_k(\varepsilon_k) \cdot B^{K-1} \leq \varepsilon, \quad (\text{S12})$$

for  $\varepsilon_k \leq \varepsilon/(K \cdot B^{K-1})$ . For shift-variant operators (Tier-2), the per-factor bound  $\varepsilon_k$  is achieved by partitioning the spatial domain into  $L$  isoplanatic patches on each of which the operator is shift-invariant to within  $\varepsilon_k$ ; the Convolution primitive  $C$  applied per patch with interpolation across patch boundaries yields the required approximation. The patch count  $L$  is finite because the operator norm is bounded ( $\|H_k\| \leq B$ ) and spatial variation is Lipschitz-continuous at Tier-2 fidelity. Complexity:  $|V| \leq 5K \leq 5N_{\text{max}}$ ; empirically  $|V| \leq 6$  for all validated modalities.  $\square$

**Remark (tightness).** The worst-case bound  $\varepsilon = K \cdot \max_k(\varepsilon_k) \cdot B^{K-1}$  is conservative. Empirically, the measured  $e_{\text{Tier2}}$  values (Table in the main text) are 10–100 $\times$  below the worst-case bound, because: (i) most stages are exact ( $\varepsilon_k = 0$  for  $M, \Pi, F, \Sigma, S$ ); (ii) the sub-multiplicativity inequality is rarely tight when factors are near-unitary; and (iii) the isoplanatic patch partition for shift-variant operators typically requires  $L \leq 4$  patches for the validated modalities.



### 18.3 Extension Protocol: Worked Example

**Compton scatter imaging.** The forward model involves direction change (angle  $\theta$ , Klein–Nishina cross section) and energy shift ( $E_0 \rightarrow E_s$ ). Testing all 9-primitive DAGs: the best achieves  $e_{\text{Tier2}} = 0.34 \gg \varepsilon$ . None of the original primitives can represent direction change with energy shift. Introducing Scatter  $R(\sigma, \Delta\varepsilon)$ : the DAG  $M(n_e) \rightarrow R_{\text{KN}} \rightarrow D(E)$  achieves  $e_{\text{Tier2}} < 0.01$ . Scatter is needed by 5+ modalities (Compton, Raman, fluorescence, DOT, Brillouin), satisfying the multi-modality criterion. Re-running the closure test with  $\mathcal{B}' = \mathcal{B} \cup \{R\}$ : all previously decomposed modalities remain valid.

## References

- [1] Chengshuai Yang. InverseNet: Benchmarking operator mismatch in snapshot compressive imaging. In *Under review at ECCV 2026*, 2026.
- [2] Stephen Boyd, Neal Parikh, Eric Chu, Borja Peleato, and Jonathan Eckstein. Distributed optimization and statistical learning via the alternating direction method of multipliers. *Foundations and Trends in Machine Learning*, 3(1):1–122, 2011. doi: 10.1561/22000000016.
- [3] Amir Beck and Marc Teboulle. A fast iterative shrinkage-thresholding algorithm for linear inverse problems. *SIAM Journal on Imaging Sciences*, 2(1):183–202, 2009. doi: 10.1137/080716542.
- [4] Ashwin A. Wagadarikar, Renu John, Rebecca Willett, and David J. Brady. Single disperser design for coded aperture snapshot spectral imaging. *Applied Optics*, 47(10):B44–B51, 2008. doi: 10.1364/AO.47.000B44.
- [5] Martin Uecker, Peng Lai, Mark J. Murphy, Patrick Virtue, Michael Elad, John M. Pauly, Shreyas S. Vasanawala, and Michael Lustig. ESPIRiT — an eigenvalue approach to autocalibrating parallel MRI: where SENSE meets GRAPPA. *Magnetic Resonance in Medicine*, 71(3):990–1001, 2014. doi: 10.1002/mrm.24751.
- [6] Andrew M. Maiden and John M. Rodenburg. An improved ptychographical phase retrieval algorithm for diffractive imaging. *Ultramicroscopy*, 109(10):1256–1262, 2009. doi: 10.1016/j.ultramicro.2009.05.012.
- [7] Chengshuai Yang. The finite primitive basis theorem for computational imaging: Formal foundations of the OperatorGraph representation. *arXiv preprint*, 2026. In preparation. Full formal definitions, primitive semantics, typed DAG denotation, and constructive proof.