

Designing Any Imaging System from Natural Language: Agent-Constrained Composition over a Finite Primitive Basis

Chengshuai Yang

NextGen PlatformAI C Corp, USA

Correspondence: integrityyang@gmail.com

March 13, 2026

Summary. Designing a computational imaging system—selecting the forward model, validating physical feasibility, and predicting reconstruction quality—today demands weeks of specialist effort per modality, concentrating the field within a small number of expert groups [1]. We previously established that 11 canonical primitives suffice to represent any imaging forward model across 170 modalities [2], but this Finite Primitive Basis (FPB) provides no design automation, no formal validation, and no link from natural-language intent to reconstruction error. Here we introduce `spec.md`—a structured 7-field specification format analogous to FASTA for sequences—together with three agents (Plan, Judge, Performance) that translate a natural-language description into a validated, executable forward model with formal error guarantees. We prove that the design-to-real reconstruction error decomposes into five independently bounded, independently measurable terms (Theorem 2), each linked to a specific corrective action, and that all five terms are explicitly computable from the 6-gate compiler output. On 6 real-data modalities spanning all 5 carrier families—clinical CT [3], multi-coil MRI, coded-aperture spectral imaging (CASSI), snapshot compressive video (CACTI), in-vivo ultrasound (PICMUS), and 4D-STEM electron ptychography—the framework achieves $98.1 \pm 0.6\%$ of expert-tuned quality with up to $420\times$ reduction in setup time; across 36 validated modalities, the Judge rejects ill-posed designs with 96% precision and the Performance Agent predicts PSNR within ± 1.8 dB ($R^2 = 0.91$). A tier-lifting protocol demonstrated on diffuse-optical and coherent-imaging cross-tier combinations enables hybrid reconstruction across the four-level fidelity hierarchy, reducing unmodeled-physics error to machine precision on selected cases.

Introduction

Every computational imaging system obeys the same mathematical structure: a forward model A maps an unknown object x to measurements $y = A(x) + \text{noise}$, and reconstruction inverts this mapping [7]. Designing A —selecting physical operators (Radon projection for CT, k-space encoding for MRI, coded-aperture modulation for spectral imaging), setting dozens of parameters, and validating physical consistency—is entirely manual, taking weeks per modality. This expertise bottleneck concentrates new system development within a small number of specialist groups and excludes the broader scientific community from prototyping imaging systems for their own domains.

The mathematical foundation for automation now exists. We previously showed [2] that 11 canonical primitives—Propagate (P), Modulate (M), Project (Π), Encode (F), Convolve (C), Accumulate (Σ), Detect (D), Sample (S), Disperse (W), Scatter (R), and Transform (Λ)—suffice to express any imaging forward model across 170 modalities, 22 application domains, and 5 carrier families. Each primitive is implemented across a four-level *fidelity hierarchy*: Tier-0 (geometric

optics, ray transforms), Tier-1 (Fourier/convolution, paraxial models), Tier-2 (shift-variant linear operators: Radon, k-space, coded aperture), and Tier-3 (full-physics: FDTD, Monte Carlo, learned surrogates). This *Finite Primitive Basis* (FPB) is the representation language—but it does not tell a user which primitives to compose, what parameters to set, whether the result is physically valid, or how close the designed system will be to reality.

We make the scope of ‘designable’ precise:

Definition 1 (Designable Imaging System). *An imaging system A is designable within this framework if:*

1. A admits a finite DAG decomposition over the 11 FPB primitives with representation error $\varepsilon_{FPB} < 0.01$;
2. The carrier belongs to one of 5 families: photon, X-ray, electron, acoustic, or spin/particle;
3. The system operates at one of the four fidelity tiers (Tier-0 through Tier-3);
4. All parameters are finitely specifiable with known physical bounds.

This definition encompasses 170 modalities across 22 domains (Figure 1, Extended Data Table 6).

We build the *design layer* on top of the FPB. Our contributions are:

1. **spec.md**: a structured 7-field specification format that decouples the imaging problem (*what*) from the numerical method (*how*), serving as a shareable, peer-reviewable scientific artifact (Section).
2. **Three-agent pipeline**: Plan ($\mathcal{A}_P$), Judge ($\mathcal{A}_J$), and Performance ($\mathcal{A}_\Pi$) agents that generate, formally verify, and evaluate specifications over the FPB (Section).
3. **Design-to-real error theorem**: the total reconstruction error decomposes into five independently bounded, independently measurable terms, each linked to a compiler gate and a corrective action (Theorem 2).

Scope and honesty. Of the 170 modalities, 36 have fully validated canonical decompositions with quantitative reconstruction; 6 are validated on real measured data spanning all 5 carrier families (CT, MRI, CASSI, CACTI, ultrasound, e-ptychography). The remaining ~ 134 compile through all 6 gates but lack per-modality reconstruction benchmarks—compilation alone does not guarantee reconstruction quality. Most modalities are validated at Tier-2; tier-lifting (Section) is demonstrated on two cross-tier combinations but not yet validated across all Tier-3 modalities. The agents depend on current LLM capabilities and may misspecify novel or under-documented modalities.




Figure 1. Validation pyramid. Of 170 designable modalities, all compile; 36 have quantitative reconstruction benchmarks; 6 are validated on real measured data spanning all 5 carrier families; 3 have ground-truth PSNR comparisons (CT, MRI, CASSI). Each tier is a strict subset.

Results

The spec.md Specification Format

The framework’s central design choice is that `spec.md`—not code—is the canonical scientific artifact. Just as FASTA standardized biological sequences, `spec.md` standardizes imaging system representation. A specification separates the imaging problem (*what*) from the numerical method (*how*), enabling sharing, version control, and peer review independently of implementation. A journal could require a `spec.md` alongside a methods section, analogous to how genomics journals require accession numbers. Compared to alternatives—prose descriptions (ambiguous, non-executable), code (implementation-coupled, non-portable), block diagrams (unstructured, unversioned)—`spec.md` is simultaneously human-readable, machine-parseable, and formally verifiable through the 6-gate compiler.

A `spec.md` has seven mandatory fields (Figure 2). The Plan Agent generates it from natural language; the Judge verifies it; the Performance Agent predicts reconstruction quality.

Seven-field schema: `modality | carrier | geometry | object | forward_model | noise | target`

(a) Clinical CT

```
modality: computed_tomography  carrier: xray
geometry: parallel_beam, n_angles=128
object: 128×128 image, non-negative
forward_model: Radon(I) → Detect(D, intensity)
noise: Poisson, I_0=1e4  target: PSNR ≥ 30 dB
```

(b) CASSI (spectral)

```
modality: cassi  carrier: photon
geometry: coded_aperture + disperser, λ=[400,700] nm
object: 256×256×28 spectral cube
forward_model: Modulate(M) → Disperse(W) → Accumulate(Σ) → Detect(D)
noise: Gaussian, σ=0.01  target: PSNR ≥ 28 dB
```

(c) Accelerated MRI

```
modality: mri  carrier: spin
geometry: cartesian_kspace, acceleration=4x
object: 256×256 complex image
forward_model: Modulate(M, coil) → Encode(F, kspace) → Sample(S) → Detect(D)
noise: Gaussian, SNR=30 dB  target: SSIM ≥ 0.9
```

Figure 2. The spec.md specification format. Seven mandatory fields fully determine an imaging system design. Three examples span X-ray (CT), optical (CASSI), and spin (MRI) carriers. Each `forward_model` field is a primitive chain over the 11 FPB operators.

The Three-Agent Pipeline

Plan Agent ($\mathcal{A}_P$). Accepts a natural-language description and generates a `spec.md`, selecting FPB primitives, carrier type, geometry, noise model, and parameters from the 170-modality registry and 36 validated canonical decompositions.

Judge Agent ($\mathcal{A}_J$). Validates the specification via the 6-gate compiler: (1) DAG acyclicity, (2) canonical chain fidelity, (3) complexity bounds ($N \leq 20$, $D \leq 10$), (4) nonlinear family constraints (VC dimension ≤ 3), (5) adjoint consistency, and (6) representation error. The Judge additionally checks carrier-geometry compatibility, noise model consistency, and reconstruction




Figure 3. Specification-driven design pipeline. A natural-language prompt is the sole user input. $\mathcal{A}_P$ generates a `spec.md`; $\mathcal{A}_J$ validates it via a 6-gate compiler (Gates 1–6 map to Theorem 2 error terms: $\text{DAG} \rightarrow \varepsilon_{\text{spec}}$, $\text{Chain} \rightarrow \varepsilon_{\text{trans}}$, $\text{Bounds} \rightarrow \text{complexity}$, $\text{NL} \rightarrow \varepsilon_{\text{param}}$, $\text{Adj.} \rightarrow \text{consistency}$, $\text{Err.} \rightarrow \varepsilon_{\text{FPB}}$); rejected designs trigger redesign (dashed blue, up to 3 rounds). $\mathcal{A}_\Pi$ predicts reconstruction quality and outputs $\hat{A}$ with an explicit ε -bound report. See Figure 4 for per-modality error breakdown.

feasibility (condition number, null-space analysis). Compiler failures trigger rejection with a formal diagnosis; the Plan Agent regenerates (up to 3 rounds).

Performance Agent ($\mathcal{A}_\Pi$). Predicts reconstruction quality (PSNR, SSIM), computational cost, and measurement SNR from analytical estimates (condition number, noise amplification factor) combined with calibrated lookup from the performance database.

Formal verification. The Judge does not *predict* feasibility; it *executes deterministic verification*. If the Judge certifies a design, every term in the five-part error decomposition (Theorem 2) is explicitly bounded—a contract that learned design tools cannot provide.

Design-to-Real Error Bound

The following theorem is the paper’s central result. It converts the implicit, unmeasured gap between a designed model and reality into an explicit five-term error budget.

Theorem 2 (Design-to-Real Error Decomposition). *Let A_{true} be the true imaging forward*

model and A_{agent} the agent-designed model for a designable system (Definition 1). Let $\hat{x}$ be the regularized reconstruction from A_{agent} and x^* the true object. Under Assumptions A1–A4 (Methods), the reconstruction error satisfies:

$$\|\hat{x} - x^*\| \leq \underbrace{\varepsilon_{\text{FPB}}}_{\text{basis}} + \underbrace{\varepsilon_{\text{spec}}}_{\text{specification}} + \underbrace{\varepsilon_{\text{trans}}}_{\text{translation}} + \underbrace{\varepsilon_{\text{param}}}_{\text{parameter}} + \underbrace{\varepsilon_{\text{unmod}}}_{\text{unmodeled}} + \varepsilon_{\text{recon}} \quad (1)$$

where each design-error term is independently bounded and independently measurable:

- $\varepsilon_{\text{FPB}} < 0.01$: FPB representation error. Measured by: Gate 6. Guaranteed by: the basis theorem [2].
- $\varepsilon_{\text{spec}} = 0$ when all 6 compiler gates pass. Bounded by: $C_1 \cdot \delta_{\text{spec}}$ otherwise.
- $\varepsilon_{\text{trans}} = 0$ when the canonical chain matches the 36-modality registry (31/36 = 86%). Bounded by: $C_2 \cdot \delta_{\text{chain}}$ otherwise.
- $\varepsilon_{\text{param}} \leq \frac{2B}{\alpha + \lambda} L_A \|\theta_{\text{agent}} - \theta_{\text{true}}\|$, where L_A is the operator Lipschitz constant (computed per primitive by Gate 4), α is the coercivity constant, λ the regularization weight, and B a signal bound. Reducible by: calibration.
- $\varepsilon_{\text{unmod}} = 0$ at Tier-3; otherwise reducible by tier-lifting (Section).
- $\varepsilon_{\text{recon}}$: irreducible reconstruction error (noise floor + regularization bias), independent of the design.

Compiler–theorem link. When all 6 gates pass: $\varepsilon_{\text{FPB}} < 0.01$ (Gate 6), $\varepsilon_{\text{spec}} = 0$ (Gates 1–5 jointly), and $\varepsilon_{\text{param}}$ is bounded by the Lipschitz constant from Gate 4. If additionally the canonical chain matches the registry: $\varepsilon_{\text{trans}} = 0$. At Tier-3: $\varepsilon_{\text{unmod}} = 0$. The total design error then reduces to $L_A \|\Delta\theta\| \cdot 2B/(\alpha + \lambda) + 0.01$, dominated by parameter mismatch alone—which is reducible by calibration.

Diagnostic protocol. Because each term is independently measurable, a practitioner can determine which term dominates a given reconstruction residual and take the corresponding corrective action: calibration for $\varepsilon_{\text{param}}$, specification refinement for $\varepsilon_{\text{spec}}$, tier-lifting for $\varepsilon_{\text{unmod}}$. In practice, ε_{FPB} and $\varepsilon_{\text{spec}}$ are negligible when the compiler certifies the design. The full proof with explicit assumptions appears in Methods.

Cross-Modality Validation

We validated the framework on all 170 modalities. For each, the Plan Agent generated a `spec.md` from a one-sentence description, the Judge validated it, and the Performance Agent predicted reconstruction quality. Table 1 reports 14 representative modalities spanning all 5 carrier families (6 with real measured data), grouped by carrier; full results for all 36 validated modalities appear in Methods.

The recovery ratio ρ_{recov} is computed from a 4-scenario protocol validated on 6 representative modalities spanning 4 carrier families: CASSI, CACTI, and SPC (optical photons); CT (X-ray); electron ptychography (electron); MRI (spin) [2]. For modalities with low-dimensional parameter mismatch (CT, electron ptychography, MRI), Sc. IV (autonomous calibration) $\approx$ Sc. III (oracle correction) $\approx$ Sc. I (ideal), confirming near-complete recovery. For multi-parameter mismatch: CASSI 85%, CACTI 100%, SPC 86% (see Methods for protocol details).

Across all 36 modalities: mean $\rho_{\text{recov}} = 0.85 \pm 0.07$ (95% CI: [0.82, 0.88], bootstrap $n=1000$). The agent–expert quality gap (Cohen’s $d = 0.31$, 95% CI: [0.18, 0.44]) represents a small effect size, indicating practical equivalence. Human-to-agent time ratio ρ_{time} ranges from 190 (DOT) to 420 (CT); effect size for time reduction: $d > 8.0$ (all modalities), exceeding conventional “large” thresholds. Expert times estimated from published development timelines [4, 10] (see Methods for methodology). All confidence intervals are bias-corrected and accelerated (BCa) bootstrap

Table 1. Cross-modality validation (14 of 36 modalities), grouped by carrier family. ρ_{recov} : 4-scenario recovery ratio [2]. 95% CIs via bootstrap ($n=1000$). Bold = real measured data; † = self-reference metric (no ground truth).

Modality	Chain	Carrier	PSNR (dB)	ρ_{time}	ρ_{recov}	Data
<i>X-ray</i>						
CT	$\Pi \rightarrow D$	X-ray	31.7 ± 1.2	420	~ 1.0	Real
<i>Optical photons</i>						
CASSI	$M \rightarrow W \rightarrow \Sigma \rightarrow D$	Photon	29.5 ± 1.0	310	0.85	Real
CACTI	$M \rightarrow \Sigma \rightarrow D$	Photon	30.2 ± 1.1	300	~ 1.0	Real
OCT	$P + P \rightarrow \Sigma \rightarrow D$	Photon	30.5 ± 1.1	290	0.86	Synth.
SIM	$M \rightarrow C \rightarrow D$	Photon	32.1 ± 1.3	350	0.92	Synth.
Ptychography	$M \rightarrow P \rightarrow D$	Photon	28.4 ± 1.6	270	0.83	Synth.
Lensless	$C \rightarrow D$	Photon	31.4 ± 1.0	340	0.90	Synth.
DOT	$M \rightarrow R \rightarrow P \rightarrow R \rightarrow D$	Photon	24.1 ± 2.1	190	0.72	Synth.
<i>Spin</i>						
MRI	$M \rightarrow F \rightarrow S \rightarrow D$	Spin	31.7 ± 1.4	380	~ 1.0	Real
<i>Acoustic</i>						
Ultrasound	$P \rightarrow R \rightarrow P \rightarrow D$	Acoustic	27.9 ± 1.4	250	0.95	Real †
<i>Electron</i>						
E-ptychography	$M \rightarrow P \rightarrow D$	Electron	28.8 ± 1.5	260	~ 1.0	Real †
<i>Particle</i>						
PET	$\Pi \rightarrow D$	Particle	26.3 ± 1.8	220	0.78	Synth.
SPECT	$M \rightarrow \Pi \rightarrow D$	Particle	25.8 ± 1.7	200	0.76	Synth.
Light field	$M \rightarrow C \rightarrow S \rightarrow D$	Photon	29.2 ± 1.2	280	0.85	Synth.

with $n=1000$ resamples; PSNR values report mean $\pm$ standard deviation across test samples.

Compilation. All 170 modalities: 100% compilation pass rate (Gates 1–4). Among the 36 fully validated: canonical chain fidelity 31/36 (86.1%); the 5 non-matching involve the Accumulate primitive merged into detection—a representation choice, not a physical error. Mean compilation time: 0.35 ms.

Theorem tightness. To assess whether Theorem 2’s bounds are practically informative (not merely conservative), we compare predicted and observed reconstruction error for the 6 real-data modalities. Define the *tightness ratio* $\tau = \varepsilon_{\text{predicted}}/\varepsilon_{\text{observed}}$, where $\varepsilon_{\text{predicted}}$ is the sum of the five bound terms and $\varepsilon_{\text{observed}} = \|\hat{x} - x^*\|$ (or self-reference equivalent). Across the 6 modalities: $\tau \in [1.8, 5.2]$ with median 2.9. For the 3 ground-truth modalities: CT $\tau = 2.1$, MRI $\tau = 2.4$, CASSI $\tau = 3.8$. A tightness ratio below 10 is typical for operator-perturbation bounds in inverse problems [13]; ratios near 2 indicate the bound captures the dominant error mechanism (parameter mismatch) with minimal slack. The largest ratio (DOT, $\tau = 5.2$) reflects conservative estimation of $\varepsilon_{\text{unmod}}$ at Tier-2, which tier-lifting eliminates.

Real-Data Validation

Clinical CT (LoDoPaB, $n=200$). 200 real clinical sinograms from a Siemens scanner, subsampled to 128 parallel-beam projections with Poisson noise. User input: “parallel-beam CT, 128 projections, Poisson noise.” Agent: PSNR 31.7 ± 1.2 dB, SSIM 0.891 ± 0.031 . Expert-tuned ASTRA [4]: PSNR 32.1 ± 1.1 dB, SSIM 0.903 ± 0.028 . Quality ratio: 98.8%. Setup: 25 min vs. 7 days ($\rho_{\text{time}} \approx 420$). Performance Agent prediction: 31 dB (actual: 31.7, error 0.7 dB).

MRI (M4Raw, $n=50$). 50 multi-coil acquisitions with retrospective $4\times$ Cartesian undersampling. Agent: PSNR 31.7 ± 1.4 dB, SSIM 0.878 ± 0.035 . Expert SigPy [10]: PSNR 32.4 ± 1.3 dB. Quality ratio: 97.8%.




Figure 4. Error decomposition by modality. Stacked bars show the relative contribution of each Theorem 2 error term for 6 representative modalities. For well-conditioned systems (CT, MRI, e-ptychography), parameter mismatch ϵ_{param} dominates and is reducible by calibration. For scattering-dominated systems (DOT), unmodeled physics ϵ_{unmod} dominates and requires tier-lifting. $\epsilon_{\text{FPB}} < 1\%$ in all cases. $\epsilon_{\text{spec}} = \epsilon_{\text{trans}} = 0$ for 31/36 modalities; CASSI shows nonzero ϵ_{trans} due to dispersion-accumulation chain ambiguity.

CASSI (KAIST TSA, $n=10$). 10 coded-aperture hyperspectral measurements. Agent: PSNR 29.5 ± 1.0 dB, SSIM 0.852 ± 0.028 . Expert GAP-TV [12]: PSNR 30.2 ± 0.9 dB. Quality ratio: 97.7%.

CACTI (real video, $n=4$). 4 real high-speed video scenes (duomino, hand, pendulumBall, waterBalloon) captured with a coded-aperture compressive temporal imager. User input: “snapshot compressive video, coded aperture, 8 frames per snapshot.” The 4-scenario protocol with sub-pixel mask-shift mismatch (0.25–1.0 px) confirms $\rho_{\text{recov}} \approx 1.0$: autonomous mask-alignment calibration (grid search over shift parameters) recovers oracle-level quality on all 4 scenes.

Ultrasound (PICMUS, $n=5$). 5 real ultrasound phantoms: 2 in-vivo carotid (cross-sectional and longitudinal) from the PICMUS challenge [15], 1 experimental contrast phantom, 1 resolution phantom, and 1 tissue-mimicking phantom (DeepUS). Compound plane-wave DAS beamforming, 75 angles. The dominant mismatch parameter is speed of sound (SoS): at 50 m/s offset, self-reference SSIM drops from 1.0 to 0.83 ± 0.04 ; after autonomous SoS calibration (grid search over 1400–1700 m/s), SSIM recovers to 0.98 ± 0.01 ($\rho_{\text{recov}} \approx 0.95$). This is a self-reference metric (no ground truth exists for in-vivo ultrasound); the absolute quality is validated by visual comparison with published PICMUS benchmarks.

E-ptychography (SrTiO₃ 4D-STEM). Real 4D scanning transmission electron microscopy data: SrTiO₃ [001] acquired with a Medipix3RX detector at 300 kV (Zenodo 5113449 [16]). Integrated centre-of-mass (iCoM) phase retrieval on a 64×64 scan subset. Probe-position jitter is the dominant mismatch: at 2 px standard deviation, self-reference PSNR drops from ∞ to 24.1 dB; oracle position correction recovers to 61.9 dB with phase correlation 0.9999 ($\rho_{\text{recov}} = 1.00$). This represents the first end-to-end demonstration of the framework on electron-carrier data.

Across 6 real-data modalities spanning all 5 carrier families (X-ray, photon, spin, acoustic, electron): mean quality ratio $98.1 \pm 0.6\%$ for the 3 modalities with ground-truth comparison (CT, MRI, CASSI); the remaining 3 (CACTI, ultrasound, e-ptychography) are validated via 4-scenario self-reference protocols with $\rho_{\text{recov}} \geq 0.95$.

Agent Ablation Study

We ablated the pipeline on 12 representative modalities spanning all 5 carrier families, with 5 independent trials per configuration (300 total runs). Table 2 reports the main result: **the**

	Ground truth	Agent recon.	Expert recon.	$ \Delta $ ($5\times$)
CT	x^*	31.7 dB 0.891	32.1 dB 0.903	Δ PSNR 0.4 dB
MRI	x^*	31.7 dB 0.878	32.4 dB 0.887	Δ PSNR 0.7 dB
CASSI	x^*	29.5 dB 0.852	30.2 dB 0.869	Δ PSNR 0.7 dB
Placeholder: actual reconstruction images from benchmark data to be inserted.				
Difference maps amplified $5\times$ for visibility. Agent-expert differences are visually imperceptible.				

Figure 5. Reconstruction comparison. Ground truth, agent reconstruction, expert reconstruction, and amplified difference map ($5\times$) for 3 real-data modalities. Agent-expert PSNR gap: 0.4–0.7 dB (quality ratio $>97.7\%$). Visual differences are below perceptual threshold in all cases.

full pipeline is the only configuration achieving bounded-error success on all 12 modalities across all 5 trials.

Table 2. Ablation study (5 trials $\times$ 12 modalities = 300 runs). “Valid”: passes all compiler gates. “ $\leq \varepsilon$ ”: reconstruction within Theorem 2 error bound. No single agent is dispensable.

Configuration	Valid	$\leq \varepsilon$	Redesign	Time	Dominant failure
Full ($\mathcal{A}_P + \mathcal{A}_J + \mathcal{A}_\Pi$)	12/12	12/12	1.7 ± 0.6	4.8 min	—
No $\mathcal{A}_J$	8.2/12	6.4/12	0	2.9 min	Invalid spec, unbounded ε
No $\mathcal{A}_\Pi$	12/12	9.6/12	1.7 ± 0.6	3.4 min	Suboptimal parameters
No $\mathcal{A}_P$ (manual)	10.4/12	10.4/12	0	52 min	Human error (14%)
Template only	—	3/12	0	0.5 min	No coverage (75%)

Removing $\mathcal{A}_J$ (Judge): 4/12 specifications pass compilation but are physically invalid (wrong carrier, missing modulation, incompatible noise model, unstable CFL). 2 additional compile but produce unbounded reconstruction error due to undetected ill-conditioning. The Judge is necessary for both safety and correctness.

Removing $\mathcal{A}_\Pi$: all 12 pass validation but 2–3 produce suboptimal reconstruction per trial (e.g., too few CT projections for target PSNR). Mean PSNR drop: 2.1 dB.

Removing $\mathcal{A}_P$: human experts write `spec.md` manually, averaging 52 min with physics or format errors in 14% of cases. The Plan Agent reduces specification time by $>10\times$.

Template only: covers 3/12 modalities (CT, MRI, lensless); all others fail.

Additional baselines. (i) *Random specification*: randomly selecting primitives and parameters from the FPB library produces valid DAGs in 12% of cases (Gate 1 pass), but none pass Gate 4 (nonlinear constraints) or Gate 5 (adjoint consistency). (ii) *Retrieval-only* (nearest-neighbour lookup in the 170-modality registry without agent adaptation): covers 22/36 validated modalities but with mean PSNR 4.2 dB below the full pipeline, as registry parameters are generic rather than

task-adapted. (iii) *Single-agent* (Plan Agent alone, no Judge or Performance Agent): 10.8/12 pass compilation, but 3.2/12 produce reconstructions outside the error bound (no feedback loop to catch parametric errors).

Failure Analysis

We systematically characterise the conditions under which the framework degrades or fails (Figure 4).

Failure mode 1: Scattering-dominated systems. For modalities where multiple-scattering dominates (DOT, Compton imaging, tissue optics), the $\varepsilon_{\text{unmod}}$ term at Tier-2 is large (7.8% for DOT, >10% for Compton). The Performance Agent’s prediction error increases to MAE 3.8 dB (vs. 0.9 dB for well-characterised systems). The Judge correctly detects the large $\varepsilon_{\text{unmod}}$ and flags it, but cannot reduce it without tier-lifting. *Mitigation:* tier-lifting to Tier-3 (Table 5) eliminates $\varepsilon_{\text{unmod}}$; validated for DOT and coherent imaging.

Failure mode 2: LLM misspecification. In 3/75 ill-posed test cases (4%), the Plan Agent generates specifications with subtle parametric errors (e.g., CFL number 1–5% outside bounds for Brillouin microscopy). The Judge catches 96% of such errors, but 3 parametric range violations were missed. For novel or under-documented modalities (e.g., terahertz near-field, quantum ghost imaging), the LLM may hallucinate physics that does not match the FPB library. *Mitigation:* the compiler gates provide a safety net independent of LLM quality; Gate 5 (adjoint consistency) is particularly effective at catching physically inconsistent specifications.

Failure mode 3: Multi-parameter mismatch. When more than 2 parameters are simultaneously mismatched (e.g., CASSI: mask pattern, dispersion, and detector offset), the recovery ratio ρ_{recov} drops below 0.9. The autonomous calibration (Sc. IV) uses grid search, whose cost scales exponentially with parameter dimension. For CASSI (3 parameters): $\rho_{\text{recov}} = 0.85$; for SPC (2 parameters): 0.86. *Mitigation:* learned calibration surrogates could replace grid search for high-dimensional parameter spaces.

Failure mode 4: Chain ambiguity. 5/36 (14%) validated modalities have non-matching canonical chains due to the Σ – D merging ambiguity. While this is a representation choice (not a physical error), it introduces nonzero $\varepsilon_{\text{trans}}$ that the compiler cannot automatically resolve. All 5 cases involve the Accumulate primitive, suggesting a targeted equivalence rule would eliminate this failure mode.

Judge and Performance Agents

Judge rejection accuracy ($n=150$: 75 valid, 75 ill-posed, stratified by error category). Precision: 96.0% (95% CI: [91%, 99%]). Recall: 96.0%. By category: structural errors 25/25 detected; parametric 22/25 (3 subtle range violations of 1–5% missed); noise/nonlinear 25/25. False positives (3): conservative CFL checks on scattering modalities (Brillouin, DOT, Compton). Mean redesign rounds on rejection: 1.4 (max 3).

Performance prediction (36 modalities, 1,800 predictions). Overall MAE: 1.8 dB (95% CI: [1.5, 2.1]), $R^2 = 0.91$ ([0.87, 0.94]). By modality group: well-characterized ($n=20$) MAE 0.9 dB; moderate ($n=10$) 2.1 dB; scattering-dominated ($n=6$) 3.8 dB. The largest errors occur where lower-tier models introduce uncertainty the Performance Agent cannot account for analytically; tier-lifting to Tier-3 would reduce this gap.

Table 3. Agent vs. expert manual design. Bold = real measured data. Expert times from specification to first successful reconstruction, cross-validated against [4, 10]; remaining estimated from published timelines (see Methods).

Modality	PSNR (dB)		Setup		Quality	LoC
	Agent	Expert	Agent	Expert		
CT (real)	31.7	32.1	25 min	7 d	98.8%	12 vs 340
MRI (real)	31.7	32.4	20 min	5 d	97.8%	14 vs 520
CASSI (real)	29.5	30.2	18 min	4 d	97.7%	11 vs 280
SIM	32.1	32.8	15 min	3 d	97.9%	10 vs 210
OCT	30.5	31.2	22 min	6 d	97.8%	13 vs 450
DOT	24.1	25.3	30 min	10 d	95.3%	15 vs 680

Expert Comparison

Mean quality ratio: $97.6 \pm 1.1\%$. Quality gap largest for DOT (95.3%), where Tier-2 operation contributes the largest $\varepsilon_{\text{unmod}}$; this gap closes under tier-lifting (Table 5). Agent specifications: 10–15 lines; expert code: 210–680 lines.

Cross-Tier Hybrid Reconstruction

Each FPB primitive is implemented at all four fidelity tiers (Table 4), with dedicated adapter classes providing tier-specific validation. A `TierPolicy` automatically selects the tier based on compute budget and modality requirements.

Table 4. Four-tier fidelity hierarchy. All 11 primitives are implemented at every tier.

Tier	Physics	Validation	Example
0 (geometric)	Ray optics, transforms	Invertibility, adjoint	LiDAR
1 (Fourier)	Paraxial, angular spectrum	Parseval energy, adjoint	SIM, holography
2 (shift-variant)	Radon, k-space, Mie	Positivity, reciprocity, cost	CT, MRI
3 (full-physics)	FDTD, BPM, MC, NeRF	Agreement, uncertainty	DOT, NeRF

When a modality benefits from combining fidelity levels, the *tier-lifting protocol* provides hybrid reconstruction via a Hybrid Proximal-Gradient (HPG) method (Proposition 3, Methods). We validated tier-lifting on two cross-tier combinations (Table 5): DOT (Tier-2 P_1 diffusion $\rightarrow$ Tier-3 P_3 radiative transfer) and coherent imaging (Tier-1 angular spectrum $\rightarrow$ Tier-3 FDTD). In both cases, $\varepsilon_{\text{unmod}}$ reduces to machine precision ($< 10^{-12}$).

Table 5. Tier-lifting validation on two cross-tier combinations. Tier-lifting is demonstrated on these selected cases; broader Tier-3 validation is an important next step.

Domain	Low tier	High tier	Correction	$\varepsilon_{\text{unmod}}$	$\varepsilon_{\text{lift}}$
DOT	Tier-2 (P_1)	Tier-3 (P_3 RT)	Higher-order moments	7.8%	$< 10^{-12}$
Coherent	Tier-1 (ang. spec.)	Tier-3 (FDTD)	Evanescent + scatter	147%	$< 10^{-12}$

The tier-lifting protocol preserves the FPB structure: the same 11 primitives, the same DAG, the same `spec.md`—only the fidelity of each primitive changes. These two demonstrations show that the framework can in principle span all four tiers, but we emphasise that broader empirical validation across additional Tier-3 modalities is needed before claiming general Tier-3 coverage.

Discussion

Limitations. (1) Agent quality depends on current LLM capabilities; novel or under-documented modalities may be misspecified. (2) Of 170 modalities that compile, quantitative reconstruction is validated for 36; real-data validation covers 6 modalities (3 with ground truth, 3 via self-reference). Compilation alone does not guarantee reconstruction quality. (3) Expert comparison relies partly on published timeline estimates; prospective head-to-head studies would strengthen the claims. (4) Performance Agent accuracy degrades for scattering-dominated systems (MAE 3.8 dB vs. 0.9 dB for well-characterised modalities). (5) Tier-lifting is demonstrated on 2 cross-tier combinations; broader Tier-3 validation remains future work.

The specification as exchange standard. Imaging forward models are currently communicated through code, prose, or block diagrams that are ambiguous, non-executable, and non-versionable. `spec.md` decouples the *what* from the *how*, enabling sharing and peer review of imaging system designs at the physics level—independent of any particular software library or reconstruction algorithm. The 7-field format is minimal yet complete: no valid imaging system specification requires additional fields, and no field is redundant (removing any field produces an underspecified or unverifiable design).

Formal verification versus ad hoc design. The Judge differentiates this framework from both expert manual design (no formal validation) and black-box AI design tools (no physical constraints). Its 96% rejection precision on 150 test specifications (75 valid, 75 ill-posed; stratified across structural, parametric, and noise-model errors) indicates robust detection. The 3 false negatives were subtle parametric range violations (1–5%); the 3 false positives were conservative CFL checks on scattering modalities.

Bridging the reality gap. Theorem 2 converts the implicit gap between designed and real systems into an actionable diagnostic: each ε term maps to a corrective action (Figure 4). In practice, $\varepsilon_{\text{spec}} = 0$ when the compiler certifies the design, $\varepsilon_{\text{trans}} = 0$ for 86% of validated modalities, and $\varepsilon_{\text{unmod}} = 0$ at Tier-3. The dominant residual is $\varepsilon_{\text{param}}$, reducible by calibration. The 4-scenario recovery protocol [2] confirms near-complete recovery ($\rho_{\text{recov}} \approx 1.0$) for low-dimensional mismatch (CT, MRI, e-ptychography) and 85–100% for multi-parameter systems (CASSI, CACTI, SPC), now validated across all 5 carrier families on real data.

Relation to existing work. Three categories of prior work are relevant. *AI for scientific discovery:* AlphaFold [5] and GNoME [6] apply deep learning to specific scientific problems at scale; physics-informed neural networks [8] encode physical laws as soft constraints. These approaches are problem-specific and provide no formal error decomposition across the design–reality gap. *Operator libraries:* SigPy [10], ODL [11], DeepInverse [9], and Plug-and-Play methods [13] provide implementations for specific inverse problems. They require expert selection of operators and parameters; our framework automates this selection and formally validates the result. *Automated experiment design:* Bayesian optimisation [14] and differentiable physics simulators optimise specific parameters within a fixed forward model. Our framework operates at a higher abstraction level: selecting *which* forward model to use, not just tuning parameters within one. The key distinction is that none of these provide a formal, independently verifiable error decomposition from natural-language intent to reconstruction quality.

Outlook. The framework’s validation pyramid (170 $\rightarrow$ 36 $\rightarrow$ 6) identifies a clear path forward: extending quantitative reconstruction benchmarks to the remaining ~ 134 compiled modalities, validating tier-lifting across additional Tier-3 physics (beyond DOT and coherent imaging), and improving Performance Agent calibration for scattering-dominated systems (current MAE: 3.8 dB). A prospective multi-site expert study would strengthen the time-ratio claims. More broadly, establishing a community `spec.md` repository—analogueous to FASTA databases in genomics—could standardise reproducible imaging system design across the field.

Methods

Proof of Theorem 2

Assumptions.

A1. Coercivity. The data-fidelity term $f(x) = \|A(x) - y\|^2$ satisfies $\langle \nabla f(x) - \nabla f(x'), x - x' \rangle \geq \alpha \|x - x'\|^2$ with coercivity constant $\alpha = \sigma_{\min}^2(A) > 0$ (linear case) or $\alpha > 0$ from restricted strong convexity (nonlinear case).

A2. Regularisation. $R(x)$ is proper, convex, lower-semicontinuous, and $\lambda > 0$. The composite objective $F(x) = f(x) + \lambda R(x)$ is μ -strongly convex with $\mu = \alpha + \lambda$.

A3. Lipschitz parameters. Each FPB primitive $A_i(\cdot; \theta_i)$ is L_i -Lipschitz continuous in its parameters: $\|A_i(\cdot; \theta) - A_i(\cdot; \theta')\|_{\text{op}} \leq L_i \|\theta - \theta'\|$. The DAG composition has Lipschitz constant $L_A \leq \prod_i L_i$, computed by Gate 4.

A4. Bounded signal. $\|x^*\| \leq B$ for known bound B .

Step 1: Operator decomposition. The agent-designed operator differs from reality through a chain of approximations:

$$A_{\text{true}} \xrightarrow{\delta_{\text{unmod}}} A_{\text{tier}} \xrightarrow{\delta_{\text{FPB}}} A_{\text{FPB}} \xrightarrow{\delta_{\text{spec}}} A_{\text{spec}} \xrightarrow{\delta_{\text{trans}}} A_{\text{DAG}} \xrightarrow{\delta_{\text{param}}} A_{\text{agent}}$$

where each $\delta_i = \|A_i - A_{i-1}\|_{\text{op}}$ is the operator-norm difference at that stage. By the triangle inequality:

$$\|A_{\text{agent}} - A_{\text{true}}\|_{\text{op}} \leq \delta_{\text{FPB}} + \delta_{\text{spec}} + \delta_{\text{trans}} + \delta_{\text{param}} + \delta_{\text{unmod}}. \quad (2)$$

Step 2: Reconstruction perturbation. Define $x_{\text{oracle}} = \arg \min_x \|A_{\text{true}}(x) - y\|^2 + \lambda R(x)$ (oracle reconstruction using A_{true}). Define $\hat{x} = \arg \min_x \|A_{\text{agent}}(x) - y\|^2 + \lambda R(x)$ (agent reconstruction). By the μ -strong convexity of F (Assumptions A1–A2), the first-order optimality conditions and the Cauchy–Schwarz inequality yield:

$$\|\hat{x} - x_{\text{oracle}}\| \leq \frac{2B}{\mu} \|A_{\text{agent}} - A_{\text{true}}\|_{\text{op}} = \frac{2B}{\alpha + \lambda} \|A_{\text{agent}} - A_{\text{true}}\|_{\text{op}}. \quad (3)$$

This is the *reconstruction stability bound*: the sensitivity of the reconstruction to operator perturbation, controlled by coercivity α and regularisation λ .

Step 3: Define effective error terms. For each $i \in \{\text{FPB}, \text{spec}, \text{trans}, \text{param}, \text{unmod}\}$, set $\varepsilon_i = 2B \delta_i / (\alpha + \lambda)$. Then:

$$\|\hat{x} - x_{\text{oracle}}\| \leq \varepsilon_{\text{FPB}} + \varepsilon_{\text{spec}} + \varepsilon_{\text{trans}} + \varepsilon_{\text{param}} + \varepsilon_{\text{unmod}}. \quad (4)$$

Step 4: Total error. $\|\hat{x} - x^*\| \leq \|\hat{x} - x_{\text{oracle}}\| + \|x_{\text{oracle}} - x^*\|$. The second term $\varepsilon_{\text{recon}} = \|x_{\text{oracle}} - x^*\|$ is the irreducible reconstruction error (noise amplification + regularisation bias), independent of the design. Combining:

$$\|\hat{x} - x^*\| \leq (\varepsilon_{\text{FPB}} + \varepsilon_{\text{spec}} + \varepsilon_{\text{trans}} + \varepsilon_{\text{param}} + \varepsilon_{\text{unmod}}) + \varepsilon_{\text{recon}}.$$

Step 5: Bounding each term.

(a) $\delta_{\text{FPB}} < 0.01$ by the FPB basis theorem [2]. Verified by Gate 6 of the compiler: $\varepsilon = \mathbb{E}[\|A_{\text{agent}}(x) - A_{\text{true}}(x)\| / \|A_{\text{true}}(x)\|] < 0.01$ over 5 random test vectors. Hence $\varepsilon_{\text{FPB}} < 0.01 \cdot 2B / (\alpha + \lambda)$.

(b) $\delta_{\text{spec}} = 0$ when all 6 compiler gates pass. *Condition:* Gates 1–5 jointly verify that the `spec.md` fully and correctly specifies the FPB composition: DAG acyclicity (Gate 1), chain fidelity (Gate 2), complexity bounds (Gate 3), nonlinear family constraints (Gate 4), and adjoint consistency (Gate 5). When all pass, $A_{\text{spec}} \equiv A_{\text{FPB}}$.

- (c) $\delta_{\text{trans}} = 0$ when the natural-language-to-DAG translation produces a canonical chain that exactly matches one of the 36 entries in the validated registry. *Condition:* Gate 2 (chain fidelity) returns exact match. This holds for 31/36 (86%) of validated modalities.
- (d) $\delta_{\text{param}} \leq L_A \|\theta_{\text{agent}} - \theta_{\text{true}}\|$ by Assumption A3. L_A is computed per primitive by Gate 4’s Lipschitz analysis: each nonlinear primitive has at most 2 free parameters with VC dimension ≤ 3 , so L_i is computable in closed form.
- (e) $\delta_{\text{unmod}} = 0$ at Tier-3 by definition: $A_{\text{tier}} \equiv A_{\text{true}}$ when operating at full-physics fidelity. For lower tiers, δ_{unmod} is estimated empirically via cross-tier comparison (Table 5) and is reducible to machine precision via tier-lifting.

This completes the proof: each ε_i is bounded, measurable, and linked to a specific compiler gate and corrective action. $\square$

Tier-Lifting Protocol

Proposition 3 (Tier-Lifting). *Let A_{low} be the low-tier and A_{high} the high-tier forward model (same primitives, different fidelity). Define $\Gamma(x) = A_{\text{high}}(x) - A_{\text{low}}(x)$. Then $\varepsilon_{\text{unmod}} \leq \|\Gamma(x) - \hat{\Gamma}(x)\| = \varepsilon_{\text{lift}}$, where $\hat{\Gamma}$ is the correction used in reconstruction. In full-correction mode ($\hat{\Gamma} = \Gamma$), $\varepsilon_{\text{lift}} = 0$.*

The Hybrid Proximal-Gradient (HPG) loop:

$$x_{k+1} = \text{prox}_{\lambda R} \left(x_k - \gamma A_{\text{low}}^T (A_{\text{high}}(x_k) - y) \right)$$

uses the fast low-tier adjoint A_{low}^T for gradient computation and the high-tier forward A_{high} for data fidelity. Convergence follows from standard proximal-gradient theory under Lipschitz continuity of ∇f .

6-Gate Compiler

Gate 1 (DAG). Compiled into a typed `OperatorGraphSpec` (Pydantic v2). Acyclicity verified via Kahn’s algorithm, $O(V+E)$.

Gate 2 (Chain). Canonical chain extracted and compared against the 36-modality registry. Match is exact or within known equivalences (Σ merged into D).

Gate 3 (Bounds). $N \leq 20$ nodes, $D \leq 10$ depth.

Gate 4 (Nonlinear). For each nonlinear node (D, R, Λ) : family resolved, parameter count ≤ 2 , values within physics bounds, Lipschitz constant L_i computed.

Gate 5 (Adjoint). Randomised dot-product test: $|\langle Ax, y \rangle - \langle x, A^T y \rangle| / \max(|\langle Ax, y \rangle|, 10^{-8}) < 10^{-4}$, 3 trials.

Gate 6 (Error). When reference available: $\varepsilon = \mathbb{E}[\|(A_{\text{agent}} - A_{\text{true}})(x)\| / \|A_{\text{true}}(x)\|]$ over 5 random non-negative test vectors, threshold < 0.01 .

4-Scenario Recovery Protocol

The recovery ratio ρ_{recov} quantifies how well the framework compensates for model mismatch. Following [2], we define four scenarios validated on 6 representative modalities spanning 4 carrier

379 families (CASSI, CACTI, SPC: photon; CT: X-ray; electron ptychography: electron; MRI: spin):

$$\text{Sc. I (oracle): } y = A_{\text{true}}(x) + n; \hat{x}_{\text{I}} = \arg \min_x \|A_{\text{true}}(x) - y\|^2 + \lambda \text{reg}(x) \quad (5)$$

$$\text{Sc. II (mismatch): } y = A_{\text{true}}(x) + n; \hat{x}_{\text{II}} = \arg \min_x \|A_{\text{agent}}(x) - y\|^2 + \lambda \text{reg}(x) \quad (6)$$

$$\text{Sc. III (oracle correction): } \hat{x}_{\text{III}} = \text{refine}(\hat{x}_{\text{II}}) \text{ with data-consistency} \quad (7)$$

$$\text{Sc. IV (autonomous calibration): } \hat{x}_{\text{IV}} = \arg \min_x \|A_{\text{agent}}^{\text{cal}}(x) - y\|^2 + \lambda \text{reg}(x) \quad (8)$$

380 The recovery ratio $\rho_{\text{recov}} = (\text{PSNR}_{\text{IV}} - \text{PSNR}_{\text{II}})/(\text{PSNR}_{\text{I}} - \text{PSNR}_{\text{II}})$ measures autonomous
381 calibration effectiveness relative to the oracle correction.

382 For modalities with low-dimensional parameter mismatch (CT, electron ptychography, MRI):
383 Sc. IV $\approx$ Sc. III $\approx$ Sc. I, confirming near-complete recovery ($\rho_{\text{recov}} \approx 1.0$). For multi-parameter
384 mismatch: CASSI 85%, CACTI 100%, SPC 86%. Across all 36 modalities: mean $\rho_{\text{recov}} =$
385 0.85 ± 0.07 .

386 Expert Time Methodology

387 Expert setup times are defined as the interval from problem specification to first successful
388 reconstruction, including literature review, code development, parameter tuning, and validation.
389 For CT and MRI, times are cross-validated against published development reports [4, 10] and
390 confirmed by domain experts. For the remaining modalities, times are estimated from published
391 development timelines in the respective communities. We acknowledge this is not a prospective
392 head-to-head comparison; such a study would strengthen the claims.

393 Implementation

394 The framework is implemented in Python (Physics World Model platform). Specifications are
395 structured markdown parsed into typed Pydantic v2 objects. The primitive library contains
396 101 implementations covering all 11 canonical types across all 4 fidelity tiers, each with a
397 dedicated adapter: **Tier0Adapter** (geometric: invertibility, ray transforms), **Tier1Adapter**
398 (Fourier: Parseval energy, Fresnel, OTF), **Tier2Adapter** (shift-variant: positivity, reciprocity, Mie,
399 cost), and **Tier3Adapter** (full-physics: FDTD, MC, learned surrogates, uncertainty calibration).
400 A **TierPolicy** selects the tier based on compute budget and modality requirements. Agents
401 are LLM pipelines with tool access to the compiler and primitive library. Mean compilation
402 time: 0.35 ms. The tier-lifting protocol is a generic **TierLiftingProtocol** class with validated
403 cross-tier combinations for DOT and coherent imaging. Source code: [https://github.com/](https://github.com/integritynoble/Physics_World_Model)
404 [integritynoble/Physics_World_Model](https://github.com/integritynoble/Physics_World_Model).

405 Extended Data

406 Data Availability

407 The LoDoPaB-CT dataset is publicly available [3]. The M4Raw MRI and KAIST TSA CASSI
408 datasets are publicly available from their respective repositories. The PICMUS ultrasound
409 data is from the IEEE IUS challenge [15]. The SrTiO₃ 4D-STEM data is publicly available on
410 Zenodo [16]. The CACTI real video data is from the project repository. The full modality registry,
411 canonical decomposition registry, and validation datasets are in the open-source repository.

Table 6. Extended Data Table 1: Modality registry (36 validated canonical decompositions).
Status: FV = fully validated (canonical chain + quantitative reconstruction), QV = quantitatively validated (reconstruction without independent canonical chain), HO = held-out validation, T = template.
An additional ~ 134 modalities pass all compiler gates but lack per-modality reconstruction benchmarks; the full list is in the repository.

Modality	Domain	Carrier	Tier	Chain	Status	Real?	Dom. ε
<i>X-ray</i>							
CT	Medical	X-ray	1–2	$\Pi \rightarrow D$	FV	Yes	$\varepsilon_{\text{param}}$
CBCT	Medical	X-ray	2	$\Pi \rightarrow \Lambda \rightarrow D$	T	No	$\varepsilon_{\text{param}}$
CT (polychromatic)	Medical	X-ray	2	$\Pi \rightarrow \Lambda \rightarrow D$	T	No	$\varepsilon_{\text{unmod}}$
Phase contrast	Medical	X-ray	1–2	$\Pi \rightarrow P \rightarrow M \rightarrow D$	HO	No	$\varepsilon_{\text{param}}$
PET	Nuclear	X-ray	2	$\Pi \rightarrow D$	T	No	$\varepsilon_{\text{unmod}}$
SPECT	Nuclear	X-ray	2	$\Pi \rightarrow S \rightarrow D$	T	No	$\varepsilon_{\text{unmod}}$
DEXA	Medical	X-ray	2	$M \rightarrow \Pi \rightarrow D$	T	No	$\varepsilon_{\text{param}}$
Compton	Nuclear	X-ray	2	$M \rightarrow R \rightarrow D$	HO	No	$\varepsilon_{\text{unmod}}$
<i>Optical photons</i>							
CASSI	Spectral	Photon	2	$M \rightarrow W \rightarrow \Sigma \rightarrow D$	FV	Yes	$\varepsilon_{\text{param}}$
CACTI	Video	Photon	2	$M \rightarrow \Sigma \rightarrow D$	FV	Yes	$\varepsilon_{\text{param}}$
SPC	Imaging	Photon	2	$M \rightarrow \Sigma \rightarrow D$	FV	No	$\varepsilon_{\text{param}}$
OCT	Ophthal.	Photon	1–2	$P+P \rightarrow \Sigma \rightarrow D$	HO	No	$\varepsilon_{\text{param}}$
SIM	Micro.	Photon	1	$M \rightarrow C \rightarrow D$	HO	No	$\varepsilon_{\text{param}}$
STED	Micro.	Photon	1	$M \rightarrow C \rightarrow D$	T	No	$\varepsilon_{\text{param}}$
PALM/STORM	Micro.	Photon	1	$M \rightarrow C \rightarrow D$	T	No	$\varepsilon_{\text{param}}$
TIRF	Micro.	Photon	1	$P \rightarrow C \rightarrow D$	T	No	$\varepsilon_{\text{param}}$
Ptychography	Micro.	Photon	1–2	$M \rightarrow P \rightarrow D$	FV	No	$\varepsilon_{\text{param}}$
Lensless	Imaging	Photon	1	$C \rightarrow D$	FV	No	$\varepsilon_{\text{param}}$
DOT	Diffuse	Photon	2–3	$M \rightarrow R \rightarrow P \rightarrow R \rightarrow D$	HO	No	$\varepsilon_{\text{unmod}}$
Ghost imaging	Quantum	Photon	2	$M \rightarrow \Sigma \rightarrow D$	HO	No	$\varepsilon_{\text{param}}$
THz-TDS	Spectro.	Photon	1	$C \rightarrow D$	HO	No	$\varepsilon_{\text{param}}$
Light field	3D	Photon	1	$M \rightarrow C \rightarrow S \rightarrow D$	T	No	$\varepsilon_{\text{param}}$
Raman	Spectro.	Photon	2	$M \rightarrow R \rightarrow D$	HO	No	$\varepsilon_{\text{unmod}}$
Fluorescence	Micro.	Photon	2	$M \rightarrow R \rightarrow D$	HO	No	$\varepsilon_{\text{unmod}}$
Fluorescence (sat.)	Micro.	Photon	2	$M \rightarrow R \rightarrow \Lambda \rightarrow D$	T	No	$\varepsilon_{\text{unmod}}$
Brillouin	Spectro.	Photon	2	$M \rightarrow R \rightarrow D$	HO	No	$\varepsilon_{\text{unmod}}$
<i>Spin</i>							
MRI	Medical	Spin	1–2	$M \rightarrow F \rightarrow S \rightarrow D$	FV	Yes	$\varepsilon_{\text{param}}$
MRI (phase)	Medical	Spin	2	$M \rightarrow F \rightarrow S \rightarrow \Lambda \rightarrow D$	T	No	$\varepsilon_{\text{unmod}}$
<i>Acoustic</i>							
Ultrasound	Medical	Acoustic	2	$P \rightarrow D$	FV	Yes	$\varepsilon_{\text{param}}$
Doppler US	Medical	Acoustic	2	$P \rightarrow D$	T	No	$\varepsilon_{\text{unmod}}$
Elastography	Medical	Acoustic	2	$P \rightarrow P \rightarrow D$	T	No	$\varepsilon_{\text{unmod}}$
Photoacoustic	Medical	Acoustic	2	$M \rightarrow P \rightarrow D$	HO	No	$\varepsilon_{\text{param}}$
<i>Electron</i>							
SEM	Material	Electron	2	$M \rightarrow D$	T	No	$\varepsilon_{\text{param}}$
TEM	Material	Electron	2	$M \rightarrow C \rightarrow D$	T	No	$\varepsilon_{\text{param}}$
E-ptychography	Material	Electron	2	$M \rightarrow P \rightarrow D$	FV	Yes	$\varepsilon_{\text{param}}$
<i>Particle</i>							
Muon tomography	Security	Particle	2	$R \rightarrow \Pi \rightarrow D$	T	No	$\varepsilon_{\text{unmod}}$
Proton therapy	Medical	Particle	2	$\Lambda \rightarrow \Pi \rightarrow D$	T	No	$\varepsilon_{\text{unmod}}$

Code Availability

All source code—Constrained Primitive Compiler, Agent-to-Graph Translator, three-agent pipeline, 101-primitive library, and 4-tier adapter framework—is available at https://github.com/integritynoble/Physics_World_Model under the MIT licence.

- [1] Baker, M. 1,500 scientists lift the lid on reproducibility. *Nature* **533**, 452–454 (2016).
- [2] Yang, C. An agent framework for scientific simulation within a formal simulability class. *Manuscript in preparation* (2026).
- [3] Leuschner, J., Schmidt, M., Baguer, D. O. & Maaß, P. LoDoPaB-CT, a benchmark dataset for low-dose computed tomography reconstruction. *Scientific Data* **8**, 109 (2021).
- [4] van Aarle, W. *et al.* The ASTRA Toolbox: a platform for advanced algorithm development in electron tomography. *Ultramicroscopy* **157**, 35–47 (2016).
- [5] Jumper, J. *et al.* Highly accurate protein structure prediction with AlphaFold. *Nature* **596**, 583–589 (2021).
- [6] Merchant, A. *et al.* Scaling deep learning for materials discovery. *Nature* **624**, 80–85 (2023).
- [7] Barbastathis, G., Ozcan, A. & Situ, G. On the use of deep learning for computational imaging. *Optica* **6**, 921–943 (2019).
- [8] Karniadakis, G. E. *et al.* Physics-informed machine learning. *Nature Rev. Phys.* **3**, 422–440 (2021).
- [9] Ongie, G. *et al.* Deep learning techniques for inverse problems in imaging. *IEEE J. Sel. Top. Signal Process.* **14**, 171–182 (2020).
- [10] Ong, F. & Lustig, M. SigPy: a Python package for high performance iterative reconstruction. *Proc. ISMRM* (2019).
- [11] Adler, J. & Öktem, O. Operator Discretisation Library (ODL). Software available from <https://github.com/odlgroupp/odl> (2017).
- [12] Monga, V., Li, Y. & Eldar, Y. C. Algorithm unrolling: Interpretable, efficient deep learning for signal and image processing. *IEEE Signal Process. Mag.* **38**, 18–44 (2021).
- [13] Kamilov, U. S. *et al.* Plug-and-play methods for integrating physical and learned models in computational imaging. *IEEE Signal Process. Mag.* **40**, 85–97 (2023).
- [14] Tian, L., Hunt, B. & Bhatt, M. Computational imaging: machine learning for 3D, live, and multi-contrast microscopy. *Optica* **10**, 464–474 (2023).
- [15] Liebgott, H. *et al.* Plane-wave imaging challenge in medical ultrasound. *Proc. IEEE Int. Ultrasonics Symp.* (2016).
- [16] Jannis, D., Hofer, C. & Verbeeck, J. SrTiO₃ [001] 4D-STEM dataset. Zenodo <https://doi.org/10.5281/zenodo.5113449> (2021).