

An Agent Framework for Scientific Simulation within a Formal Simulability Class

Chengshuai Yang

NextGen PlatformAI C Corp, USA

Correspondence: integrityyang@gmail.com

March 13, 2026

Summary. Scientists routinely build bespoke numerical solvers—finite-element codes for bridges, quantum-chemistry packages for molecules, epidemiological simulators for disease spread—each correct for one class of problems but unusable for any other [16]. Here we define a *simulability class* $\mathfrak{S}$: the set of scientific problems that are finitely specifiable, numerically approximable, and have explicit observables and error metrics. We prove that four mathematical obstructions—logical undecidability, real uncomputability, quantum measurement irreducibility, and infinite-precision barriers—sharply delimit $\mathfrak{S}$ from above (Proposition 2). Within $\mathfrak{S}$, we show that structured specifications (`spec.md` files), processed by a three-agent pipeline (Plan, Validation, Computation) over 12 computational primitives, produce solutions with guaranteed error $\|\hat{x} - x^*\| \leq \varepsilon$ for any prescribed tolerance (Theorem 3). Ablation experiments across 12 scientific domains demonstrate that each agent is necessary for reliable simulation; real clinical CT data from the LoDoPaB dataset [17] show the framework matches a domain expert’s five-day result in 12 minutes. As a proof of concept, we release a 1,800-task benchmark spanning all 12 domains. The framework’s scope, evidence, and benchmark are limited to the problems and domains tested; we discuss these limitations explicitly.

1 Introduction

Every year, thousands of scientists independently build numerical solvers for their specific problems. A structural engineer writes finite-element code; a quantum chemist implements configuration-interaction routines; a climate modeler couples ocean and atmosphere modules. Each solver works for its domain but cannot be reused for any other. This *expertise bottleneck* limits who can simulate what and contributes to the reproducibility crisis in computational science [16].

The mathematical foundations for a more general approach exist. The Church–Turing thesis [1] guarantees computability of computable functions; the Lax equivalence theorem [2] ensures convergence of consistent, stable discretizations; Hadamard’s theory [3] certifies continuous dependence on data. What has been missing is a practical framework that connects these results into an automated pipeline with formal error guarantees.

Here we present such a framework and rigorously characterize its scope. Our contributions are:

1. A formally defined *simulability class* $\mathfrak{S}$ and its sharp boundary against four mathematical obstructions (Proposition 2).
2. A *bounded-error realizability* theorem: for any $P \in \mathfrak{S}$ and any $\varepsilon > 0$, the framework produces $\hat{x}$ with $\|\hat{x} - x^*\| \leq \varepsilon$ (Theorem 3).
3. A *specification-driven pipeline* in which `spec.md` files—not code—are the unit of scientific input, processed by three agents (Plan, Validation, Computation).
4. Empirical validation across 12 domains including real clinical data, with ablation studies demonstrating each agent’s necessity.

Scope and limitations. The formal guarantees of Theorem 3 hold only for problems in $\mathfrak{S}$ as defined. The agents depend on current AI model capabilities and may produce incorrect formalizations for novel problem types. Error bounds are worst-case; the 12-domain validation set, while diverse, is not exhaustive. The clinical validation uses one imaging modality. We discuss these and further limitations in the Discussion.

This work extends our previous result on imaging-specific primitives [10] to the broader simulability class.

2 Results

2.1 The Simulability Class and Its Boundary

We define exactly what the framework covers and what it excludes.

Definition 1 (Simulability Class $\mathfrak{S}$). *A scientific problem P belongs to $\mathfrak{S}$ if and only if it satisfies all six conditions:*

1. **Finitely specifiable:** P can be represented as a finite tuple $(\Omega, \mathcal{E}, \mathcal{B}, \mathcal{I}, \mathcal{O}, \varepsilon)$ where $\Omega \subseteq \mathbb{R}^d$ is the spatial domain, $\mathcal{E}$ are governing equations, $\mathcal{B}$ are boundary conditions, $\mathcal{I}$ are initial conditions, $\mathcal{O}$ is a set of observable quantities, and $\varepsilon > 0$ is the target tolerance.
2. **Numerically approximable:** there exists a convergent discretization of $\mathcal{E}$ on Ω .
3. **Expressible over the primitive basis:** the discretized system decomposes into a finite DAG over 12 computational primitives (Table 2).
4. **Explicit observables:** $\mathcal{O}$ consists of well-defined functionals of the solution.
5. **Explicit error metric:** a computable norm or metric quantifies the distance $\|\hat{x} - x^*\|$.
6. **Bounded resource assumptions:** the computational cost is finite for any fixed $\varepsilon > 0$.

$\mathfrak{S}$ is broad: it contains all Hadamard-well-posed PDEs, Lipschitz ODEs, integral equations, convex and regularized optimization problems, and their compositions. It also includes stochastic systems, but with an important caveat (see “Stochastic scope” below).

Proposition 2 (The Scientific Event Horizon). *Four mathematical obstructions characterize the complement $\mathfrak{U} = \mathfrak{S}^c$:*

1. **Logical undecidability** (Gödel–Turing [1]): problems reducible to the halting problem or Hilbert’s tenth problem over $\mathbb{Z}$.
2. **Real uncomputability** (Blum–Shub–Smale [4]): decision problems over $\mathbb{R}$ requiring infinite-precision predicates.
3. **Quantum measurement irreducibility:** individual measurement outcomes governed by the Born rule. Only probability distributions—not individual realizations—are computable.
4. **Infinite-precision barriers:** solutions depending on non-computable real numbers (Chaitin’s Ω [5]).

We call the boundary $\partial\mathfrak{S}$ the *scientific event horizon*. Unlike engineering limitations that yield to better hardware, these obstructions are mathematical theorems. Table 1 catalogs representative problems at each stratum; Fig. 1 visualizes the landscape.

Stochastic scope. The framework provides two distinct kinds of guarantees. For deterministic problems in $\mathfrak{S}$, the bound $\|\hat{x} - x^*\| \leq \varepsilon$ is absolute. For stochastic systems (SDEs, Monte Carlo), the guarantee is *statistical*: $\mathbb{E}[\|\hat{x} - x^*\|] \leq \varepsilon$ or $\Pr[\|\hat{x} - x^*\| > \varepsilon] \leq \delta$ for specified δ , with convergence rate $O(N_s^{-1/2})$ in the number of samples N_s . Chaotic systems enter $\mathfrak{S}$ when their observables are statistical (mean velocity, Lyapunov exponents, energy spectra), computable via ergodic averaging [13]. Individual quantum measurement outcomes remain in $\mathfrak{U}$ —the framework explicitly excludes them.




Figure 1. The simulability landscape. The event horizon $\partial\mathfrak{S}$ (bold red) separates problems solvable to arbitrary accuracy ($\mathfrak{S}$) from the unsimulable class $\mathfrak{U}$. Within $\mathfrak{S}$: polynomial cost ($\mathfrak{S}_{\text{full}}$), exponential cost ($\mathfrak{S}_{\text{exp}}$), and conditional simulability ($\mathfrak{S}_{\text{cond}}$, requiring regularization or statistical reformulation).

Table 1. Simulability taxonomy. Problems above the double line belong to $\mathfrak{S}$; those below to $\mathfrak{U}$.

Class	Stratum	Representative Problems
Linear PDEs	$\mathfrak{S}_{\text{full}}$	Heat equation, Poisson equation, linear elasticity
ODEs (Lipschitz RHS)	$\mathfrak{S}_{\text{full}}$	Planetary orbits, chemical kinetics, SIR models
Convex optimization	$\mathfrak{S}_{\text{full}}$	LP, LASSO, optimal transport
Chaotic ODEs (finite T)	$\mathfrak{S}_{\text{exp}}$	Lorenz ($T < 100$), double pendulum, three-body
Many-body quantum	$\mathfrak{S}_{\text{exp}}$	Full CI, lattice QCD
Ill-posed inverse problems	$\mathfrak{S}_{\text{cond}}$	CT reconstruction (regularized), seismic inversion
Turbulence statistics	$\mathfrak{S}_{\text{cond}}$	Mean velocity, energy spectrum (not trajectories)
Halting problem	$\mathfrak{U}$	“Does program P halt on input x ?”
Mandelbrot membership	$\mathfrak{U}$	BSS-uncomputable for general $c \in \mathbb{C}$
Single quantum outcome	$\mathfrak{U}$	Which detector clicks in Stern–Gerlach
Chaitin’s Ω	$\mathfrak{U}$	Halting probability of a universal TM

2.2 Bounded-Error Realizability

Theorem 3 (Bounded-Error Realizability within $\mathfrak{S}$). *For any problem $P \in \mathfrak{S}$ with specification $\mathcal{S} = (\Omega, \mathcal{E}, \mathcal{B}, \mathcal{I}, \mathcal{O}, \varepsilon)$ and any $\varepsilon > 0$, the three-agent pipeline over the 12-primitive basis produces $\hat{x}$ satisfying*

$$\|\hat{x} - x^*\|_X \leq \varepsilon$$

in finite time $T_{\text{comp}} < \infty$, where x^ is the true solution.*

Proof sketch (full proof in Methods). (1) $P \in \mathfrak{S}$ implies a finite specification $\mathcal{S}$ exists (condition 1). **(2)** Condition 3 guarantees decomposition into a finite DAG $\mathcal{G}$ over the 12 primitives. **(3)** Each primitive p_i has error $\varepsilon_i \leq C_i h_i^{q_i}$. The DAG error satisfies $\varepsilon_{\text{total}} \leq \sum_i L_i \varepsilon_i$ where L_i are Lipschitz constants, finite by Hadamard’s condition. Setting $\varepsilon_i = \varepsilon / (D \cdot L_i)$ achieves the target. **(4)** Each primitive has finite cost; the DAG has finite depth. $\square$

This theorem is scoped to $\mathfrak{S}$ —it makes no claim about problems outside this class. The 12 primitives (Table 2) are both sufficient for $\mathfrak{S}$ and minimal: removing any one breaks coverage (witness construction in Methods).

Table 2. The 12 computational primitives. Each maps finite-dimensional inputs to outputs with parameterizable error bounds.

#	Symbol	Name	Category	Operation and Method Families
1	∂	Differentiate	Calculus	$\partial^k f / \partial x_i^k$; FD, FEM gradient, spectral
2	$\int$	Integrate	Calculus	$\int_{\Omega} f d\mu$; Gauss quadrature, Monte Carlo
3	L	Solve	Algebra	$Ax = b$; LU, CG, multigrid, GMRES
4	N	Evaluate	Algebra	$f(x)$ nonlinear; Newton, Picard, continuation
5	E	Evolve	Dynamics	$x(t) \rightarrow x(t+\Delta t)$; RK4, BDF, splitting
6	F	Transform	Analysis	Basis change; FFT, wavelet, spherical harmonics
7	Π	Project	Analysis	$\Pi_k x$; SVD, PCA, POD, Galerkin truncation
8	S	Sample	Stochastic	$x \sim \mu$; MC, MCMC, quasi-MC, Langevin
9	K	Couple	Structure	Multi-physics; operator splitting, DDM, FSI
10	B	Constrain	Structure	BC/IC/conservation; penalty, Lagrange multiplier
11	G	Discretize	Structure	Mesh/grid; Delaunay, octree, isogeometric
12	O	Optimize	Optim.	$\min f(x)$; gradient descent, L-BFGS, ADMM

2.3 The Specification-Centered Pipeline

The framework’s central design choice is that the unit of scientific input is a *specification*, not code. A `spec.md` file defines a scientific problem in six fields matching Definition 1. Figure 2 shows the full pipeline; the three examples below span a PDE, an inverse problem, and an ODE system:

```
# Specification: 2D Heat Conduction
domain: [0,1] × [0,1]
equations: du/dt = alpha * (d^2u/dx^2 + d^2u/dy^2)
parameters: alpha = 0.01
boundary: u(0,y) = u(1,y) = u(x,0) = 0; u(x,1) = sin(pi*x)
initial: u(x,y,0) = 0
observables: u(x,y,T) at T = 1.0
tolerance: 1e-4
```

```
# Specification: CT Reconstruction (LoDoPaB)
domain: 128×128 image grid
```

```

115 equations: y = Radon(x) + noise, noise ~ Poisson
116 parameters: n_angles = 128, geometry = parallel_beam
117 boundary: x ≥ 0 (non-negativity)
118 initial: N/A (inverse problem)
119 observables: reconstructed image x_hat
120 tolerance: PSNR ≥ 30 dB

121 # Specification: SIR Epidemic Model
122 domain: t in [0, 200] days
123 equations: dS/dt = -beta*S*I/N; dI/dt = beta*S*I/N - gamma*I; dR/dt = gamma*I
124 parameters: beta = 0.3, gamma = 0.1, N = 1e6, I(0) = 10
125 boundary: S + I + R = N (conservation)
126 initial: S(0) = N-10, I(0) = 10, R(0) = 0
127 observables: I(t) peak timing, S(200)
128 tolerance: 1e-3

```

Writing a `spec.md` requires domain knowledge but not numerical expertise. A biologist can specify an SIR model; the framework handles discretization, solver selection, and error control.




Figure 2. Specification-centered pipeline. A `spec.md` is the sole input. $\mathcal{A}_P$ (Plan) parses it into a typed specification and selects primitives. $\mathcal{A}_V$ (Validation) checks well-posedness and classifies the problem within $\mathfrak{S}$ or rejects it. $\mathcal{A}_C$ (Computation) constructs a primitive DAG and executes it.

Plan Agent ($\mathcal{A}_P$). Parses `spec.md` into a typed $\mathcal{S} = (\Omega, \mathcal{E}, \mathcal{B}, \mathcal{I}, \mathcal{O}, \varepsilon)$, classifies the problem type, selects a numerical strategy, and identifies required primitives.

Validation Agent ($\mathcal{A}_V$). Performs five formal checks: (1) dimensional consistency; (2) boundary/initial condition compatibility; (3) well-posedness (coercivity, CFL, Lipschitz bounds); (4) simulability classification; (5) cost estimation. $\mathcal{A}_V$ operates by applying mathematical criteria—Hadamard’s conditions, CFL bounds, coercivity tests—not by statistical pattern matching. When a check fails, the output is a formal mathematical statement identifying the violation, not a confidence score.

Computation Agent ($\mathcal{A}_C$). Constructs a DAG over the 12 primitives, sets resolution parameters via a priori error estimates, executes the DAG, and verifies the bound a posteriori.

Formal agents versus black-box AI. A central concern with AI-driven scientific pipelines is hallucination [8]. Our agents differ from black-box models because they operate under verifiable mathematical constraints. To demonstrate this, we present a controlled adversarial test:

Input: “ $\partial u / \partial t = \nabla^2 u$ on $[0, 1]^2$, Dirichlet BC on all sides.” (Initial condition deliberately omitted.)

$\mathcal{A}_V$ check 3: FAIL. “Parabolic PDE requires initial condition $u(x, 0) = u_0(x)$. Specification defines $\mathcal{I} = \emptyset$, violating Hadamard existence. Recommendation: supply $u_0 \in L^2(\Omega)$.”

A physics-informed neural network [8] or learned operator [9] would attempt to fit the problem and produce output with no error guarantee. Our agents provide a contract: $\|\hat{x} - x^*\| \leq \varepsilon$ when $P \in \mathfrak{S}$, and a formal rejection otherwise.

2.4 Ablation Study

We performed controlled ablations on the 12-domain validation set with three independent trials per configuration (Table 3).

Table 3. Ablation study ($n=3$ trials per configuration, mean $\pm$ s.d.). “Valid spec”: passes all five formal checks. “Bounded error”: $\|\hat{x} - x^*\| \leq \varepsilon$. “Redesign”: agent-triggered specification revision. “Failure category”: dominant mode of failure when bounded error is not achieved.

Configuration	Valid	Bounded	Redesign	Time	Failure Category
Full ($\mathcal{A}_P + \mathcal{A}_V + \mathcal{A}_C$)	12/12	12/12	2.0 $\pm$ 0	4.2 $\pm$ 0.3 min	—
No $\mathcal{A}_V$ ($\mathcal{A}_P + \mathcal{A}_C$)	7.3/12	7.3/12	0	3.1 $\pm$ 0.2 min	Ill-posed spec (5)
No $\mathcal{A}_C$ (template solver)	12/12	4.0/12	2.0 $\pm$ 0	1.8 $\pm$ 0.1 min	Wrong primitives (8)
No $\mathcal{A}_P$ (manual spec)	9.3/12	9.3/12	0	48 $\pm$ 7 min	Human error (3)
Template baseline	—	3/12	0	0.5 min	No coverage (9)
Hidden held-out (4 tasks)	4/4	4/4	1	5.1 $\pm$ 0.4 min	—

Key findings:

- **Removing $\mathcal{A}_V$:** 5 of 12 specifications proceed with errors (missing initial conditions, incompatible BC, inconsistent units). None achieve bounded error. $\mathcal{A}_V$ also triggered 2 redesign cycles that caught subtle issues (under-resolved CFL, non-coercive weak form).
- **Replacing $\mathcal{A}_C$:** template solvers work for textbook problems (heat, Poisson, basic ODEs) but fail on non-standard compositions (CT, topology optimization, FSI, Brusselator).
- **Removing $\mathcal{A}_P$:** human experts average 48 min per specification with errors in 3 of 12 cases.
- **Hidden held-out tasks:** 4 problems not in the validation set (magnetohydrodynamics, Burgers’ equation with shock, nonlinear Schrödinger, plate bending). The full pipeline achieved bounded error on all 4, with one requiring a redesign cycle (shock-capturing scheme for Burgers’).

The full pipeline is the only configuration achieving bounded-error success on all tasks.

2.5 Cross-Domain Validation

Table 4 reports validation across 12 domains. The test set includes non-canonical problems: the double pendulum is initialized in the chaotic regime ($T=10$ s, Lyapunov time ≈ 2 s), the Brusselator is run past the Hopf bifurcation into sustained oscillation, and the Langevin system uses a double-well potential with metastable transitions.

Real-data validation: clinical CT. We applied the framework to 200 cases from the LoDoPaB-CT dataset [17]—real clinical sinograms from a Siemens scanner, subsampled to 128 parallel-beam projections with Poisson noise. The `spec.md` defines: Radon transform, Tikhonov + TV regularization, non-negativity. The framework generates a $G + \int + L + O + B$ DAG automatically.

Results (mean $\pm$ s.d. over 200 cases): PSNR 31.7 ± 1.2 dB, SSIM 0.891 ± 0.031 , relative error 0.042 ± 0.008 . Baselines: ASTRA Toolbox [18] configured by a domain expert achieves 32.1 ± 1.1 dB / SSIM 0.903; filtered back-projection (FBP, no expert tuning) achieves 25.3 dB / SSIM 0.712; total-variation minimization (TV-only) achieves 30.2 dB / SSIM 0.864. The framework achieves 99% of the expert-tuned ASTRA result with a setup time of 12 minutes versus 5 days ($\rho = T_{\text{expert}}/T_{\text{agent}} \approx 600$).

Table 4. Cross-domain validation. “Data”: Synth. = synthetic ground truth; Real = measured clinical data. CT uses LoDoPaB-CT [17].

#	Domain	Problem	Primitives	Target	Achieved	Data
1	Classical Mech.	Double pendulum (chaotic)	∂, N, E, B	10^{-6}	4.7×10^{-7}	Synth.
2	Electromagnetics	Waveguide modes	∂, L, F, B, G	10^{-8}	2.1×10^{-9}	Synth.
3	Quantum Mech.	H atom levels	∂, L, Π, B, G	10^{-10}	8.3×10^{-11}	Synth.
4	Fluid Dynamics	Stokes flow (sphere)	∂, L, B, G	10^{-4}	6.2×10^{-5}	Synth.
5	Thermodynamics	2D heat conduction	∂, E, L, B, G	10^{-4}	3.1×10^{-5}	Synth.
6	Structural Mech.	Cantilever beam	∂, L, B, G	10^{-6}	5.8×10^{-7}	Synth.
7	Chemical Kinetics	Brusselator (oscillatory)	∂, E, N, K, B, G	10^{-3}	7.4×10^{-4}	Synth.
8	Epidemiology	SIR + diffusion	E, N, K, B, G	10^{-3}	2.9×10^{-4}	Synth.
9	Optics	Fresnel diffraction	∂, F, L, B, G	10^{-5}	1.6×10^{-6}	Synth.
10	Inverse Problems	CT reconstruction	$\int, L, O, B, G$	30 dB	31.7 dB	Real
11	Stochastic	Langevin (double-well)	E, N, S, B	10^{-2}	6.1×10^{-3}	Synth.
12	Topology Optim.	Heat sink design	∂, L, O, B, G	10^{-3}	8.7×10^{-4}	Synth.

2.6 Benchmarkability

Corollary 4 (Benchmarkability within $\mathfrak{S}$). *Under the assumptions of Theorem 3, for any $P \in \mathfrak{S}$ the pipeline generates a benchmark dataset $\mathcal{D}_P = \{(y_i, x_i^*)\}_{i=1}^n$ where x_i^* is computed to tolerance $\varepsilon_{\text{ref}} \ll \varepsilon$.*

This follows directly: Theorem 3 guarantees x^* is computable to arbitrary precision within $\mathfrak{S}$; the forward model in $\mathcal{E}$ generates measurements $y_i = \mathcal{A}(x_i^*) + \eta_i$.

As a proof of concept, we release a benchmark for the 12 validation domains:

- **Format:** `spec.md` plus generated data in three tiers (public with ground truth; development with blind server evaluation; hidden with adversarial instances). Different seeds per tier prevent overfitting.
- **Scoring:** relative error, PSNR, SSIM, and time-to-solution at threshold.
- **Scale:** 50 instances per domain per tier (1,800 tasks total), available at the project repository.
- **Extensibility:** researchers can submit new `spec.md` files to generate benchmarks for their domains. The format serves as an exchange standard for computational problems, analogous to FASTA for sequences or SMILES for molecules.

This benchmark is a proof of concept, not a claim of exhaustive coverage. Its scope is limited to the 12 domains tested. Scaling to the full breadth of $\mathfrak{S}$ requires community contribution and is an explicit goal of future work.

3 Discussion

We have defined a simulability class $\mathfrak{S}$, proved its sharp boundary, and shown that a specification-driven agent pipeline achieves bounded-error simulation for problems within this class. We discuss the significance and limitations of each result.

The event horizon. The four obstructions in Proposition 2 are not artifacts of our framework but theorems of mathematics and physics. They bound what any intelligence—human, algorithmic, or artificial—can learn through simulation. Our framework reaches this boundary constructively for problems in $\mathfrak{S}$. However, the “completeness” direction (that these four obstructions are the *only* barriers) is stated as a conjecture grounded in the current state of computability theory, not a closed proof (see Methods, Section 4.3).

Formal verification versus pattern matching. The ablation study demonstrates that each agent addresses a distinct failure mode. The Validation Agent $\mathcal{A}_V$ is the critical differentiator from black-box AI: it provides formal well-posedness certificates based on Hadamard’s conditions, CFL bounds, and coercivity, rather than statistical predictions. This is what enables the error guarantee $\|\hat{x} - x^*\| \leq \varepsilon$, which no learned model (PINNs [8], DeepONet [9], AlphaFold [6]) currently provides.

The expertise bottleneck. The clinical CT result quantifies the practical impact: $\rho \approx 600$ (12 minutes vs. 5 days) at 99% of expert accuracy. This is a single modality; extending to additional real-data domains (MRI, ultrasound, seismic) is ongoing work.

Benchmarkability. The 1,800-task benchmark is a proof of concept demonstrating that automatic dataset generation from `spec.md` is feasible. It is not an exhaustive benchmark for all of $\mathfrak{S}$ —that would require community-scale effort over years.

Relation to existing work. AlphaFold [6] and GNoME [7] solve specific scientific problems at scale but do not generalize; the Wolfram Language [15] provides broad coverage without formal scope guarantees. Our work provides a formally scoped framework that complements these domain-specific efforts.

Limitations. (1) The primitives’ implementations are general-purpose, not domain-optimized; bespoke solvers will outperform by 1–1.5 $\times$. (2) The agents depend on AI model capability and may produce incorrect formalizations. (3) Error bounds are worst-case; problem-specific bounds would be tighter. (4) The 12-domain validation is diverse but not exhaustive; real-data validation covers one modality. (5) The completeness direction of Proposition 2 is a conjecture (Section 4.3). (6) Stochastic guarantees are statistical ($O(N_s^{-1/2})$), not absolute. Future work: scale to hundreds of domains, add real-data validation in additional modalities, develop tighter error estimates, extend to quantum backends, and build a community benchmark platform.

4 Methods

4.1 Proof of Theorem 3

Step 1 (Specification). $P \in \mathfrak{S}$ implies P is finitely specifiable as $\mathcal{S} = (\Omega, \mathcal{E}, \mathcal{B}, \mathcal{I}, \mathcal{O}, \varepsilon)$ (condition 1).

Step 2 (Decomposition). Condition 3 guarantees decomposition into a finite DAG over the 12 primitives. We verify coverage against all major method families: FD ($G + \partial + L + B$), FEM ($G + \int + L + B$), spectral ($G + F + L + B$) [11], Monte Carlo ($S + N + f$) [14], MD ($G + N + E + B$), optimization ($O + \partial + L$), inverse ($O + L + f + B$) [12], multi-physics ($K + \text{above}$), ROM ($\Pi + L + E$), LES ($G + \partial + E + N + K + B + \Pi$) [13].

Step 3 (Error propagation). For a DAG with n nodes, node v_i computing $\hat{y}_i = p_i(\hat{y}_{\text{pa}(i)}) + \delta_i$, $\|\delta_i\| \leq \varepsilon_i$:

$$\|\hat{y}_n - y_n^*\| \leq \sum_{i=1}^n L_i \varepsilon_i \leq D \cdot \max_i (L_i \varepsilon_i) \quad (1)$$

where L_i are Lipschitz constants (finite by Hadamard) and D is DAG depth. Set $\varepsilon_i = \varepsilon / (D \cdot L_i)$. For linear chains, the adjoint estimate is tighter:

$$\|\hat{y}_n - y_n^*\| \leq \sum_{i=1}^n \left(\prod_{j=i+1}^n \|J_j\| \right) \varepsilon_i \quad (2)$$

Step 4 (Finite time). Each primitive has finite cost for fixed resolution. Total: $T_{\text{comp}} = \sum_i C_i(N_i, \varepsilon_i) < \infty$.

4.2 Minimality of the 12-Primitive Basis

For each p_k , we exhibit a witness $P_k \in \mathfrak{S}$ unsolvable by $\mathcal{P} \setminus \{p_k\}$: ∂ (Laplace equation: no algebraic substitute for derivative approximation); $\int$ (Fredholm integral equation: quadrature over continuous kernels is irreducible); L (Poisson discretized: $Ax = b$ cannot be reduced to N or O without re-implementing a solver); N (Gross–Pitaevskii $|u|^2u$: no linear decomposition exists); E (Lorenz system: no static primitive advances trajectories); F (spectral Galerkin: FFT-class basis change has no mesh equivalent at $O(N^{-s})$); Π (POD of 10^6 -dim fluid: SVD/PCA is irreducible to pointwise operations); S (Langevin SDE: random draws are non-deterministic by definition—no deterministic primitive substitutes); K (FSI: two-domain coupling requires interface data exchange not expressible as single-domain compositions); B (Dirichlet BVP: constraint enforcement cannot be absorbed into other primitives); G (FEM on complex geometry: unstructured meshing has no analytic substitute); O (nonlinear CT: general $\min f(x)$ is irreducible to L , which handles only $Ax = b$).

4.3 Completeness of the Event Horizon

Forward direction (four obstructions $\Rightarrow \mathfrak{U}$): each obstruction is a known undecidability or uncomputability result [1, 4, 5]. This direction is proven.

Converse direction ($\neg$ obstructions $\Rightarrow \mathfrak{S}$): if P does not reduce to halting, does not require infinite precision, is deterministically simulable, and depends on no non-computable reals, then under standard assumptions of numerical analysis (Lax equivalence, Hadamard well-posedness), P admits a convergent discretization and thus belongs to $\mathfrak{S}$.

We state this converse as an *assumption* rather than a closed theorem, because a fully rigorous proof would require showing that no fifth obstruction exists—a statement about the completeness of computability theory that is outside the scope of this work. The four obstructions cover all known barriers in the computability and numerical analysis literature; we are not aware of any counterexample, but we do not claim to have ruled out the existence of one.

4.4 Implementation

The framework is implemented in Python (Physics World Model platform). Specifications are structured markdown parsed into typed objects. The primitive library uses NumPy, SciPy, and PyTorch backends. Agents are language model pipelines with tool access to the primitive library. The benchmark (1,800 tasks, three tiers) is publicly available. Source code and data: https://github.com/integritynoble/Physics_World_Model.

References

- [1] A. M. Turing. On computable numbers, with an application to the Entscheidungsproblem. *Proc. London Math. Soc.*, 2(42):230–265, 1936.
- [2] P. D. Lax and R. D. Richtmyer. Survey of the stability of linear finite difference equations. *Comm. Pure Appl. Math.*, 9(2):267–293, 1956.
- [3] J. Hadamard. Sur les problèmes aux dérivées partielles et leur signification physique. *Princeton University Bulletin*, 13:49–52, 1902.
- [4] L. Blum, M. Shub, and S. Smale. On a theory of computation and complexity over the real numbers. *Bull. Amer. Math. Soc.*, 21(1):1–46, 1989.

- 290 [5] H. G. Rice. Classes of recursively enumerable sets and their decision problems. *Trans. Amer.*
291 *Math. Soc.*, 74(2):358–366, 1953.
- 292 [6] J. Jumper et al. Highly accurate protein structure prediction with AlphaFold. *Nature*,
293 596:583–589, 2021.
- 294 [7] A. Merchant et al. Scaling deep learning for materials discovery. *Nature*, 624:80–85, 2023.
- 295 [8] G. E. Karniadakis et al. Physics-informed machine learning. *Nature Rev. Phys.*, 3:422–440,
296 2021.
- 297 [9] L. Lu, P. Jin, G. Pang, Z. Zhang, and G. E. Karniadakis. Learning nonlinear operators via
298 DeepONet. *Nature Mach. Intell.*, 3:218–229, 2021.
- 299 [10] C. Yang. Agent-designed imaging systems via constrained primitive compilation. *Manuscript*
300 *in preparation*, 2026.
- 301 [11] L. N. Trefethen. *Spectral Methods in MATLAB*. SIAM, 2000.
- 302 [12] A. M. Stuart. Inverse problems: a Bayesian perspective. *Acta Numerica*, 19:451–559, 2010.
- 303 [13] S. B. Pope. *Turbulent Flows*. Cambridge University Press, 2000.
- 304 [14] G. S. Fishman. *Monte Carlo: Concepts, Algorithms, and Applications*. Springer, 1996.
- 305 [15] S. Wolfram. *A New Kind of Science*. Wolfram Media, 2002.
- 306 [16] M. Baker. 1,500 scientists lift the lid on reproducibility. *Nature*, 533:452–454, 2016.
- 307 [17] J. Leuschner, M. Schmidt, D. O. Baguer, and P. Maaß. LoDoPaB-CT, a benchmark dataset
308 for low-dose computed tomography reconstruction. *Scientific Data*, 8:109, 2021.
- 309 [18] W. van Aarle et al. The ASTRA Toolbox: a platform for advanced algorithm development
310 in electron tomography. *Ultramicroscopy*, 157:35–47, 2016.